

# Database Processing

## CS 451 / 551

### Lecture 5: Searching and Indexing: Part 2



**Suyash Gupta**

Assistant Professor

Distopia Labs and ORNG

Dept. of Computer Science

(E) [suyash@uoregon.edu](mailto:suyash@uoregon.edu)

(W) [gupta-suyash.github.io](https://github.com/suyashgupta)



**Assignment 1 is Out!**  
**Deadline:** Oct 28, 2025 at 11:59pm

**Start collaborating with your groups!**

**Quiz 1:** Oct 16, 2025 (in class)

# Last Class

- We discussed sequential indexes: sparse, dense, multi-level.
- What are the challenges with these indexes?
  - A lot of file reorganization is needed when adding or deleting a record.
  - Can we avoid the reorganization? Yes, but
  - Then records are no longer mapped sequentially on the disk.
- Can we do better?

# How to determine a Good Index?

- A good index should help to search a record fast!
- Characteristics of a good index:
  - **Access Types:** Supports accessing a particular record (point query) and/or records within a specified range (range query).
  - **Access Time:** Time to find a particular record.
  - **Insertion Time:** Time to insert a new record in the index (includes time to find the right place to insert).
  - **Deletion Time:** Time to delete a new record in the index (includes time to find the item to be deleted).
  - **Space Overhead:** The space consumed by the index.

# A More desirable Index Structure

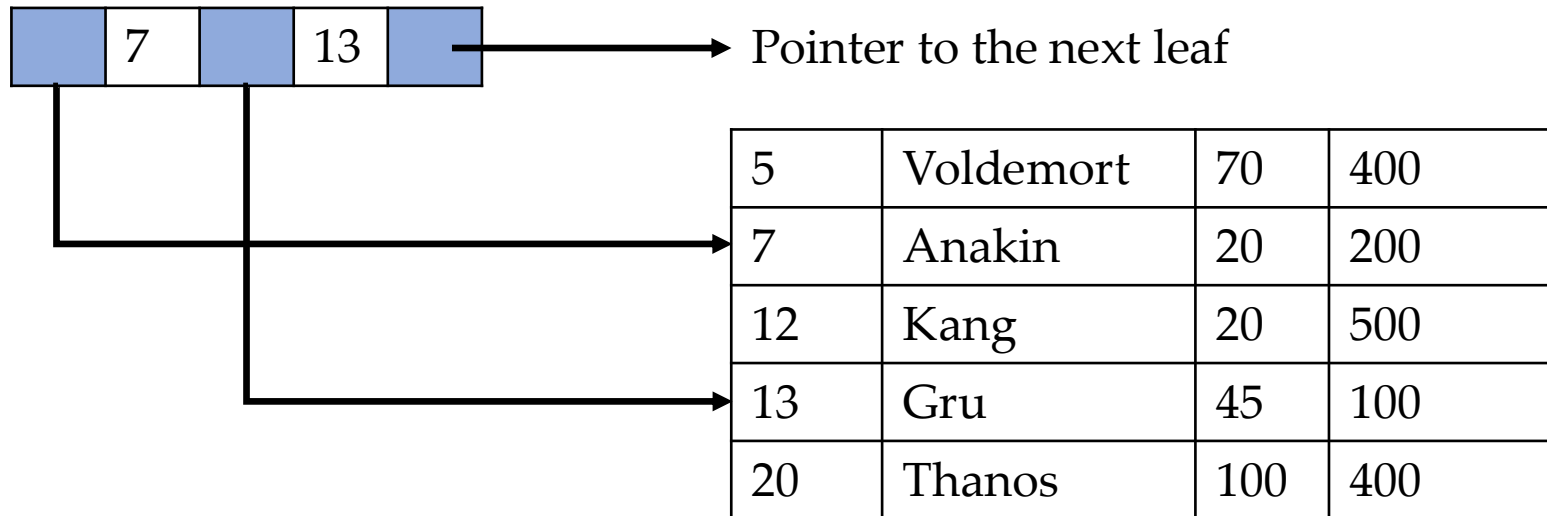
- Should ensure minimal reorganization.
- Should support sequential data access from disk.

# B<sup>+</sup>-Tree

# B<sup>+</sup>-Tree

- Another tree from the family of Balanced Trees.
- Three types of nodes: root, internal nodes, and leaf nodes.
- Every leaf node is at the same height.
- Give a value  $n$ , each internal node has:
  - $k$  children
  - $k - 1$  search keys
  - where,  $k$  is between  $\lceil n/2 \rceil$  to  $n$ .
- Root can have less than  $\lceil n/2 \rceil$  children but should have at least 2 children if there are more than one node in the tree.

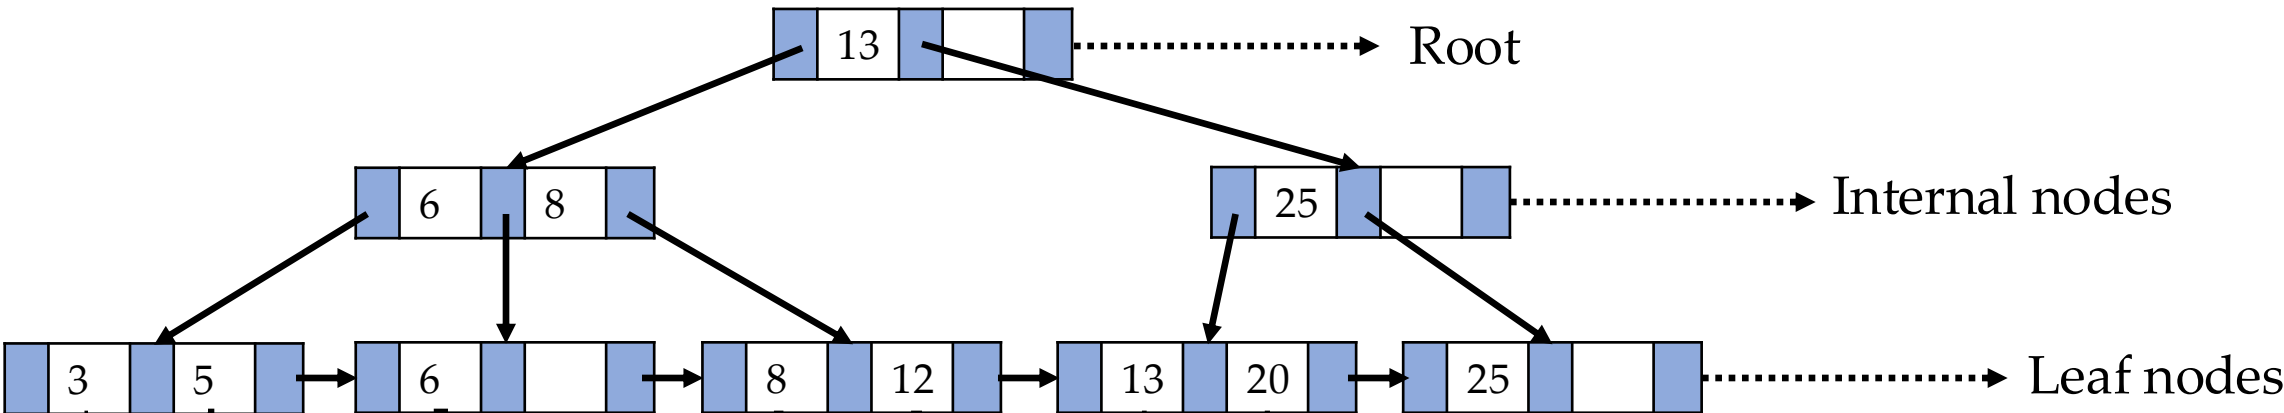
# B<sup>+</sup>-Tree Leaf Node Structure



Internal nodes also have similar structure, except they point to other tree nodes



# B<sup>+</sup>-Tree At a Glance



B<sup>+</sup> Tree with ID used for indexing

3	Joffrey	18	600
5	Voldemort	70	400
6	Scarecrow	30	150
8	Anakin	20	200
12	Kang	20	500
13	Gru	45	100
20	Thanos	100	400
25	Joker	66	200


File

# Searching data in B<sup>+</sup>-Tree

- Notice that the keys are stored in B+ tree in a **sorted manner**.
- We claim that the data is stored in B+ tree in sorted order because if you perform an in-order traversal, then you will get a sorted list.


# Searching data in B<sup>+</sup>-Tree

- Let's run an **in-order traversal**, where we will only output data in the leaf nodes.




# Searching data in B<sup>+</sup>-Tree

- Let's run an in-order traversal, where we will only output data in the leaf nodes.




# Searching data in B<sup>+</sup>-Tree

- Let's run an in-order traversal, where we will only output data in the leaf nodes.




# Searching data in B<sup>+</sup>-Tree

- Let's run an in-order traversal, where we will only output data in the leaf nodes.




# Searching data in B<sup>+</sup>-Tree

- Let's run an in-order traversal, where we will only output data in the leaf nodes.



# Searching data in B<sup>+</sup>-Tree


- Let's run an in-order traversal, where we will only output data in the leaf nodes.






# Searching data in B<sup>+</sup>-Tree

- Let's run an in-order traversal, where we will only output data in the leaf nodes.




# Searching data in B<sup>+</sup>-Tree

- Let's run an in-order traversal, where we will only output data in the leaf nodes.




# Searching data in B<sup>+</sup>-Tree

- Let's run an in-order traversal, where we will only output data in the leaf nodes.




# Searching data in B<sup>+</sup>-Tree

- Let's run an in-order traversal, where we will only output data in the leaf nodes.




# Searching data in B<sup>+</sup>-Tree

- Let's run an in-order traversal, where we will only output data in the leaf nodes.




# Searching data in B<sup>+</sup>-Tree

- Let's run an in-order traversal, where we will only output data in the leaf nodes.




# Searching data in B<sup>+</sup>-Tree

- Let's run an in-order traversal, where we will only output data in the leaf nodes.



# Searching data in B<sup>+</sup>-Tree


- Let's run an in-order traversal, where we will only output data in the leaf nodes.






# Searching data in B<sup>+</sup>-Tree

- Let's run an in-order traversal, where we will only output data in the leaf nodes.




# Searching data in B<sup>+</sup>-Tree

- Let's run an in-order traversal, where we will only output data in the leaf nodes.




# Searching data in B<sup>+</sup>-Tree

- Let's run an in-order traversal, where we will only output data in the leaf nodes.




# Searching data in B<sup>+</sup>-Tree

- Let's run an in-order traversal, where we will only output data in the leaf nodes.




# Searching data in B<sup>+</sup>-Tree

- Let's run an in-order traversal, where we will only output data in the leaf nodes.




# Searching data in B<sup>+</sup>-Tree

- Now, let's try to search a key → Say we want to search key 12.
  - We need to traverse the tree in the in-order fashion.
  - Stop traversing if one of the following three cases occur:
    - Key is found!
    - You encounter a Key greater than the search key.
    - You have reached the last key or leaf node of the tree.




# Searching data in B<sup>+</sup>-Tree

- Now, let's try to **search a key** → Say we want to **search key 12**.
  - We need to traverse the tree in the in-order fashion.
  - **Stop traversing** if one of the following three cases occur:
    - Key is found!
    - You encounter a Key greater than the search key.
    - You have reached the last key or leaf node of the tree.



# Searching data in B<sup>+</sup>-Tree


- Now, let's try to search a key → Say we want to search key 12.
  - We need to traverse the tree in the in-order fashion.
  - Stop traversing if one of the following three cases occur:
    - Key is found!
    - You encounter a Key greater than the search key.
    - You have reached the last key or leaf node of the tree.





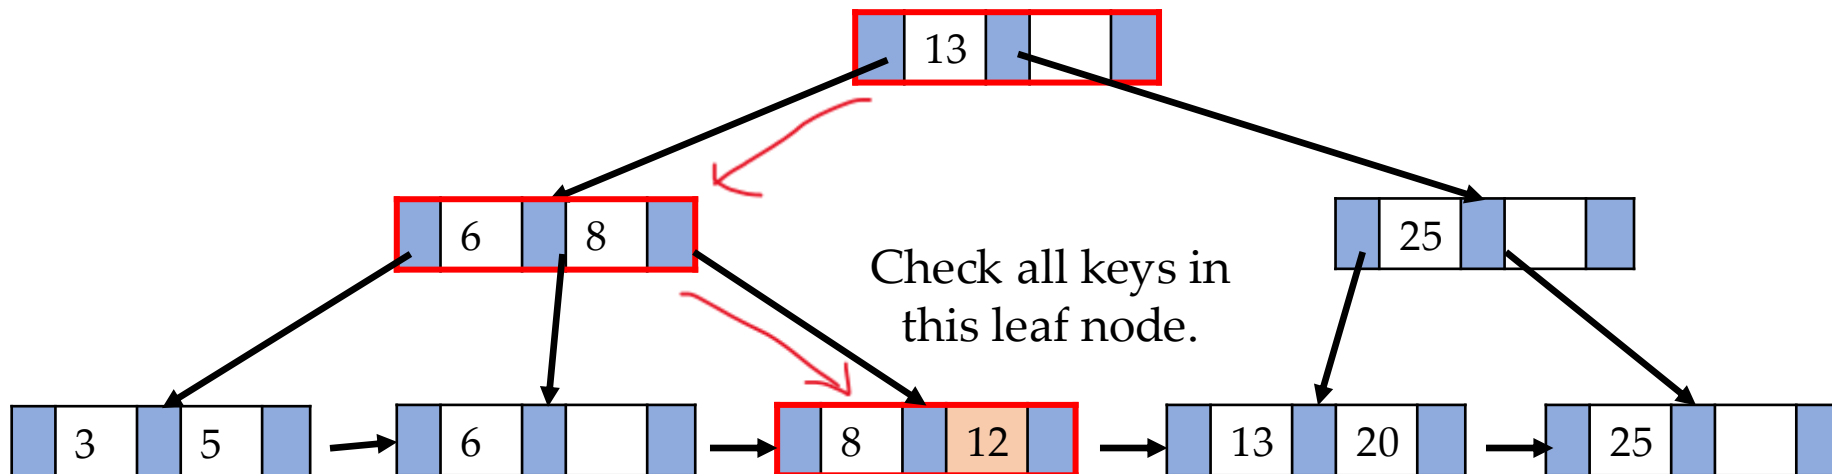
# Searching data in B<sup>+</sup>-Tree

- Now, let's try to search a key → Say we want to search key 12.
  - We need to traverse the tree in the in-order fashion.
  - Stop traversing if one of the following three cases occur:
    - Key is found!
    - You encounter a Key greater than the search key.
    - You have reached the last key or leaf node of the tree.




# Searching data in B<sup>+</sup>-Tree

- Now, let's try to search a key → Say we want to search key 12.
  - We need to traverse the tree in the in-order fashion.
  - Stop traversing if one of the following three cases occur:
    - Key is found!
    - You encounter a Key greater than the search key.
    - You have reached the last key or leaf node of the tree.




# Searching data in B<sup>+</sup>-Tree

- What we just did was a **Point Query**, where I wanted to search a specific item.
- Say we want to search a range of keys (**Range Query**) → **Keys from 12 to 24**.
  - We need to traverse the tree in the in-order fashion to reach the first key in the range, that is, first leaf node.
  - Then, perform linearly scan → follow the leaf pointers till you hit the last key or a key greater than the range.




# Searching data in B<sup>+</sup>-Tree

- What we just did was a **Point Query**, where I wanted to search a specific item.
- Say we want to search a range of keys (**Range Query**) → **Keys from 12 to 24**.
  - We need to traverse the tree in the in-order fashion to reach the first key in the range, that is, first leaf node.
  - Then, perform linearly scan → follow the leaf pointers till you hit the last key or a key greater than the range.




# Searching data in B<sup>+</sup>-Tree

- What we just did was a **Point Query**, where I wanted to search a specific item.
- Say we want to search a range of keys (**Range Query**) → **Keys from 12 to 24**.
  - We need to traverse the tree in the in-order fashion to reach the first key in the range, that is, first leaf node.
  - Then, perform linearly scan → follow the leaf pointers till you hit the last key or a key greater than the range.




# Searching data in B<sup>+</sup>-Tree

- What we just did was a **Point Query**, where I wanted to search a specific item.
- Say we want to search a range of keys (**Range Query**) → **Keys from 12 to 24**.
  - We need to traverse the tree in the in-order fashion to reach the first key in the range, that is, first leaf node.
  - Then, perform linearly scan → follow the leaf pointers till you hit the last key or a key greater than the range.




# Searching data in B<sup>+</sup>-Tree

- What we just did was a **Point Query**, where I wanted to search a specific item.
- Say we want to search a range of keys (**Range Query**) → **Keys from 12 to 24**.
  - We need to traverse the tree in the in-order fashion to reach the first key in the range, that is, first leaf node.
  - Then, perform linearly scan → follow the leaf pointers till you hit the last key or a key greater than the range.



# Searching data in B<sup>+</sup>-Tree


- What we just did was a **Point Query**, where I wanted to search a specific item.
- Say we want to search a range of keys (**Range Query**) → **Keys from 12 to 24**.
  - We need to traverse the tree in the in-order fashion to reach the first key in the range, that is, first leaf node.
  - Then, perform linearly scan → follow the leaf pointers till you hit the last key or a key greater than the range.






# Searching data in B<sup>+</sup>-Tree

- What we just did was a **Point Query**, where I wanted to search a specific item.
- Say we want to search a range of keys (**Range Query**) → **Keys from 12 to 24**.
  - We need to traverse the tree in the in-order fashion to reach the first key in the range, that is, first leaf node.
  - Then, perform linearly scan → follow the leaf pointers till you hit the last key or a key greater than the range.




# Searching data in B<sup>+</sup>-Tree

- What we just did was a **Point Query**, where I wanted to search a specific item.
- Say we want to search a range of keys (**Range Query**) → **Keys from 12 to 24**.
  - We need to traverse the tree in the in-order fashion to reach the first key in the range, that is, first leaf node.
  - Then, perform linearly scan → follow the leaf pointers till you hit the last key or a key greater than the range.




# Searching data in B<sup>+</sup>-Tree

- What we just did was a **Point Query**, where I wanted to search a specific item.
- Say we want to search a range of keys (**Range Query**) → **Keys from 12 to 24**.
  - We need to traverse the tree in the in-order fashion to reach the first key in the range, that is, first leaf node.
  - Then, perform linearly scan → follow the leaf pointers till you hit the last key or a key greater than the range.



# Searching data in B<sup>+</sup>-Tree

- What we just did was a **Point Query**, where I wanted to search a specific item.
- Say we want to search a range of keys (**Range Query**) → **Keys from 12 to 24**.
  - We need to traverse the tree in the in-order fashion to reach the first key in the range, that is, first leaf node.
  - Then, perform linearly scan → follow the leaf pointers till you hit the last key or a key greater than the range.



But 25 is out of range,  
So terminate search

# Search Complexity of B<sup>+</sup>-Tree

# Search Complexity of B<sup>+</sup>-Tree

- If each node can have  $n$  search keys (pointers), and total  $N$  records in the tree,
  - $O(\log_{n/2} N)$  is the length of the path.
- Ex: If  $n = 100$  and  $N = 1,000,000$ , only 4 nodes need to be accessed.
  - Only 4 blocks need to be read from disk.
- This is also an important **distinction between B<sup>+</sup>-tree and Binary trees**.
  - We can design B<sup>+</sup>-tree, where node size is large enough to be block size.
  - So one block fetch gives access to one node of the B<sup>+</sup>-tree.
- Notice that **root is most frequently accessed**.
  - Place it in your database buffer, which will save lookup cost.

# Insertions and Deletions in B<sup>+</sup>-Tree

# Insertions and Deletions in B<sup>+</sup>-Tree

- Insertions and Deletions are slightly more complex.
- You may need to **split a node** or **merge two nodes**.
- Split and merge operations can be avoided if there is a space, or you are not violating the B<sup>+</sup>-tree conditions.
- Remember; Give a value ***n***, each internal node has:
  - ***k*** children
  - ***k* – 1** search keys
  - where, ***k*** is between  **$\lceil n/2 \rceil$**  to ***n***.
- Lets look at a live demonstration.



# Insertions and Deletions in B<sup>+</sup>-Tree

30, 12, 56, **45**, 18, **16**, 10, **14**, 8, **6**, 90, 83, 67, 76, 49, 78,

56, 49, **67**, 83, 78, 90, 18, 30, 76,

# Insertions and Deletions Complexity

- If each node can have  $n$  search keys (pointers), and total  $N$  records in the tree,
  - $O(\log_{n/2} N)$  is the number of I/O operations needed.
- Notice that insertion and deletion complexity is still same as search!
- This is the worst case complexity, on average fewer I/O operations are required.

# Can we use B<sup>+</sup>-tree for File organization?

- Till now, we used B<sup>+</sup>-tree for designing an index for our file.
- How about we use it to even organize our files.
- The leaf nodes of the B<sup>+</sup>-tree can store actual records.
- If each leaf has same size as the disk block, then one disk block I/O fetches necessary records.

# Self Reading Task

- Difference between **B-tree** and **B<sup>+</sup>-tree**.
- Disadvantages of B-tree when compared to B<sup>+</sup>-tree?

# Special Indices?

- Often, some attributes have only a small set of possible values.
  - Course with grades Pass or Fail.
  - Daily Attendance: Present or Absent
- For some attributes, we can create groups for their values.
  - Faculty Title: Assistant Prof., Associate Prof., Professor
  - Salary Payscale: L1 ( < 100); L2 ( 100 - 300); L3 (300 – 500); L4 ( > 500)

# Bitmap Indices


- Often, some attributes have only a small set of possible values.
  - Course with grades Pass or Fail.
  - Daily Attendance: Present or Absent
- For some attributes, we can create groups for their values.
  - Faculty Title: Assistant Prof., Associate Prof., Professor
  - Salary Payscale: L1 ( < 100); L2 ( 100 - 300); L3 (300 – 500); L4 ( > 500)
- Constructing Bitmap indices is useful for such attributes.
- Each value is represented with the help of a bitmap.
  - Each record needs a sequential identifier.
  - The size of the bitmap is equal to number of records.
- One bitmap for each value!

# Bitmap Indices


- Assume that the following is our table:
  - We can construct bitmaps for Grade and Payscale.

ID	Name	Grade	Payscale
1	Voldemort	P	L1
2	Anakin	P	L2
3	Kang	F	L1
4	Gru	F	L2
5	Thanos	P	L4
6	Joker	F	L3
7	Jeoffrey	P	L1

Bitmap for Grade



Bitmap for Payscale



# Bitmap Indices

- Assume that the following is our table:
  - We can construct bitmaps for Grade and Payscale.

ID	Name	Grade	Payscale
1	Voldemort	P	L1
2	Anakin	P	L2
3	Kang	F	L1
4	Gru	F	L2
5	Thanos	P	L4
6	Joker	F	L3
7	Jeoffrey	P	L1

Bitmap for Grade

P 1 

--	--	--	--	--	--	--

F 0 

--	--	--	--	--	--	--

Bitmap for Payscale

L1 

1						
---	--	--	--	--	--	--

L2 

0						
---	--	--	--	--	--	--

L3 

0						
---	--	--	--	--	--	--

L4 

0						
---	--	--	--	--	--	--



# Bitmap Indices

- Assume that the following is our table:
  - We can construct bitmaps for Grade and Payscale.

ID	Name	Grade	Payscale
1	Voldemort	P	L1
2	Anakin	P	L2
3	Kang	F	L1
4	Gru	F	L2
5	Thanos	P	L4
6	Joker	F	L3
7	Jeoffrey	P	L1

Bitmap for Grade

P	1	1					
---	---	---	--	--	--	--	--

F	0	0					
---	---	---	--	--	--	--	--

Bitmap for Payscale

L1	1	0					
----	---	---	--	--	--	--	--

L2	0	1					
----	---	---	--	--	--	--	--

L3	0	0					
----	---	---	--	--	--	--	--

L4	0	0					
----	---	---	--	--	--	--	--

# Bitmap Indices

- Assume that the following is our table:
  - We can construct bitmaps for Grade and Payscale.

ID	Name	Grade	Payscale
1	Voldemort	P	L1
2	Anakin	P	L2
3	Kang	F	L1
4	Gru	F	L2
5	Thanos	P	L4
6	Joker	F	L3
7	Jeoffrey	P	L1

Bitmap for Grade

P 

1	1	0				
---	---	---	--	--	--	--

F 

0	0	1				
---	---	---	--	--	--	--

Bitmap for Payscale

L1 

1	0	1				
---	---	---	--	--	--	--

L2 

0	1	0				
---	---	---	--	--	--	--

L3 

0	0	0				
---	---	---	--	--	--	--

L4 

0	0	0				
---	---	---	--	--	--	--

# Bitmap Indices

- Assume that the following is our table:
  - We can construct bitmaps for Grade and Payscale.

ID	Name	Grade	Payscale
1	Voldemort	P	L1
2	Anakin	P	L2
3	Kang	F	L1
4	Gru	F	L2
5	Thanos	P	L4
6	Joker	F	L3
7	Jeoffrey	P	L1

Bitmap for Grade

<b>P</b>	1	1	0	0			
----------	---	---	---	---	--	--	--

<b>F</b>	0	0	1	1			
----------	---	---	---	---	--	--	--

Bitmap for Payscale

<b>L1</b>	1	0	1	0			
-----------	---	---	---	---	--	--	--

<b>L2</b>	0	1	0	1			
-----------	---	---	---	---	--	--	--

<b>L3</b>	0	0	0	0			
-----------	---	---	---	---	--	--	--

<b>L4</b>	0	0	0	0			
-----------	---	---	---	---	--	--	--

# Bitmap Indices

- Assume that the following is our table:
  - We can construct bitmaps for Grade and Payscale.

ID	Name	Grade	Payscale
1	Voldemort	P	L1
2	Anakin	P	L2
3	Kang	F	L1
4	Gru	F	L2
5	Thanos	P	L4
6	Joker	F	L3
7	Jeoffrey	P	L1

Bitmap for Grade

<b>P</b>	1	1	0	0	1		
----------	---	---	---	---	---	--	--

<b>F</b>	0	0	1	1	0		
----------	---	---	---	---	---	--	--

Bitmap for Payscale

<b>L1</b>	1	0	1	0	0		
-----------	---	---	---	---	---	--	--

<b>L2</b>	0	1	0	1	0		
-----------	---	---	---	---	---	--	--

<b>L3</b>	0	0	0	0	0		
-----------	---	---	---	---	---	--	--

<b>L4</b>	0	0	0	0	1		
-----------	---	---	---	---	---	--	--

# Bitmap Indices

- Assume that the following is our table:
  - We can construct bitmaps for Grade and Payscale.

ID	Name	Grade	Payscale
1	Voldemort	P	L1
2	Anakin	P	L2
3	Kang	F	L1
4	Gru	F	L2
5	Thanos	P	L4
6	Joker	F	L3
7	Jeoffrey	P	L1

Bitmap for Grade

<b>P</b>	1	1	0	0	1	0	
----------	---	---	---	---	---	---	--

<b>F</b>	0	0	1	1	0	1	
----------	---	---	---	---	---	---	--

Bitmap for Payscale

<b>L1</b>	1	0	1	0	0	0	
-----------	---	---	---	---	---	---	--

<b>L2</b>	0	1	0	1	0	0	
-----------	---	---	---	---	---	---	--

<b>L3</b>	0	0	0	0	0	1	
-----------	---	---	---	---	---	---	--

<b>L4</b>	0	0	0	0	1	0	
-----------	---	---	---	---	---	---	--

# Bitmap Indices

- Assume that the following is our table:
  - We can construct bitmaps for Grade and Payscale.

ID	Name	Grade	Payscale
1	Voldemort	P	L1
2	Anakin	P	L2
3	Kang	F	L1
4	Gru	F	L2
5	Thanos	P	L4
6	Joker	F	L3
7	Jeoffrey	P	L1

Bitmap for Grade

P 

1	1	0	0	1	0	1
---	---	---	---	---	---	---

F 

0	0	1	1	0	1	0
---	---	---	---	---	---	---

Bitmap for Payscale

L1 

1	0	1	0	0	0	1
---	---	---	---	---	---	---

L2 

0	1	0	1	0	0	0
---	---	---	---	---	---	---

L3 

0	0	0	0	0	1	0
---	---	---	---	---	---	---

L4 

0	0	0	0	1	0	0
---	---	---	---	---	---	---

# When are Bitmap Indices useful?

- Say we have the following query:

```
select * from cs_employees  
where grade = 'P';
```

- Is Bitmap index useful for this query?
- Not much,
  - You will scan the bitmap index.
  - For every record where grade is equal to P, you will fetch it from the disk.
  - So, you did not have to fetch every record.
  - However, records are stored sequentially in blocks on the disk, so you may end up fetching a lot of blocks with not required blocks!

# When are Bitmap Indices useful?

- Say we have another query:

```
select * from cs_employees
where grade = 'P' and payscale = 'L1';
```

- Is Bitmap index useful for this query?
- Significantly more,
  - We have bitmap indices on both grade and pay attributes.
  - So first, we will take an intersection of these bitmaps and then fetch!

Bitmap for Grade

P	1	1	0	0	1	0	1
---	---	---	---	---	---	---	---

Bitmap for Payscale

L1	1	0	1	0	0	0	1
----	---	---	---	---	---	---	---

Intersection Bitmap

1	0	0	0	0	0	1
---	---	---	---	---	---	---

So only two records  
are fetched!