

# Database Processing

## CS 451 / 551

### Lecture 6: Hashing



**Suyash Gupta**

Assistant Professor

Distopia Labs and ORNG

Dept. of Computer Science

(E) [suyash@uoregon.edu](mailto:suyash@uoregon.edu)

(W) [gupta-suyash.github.io](https://github.com/suyashgupta)



**Assignment 1 is Out!**  
**Deadline: Oct 28, 2025 at 11:59pm**

**Start collaborating with your groups!**

# Term Paper for Graduate Students

- Select one area.
- **Select one paper published** in 2025 from the following 4 conferences:
  - **No two students can select the same paper.**
  - Your selected paper **needs my approval.**
- VLDB, SIGMOD, OSDI, SOSP.
- Describe the following in 4-page style ACM Sigmod double-column style.
  - What is the paper's goal?
  - How is it meeting its goal?
  - What are the disadvantages of the proposed design and advantages of the proposed design?
  - Explain how can you improve the proposed design?
  - What architectural changes you need to do?
  - How to provide support for queries, say Natural Join?
- Topics:
  - Federated Learning, Vector Databases, Graph Databases, Privacy-Preserving Databases


# Unordered Indexing

- Until now, we studied ordered indexes, such as clustered indexes and trees.
- Next, we will look at unordered indexes → Hash indexes.

# Hashing

# Hashing

- Three key components of a hash index:
  - A **hash table**, which stores all the keys.
  - A **function** that helps to map the key to hash table.
  - An **hashing algorithm**



# Types of Hashing

# Types of Hashing

- Two types of hashing schemes:
  - **Static Hashing** → Size of hash map is fixed; cannot be increased.
  - **Dynamic Hashing** → Size of hash map can increase as needed.
    - Essentially as your databases increases over time, you can accommodate more data.



# Complexity of Hashing

- As hashed indexes are unordered, they do not force maintaining any specific order.
- The position of a key in the hash table is dictated by the hash function.
- **Average case complexity** for insertion, deletion, and search  $\rightarrow O(1)$ 
  - But, there are constants, which matter.
- **Worst case complexity**, given  $n$  keys  $\rightarrow O(n)$
- Hash tables support **random access**, unlike earlier indexes, which support sequential access.

# Static Hashing

- Say, we know that in our database there will be **5 records**.
- So, we select a hash function and create a **hash table (array) of size 5**.

13	Gru	45	100
----	-----	----	-----

$$h(x) = \textit{key} \% n = (13) \% n$$

Hash Table

0	
1	
2	
3	
4	


$n = 5$

File Storage

13	Gru	45	100
----	-----	----	-----


# Static Hashing

- Say, we know in our database there will be 5 records.
- So, we select a hash function and create a **hash table (array)** of size 5.



# Static Hashing

- Say, we know in our database there will be 5 records.
- So, we select a hash function and create a **hash table (array) of size 5**.



# Challenges for Static Hashing

# Challenges for Static Hashing

- Fixed number of Keys
- Duplicate Keys
- Collisions
- Disk Access Cost

# Challenges for Static Hashing

- **Fixed number of Keys** → You should know the total size of the database in the future, and it cannot grow any further!
- For example, this hash table can only store 5 keys and if in the future your database gets a **6<sup>th</sup> record**, you need to **reorganize** → **change hash table** → **too expensive!**

Hash Table

0	
1	
2	
3	
4	


$$n = 5$$

# Challenges for Static Hashing

- **Unique Keys** → How to store and search for duplicate keys?
- Hash function would map duplicate keys to the same location.
  - Overwrite pointer to existing record?
  - How do you search for an existing record with duplicate keys?

20	Scarecrow	30	200
----	-----------	----	-----

$(20) \% n$



File Storage

13	Gru	45	100
4	Voldemort	70	400
20	Thanos	100	400



# Challenges for Static Hashing

- **No Collisions** → Perfect hashing function that ensures there are no collisions.
- Hash function may end up assigning the same location to two or more records.

5	Jeoffrey	18	600
---	----------	----	-----

$(5) \% n$

0	
1	
2	
3	
4	

$n = 5$

File Storage

13	Gru	45	100
4	Voldemort	70	400
20	Thanos	100	400

# Challenges for Static Hashing

- **Disk Access Cost** and Lack of opportunities for **Pre-fetching**.
- Fetching a single record (point query) is fast. But, say I want to fetch a range of records. These records could be spread **across the disk → multiple blocks!**
- No longer sequential access. Moreover, **File Manager cannot even predict!**

# Design Decisions for Static Hashing

# Design Decisions for Static Hashing

- **Good Hash Function:**
  - Maps a large set of keys to a small array.
  - Dilemma b/w using a fast hash function vs. a hash function with low collisions.
- **Hashing Algorithm:**
  - How to handle key collisions when they occur?
  - Dilemma b/w allocating a large table to prevent collisions vs. setting up rules that allow storing duplicate and colliding keys!

# Hash Functions

- Given an input key, it return an integer representation of that key.
  - Essentially, you can use hash function to convert an arbitrary byte array into a fixed-length code.
- We want a hash function that is both **fast** and has a **low collision rate**.
- Notice that we are allowing collisions as we desire fast hashing!
- Alternatively, you can use a cryptographic hash function, like SHA256.
  - No collisions!
  - Extremely secure → NIST recommended
  - Extremely slow!

# Hash Functions

- Fortunately, we don't have to create a hash functions!
- [CRC-64](#) (1975)
  - Used in networks for error detection
- [MurmurHash](#) (2008)
  - Fast, general-purpose hash function.
- [Google CityHash](#) (2011)
  - Fast for keys of short length.
- [Facebook XXHash](#) (2012)
  - State-of-the-art
- [Google FarmHash](#) (2014)
  - Better version of CityHash; reduced collisions

# Hash Schemes Performance

If you want to test the performance of various hash functions, or play with different hash functions, check out [SMHasher](#).

# Static Hashing Algorithms

- We will be looking at two common algorithms:
  - **Linear Probe Hashing**
  - **Cuckoo Hashing**
- These algorithms are also termed as **open addressing**:
  - Essentially, the key may not be in the location where the hash function points.
- More advanced algorithms (**not part of this course**)
  - Robinhood hashing
  - Hopscotch hashing
  - Swiss Tables




# Linear Probe Hashing


# Linear Probe Hashing

- Simplest hashing algorithm → resolves collision by searching for next empty **slot**.
- Requires a **fixed-size giant array** (smaller the size, more collisions).
  - Hash table's **load factor** (like a threshold) determines when the **table is too full**.
  - No new key should be added, otherwise collisions → **allocate new table!**
- **Inserting a key:**
  - Use your hash function to find a **slot** (position).
  - If the location is empty, store the key in that slot.
  - Otherwise, start sequential scanning from that location.
  - When you find an empty slot, insert your key in that slot.
- **Deletion and Search:**
  - Same as insertion.


# Linear Probe Hashing



# Linear Probe Hashing



# Linear Probe Hashing



# Linear Probe Hashing

17	Anakin	45	300
----	--------	----	-----

$$(17) \% 7$$



Hash Table

0	7
1	
2	
3	10
4	
5	
6	

$$n = 7$$

File Storage

10	Gru	45	100
7	Voldemort	70	400

# Linear Probe Hashing

17	Anakin	45	300
----	--------	----	-----

$$(17) \% 7$$

Hash Table

0	7
1	
2	
3	10
4	
5	
6	

$$n = 7$$

File Storage

10	Gru	45	100
7	Voldemort	70	400

# Linear Probe Hashing

17	Anakin	45	300
----	--------	----	-----

$(17) \% 7$

Hash Table

0	7
1	
2	
3	10
4	17
5	
6	

$n = 7$

File Storage

10	Gru	45	100
7	Voldemort	70	400
17	Anakin	45	300



# Linear Probe Hashing

24	Joker	60	300
----	-------	----	-----

$$(24) \% 7$$



Hash Table

0	7
1	
2	
3	10
4	17
5	
6	

$$n = 7$$

File Storage

10	Gru	45	100
7	Voldemort	70	400
17	Anakin	45	300

# Linear Probe Hashing

24	Joker	60	300
----	-------	----	-----

$$(24) \% 7$$

Hash Table

0	7
1	
2	
3	10
4	17
5	
6	

$$n = 7$$

File Storage

10	Gru	45	100
7	Voldemort	70	400
17	Anakin	45	300

# Linear Probe Hashing

24	Joker	60	300
----	-------	----	-----

$$(24) \% 7$$

Hash Table

0	7
1	
2	
3	10
4	17
5	
6	

$$n = 7$$

File Storage

10	Gru	45	100
7	Voldemort	70	400
17	Anakin	45	300

# Linear Probe Hashing

24	Joker	60	300
----	-------	----	-----

$$(24) \% 7$$

Hash Table

0	7
1	
2	
3	10
4	17
5	24
6	

$$n = 7$$

File Storage

10	Gru	45	100
7	Voldemort	70	400
17	Anakin	45	300
24	Joker	60	300

# Linear Probe Hashing

5	Thanos	100	500
---	--------	-----	-----

$$(5) \% 7$$

Hash Table

0	7
1	
2	
3	10
4	17
5	24
6	

$$n = 7$$

File Storage

10	Gru	45	100
7	Voldemort	70	400
17	Anakin	45	300
24	Joker	60	300

# Linear Probe Hashing

5	Thanos	100	500
---	--------	-----	-----

$$(5) \% 7$$

Hash Table

0	7
1	
2	
3	10
4	17
5	24
6	

$$n = 7$$

File Storage

10	Gru	45	100
7	Voldemort	70	400
17	Anakin	45	300
24	Joker	60	300

# Linear Probe Hashing

5	Thanos	100	500
---	--------	-----	-----

$(5) \% 7$

Hash Table

0	7
1	
2	
3	10
4	17
5	24
6	5

$n = 7$

File Storage

10	Gru	45	100
7	Voldemort	70	400
17	Anakin	45	300
24	Joker	60	300
5	Thanos	100	500

# Searching in Linear Probe Hashing

- Follow the same algorithm as you are trying to insert.
  - If the slot is empty, key not found.
  - If the slot is full, then continue to next slot.
  - Stop when you reach an empty slot or have covered all the slots.



# Deleting in Linear Probe Hashing

- How can we delete a record?
- Say, we want to delete the **record 10**, which maps to **slot 3**.

10	Gru	45	100
----	-----	----	-----

$(10) \% 7$

Hash Table

0	7
1	
2	
3	10
4	17
5	24
6	5

$n = 7$

# Deleting in Linear Probe Hashing

- How can we delete a record?
- Say, we want to delete the **record 10**, which maps to **slot 3**.
- Can we set slot 3 to **empty**?

10	Gru	45	100
----	-----	----	-----

$(10) \% 7$

Hash Table

0	7
1	
2	
3	
4	17
5	24
6	5

$n = 7$

# Deleting in Linear Probe Hashing

- On deleting a record, setting a slot to empty is **dangerous!**
  - Other keys could have also mapped to the same slot, but due to the slot being full, they were in subsequent locations.
  - By emptying the slot, you are indicating that other keys also do not exist!
- Two possible solutions:
  - Rearrangement
  - Tombstones

# Deletions: Rearrangement

- Once a key is deleted, you can rehash all the keys again.
- Any key that was supposed to be mapped to the same slot can now take place.
- **Too expensive! No database does this.**

Hash Table

0	7
1	
2	
3	10
4	17
5	24
6	5

$n = 7$



Hash Table

0	7
1	
2	
3	17
4	24
5	5
6	

$n = 7$

# Deletions: Tombstones

- Once a key is deleted, you place a **tombstone** for that key in that slot.
- Tombstone informs any future query that the specific key does not exist.
- However, other keys may still exist!


Hash Table

0	7
1	
2	
3	10
4	17
5	24
6	5

$n = 7$



Hash Table

0	7
1	
2	
3	
4	17
5	24
6	5

$n = 7$

For each tombstone, you need to maintain the list of keys that have been deleted!

# Duplicate Keys in Linear Probe Hashing

- How do you handle **duplicate (non-unique)** keys?
- Two ways:
  - Maintain a list of values
  - Just simply allow adding redundant keys

# Duplicate Keys: List of Values

Hash Table

0	7
1	
2	
3	10
4	
5	5
6	

$n = 7$

Lists of Values

10	Gru	45	100
10	Voldemort	70	400

5	Anakin	45	300
5	Joker	60	300
5	Thanos	100	500

# Duplicate Keys: Allow Redundant Keys

Hash Table

0	7
1	5
2	
3	10
4	5
5	5
6	10


$$n = 7$$



# Cuckoo Hashing


- Why the name cuckoo?
- Like the bird cuckoo, if we do not find a free slot for a key, we may kick out an existing key!
- In cuckoo hashing, we use multiple **hash functions** to find free slots to store the key.
  - Each hash function may give us a slot to place and if any of those slots is free, we store the key!
- If no slot is free, evict an existing key!

# Cuckoo Hashing: Insertion



# Cuckoo Hashing: Insertion

10	Gru	45	100
----	-----	----	-----



Randomly select a slot, say it selects slot 3 for storing key 10.

# Cuckoo Hashing: Insertion

10	Gru	45	100
----	-----	----	-----

$h1(10)$

$h2(10)$

Hash Table

0	
1	
2	
3	10
4	
5	
6	

$n = 7$

# Cuckoo Hashing: Insertion

5	Anakin	25	400
---	--------	----	-----

$h1(10)$

$h2(10)$

$h1(5)$

$h2(5)$

Hash Table

0	
1	
2	
3	10
4	
5	
6	

$n = 7$

# Cuckoo Hashing: Insertion

5	Anakin	25	400
---	--------	----	-----

$h1(10)$

$h2(10)$

$h1(5)$

$h2(5)$

Hash Table

0	
1	
2	
3	10
4	
5	5
6	

$n = 7$

As slot 3 is occupied, we select slot 5 to store key 5.

# Cuckoo Hashing: Insertion

18	Joker	66	300
----	-------	----	-----

$h1(10)$

$h2(10)$

$h1(5)$

$h2(5)$


$h1(18)$

$h2(18)$

Hash Table

0	
1	
2	
3	10
4	
5	5
6	

$n = 7$



# Cuckoo Hashing: Insertion

18	Joker	66	300
----	-------	----	-----

$h1(10)$

$h2(10)$

$h1(5)$

$h2(5)$

$h1(18)$

$h2(18)$

Hash Table

0	
1	
2	
3	10
4	
5	5
6	

$n = 7$

Say, it decides to kick key 5



# Cuckoo Hashing: Insertion

18	Joker	66	300
----	-------	----	-----

$h1(10)$

$h2(10)$

$h1(5)$

$h2(5)$

$h1(18)$

$h2(18)$

Hash Table

0	
1	
2	
3	10
4	
5	18
6	

$n = 7$

Say, it decides to kick key 5

# Cuckoo Hashing: Insertion

18	Joker	66	300
----	-------	----	-----

$h1(10)$

$h2(10)$

$h1(5)$

$h2(5)$

$h1(18)$

$h2(18)$

$h1(5)$

Hash Table

0	
1	
2	
3	10
4	
5	18
6	

$n = 7$

So, we need to rehash key 5,  
and only remaining slot is  
the slot occupied by key 10

# Cuckoo Hashing: Insertion

18	Joker	66	300
----	-------	----	-----

$h1(10)$

$h2(10)$

$h1(5)$

$h2(5)$

$h1(18)$

$h2(18)$

$h1(5)$

Hash Table

0	
1	
2	
3	5
4	
5	18
6	

$n = 7$

So, kick out key 10!

# Cuckoo Hashing: Insertion

18	Joker	66	300
----	-------	----	-----

**$h1(10)$**

**$h2(10)$**

**$h1(5)$**

**$h2(5)$**

**$h1(18)$**

**$h2(18)$**

**$h1(5)$**

**$h1(10)$**

Hash Table

0	10
1	
2	
3	5
4	
5	18
6	

**$n = 7$**

And now, rehash key 10

# Challenges with Cuckoo Hashing

- So what are the challenges with cuckoo hashing?
- Insertions are **expensive** → We need to do rehashing!
- We can get stuck into an **infinite loop**.
  - To exit the infinite loop, add more hash functions, or increase size of table, or maintain some list.

# Dynamic Hashing

- The biggest challenge for static hashing remains to be **fixed size of hash table**.
- Alternatively, use dynamic hashing algorithms:
  - Chained Hashing
  - Extensible Hashing
  - Linear Hashing

# Chained Hashing

# Chained Hashing

- For each slot in the hash table, there is a **linked list of buckets**.
- Essentially, collisions are resolved by **placing all keys with the same slot** into same linked list.
- Searching for a key requires scanning the linked list till you find the key or have reached end of the list.



# Chained Hashing

Simple hash function

$$h(key) = key \% n = (key) \% 7$$

Hash Table

0	
1	
2	
3	
4	
5	
6	

$$n = 7$$

# Chained Hashing

10	Gru	45	100
----	-----	----	-----

$(10) \% 7$

Hash Table

0	
1	
2	
3	
4	
5	
6	

$n = 7$


10

# Chained Hashing

7	Voldemort	70	400
---	-----------	----	-----

$(7) \% 7$

Hash Table




# Chained Hashing

17	Anakin	45	300
----	--------	----	-----

$(17) \% 7$

Hash Table




# Chained Hashing

24	Joker	60	300
----	-------	----	-----

**(24)% 7**

Hash Table



# Searching in Chained Hashing

- Use the hash function to reach the specific slot, and then scan the linked list till you find the key or have reached end of the list.
- For example, **on searching 17**, you would first reach **slot 3**, and then scan the list for **slot 3** and find it is as the **second entry in the linked list**.

# Challenges with Chained Hashing

- What is the key challenge with chained hashing?
- If a lot of keys are hashed to the same slot, then
  - You have a massively large linked list, and
  - Searching a key comes expensive → same cost as linear scan.


# Extensible Hashing

- Solves the problem of massively large linked lists.
- Requires linked lists to be split, when size crosses a threshold.
- Requires observing each key in a bit format.
- When you hash a key, you get a numeric (base-10 or base-16) representation.
  - You can convert that base-10 to binary format (base-2).
- For example: 4 can be represented as 100 in a 3-bit representation.




# Extensible Hashing

- Initially, your hash map is **1-bit**, and you have some fixed number of buckets for each bit → Say 3 buckets.




# Extensible Hashing

- Assume on passing the **key 13** through a **hash function**, the binary representation is **00011**.




# Extensible Hashing

- Another **key 7**, after passing it through a **hash function**, let the binary representation be **10011**.




# Extensible Hashing

- This way, all keys with binary representation starting from 0 go to buckets for bit 0, and vice versa for buckets for bit 1.




# Extensible Hashing

- Let's assume all the buckets for **bit 1** are full.




# Extensible Hashing

- So, now we need to **split the buckets** for bit 1. This will require expanding the bit representation from 1-bit to 2-bits.




# Extensible Hashing

- Notice that all the **2-bit** representations **starting with bit-0** continue pointing to the old buckets.



# Extensible Hashing

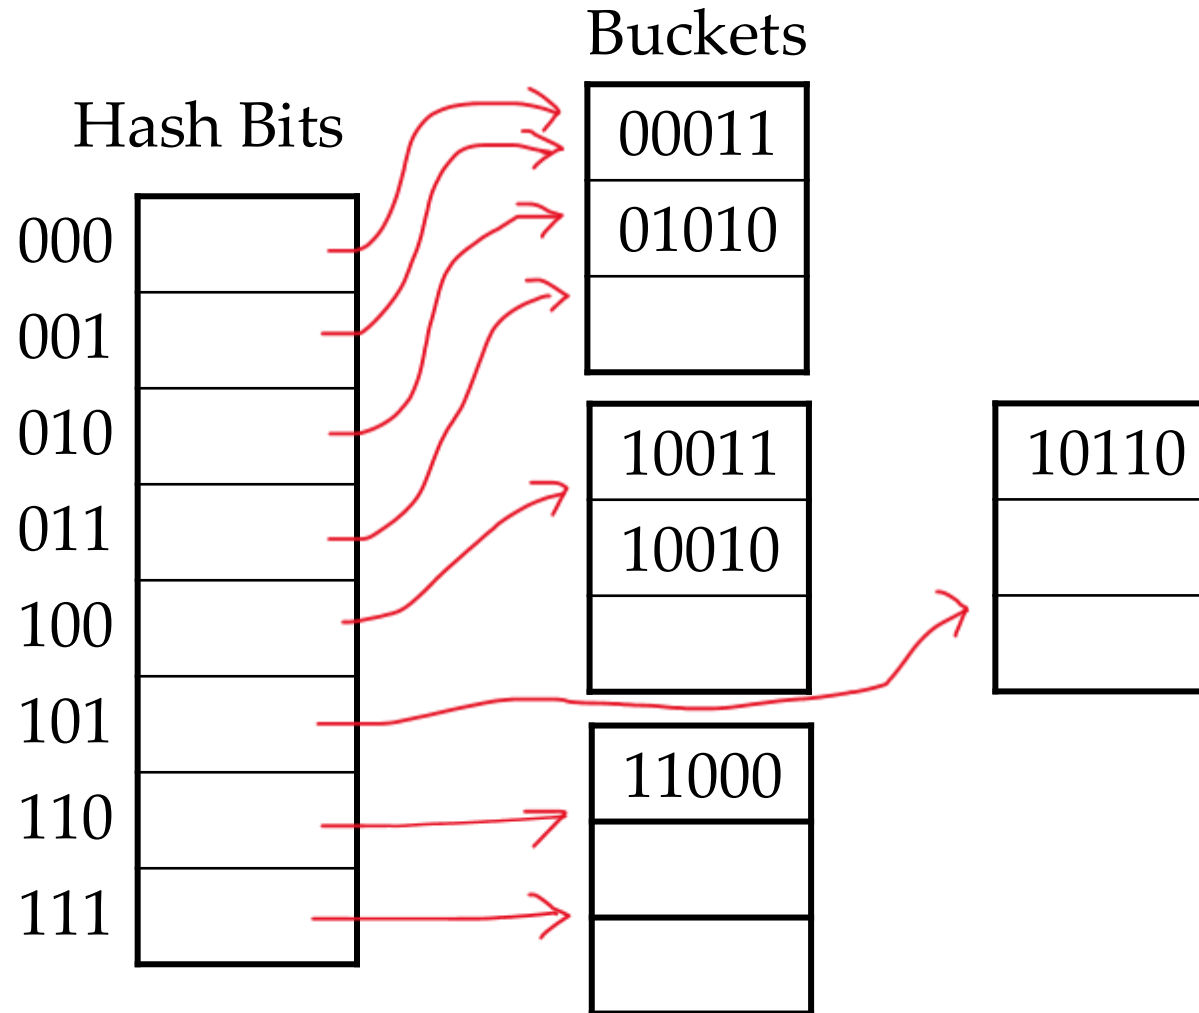
- Next, assume we received a **key 18**, and on passing it **through the hash function**, the binary representation is **10010**.





# Extensible Hashing

- Observe that all the buckets for bits 10 are full → Need to split again buckets for 10.
- Now, **3-bits**.



# Linear Hashing

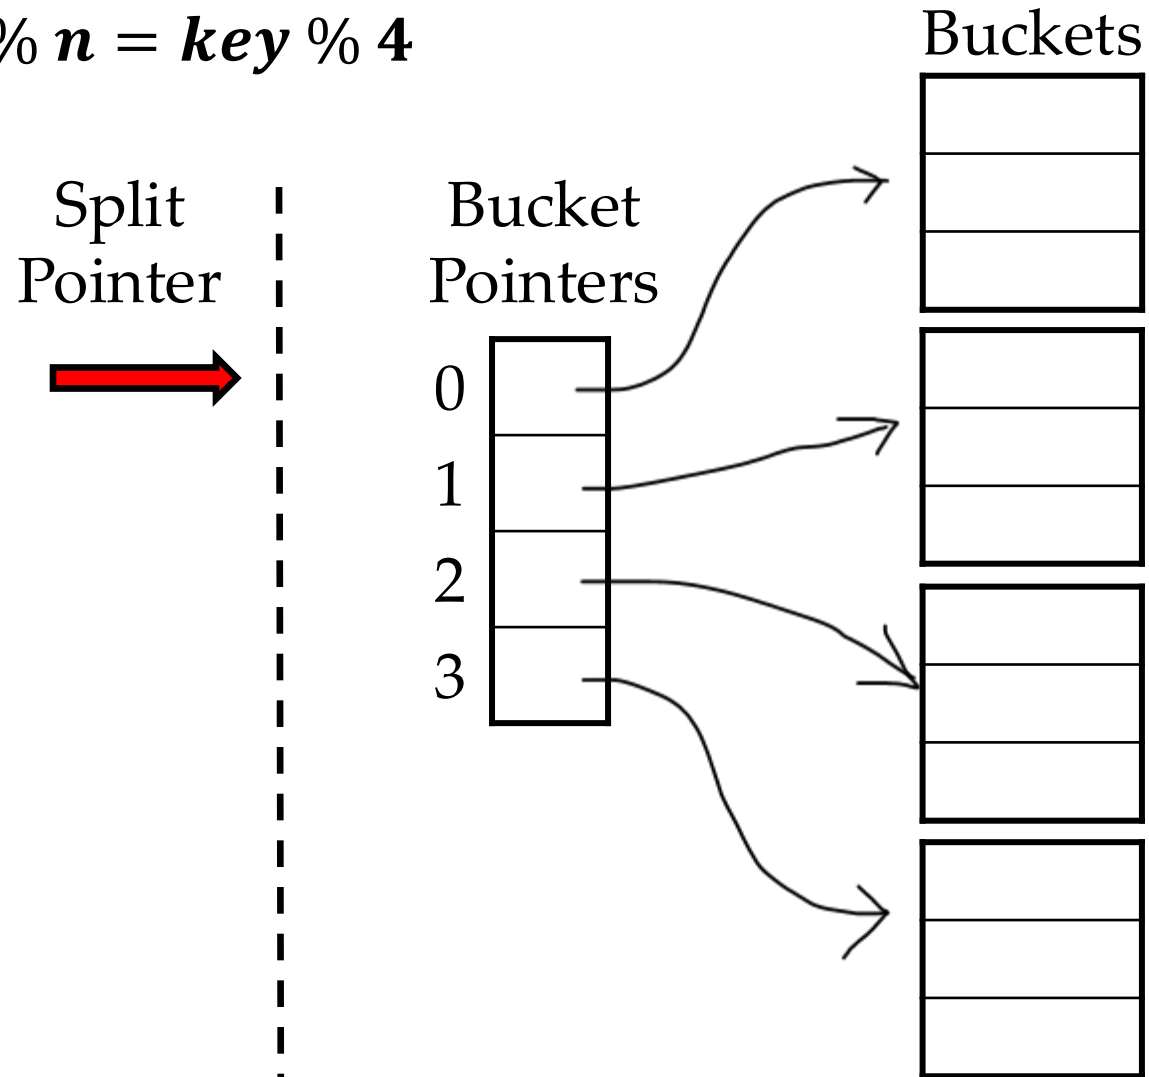
# Linear Hashing

- Extensible hashing works well, but we perform the splitting **lazily** when the buckets for some bit(s) are full.
- What if we allow splitting to happen **eagerly** in the hope that in the future we would anyways need to split.
- **Linear hashing** performs eager and random splitting.
  - We call the splitting random because you may end up splitting empty buckets.
- Note: there is no longer tracking of buckets via binary representation.
- What we need is a **split pointer** that tells where did the last split took place.
  - Every **n-th split** introduces a new hash function.

# Linear Hashing

- Initially say our hash function is:

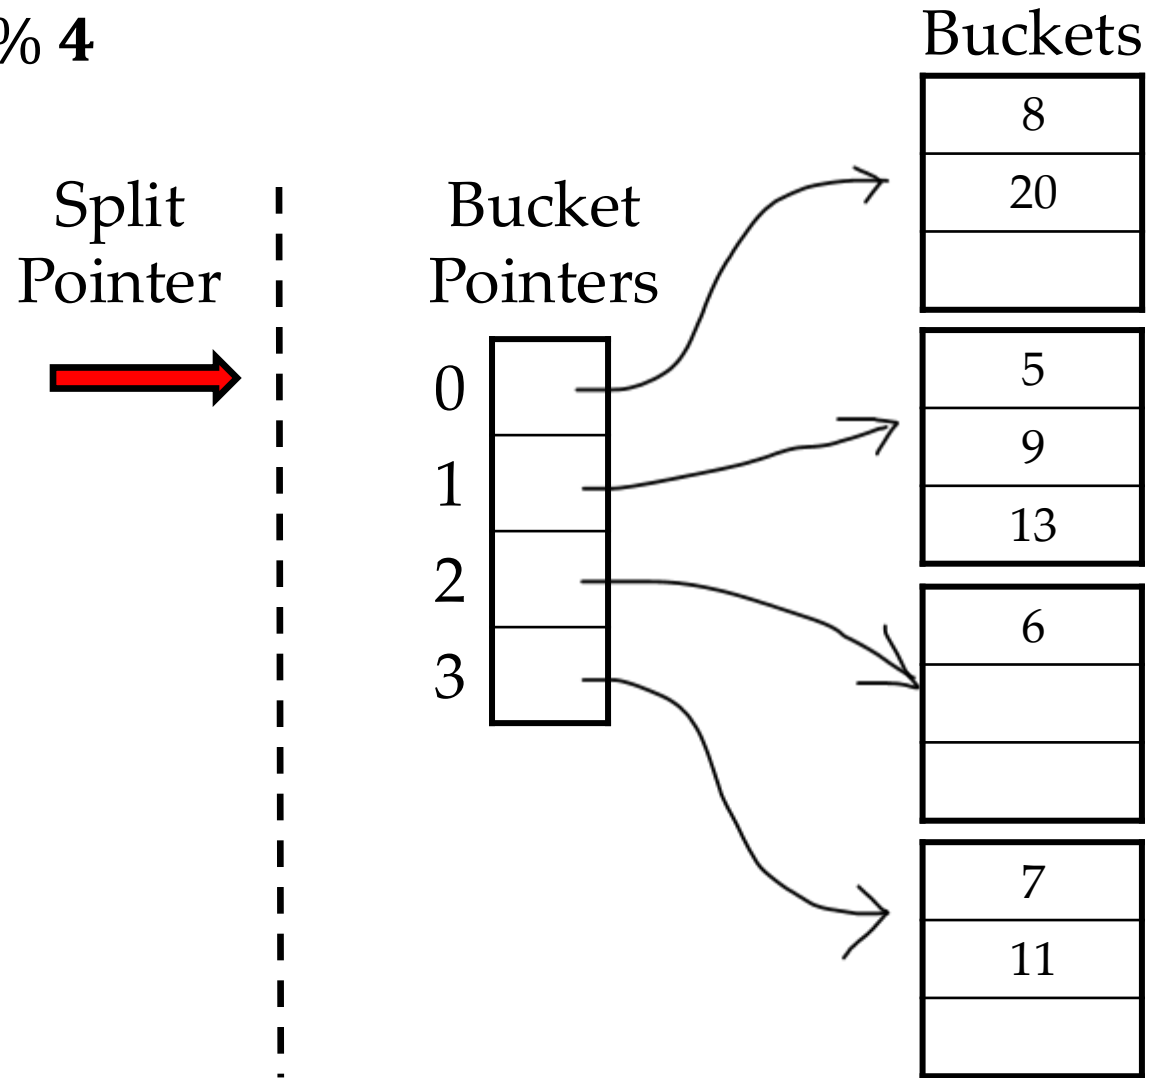
$$h_1(\textit{key}) = \textit{key} \% n = \textit{key} \% 4$$



# Linear Hashing

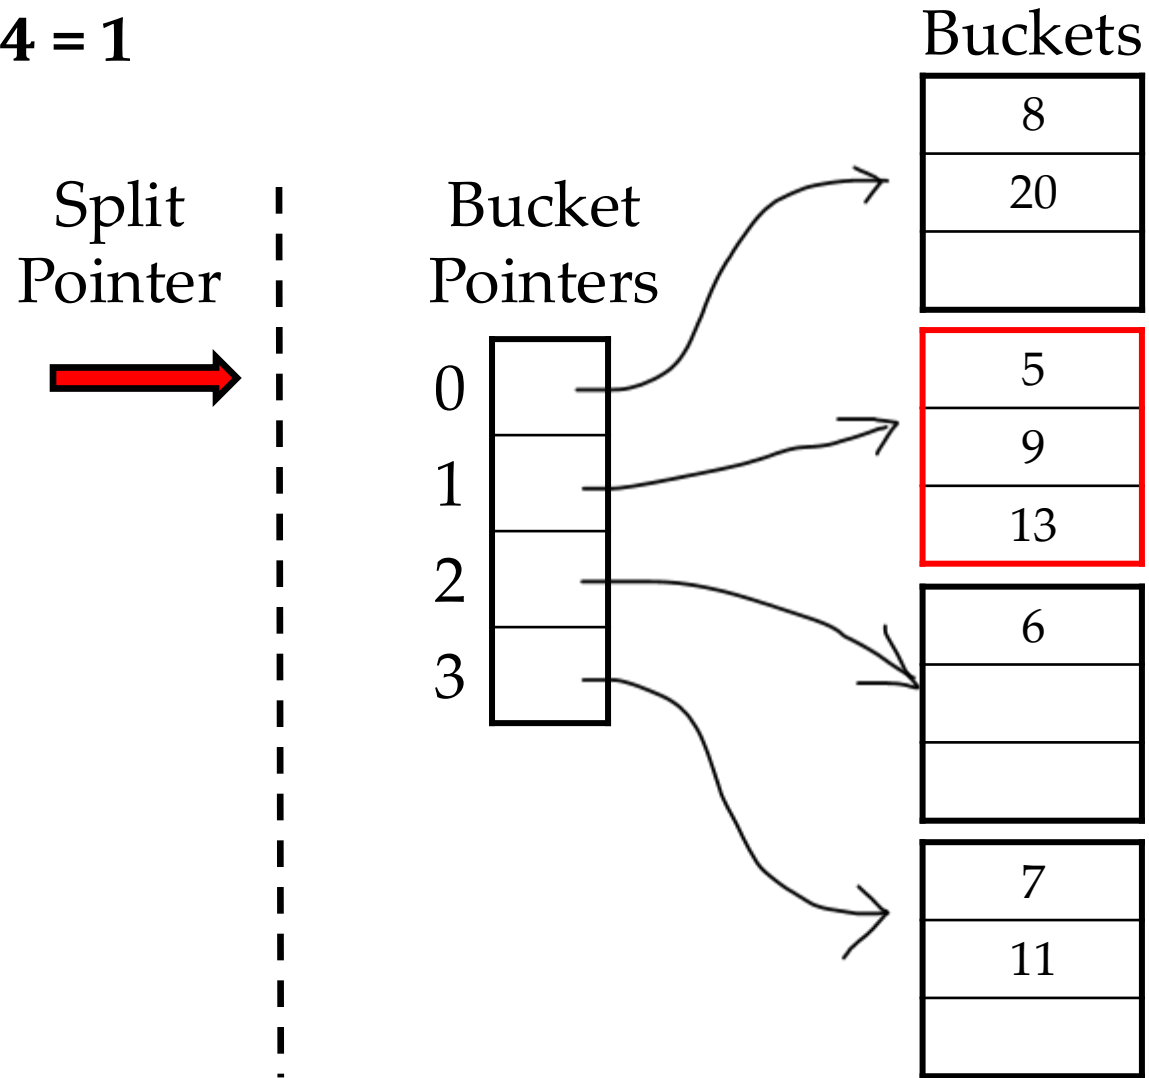
- Say our buckets look like this:

$$h_1(\textit{key}) = \textit{key} \% 4$$



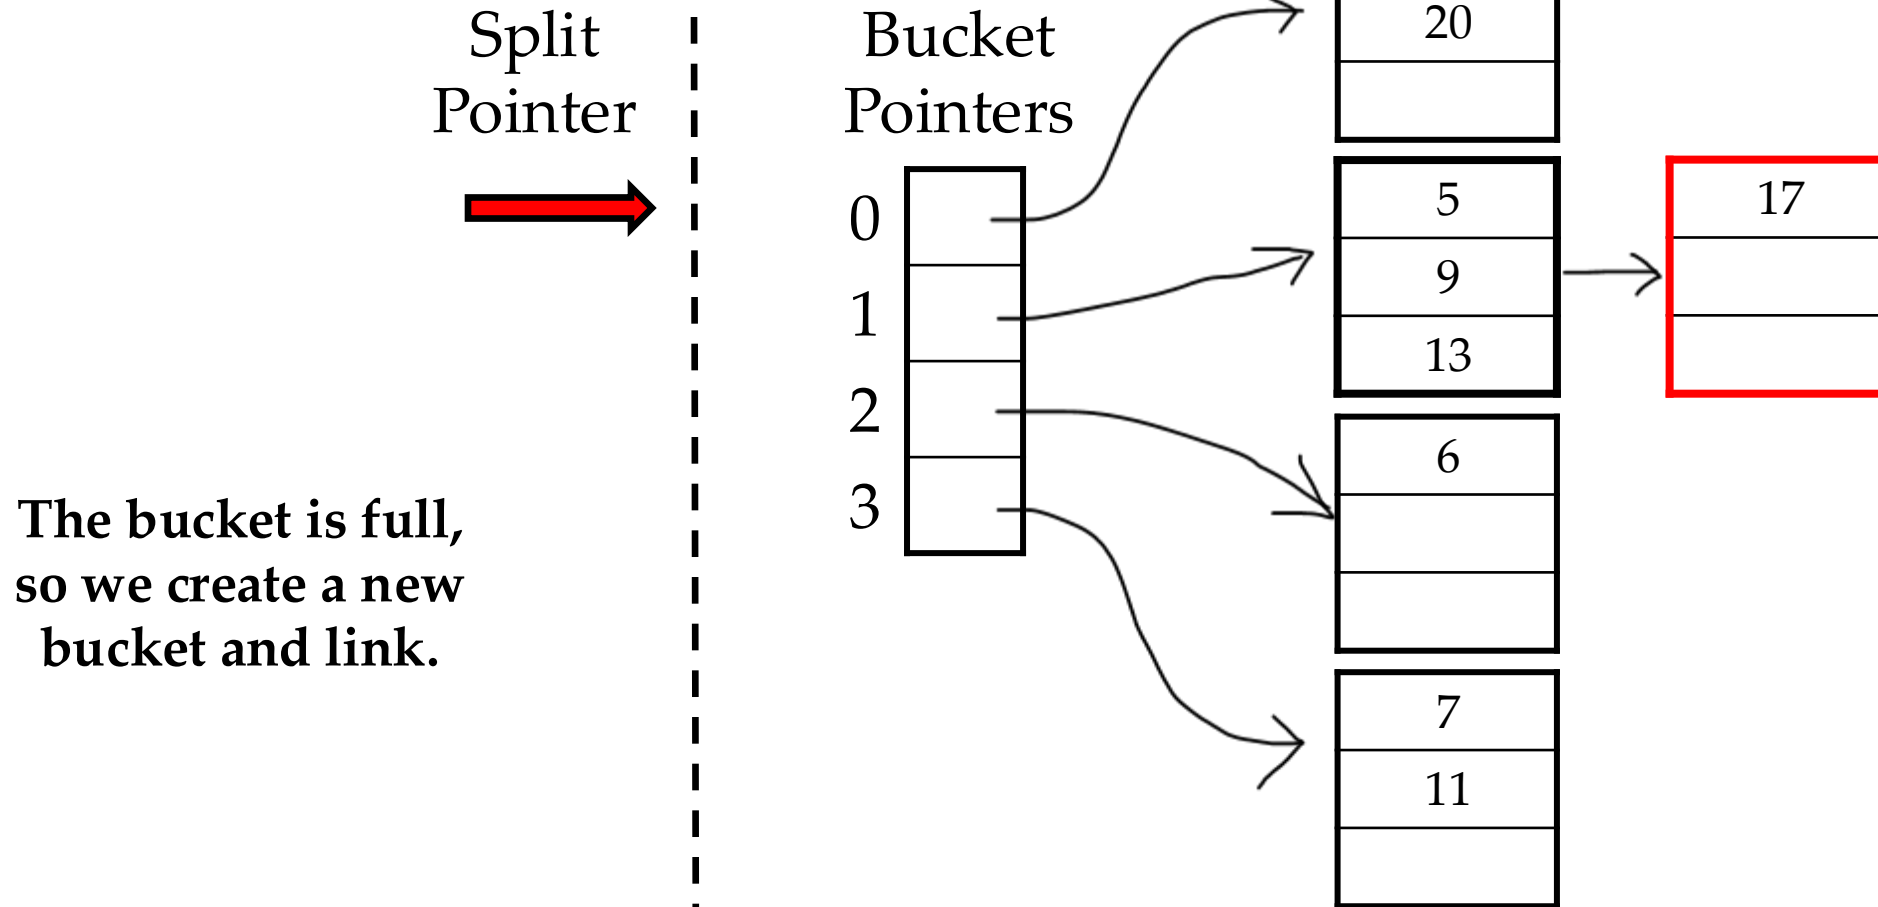
# Linear Hashing

- Let's insert a **key = 17**:  
 $h_1(17) = 17 \% 4 = 1$



# Linear Hashing

- Let's insert a **key = 17**:  
 $h_1(17) = 17 \% 4 = 1$




# Linear Hashing

- Let's insert a **key = 17**:  
 $h_1(17) = 17 \% 4 = 1$

Split  
Pointer  
→

This situation has  
caused an **overflow**,  
so we need to **split**!



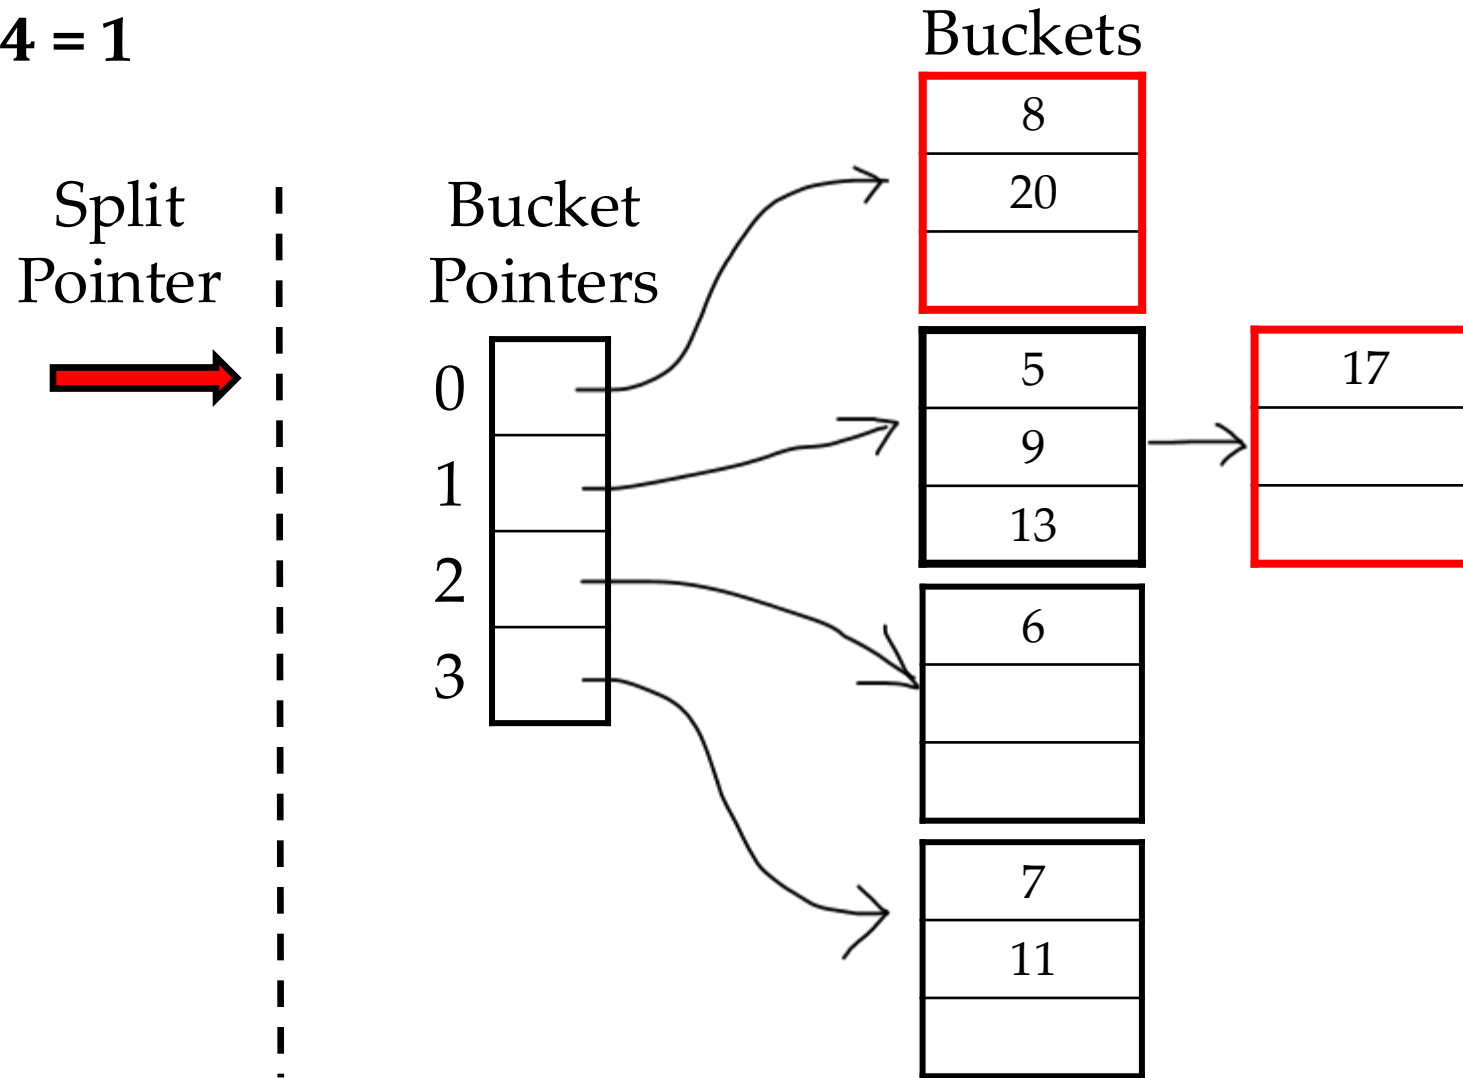


# Linear Hashing

- Let's insert a **key = 17**:

$$h_1(17) = 17 \% 4 = 1$$

My split pointer  
is at 0, so I will  
split bucket 0.




# Linear Hashing

- Let's insert a **key = 17**:  
 $h_1(17) = 17 \% 4 = 1$

Split  
Pointer  
→

My split pointer is at 0, so  
I will **split bucket 0**, and  
add a new **bucket pointer**.



# Linear Hashing

- Let's insert a **key = 17**:

$$h_1(\text{key}) = \text{key} \% 4 = 1$$

$$h_2(\text{key}) = \text{key} \% 8 = 1$$


Split  
Pointer



Introduce a new hash  
function and Rehash the  
keys in original bucket 0.

$$8 \% 8 = 0$$

$$20 \% 8 = 4$$



# Linear Hashing

- Let's insert a **key = 17**:


$$h_1(\text{key}) = \text{key} \% 4 = 1$$

$$h_2(\text{key}) = \text{key} \% 8 = 1$$

Split  
Pointer



Move the split pointer



# Linear Hashing

- Let's insert a **key = 16**:


$$h_1(16) = 16 \% 4 = 0$$

$$h_2(\text{key}) = \text{key} \% 8 = 1$$

Split  
Pointer



First try the hash  
function  $h_1(\text{key})$ .



# Linear Hashing

- Let's insert a **key = 16**:


$$h_1(16) = 16 \% 4 = 0$$

$$h_2(16) = 16 \% 8 = 0$$

Split  
Pointer



As 0 is above the split  
pointer, so we need to run  
the next hash function.



# Linear Hashing

- Let's insert a **key = 16**:


$$h_1(16) = 16 \% 4 = 0$$

$$h_2(16) = 16 \% 8 = 0$$

Split  
Pointer



As 0 is above the split  
pointer, so we need to run  
the next hash function.



# Linear Hashing

- Let's insert a **key = 12**:


$$h_1(12) = 12 \% 4 = 0$$

$$h_2(\text{key}) = \text{key} \% 8 = 1$$

Split  
Pointer



First try the hash  
function  $h_1(\text{key})$ .





# Linear Hashing

- Let's insert a **key = 12**:


$$h_1(12) = 12 \% 4 = 0$$

$$h_2(12) = 12 \% 8 = 4$$

Split  
Pointer



As 0 is above the split  
pointer, so we need to run  
the next hash function.



# Linear Hashing

- Let's insert a **key = 12**:


$$h_1(12) = 12 \% 4 = 0$$

$$h_2(12) = 12 \% 8 = 0$$

Split  
Pointer



As 0 is above the split  
pointer, so we need to run  
the next hash function.



# Linear Hashing

