

# Database Processing

## CS 451 / 551

Lecture 10:  
Sorting and Transactions



**Suyash Gupta**

Assistant Professor

Distopia Labs and ONRG

Dept. of Computer Science

(E) [suyash@uoregon.edu](mailto:suyash@uoregon.edu)

(W) [gupta-suyash.github.io](https://gupta-suyash.github.io)



**Assignment 2 is Out!**  
**Deadline: Nov 13, 2025 at 11:59pm**

**Quiz 2: Nov 6, 2025 (in class)**

# Query Processing

- Last Lecture we looked at Query Processing
  - How to measure the cost of a select operation?
  - The cost of select operation helps us to decide what type of indexes to use, what operators and attributes to access.

# Query Processing

- Last Lecture we looked at Query Processing
  - How to measure the cost of a select operation?
  - The cost of select operation helps us to decide what type of indexes to use, what operators and attributes to access.
- Next, we look at Sorting.

# Why is Sorting Interesting to Us?

# Why is Sorting Interesting to Us?

- SQL queries can require the output be sorted.
- Efficient query processing:
  - Operations like Joins and searching can be implemented.

# Ways to Sort a set of Keys

- Is it possible to sort a set of keys/records without running a sorting algorithm?

# Ways to Sort a set of Keys

- Is it possible to sort a set of keys/records without running a sorting algorithm?
- Yes, what if we have access to a B<sup>+</sup>-tree.
  - All the leaves in a B<sup>+</sup>-tree are sorted!

# Ways to Sort a set of Keys

- Is it possible to sort a set of keys/records without running a sorting algorithm?
- Yes, what if we have access to a B<sup>+</sup>-tree.
  - All the leaves in a B<sup>+</sup>-tree are sorted!
  - Can we do something more?

# Top-N Heap Sort

- Say a query contains an **ORDER BY** clause with a **LIMIT**.
  - The DBMS only needs to scan the data once to find the required number of elements.
  - Specifically, a query asks you to output only first **N** elements.

# Top-N Heap Sort

- Say a query contains an **ORDER BY** clause with a **LIMIT**.
  - The DBMS only needs to scan the data once to find the required number of elements.
  - Specifically, a query asks you to output only first **N** elements.
  - This is also an ideal candidate for **Heap Sort!**

# Top-N Heap Sort

- Say a query contains an **ORDER BY** clause with a **LIMIT**.
  - The DBMS only needs to scan the data once to find the required number of elements.
  - Specifically, a query asks you to output only first **N** elements.
  - This is also an ideal candidate for **Heap Sort!**
  - We are just going to scan data once and maintain an in-memory sorted **priority queue**.

# Top-N Heap Sort

- Say, our query is:

```
select * from enrolled order by sid  
asc fetch first 4 rows with ties
```

## Original Data

23    34    12    67    8    6    76    98    78    19    83    19

## Sorted Data

# Top-N Heap Sort

- Say, our query is:

```
select * from enrolled order by sid  
asc fetch first 4 rows with ties
```

**Original Data**

|    |    |    |    |   |   |    |    |    |    |    |    |
|----|----|----|----|---|---|----|----|----|----|----|----|
| 23 | 34 | 12 | 67 | 8 | 6 | 76 | 98 | 78 | 19 | 83 | 19 |
|----|----|----|----|---|---|----|----|----|----|----|----|

**Sorted Data**

# Top-N Heap Sort

- Say, our query is:

```
select * from enrolled order by sid  
asc fetch first 4 rows with ties
```

## Original Data

23    34    12    67    8    6    76    98    78    19    83    19

## Sorted Data

23

# Top-N Heap Sort

- Say, our query is:

```
select * from enrolled order by sid  
asc fetch first 4 rows with ties
```

## Original Data

23    34    12    67    8    6    76    98    78    19    83    19

## Sorted Data

23    34

# Top-N Heap Sort

- Say, our query is:

```
select * from enrolled order by sid  
asc fetch first 4 rows with ties
```

## Original Data

23    34    12    67    8    6    76    98    78    19    83    19

## Sorted Data

12    23    34

# Top-N Heap Sort

- Say, our query is:

```
select * from enrolled order by sid  
asc fetch first 4 rows with ties
```

## Original Data

23    34    12    67    8    6    76    98    78    19    83    19

## Sorted Data

12    23    34    67

# Top-N Heap Sort

- Say, our query is:

```
select * from enrolled order by sid  
asc fetch first 4 rows with ties
```

## Original Data

23    34    12    67    8    6    76    98    78    19    83    19



## Sorted Data

8    12    23    34

# Top-N Heap Sort

- Say, our query is:

```
select * from enrolled order by sid  
asc fetch first 4 rows with ties
```

## Original Data

23    34    12    67    8    6    76    98    78    19    83    19



## Sorted Data

6    8    12    23

# Top-N Heap Sort

- Say, our query is:

```
select * from enrolled order by sid  
asc fetch first 4 rows with ties
```

## Original Data

23    34    12    67    8    6    76    98    78    19    83    19



## Sorted Data

6    8    12    23

# Top-N Heap Sort

- Say, our query is:

```
select * from enrolled order by sid  
asc fetch first 4 rows with ties
```

## Original Data

23    34    12    67    8    6    76    **98**    78    19    83    19

## Sorted Data

6    8    12    23

# Top-N Heap Sort

- Say, our query is:

```
select * from enrolled order by sid  
asc fetch first 4 rows with ties
```

## Original Data

23    34    12    67    8    6    76    98    78    19    83    19



## Sorted Data

6    8    12    23

# Top-N Heap Sort

- Say, our query is:

```
select * from enrolled order by sid  
asc fetch first 4 rows with ties
```

## Original Data

23    34    12    67    8    6    76    98    78    **19**    83    19

## Sorted Data

6    8    12    19

# Top-N Heap Sort

- Say, our query is:

```
select * from enrolled order by sid  
asc fetch first 4 rows with ties
```

## Original Data

23    34    12    67    8    6    76    98    78    19    **83**    19

## Sorted Data

6    8    12    19

# Top-N Heap Sort

- Say, our query is:

```
select * from enrolled order by sid  
asc fetch first 4 rows with ties
```

## Original Data

23    34    12    67    8    6    76    98    78    19    83    **19**

## Sorted Data

6    8    12    19    19

**Tracks all the duplicates!**

# Ways to Sort a set of Key

- But, having a B<sup>+</sup>-tree and Top-N Heap Sort sufficient for all the sorting queries?
- What are we still missing?

# Ways to Sort a set of Key

- But, is having just a B<sup>+</sup>-tree sufficient to satisfy all the sorting queries?
- What are we still missing?
- B<sup>+</sup>-tree is just an index.
- Your data needs to be physically sorted too!
- Remember, the benefits of sequential data access only comes when data is stored in a sorted manner.

# Ways to Sort a set of Key

- Imagine you have stored the data in your disk in unsorted manner but you are now trying to fetch it in sorted manner?
- Pretty bad performance as too many disk accesses!

# Ways to Sort a set of Key

- So, depending on where your data resides (in-memory or disk, you select a sorting algorithm.
- **Why this difference?**

# Ways to Sort a set of Key

- So, depending on where your data resides (in-memory or disk, you select a sorting algorithm.
- **Why this difference?**
- Because, an in-memory sorting algorithm does not need to worry about expensive data swapping operations!

# Ways to Sort a set of Key

- So, depending on where your data resides (in-memory or disk, you select a sorting algorithm.
- **Why this difference?**
- Because, an in-memory sorting algorithm does not need to worry about expensive data swapping operations!
- In-memory swapping algorithms? → Quick Sort
- Our focus → Disk sorting algorithms.

# Ways to Sort a set of Key

- So, depending on where your data resides (in-memory or disk, you select a sorting algorithm.
- **Why this difference?**
- Because, an in-memory sorting algorithm does not need to worry about expensive data swapping operations!
- In-memory swapping algorithms? → Quick Sort
- Our focus → Disk sorting algorithms.

# External Merge-Sort

- Sorting of relations that do not fit in memory is called external sorting.
- A divide and conquer algorithm!
- Split data into parts (also called as runs).
- Sort each run individually.
- Merge the runs!

# External Merge-Sort

- What are the challenges for External Merge-Sort?

# External Merge-Sort

- What are the challenges for External Merge-Sort?
- Number of runs.
- Size of each run.
- Size of memory.
- All these factors work in conjunction to determine how fast can we perform external merge-sort.

# N-Way External Merge-Sort

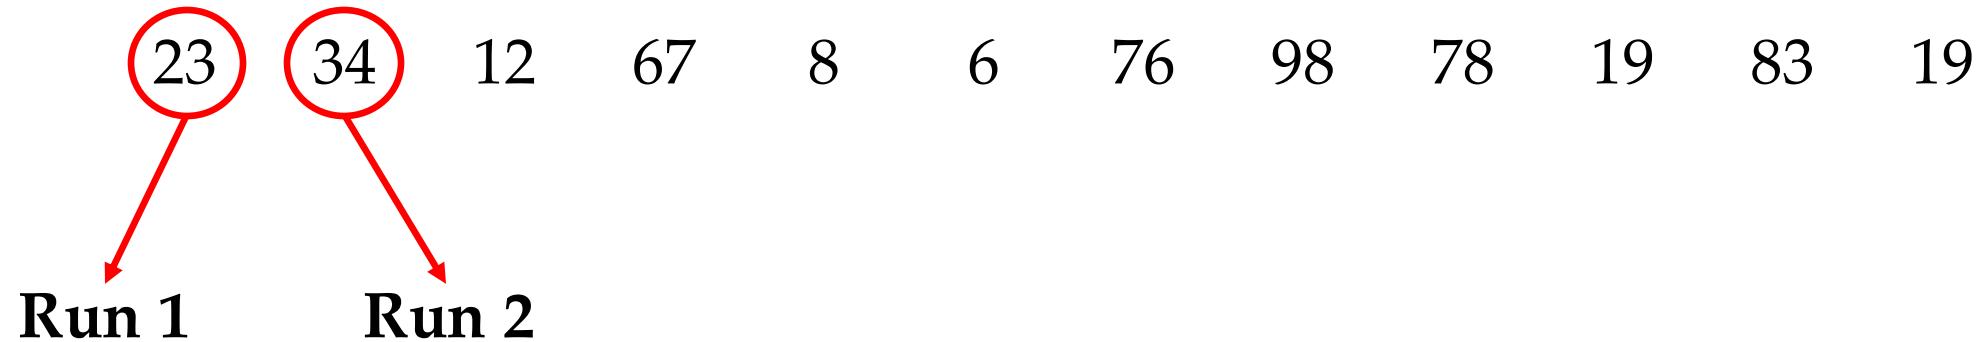
- We are going to run an external merge sort where:
  - $N$  total number of runs
  - We will assume each run consists of **M blocks**.
  - We will assume the memory can store  **$N+1$  blocks**.
  - We will fetch **one block from each run** at a time.
  - Remember → Often you would decide on the right number of runs based on how many threads you have, and how many blocks you can store.

# 2-Way External Merge-Sort → N=2

23    34    12    67    8    6    76    98    78    19    83    19

Assume that we can have only 2 runs at a time and total of 3 blocks in memory.

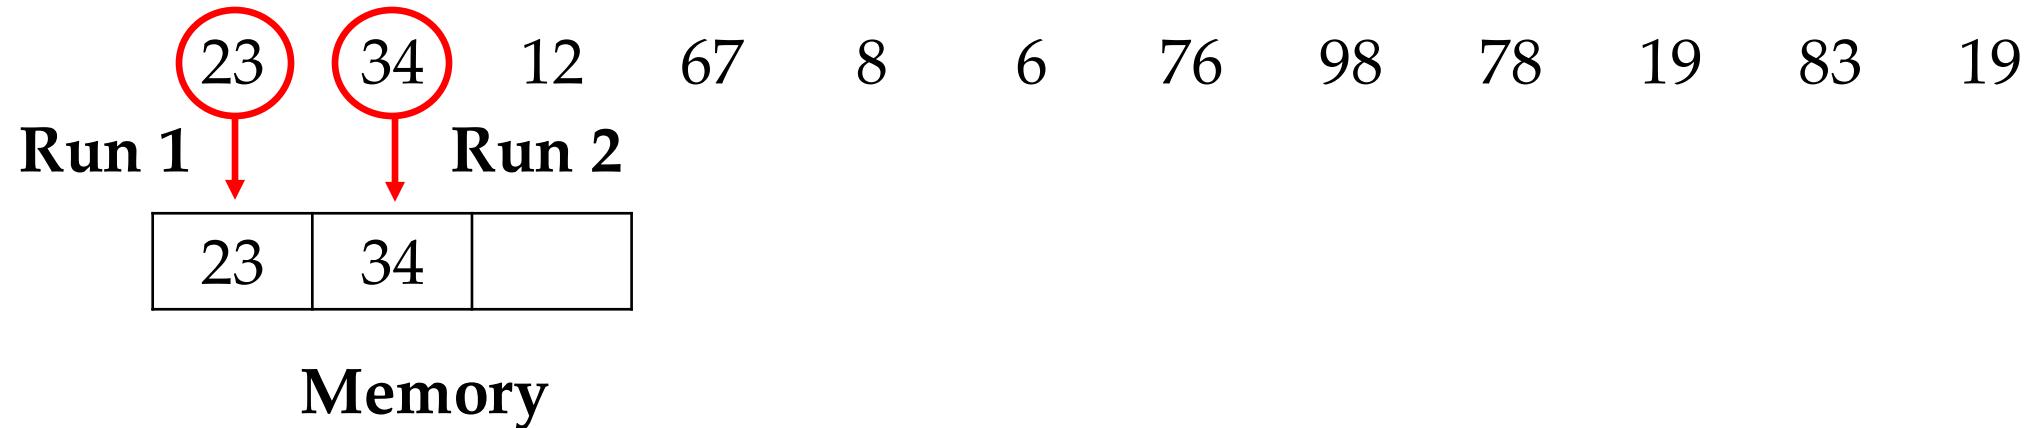
# 2-Way External Merge-Sort → N=2



Initially, Each Run is of size 1.

Each run is implicitly sorted!

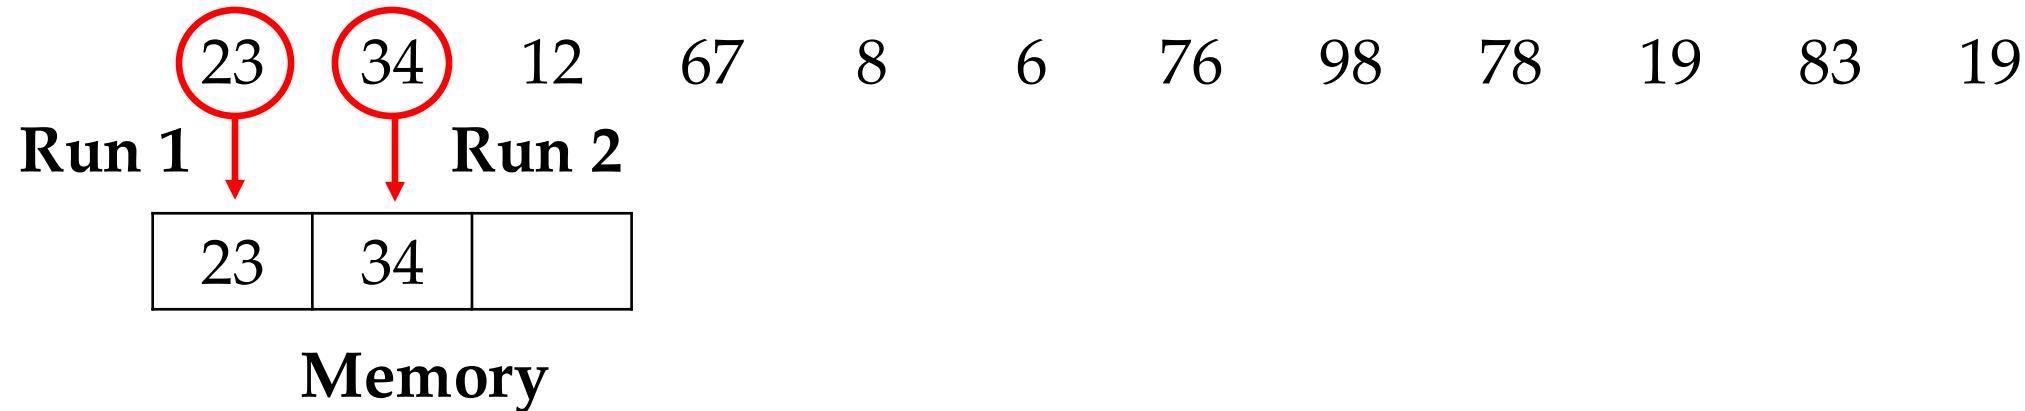
# 2-Way External Merge-Sort → N=2



Next, we merge two runs at a time.

We can have only 3 blocks in memory

# 2-Way External Merge-Sort → N=2



Use the extra block to sort.

Write back to memory

# 2-Way External Merge-Sort → N=2

23    34    12    67    8    6    76    98    78    19    83    19

|    |    |  |
|----|----|--|
| 12 | 67 |  |
|----|----|--|

**Memory**

Use the extra block to sort.

Write back to memory

# 2-Way External Merge-Sort → N=2

23    34    12    67    8    6    76    98    78    19    83    19

|   |   |  |
|---|---|--|
| 8 | 6 |  |
|---|---|--|

**Memory**

Use the extra block to sort.

Write back to memory

# 2-Way External Merge-Sort → N=2

23    34    12    67    6    6    76    98    78    19    83    19

|   |   |   |
|---|---|---|
| 8 | 6 | 8 |
|---|---|---|

**Memory**

Use the extra block to sort.

Write back to memory

# 2-Way External Merge-Sort → N=2

23    34    12    67    6    8    76    98    78    19    83    19

|   |   |   |
|---|---|---|
| 8 | 6 | 8 |
|---|---|---|

**Memory**

Use the extra block to sort.

Write back to memory

# 2-Way External Merge-Sort → N=2

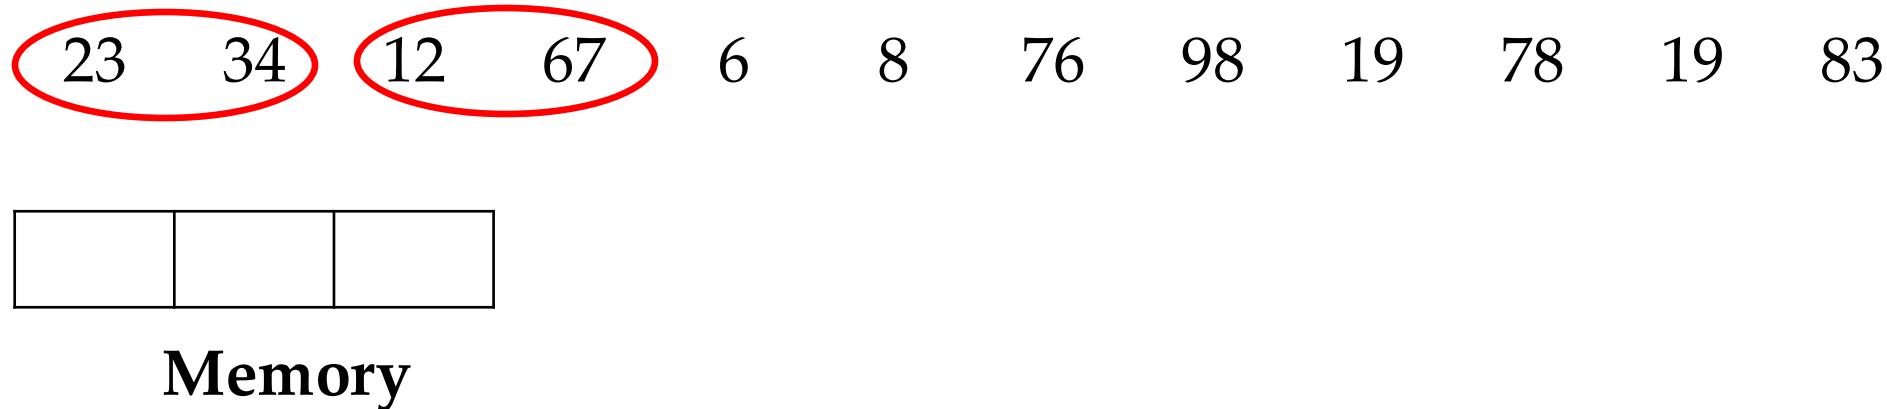
23    34    12    67    6    8    76    98    19    78    19    83



**Memory**

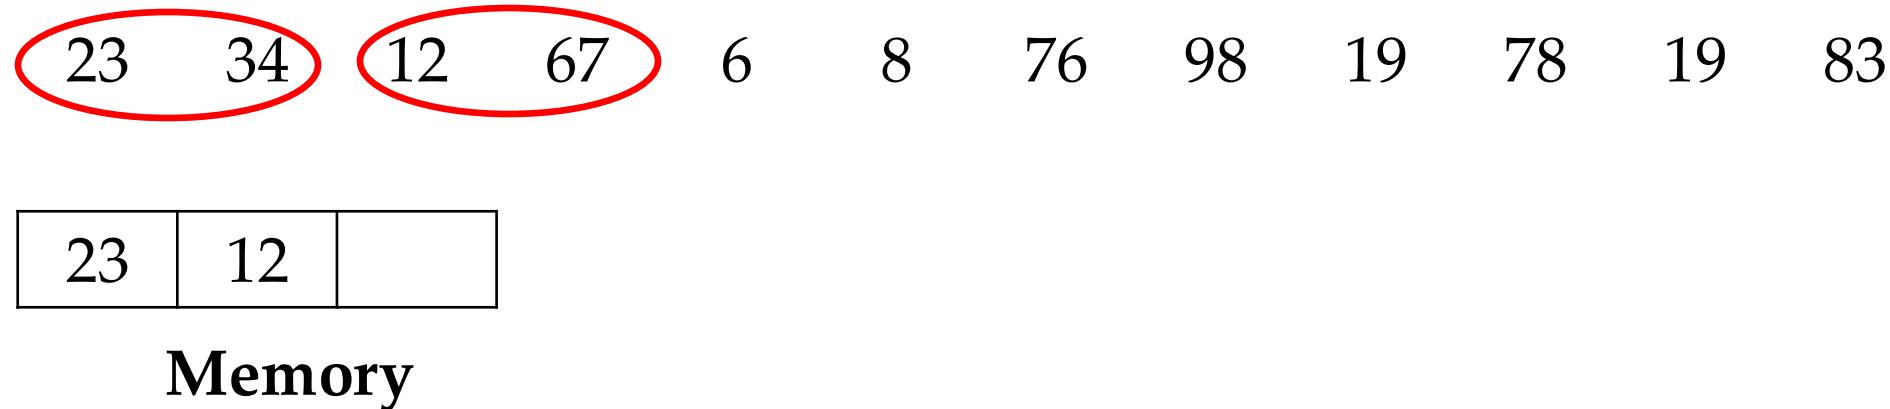
Lets continue this till the end

# 2-Way External Merge-Sort → N=2



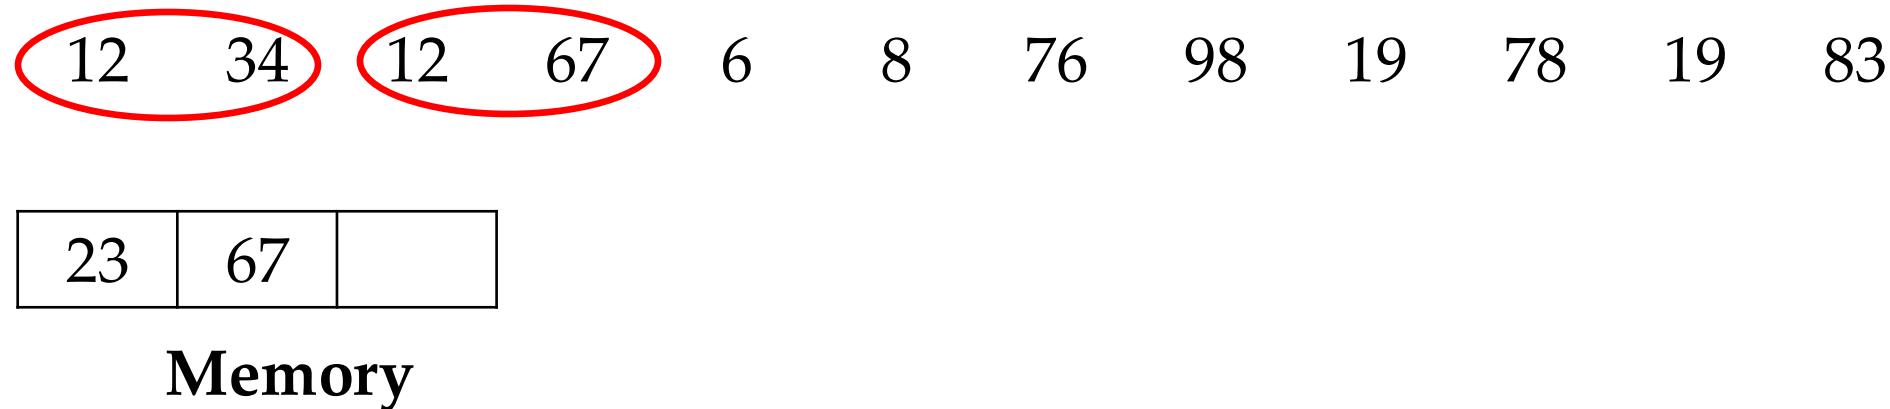
Now, the size of each run is 2 blocks

# 2-Way External Merge-Sort → N=2



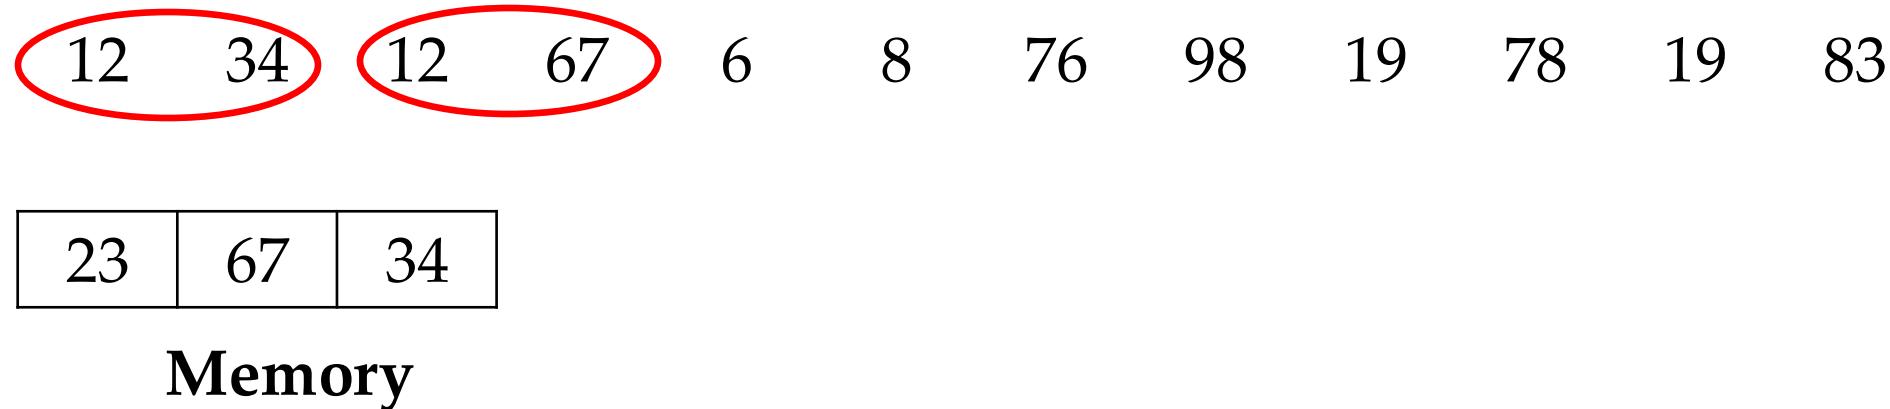
Now, the size of each run is 2 blocks

# 2-Way External Merge-Sort → N=2



Now, the size of each run is 2 blocks

# 2-Way External Merge-Sort → N=2



Now, the size of each run is 2 blocks

# 2-Way External Merge-Sort → N=2

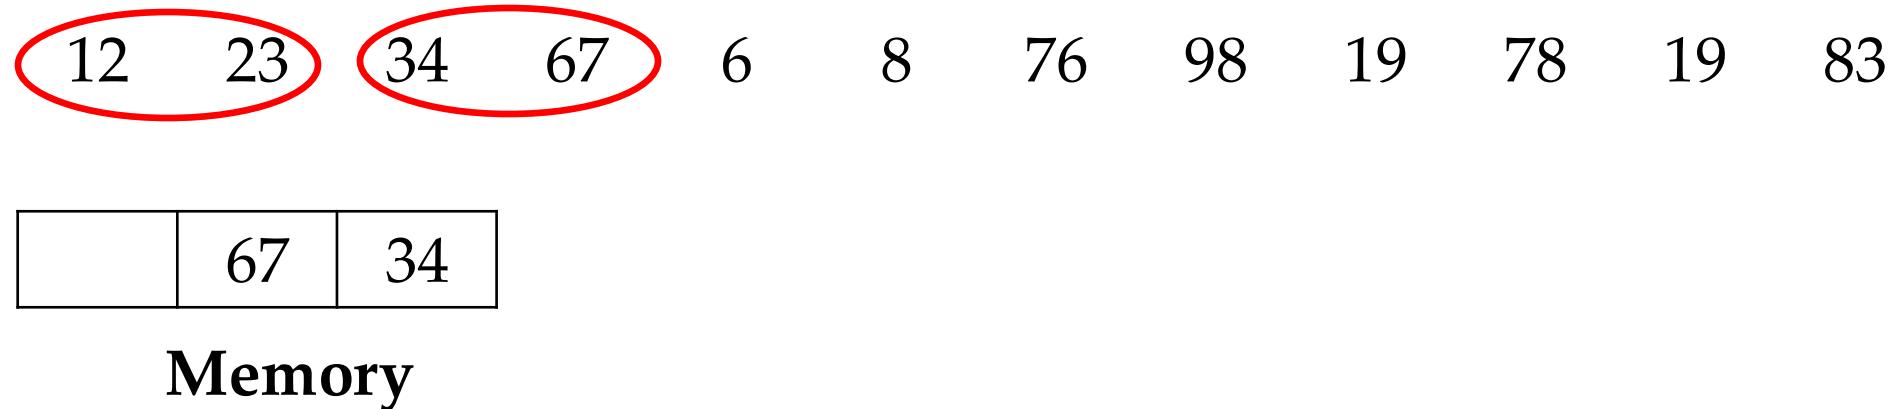
12    23    12    67    6    8    76    98    19    78    19    83

|  |    |    |
|--|----|----|
|  | 67 | 34 |
|--|----|----|

**Memory**

Now, the size of each run is 2 blocks

# 2-Way External Merge-Sort → N=2



Now, the size of each run is 2 blocks

# 2-Way External Merge-Sort → N=2

12    23    34    67    6    8    76    98    19    78    19    83

|   |    |  |
|---|----|--|
| 6 | 76 |  |
|---|----|--|

**Memory**

Now, the size of each run is 2 blocks

# 2-Way External Merge-Sort → N=2

12    23    34    67    6    8    76    98    19    78    19    83

|    |    |  |
|----|----|--|
| 19 | 19 |  |
|----|----|--|

**Memory**

Now, the size of each run is 2 blocks

# 2-Way External Merge-Sort → N=2

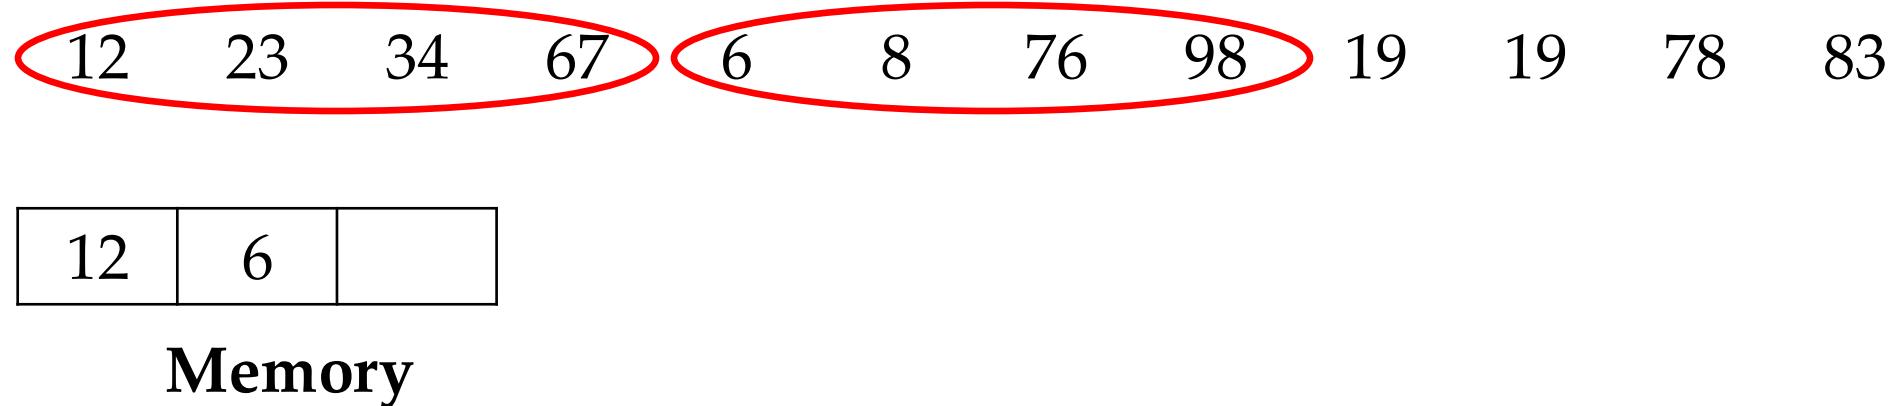
12    23    34    67    6    8    76    98    19    19    78    83

|    |    |  |
|----|----|--|
| 19 | 19 |  |
|----|----|--|

**Memory**

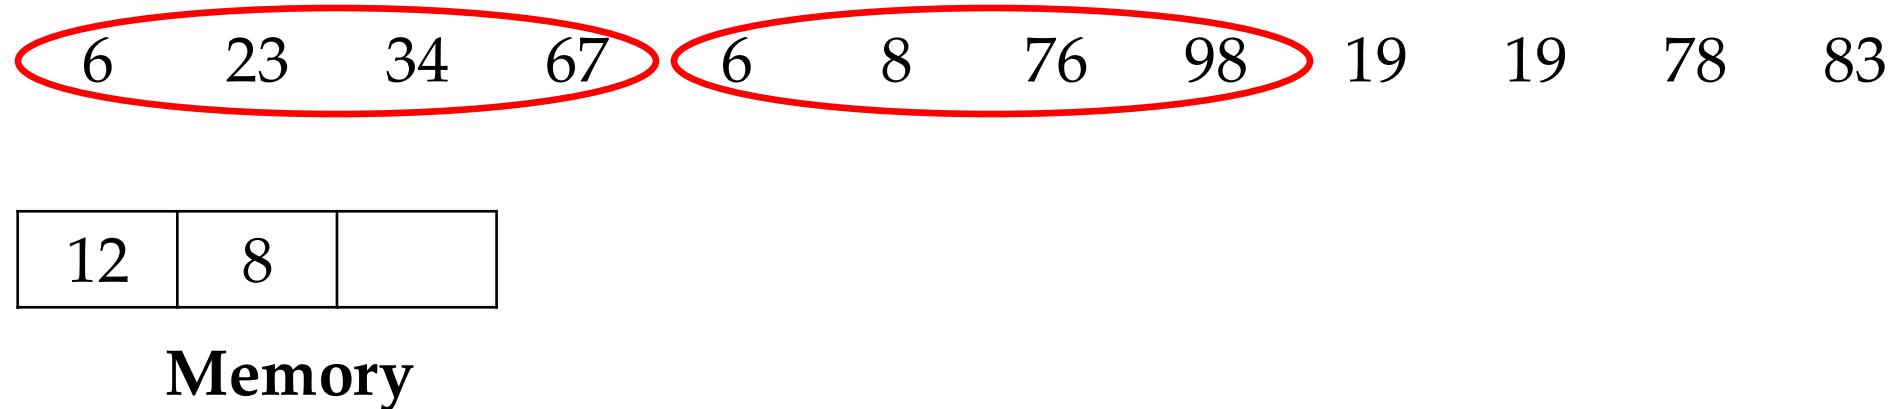
Now, the size of each run is 2 blocks

# 2-Way External Merge-Sort → N=2



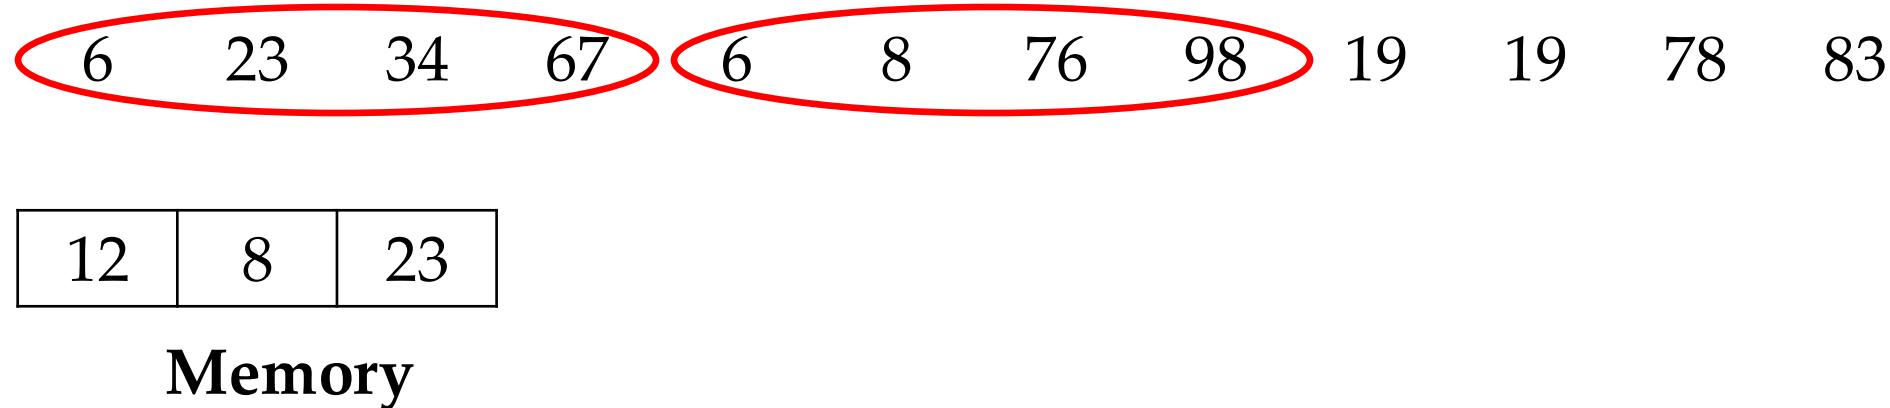
Now, the size of each run is 4 blocks

# 2-Way External Merge-Sort → N=2



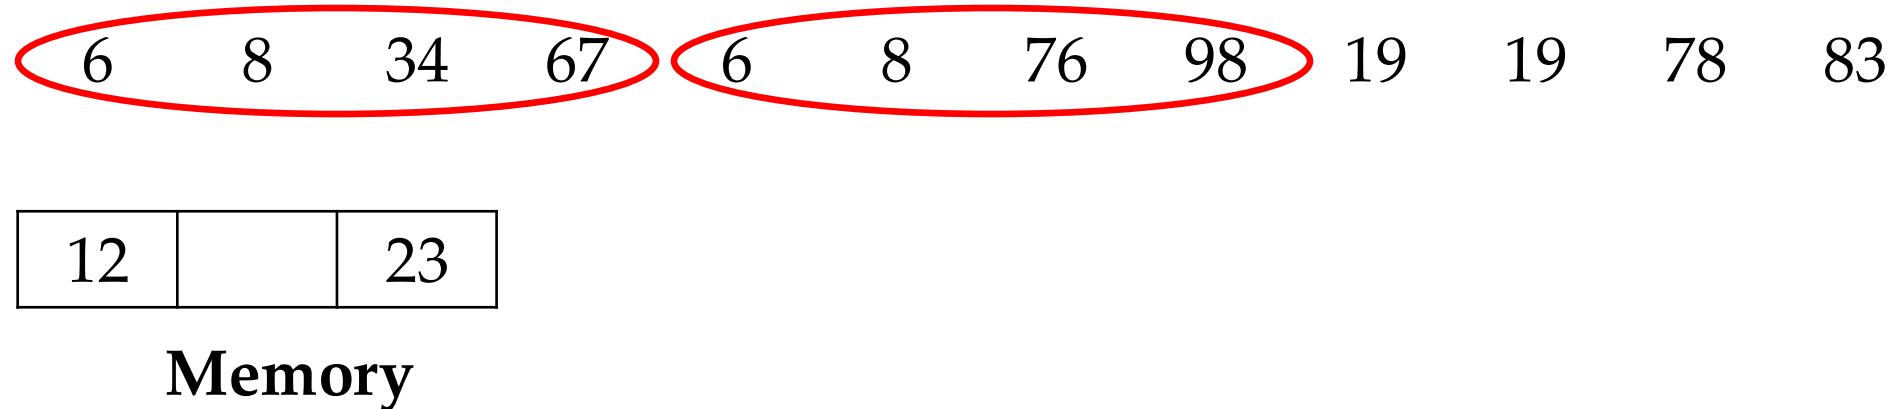
Now, the size of each run is 4 blocks

# 2-Way External Merge-Sort → N=2



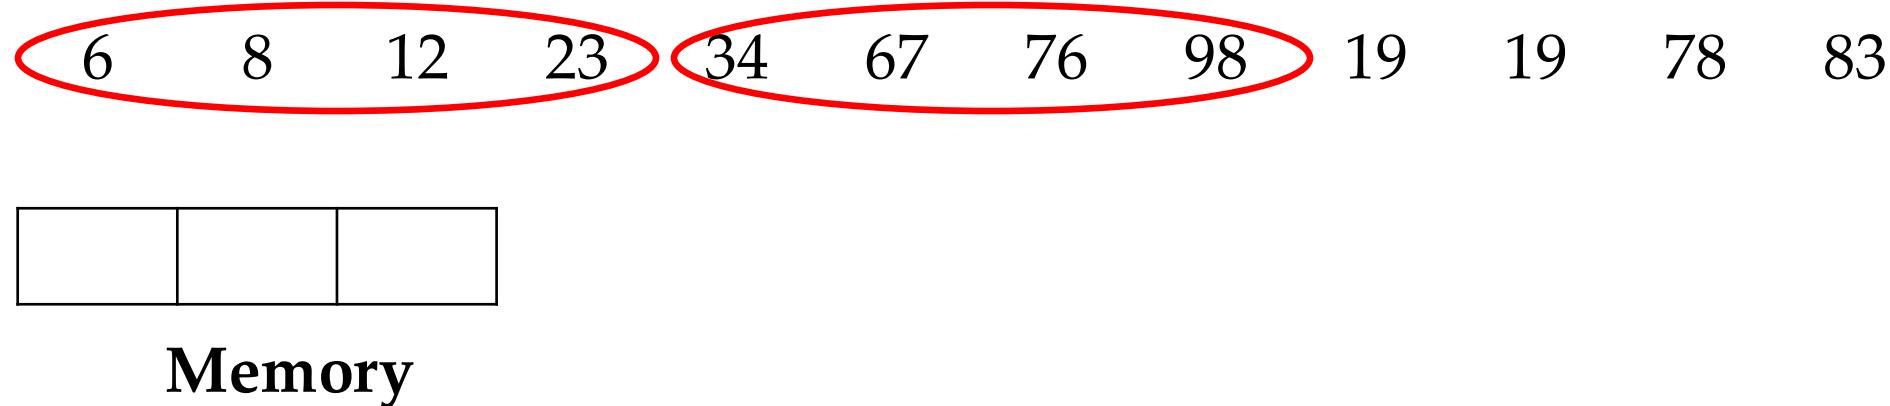
Now, the size of each run is 4 blocks

# 2-Way External Merge-Sort → N=2



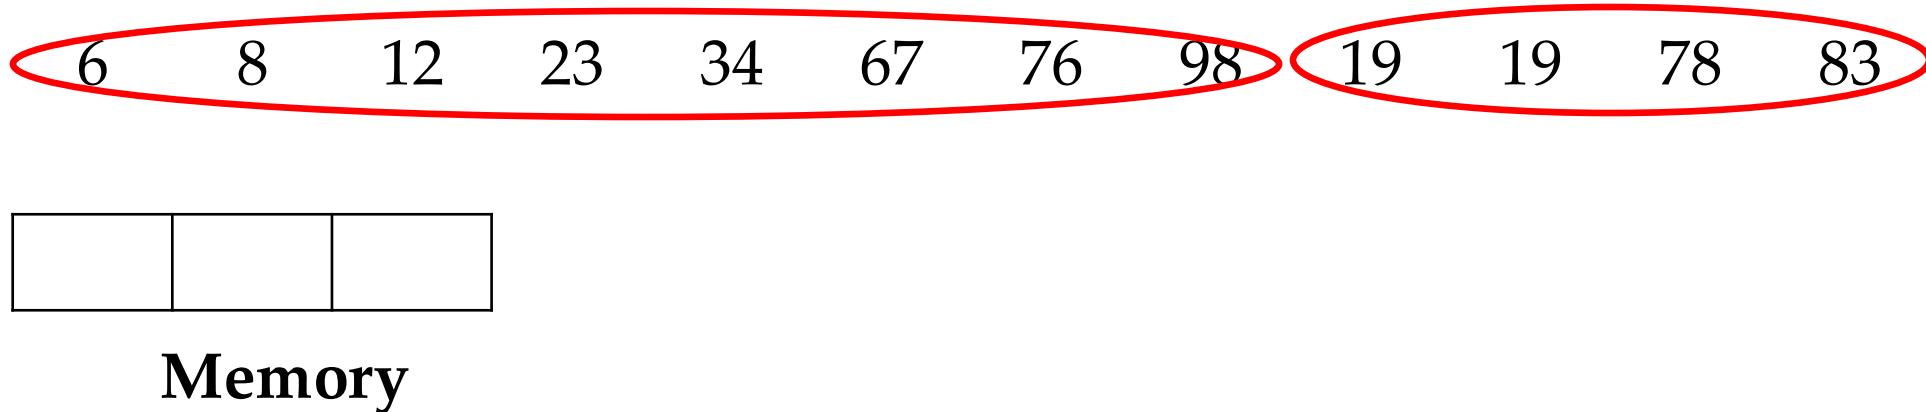
Now, the size of each run is 4 blocks

# 2-Way External Merge-Sort → N=2



Now, the size of each run is 4 blocks

# 2-Way External Merge-Sort → N=2



Now, the size of one run is 8 blocks while other is 4 blocks

# **Transactions**

# Transactions

- Transactions are ubiquitous!
- Examples: Banking, Online shopping, Trading, Social media, and so on.

# How to define a Transaction?

# How to define a Transaction?

- Transaction is a collection of operations.
- For example:
  - Moving money from one checkings account to savings account.
  - Buying a product from Amazon.

# How to define a Transaction?

- Transaction is a unit of program that reads and/or writes one or more data items.
- A common way to write a transaction in popular DBMS is by placing the body of the transaction between, “**begin transaction**” and “**end transaction**”.

# How to define a Transaction?

- Transaction is a unit of program that reads and/or writes one or more data items.
- A common way to write a transaction in popular DBMS is by placing the body of the transaction between, “**begin transaction**” and “**end transaction**”.
- Also, the reason why transaction is termed as an indivisible unit.
  - It either executes in its entirety or nothing at all.

# ACID Properties for a Transaction

?

# **ACID Properties for a Transaction**

- Each database should provide the following four properties for transactions :

# ACID Properties for a Transaction

- Each database should provide the following four properties for transactions :
- **Atomicity:** Either all operations of the transaction are reflected properly in the database, or none are.

# ACID Properties for a Transaction

- Each database should provide the following four properties for transactions:
- **Atomicity:** Either all operations of the transaction are reflected properly in the database, or none are.
- **Consistency:** Execution of a transaction in isolation (that is, with no other transaction executing concurrently) preserves the consistency of the database.

# ACID Properties for a Transaction

- Each database should provide the following four properties for transactions:
- **Atomicity:** Either all operations of the transaction are reflected properly in the database, or none are.
- **Consistency:** Execution of a transaction in isolation (that is, with no other transaction executing concurrently) preserves the consistency of the database.
  - Not referring to database consistency constraints.
- **Isolation:** For every pair of concurrent (executing at the same time) transactions  $T_i$  and  $T_j$ , either  $T_i$  finished execution before  $T_j$  started, or  $T_i$  started execution after  $T_j$  finished.
  - Transactions are unaware of other transactions executing

# ACID Properties for a Transaction

- **Atomicity:** Either all operations of the transaction are reflected properly in the database, or none are.
- **Consistency:** Execution of a transaction in isolation (that is, with no other transaction executing concurrently) preserves the consistency of the database.
  - Not referring to database consistency constraints.
- **Isolation:** For every pair of concurrent (executing at the same time) transactions  $T_i$  and  $T_j$ , either  $T_i$  finished execution before  $T_j$  started, or  $T_i$  started execution after  $T_j$  finished.
  - Transactions are unaware of other transactions executing
- **Durability:** Once a transaction completes successfully, any changes it made to the database should persist, even if there are system failures.

# ACID Properties for a Transaction

- When do ACID properties come into play?

# ACID Properties for a Transaction

- When do ACID properties come into play?
  - Concurrency!
- In a concurrent system or database, two or more transactions may attempt to fetch the same data.
- But, why is this an issue?

# ACID Properties for a Transaction

- When do ACID properties come into play?
  - Concurrency!
- In a concurrent system or database, two or more transactions may attempt to fetch the same data.
- But, why is this an issue?
  - Concurrency if not handled well can lead to ACID violations.
  - For instance. → **Race conditions!**

# Two Concurrent Transactions

T1:

```
read(A);  
A = A - 50;  
write(A);  
read(B);  
B = B + 50;  
write(B).
```

T2:

```
read(A);  
temp = A * 0.1;  
A = A - temp;  
write(A);  
read(B);  
B = B + temp;  
write(B)
```

Notice that they are accessing the same variables A and B.

# Conflicting Transactions

**Two transactions T1 and T2 if they concurrently access the same variable and at least one of that access is a write operation, then they conflict!**

# Isolation

- Users submit transactions.
- Each transaction should execute as if it were running by itself.
- But **running one transactions at a time will give poor performance.**
- With the prevalence of multi-core architecture, DBMS should take advantage of the multiple cores.
- **Concurrency permits interleaving the transaction operations.**
  - Interleaving transactions also permits running one transaction when another is waiting for some resource (I/O, user input, or fetching data from disk).
  - Need a mechanism to interleave transactions but make it appear as if they ran one-at-a-time.

# Concurrent Transactions

- Assume that these two transactions are undergoing concurrent execution.
- What are the possible interleaving of these two transactions?

T1:

```
read(A);  
A = A - 50;  
write(A);  
read(B);  
B = B + 50;  
write(B).
```

T2:

```
read(A);  
A = A * 0.95;  
write(A);  
read(B);  
B = B * 1.05 ;  
write(B)
```

# Isolation Support: Concurrency Control

- A **concurrency control protocol** lays down the mechanism for the DBMS to decide a legal/valid schedule of transactions.
- What are the **types** of concurrency control protocols?

# Isolation Support: Concurrency Control

- A concurrency control protocol lays down the mechanism for the DBMS to decide legal/valid schedule of transactions.
- What are the types of concurrency control protocols?
- **Pessimistic:** Prevent problems from arising in the first place.

# Isolation Support: Concurrency Control

- A concurrency control protocol lays down the mechanism for the DBMS to decide legal/valid schedule of transactions.
- What are the types of concurrency control protocols?
- **Pessimistic:** Prevent problems from arising in the first place.
- **Optimistic:** Assume that conflicts are rare; deal with them after they occur.

# Isolation Support: Concurrency Control

- A concurrency control protocol lays down the mechanism for the DBMS to decide legal/valid schedule of transactions.
- What are the types of concurrency control protocols?
- **Pessimistic:** Prevent problems from arising in the first place.
- **Optimistic:** Assume that conflicts are rare; deal with them after they occur.
- **But, how does a concurrency control protocol determine a valid schedule?**

# Serializable Schedules

- **Serial Schedule** → A schedule that does not interleave the operations of different transactions.
- **Equivalent Schedule** → For any database state, the effect of executing the first schedule is identical to the effect of executing the second schedule.
- **Serializable Schedule?**

# Serializable Schedules

- **Serial Schedule** → A schedule that does not interleave the operations of different transactions.
- **Equivalent Schedule** → For any database state, the effect of executing the first schedule is identical to the effect of executing the second schedule.
- **Serializable Schedule** → A schedule that is equivalent to some serial execution of the transactions (serial schedule).

# Serializable Schedules

- **Serial Schedule** → A schedule that does not interleave the operations of different transactions.
- **Equivalent Schedule** → For any database state, the effect of executing the first schedule is identical to the effect of executing the second schedule.
- **Serializable Schedule** → A schedule that is equivalent to some serial execution of the transactions (serial schedule).
- If each transaction preserves consistency, then the corresponding serializable schedule preserves consistency!

# **Isolation Levels vs. Consistency Levels**

**Isolation Levels**

**Consistency Levels**

# Isolation Levels vs. Consistency Levels

| Isolation Levels   | Consistency Levels |
|--|--------------------|
| <ul style="list-style-type: none"><li>• Correspond to the I in ACID.</li><li>• <b>Database isolation</b> is the ability of a database to allow a transaction to execute as if there are no other concurrently running transactions.</li><li>• Greater the guaranteed isolation among the transactions, lesser the system performance.</li><li>• <b>Isolation levels</b> trade off isolation guarantees for improved performance.</li></ul> |                    |

# Isolation Levels vs. Consistency Levels

## Isolation Levels

- Correspond to the I in ACID.
- **Database isolation** is the ability of a database to allow a transaction to execute as if there are no other concurrently running transactions.
- Greater the guaranteed isolation among the transactions, lesser the system performance.
- **Isolation levels** trade off isolation guarantees for improved performance.

## Consistency Levels

- Do not correspond to C in ACID.
- Unlike the C in ACID, the **database consistency** refers to the rules that make a **concurrent, distributed system** appear as a **single-threaded, centralized system**.
- **Reads** at a particular point in time **must reflect the most recently completed write** (in real-time) of that data item, no matter which server processed that write.
- **Consistency levels** trade off read results for improved performance.

# Isolation Levels vs. Consistency Levels

- More simply said:
  - Whenever you talk about transaction isolation, you will be talking about isolation levels.
  - Whenever you talk about individual operations like read/write, you will talk about consistency levels.