

Database Processing

CS 451 / 551

Lecture 13:
Two-Phase Locking and
Time-Stamp Ordering



Suyash Gupta

Assistant Professor

Distopia Labs and ONRG
Dept. of Computer Science

(E) suyash@uoregon.edu
(W) gupta-suyash.github.io



Assignment 3 is Out!
Deadline: Nov 30, 2025 at 11:59pm

Final Exam: Dec 8, 2025 at 8-10am

Syllabus → Focus on course not covered in either Quiz but you should remember indexing and storage.

Presentation

- Time slots to be released today around 11am PST.
- Each group gets a 15min time slot.
- 8 minutes to present + 7 minutes for Q/A
 - If 4 group members → 2 min for each member.
- Don't present your code.
- Present your idea.
 - How did you design?
 - Why did you select such a design?
 - Is there anything cool about your design.
 - How did each member contribute?

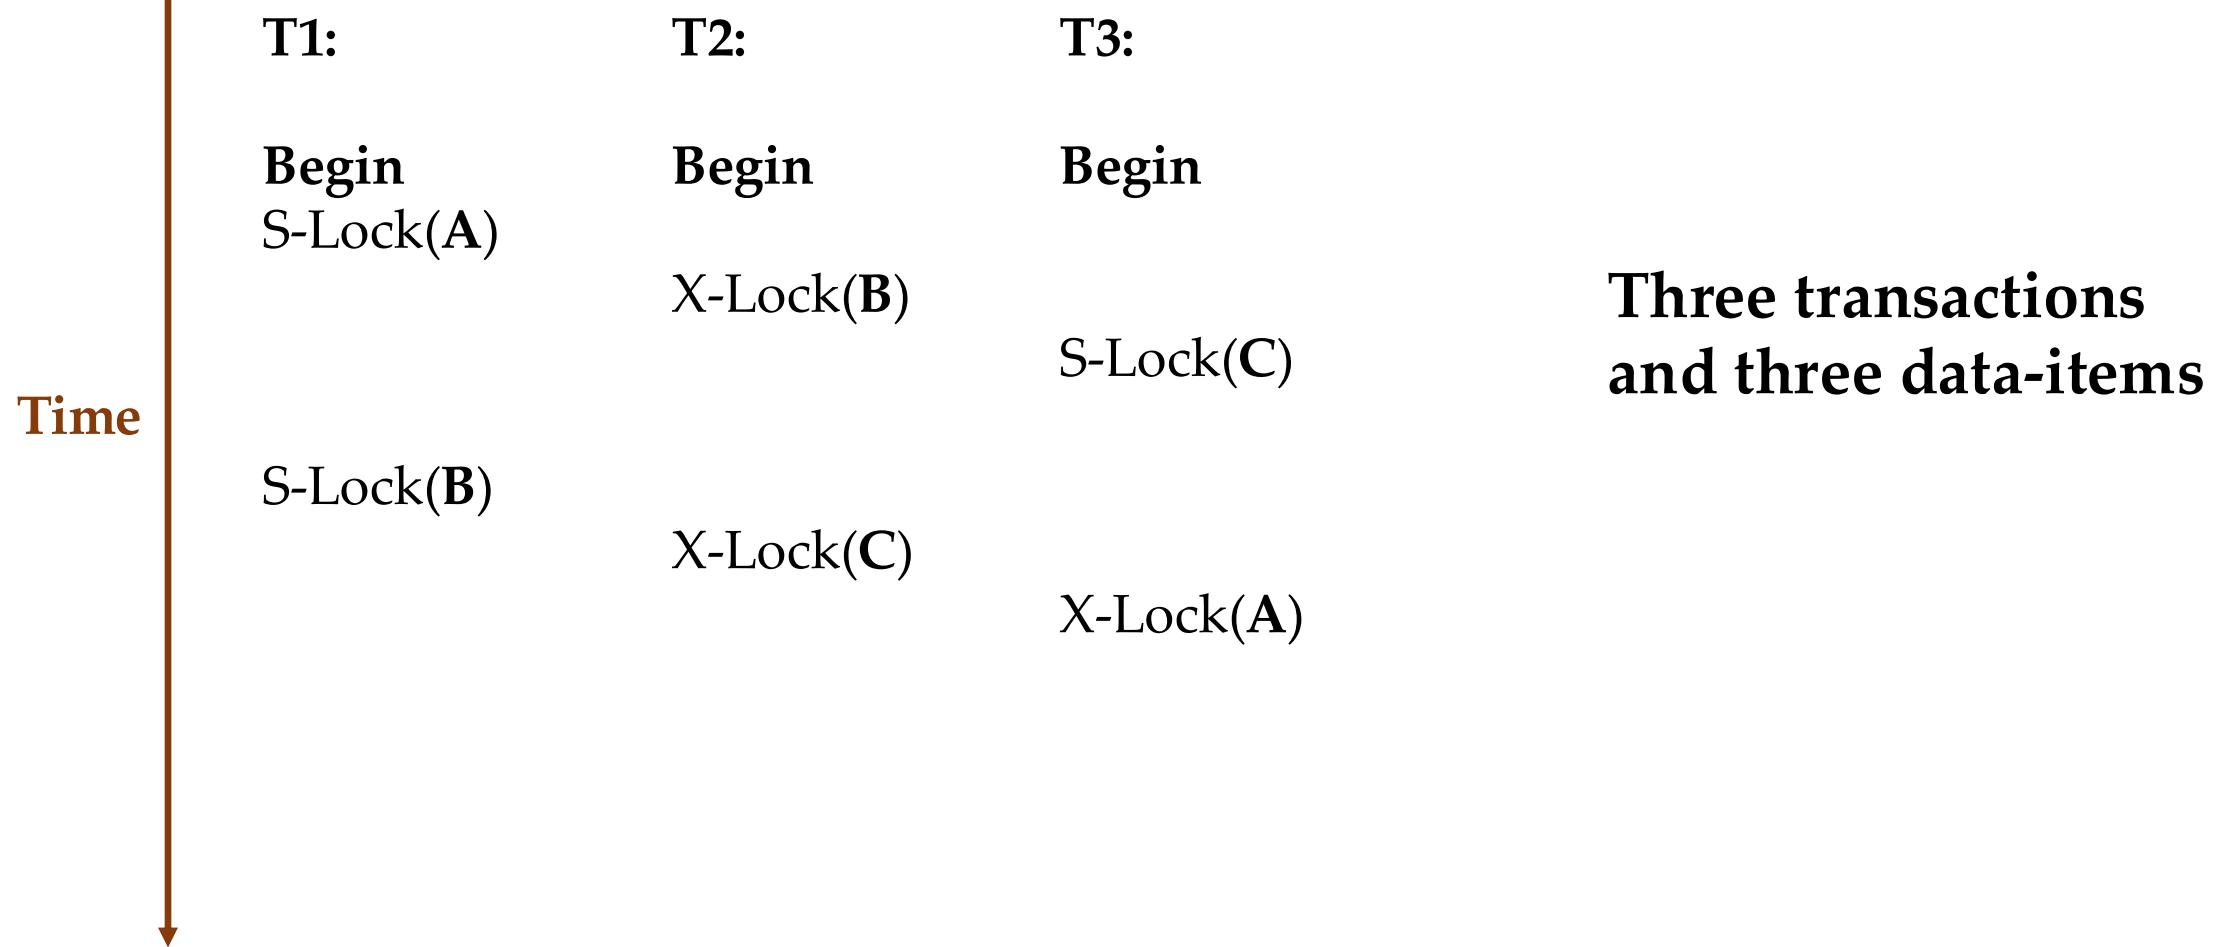
Last Class

- We discussed possibility of deadlocks in 2PL.
- There are two ways to manage deadlocks:
- **Deadlock Detection** → When deadlock occurs, detect and solve.
- **Deadlock Prevention** → Prevent deadlock from occurring in the first place.

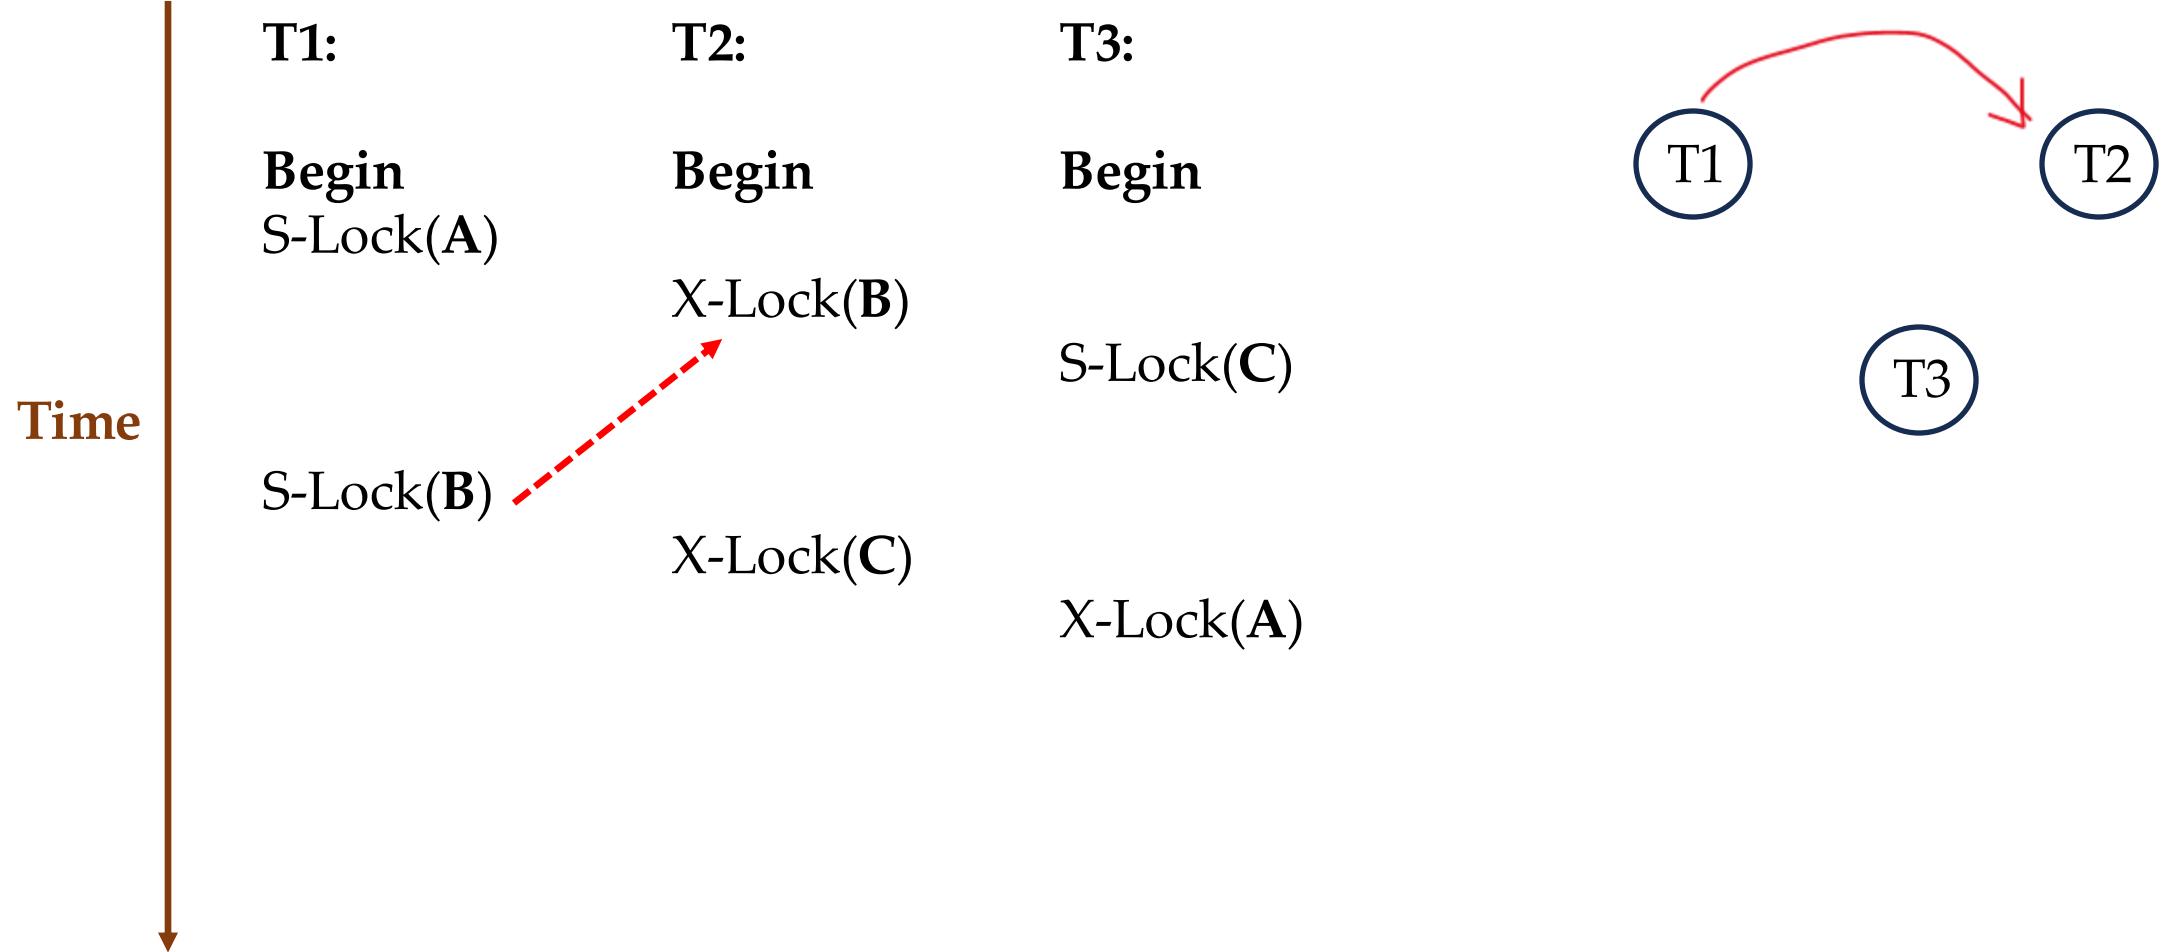
Deadlock Detection

- Create a **waits-for graph**.
- Waits-for graph keep track of what locks each transaction is waiting to acquire.
- In the wait-for graph:
 - **Nodes** are transactions
 - **Add an Edge** from transaction T_i to T_j if T_i is waiting for T_j to release a lock.
 - The system periodically **checks for cycles** in waits- for graph and then decides **how to break it**.

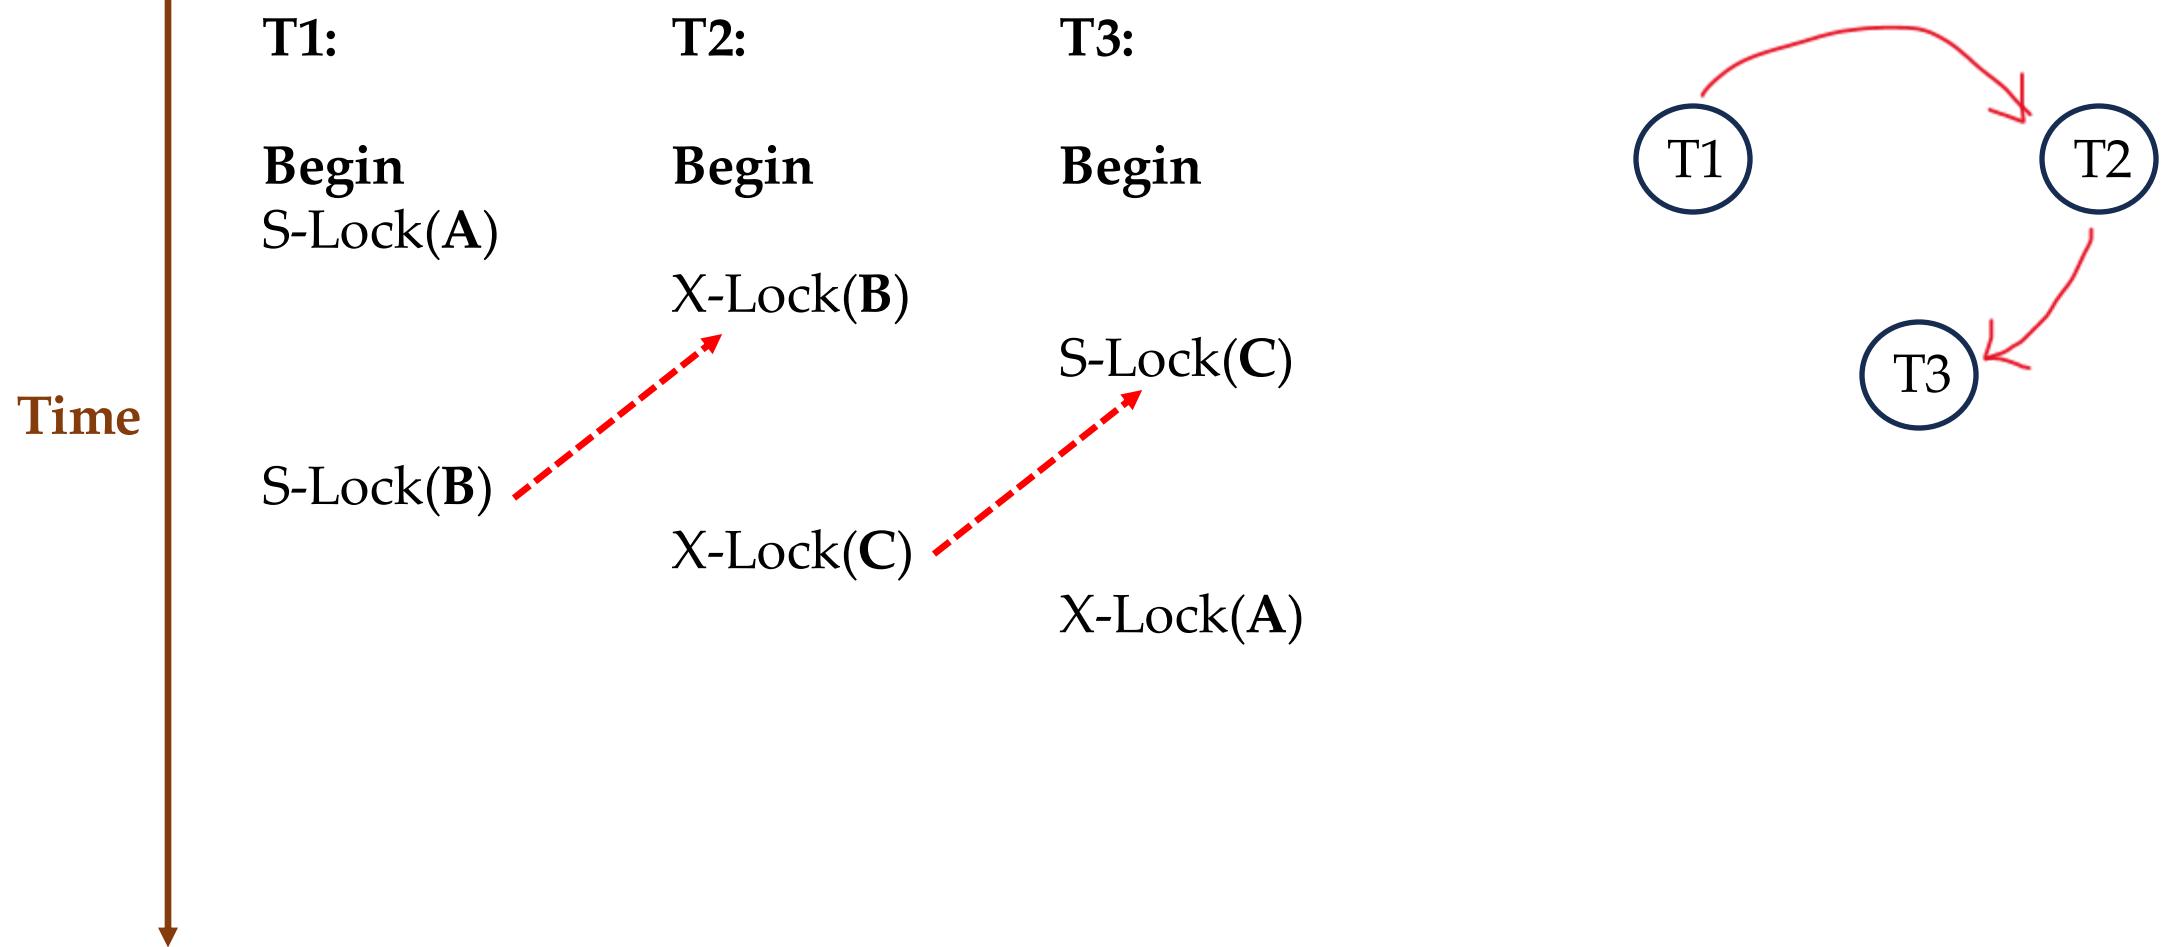
Deadlock Detection



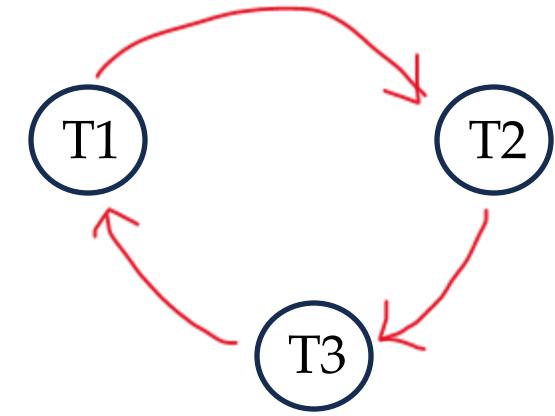
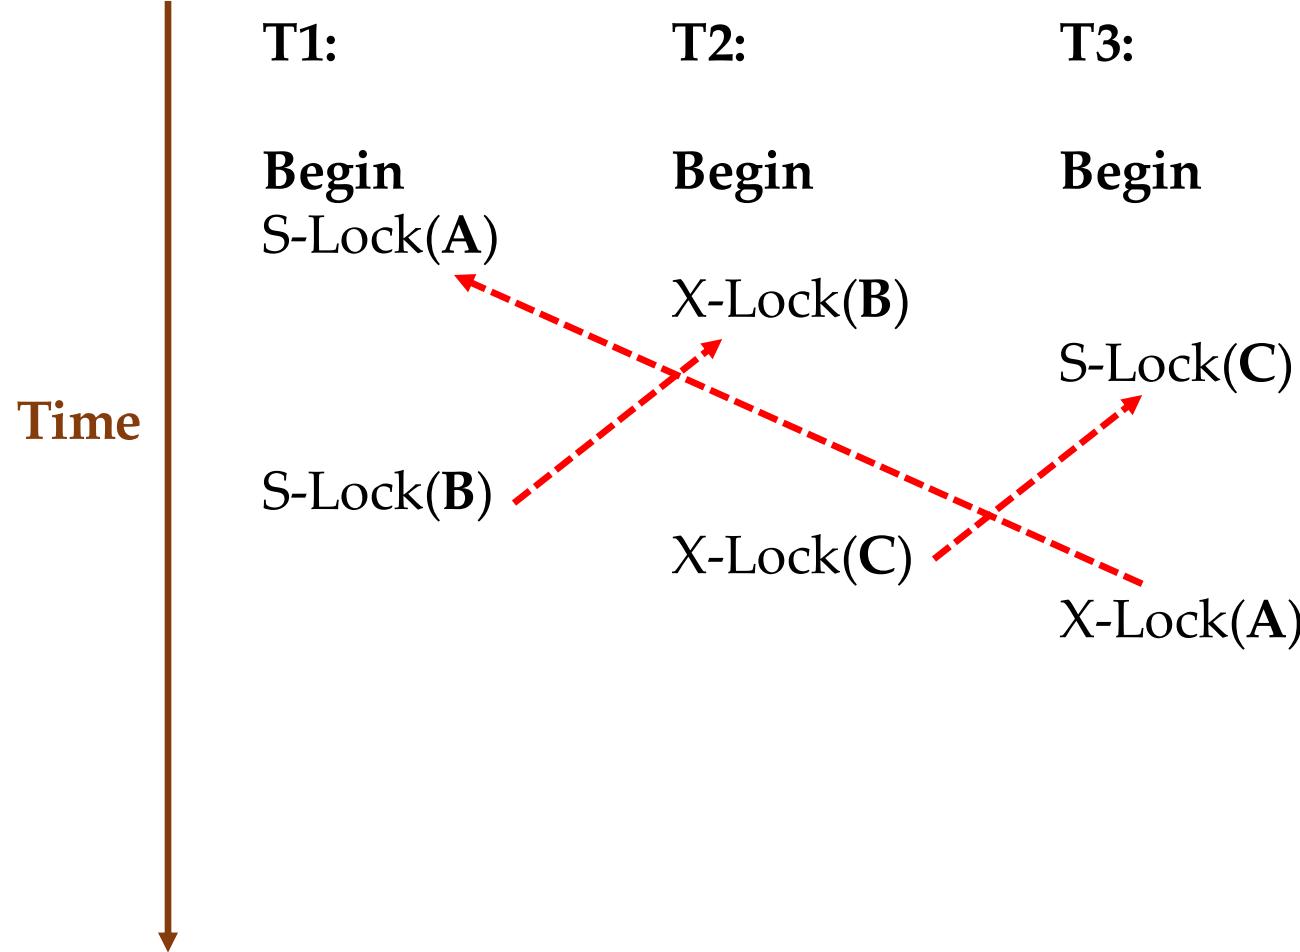
Deadlock Detection



Deadlock Detection



Deadlock Detection



Deadlock Handling

- When the DBMS detects a deadlock, it will select a **victim**.
- The victim transaction is **rolled back** to break the cycle.
- The victim transaction is either **restarted in the future** or **aborted**.
- Performance trade-off between the frequency of checking for deadlocks and the time transactions wait before deadlocks are broken.

Deadlock Handling: Victim Selection

- Selecting a victim depends on several factors:
 - Age (lowest timestamp).
 - Progress (least/most queries executed)
 - The number of items already locked.
 - The number of transactions that need to be rolled back along with it.
- Note: if your DBMS plans to restart the rolled back transaction, do take into account **starvation**.
 - The number of times a transaction has been restarted in the past.

Deadlock Prevention

Deadlock Prevention

- If a transaction **tries to acquire a lock** that is held by another transaction, **kill** one of them to prevent a deadlock.
- No need for a waits-for graph or detection algorithm.

Deadlock Prevention

- So how to achieve deadlock prevention?

Deadlock Prevention

- So how to achieve deadlock prevention?
- Assign a **timestamp** (time of arrival in the system) to each transaction.
- **Prioritize** transactions based on the value of timestamps.
 - For example: Higher timestamp, lower the priority.

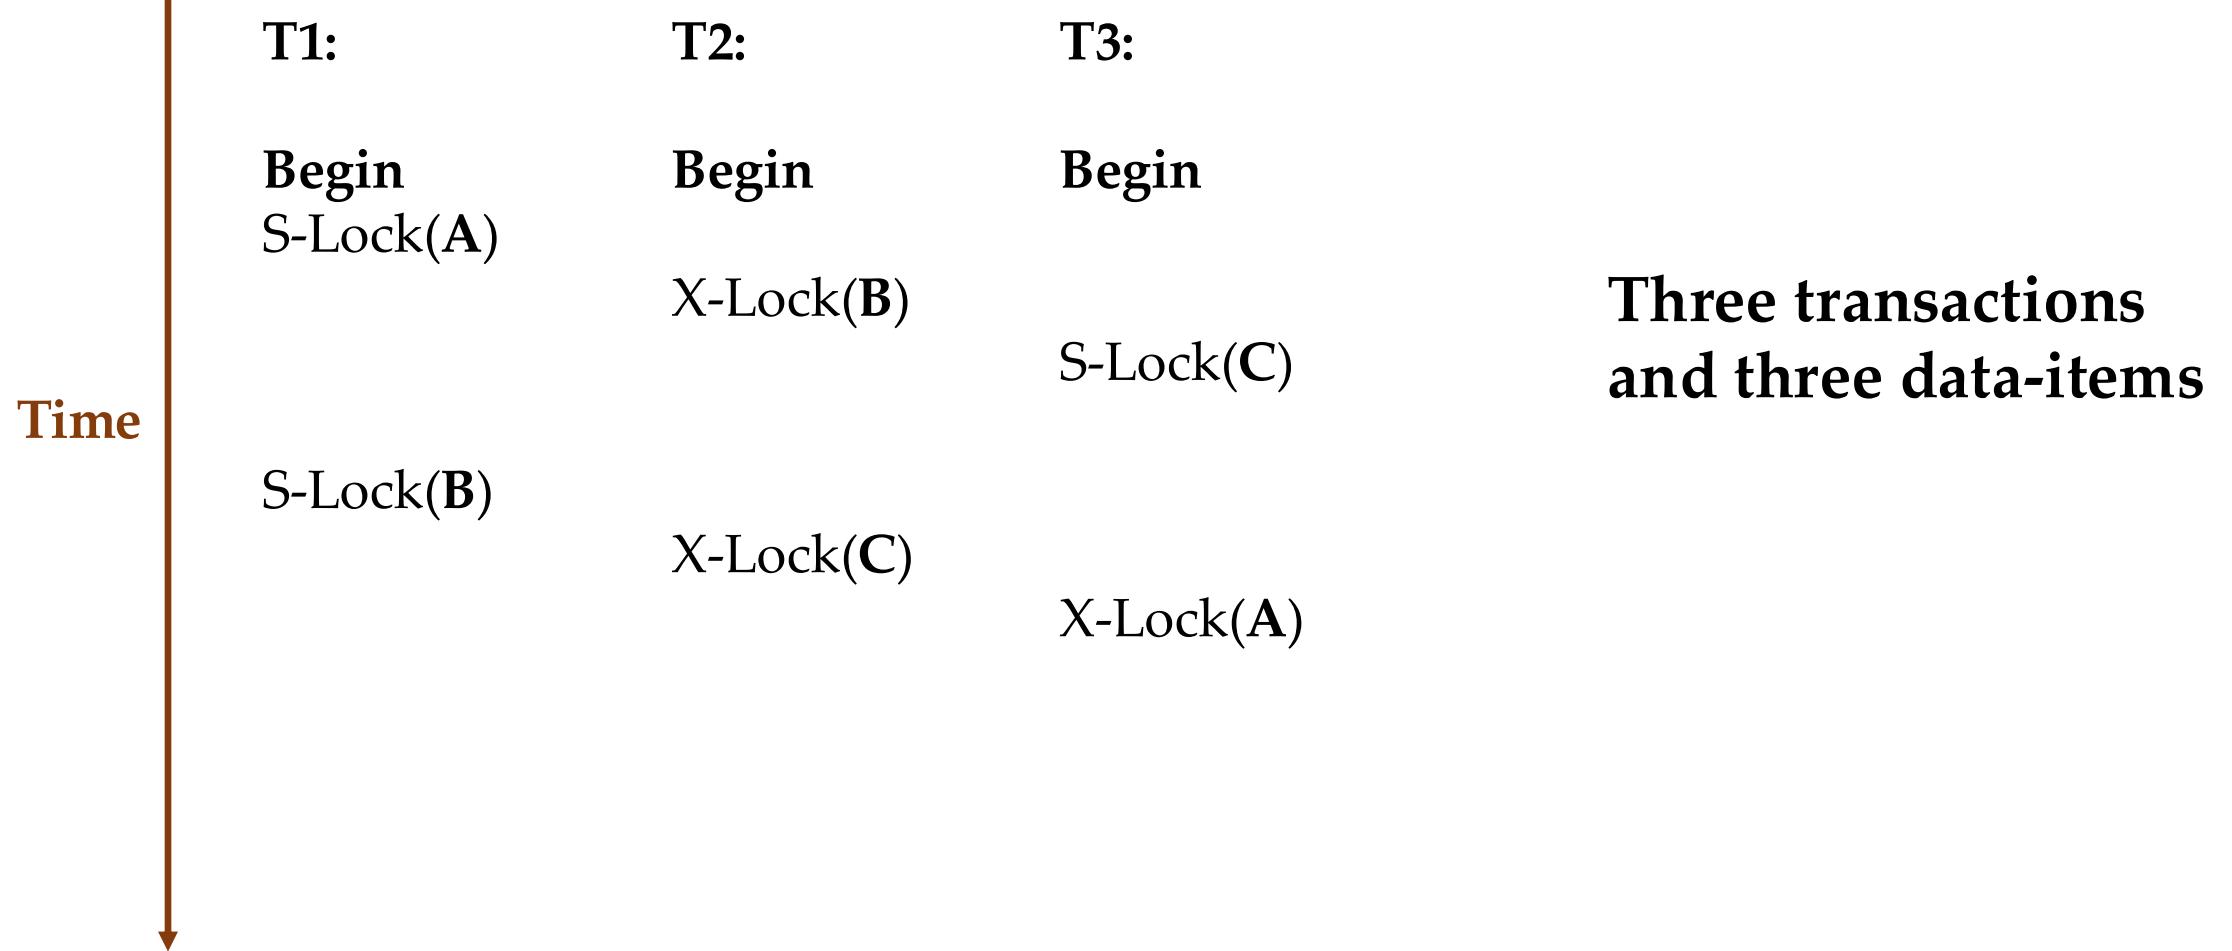
Deadlock Prevention

- So how to achieve deadlock prevention?
- Assign a **timestamp** (time of arrival in the system) to each transaction.
- **Prioritize** transactions based on the value of timestamps.
 - For example: Higher timestamp, lower the priority
- Two Designs:
 - **Wait-Die**
 - **Wound-Wait**

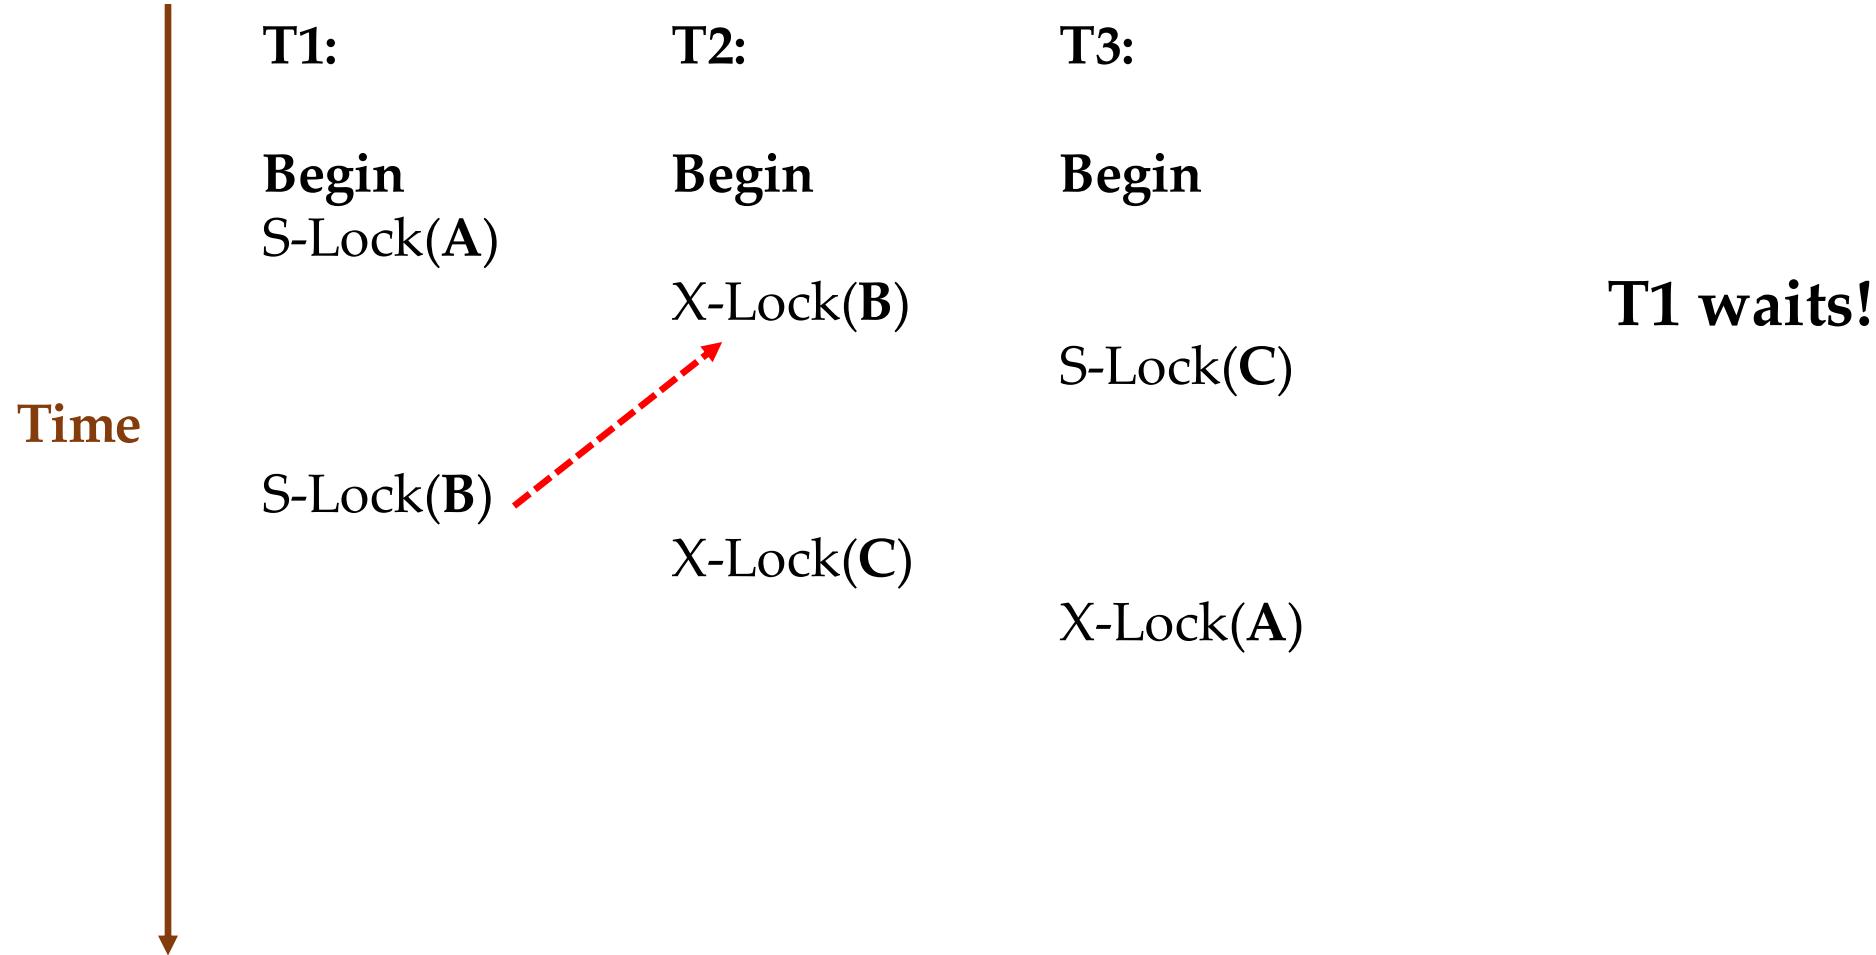
Deadlock Prevention: Wait-Die

- **Old Waits for Young**
- If requesting transaction has higher priority than holding transaction, then requesting transaction waits for holding transaction.
- Otherwise requesting transaction aborts.

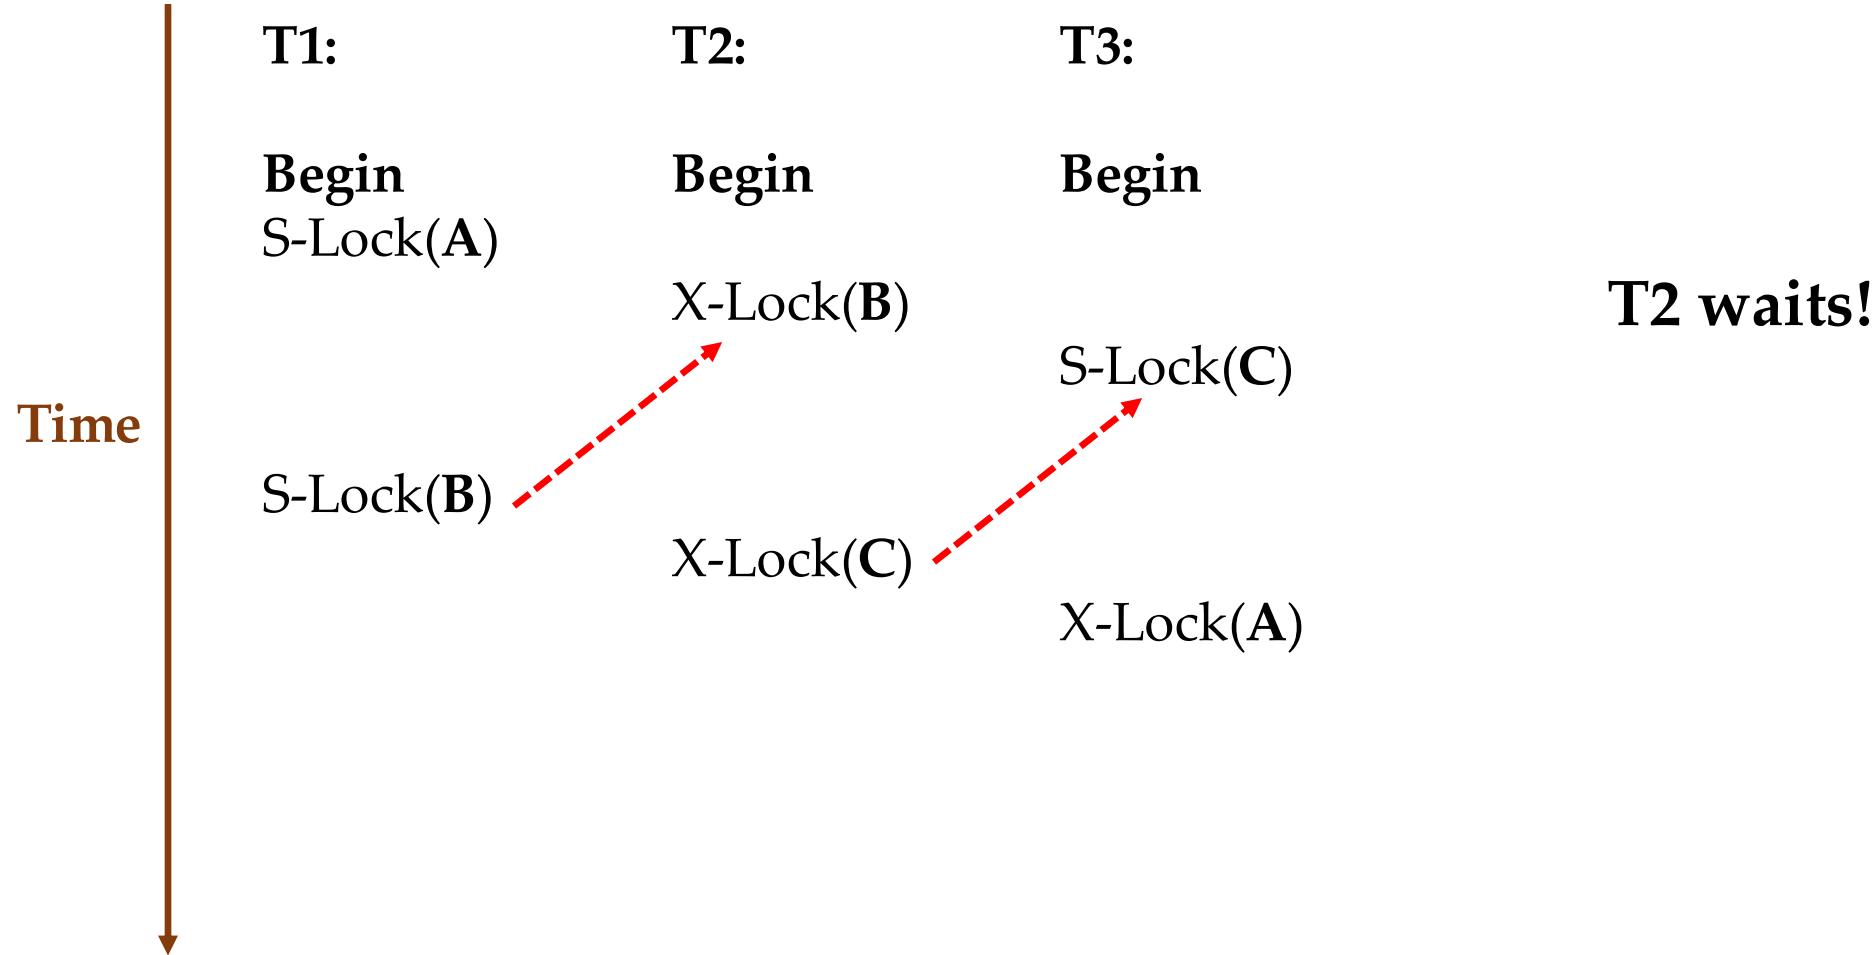
Deadlock Prevention: Wait-Die



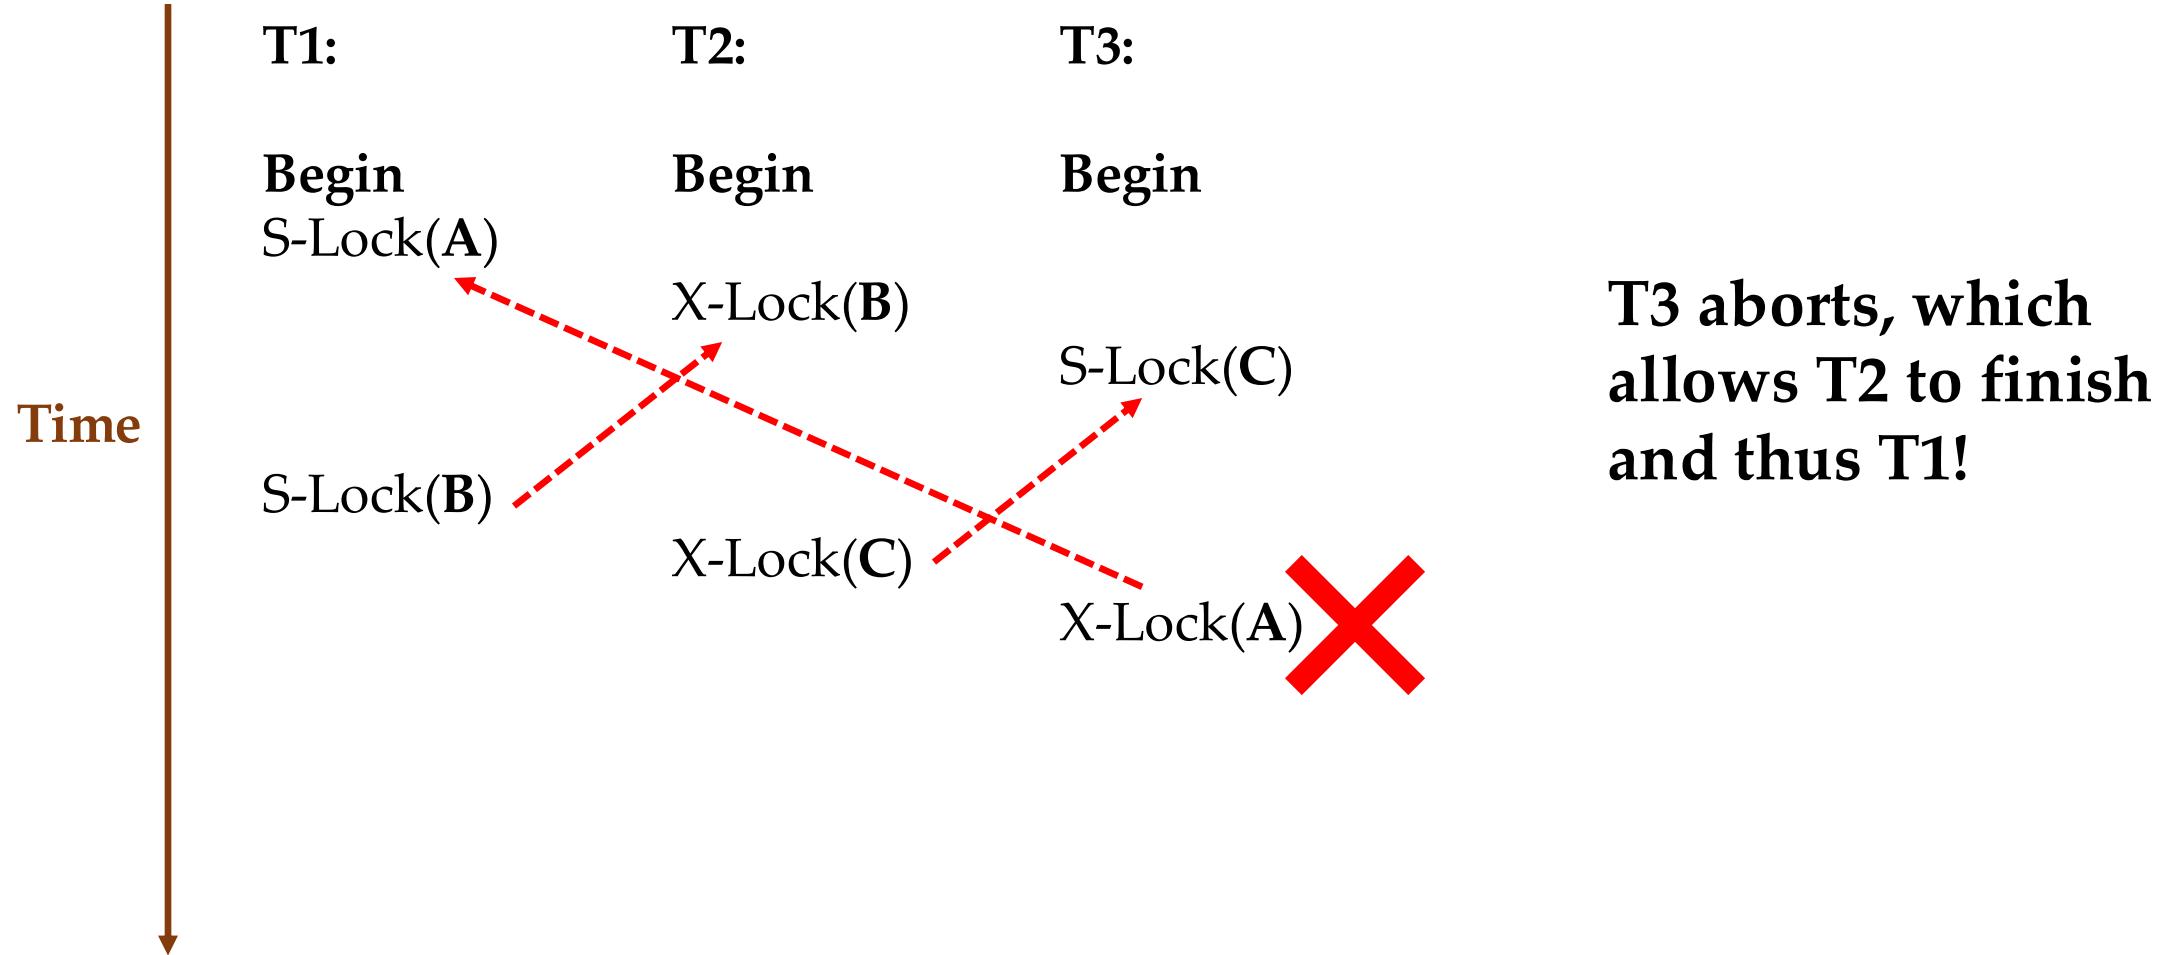
Deadlock Prevention: Wait-Die



Deadlock Prevention: Wait-Die



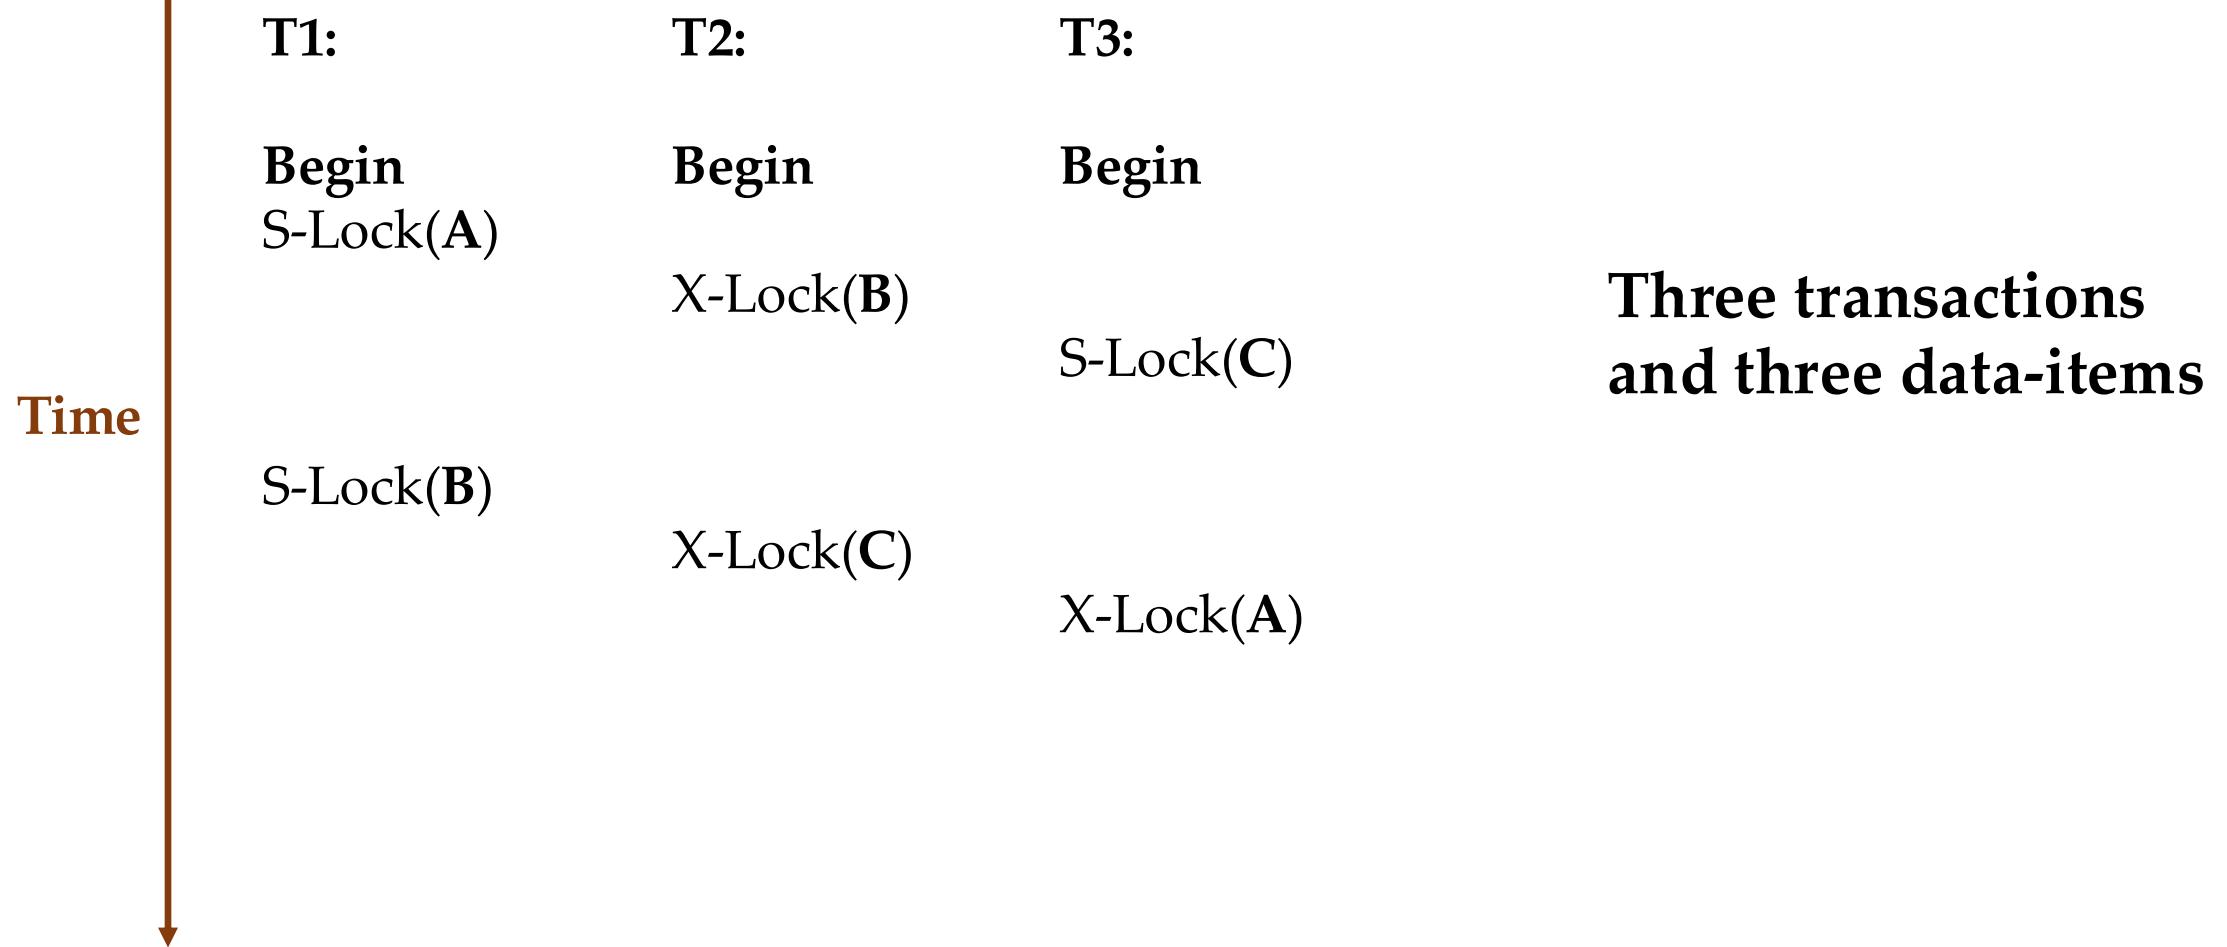
Deadlock Prevention: Wait-Die



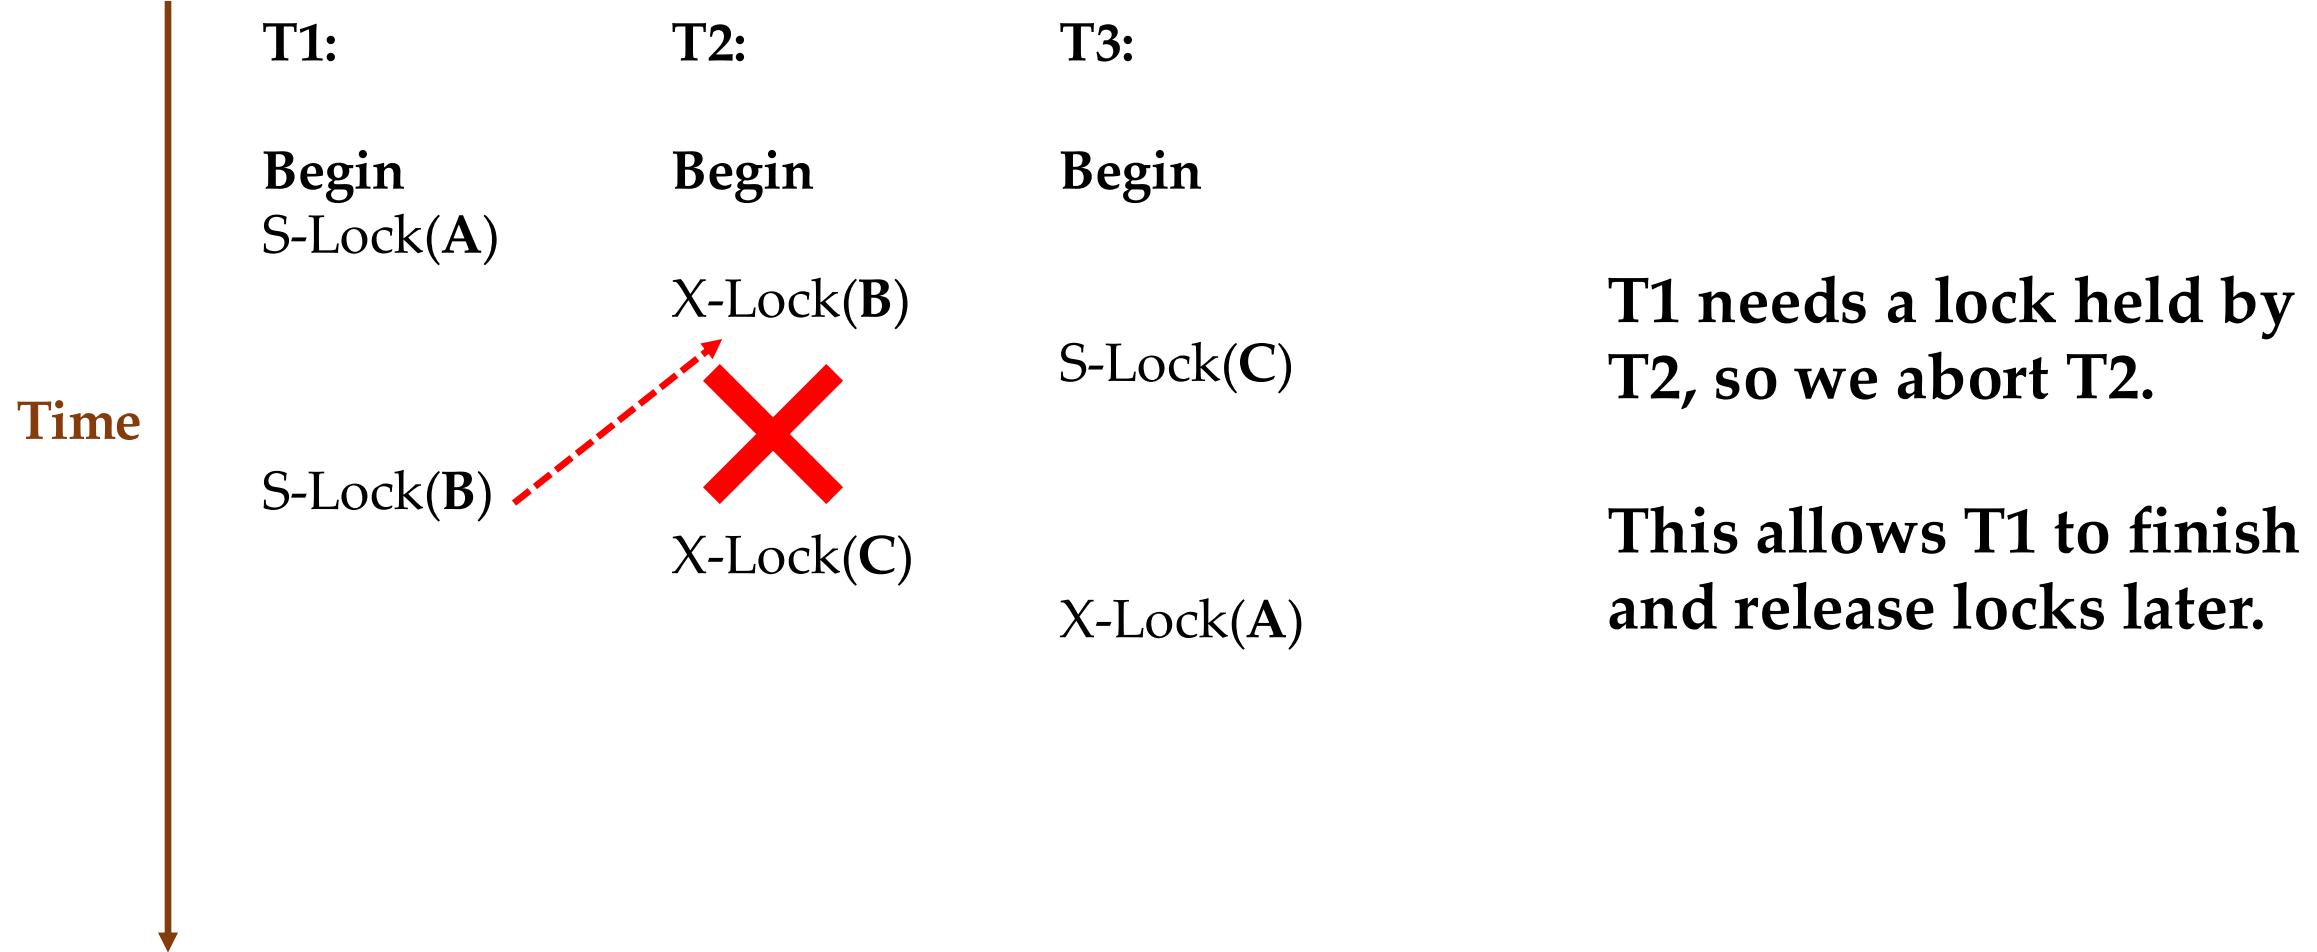
Deadlock Prevention: Wound-Wait

- **Young Waits for Old**
- If requesting transaction has higher priority than holding transaction, then holding transaction aborts and releases lock.
- Otherwise requesting transaction waits.

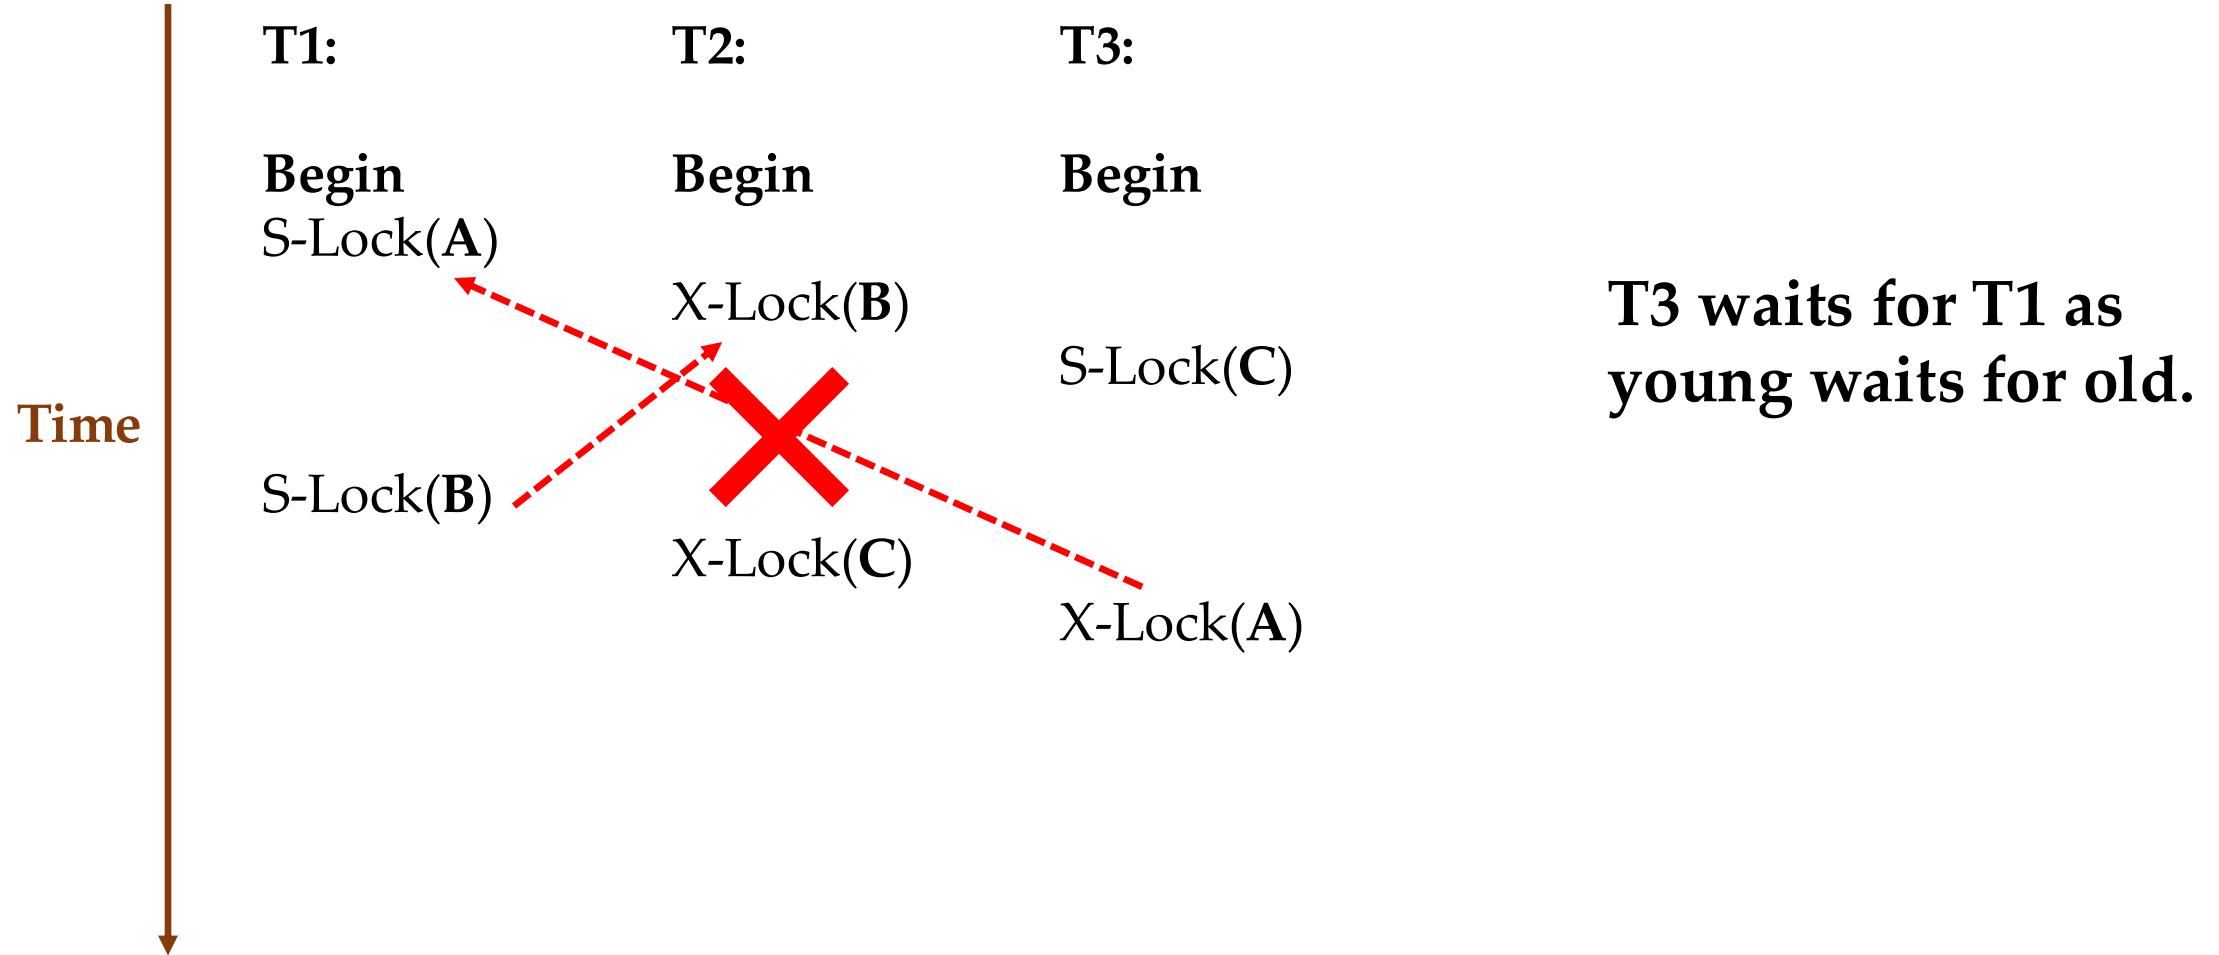
Deadlock Prevention: Wound-Wait



Deadlock Prevention: Wound-Wait



Deadlock Prevention: Wound-Wait



Deadlock Prevention

- Why do these schemes guarantee no deadlocks?

Deadlock Prevention

- Why do these schemes guarantee no deadlocks?
- Ensure that waiting for locks occur in only one.
- When a transaction restarts, what is its priority?

Deadlock Prevention

- Why do these schemes guarantee no deadlocks?
- Ensure that waiting for locks occur in only one.
- When a transaction restarts, what is its priority?
- Its original timestamp to prevent the transaction from starving.

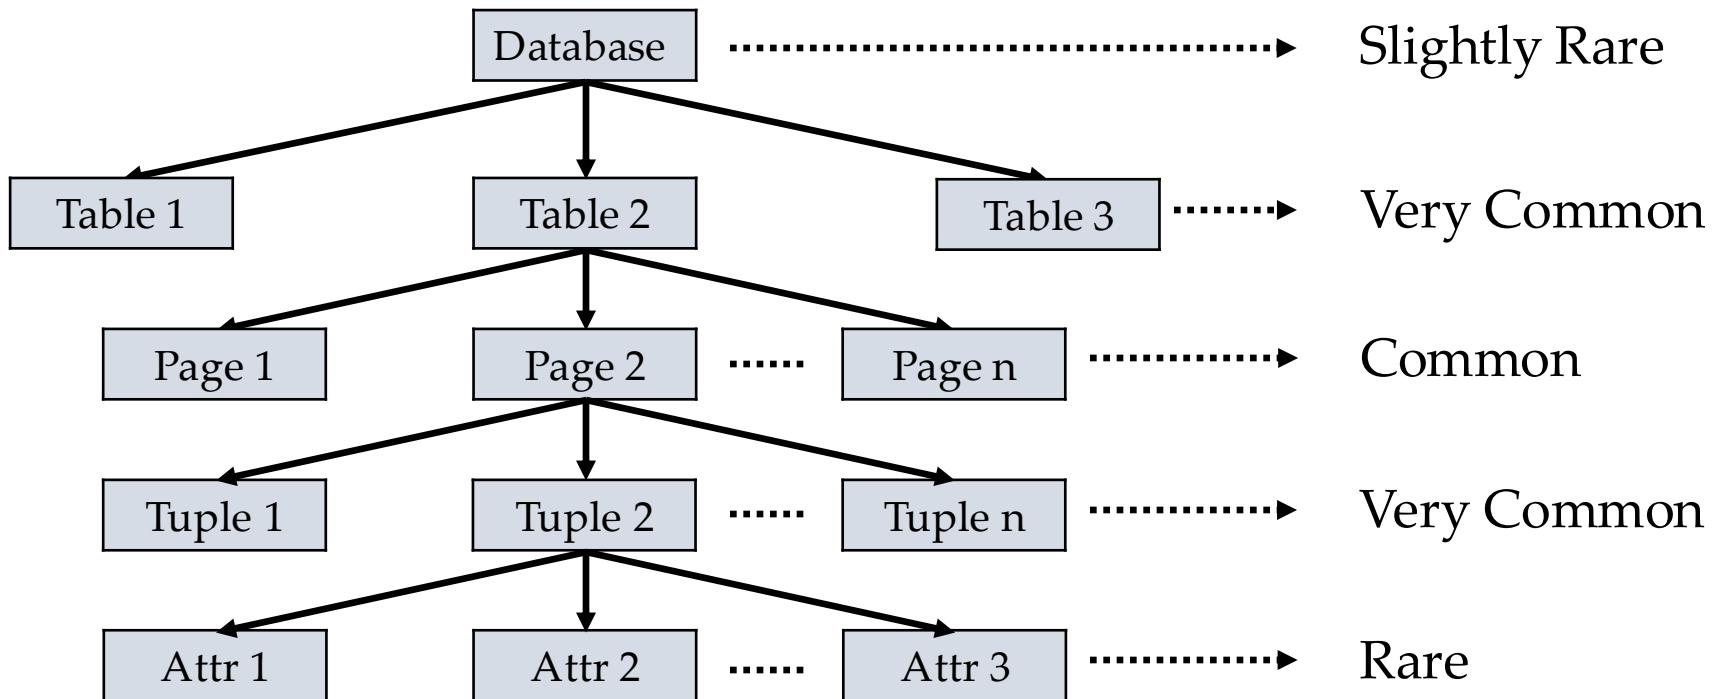
Lock Granularities

- What is the right granularity of acquiring a lock?

Lock Granularities

- What is the right granularity of acquiring a lock?
- The DBMS needs to decide the lock granularity: page, tuple, or attribute?
- Finer the lock granularity, better the performance and harder to guarantee code correctness.
- Finer the lock granularity, frequent the need to request/acquire locks.

Lock Granularities



Lock Granularities

Say a transaction T1 has locked some attribute in Table 2/ Page 1 and another transaction T2 wants to lock the full database, how to check if T2's request can / cannot be satisfied?

Intention Locks

- An **intention lock** allows locking a higher-level node in shared or exclusive mode without checking all the descendent nodes.
- If a node is locked in an **intention mode**, then some transaction has acquired a lock at the lower level in the tree.

Intention Locks

- **Intention-Shared (IS)**
 - Indicates explicit locking at lower level with S-Locks.
 - Intent to get S-Lock(s) at finer granularity.
- **Intention-Exclusive (IX)**
 - Indicates explicit locking at lower level with X-Locks.
 - Intent to get X-Lock(s) at finer granularity.
- **Shared+Intention-Exclusive (SIX)**
 - The subtree rooted by that node is locked explicitly in S mode and at a further lower level explicit locking with X-Locks.

Lock Compatibility Matrix

If a transaction T_i holds a lock, can another transaction T_j acquire a lock.

		Tj Wants				
		IS	IX	S	SIX	X
Ti Holds	IS	✓	✓	✓	✓	✗
	IX	✓	✓	✗	✗	✗
	S	✓	✗	✓	✗	✗
	SIX	✓	✗	✗	✗	✗
	X	✗	✗	✗	✗	✗

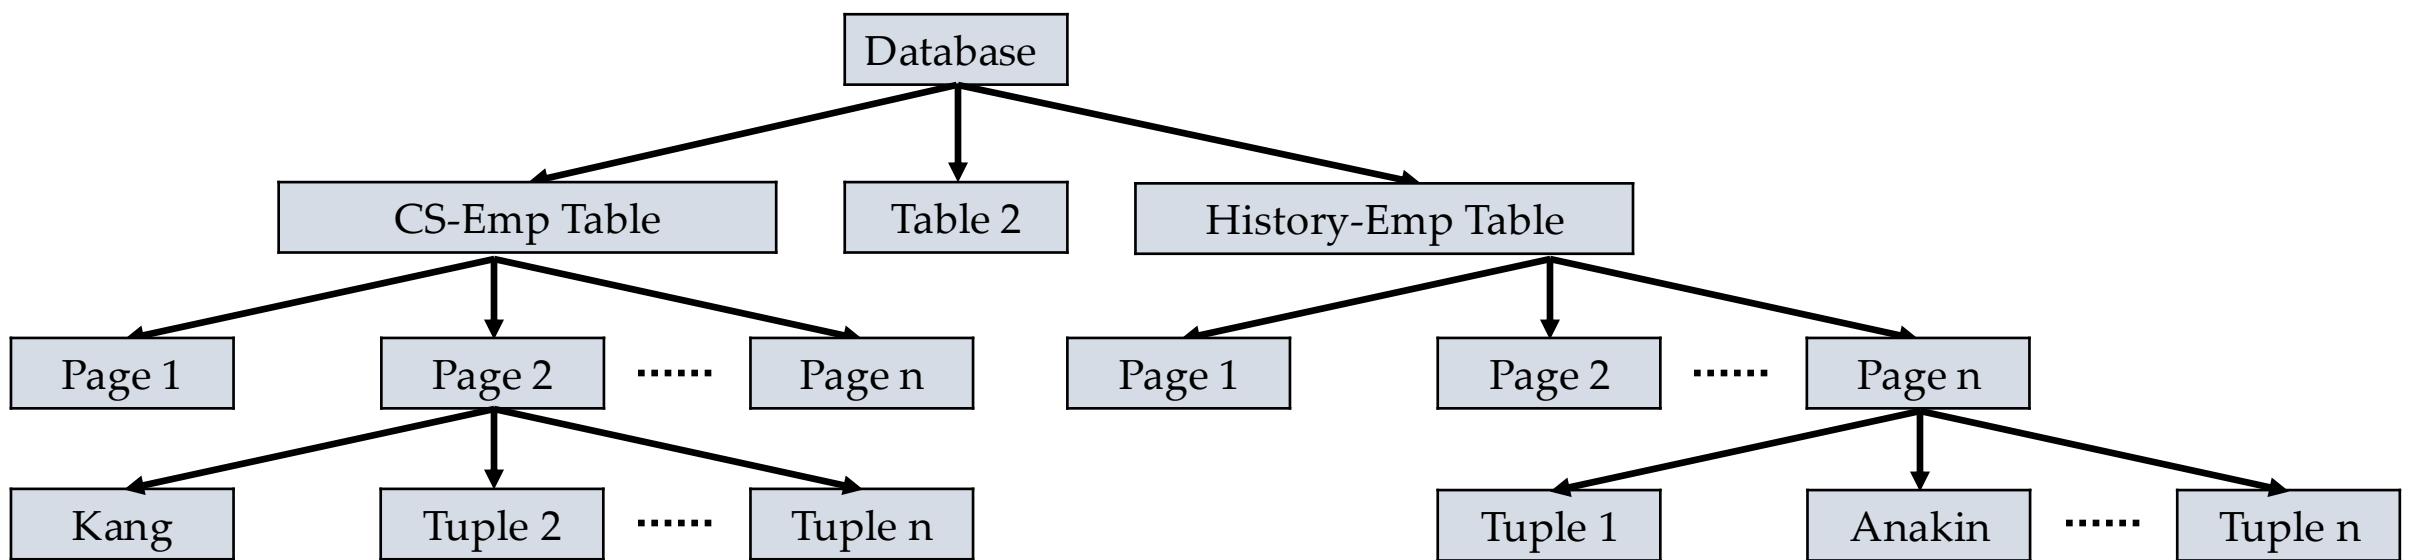
Locking Protocol

- A transaction tries to fetch an appropriate lock at highest level of the database hierarchy.
- To get **S** or **IS** lock on a node, the transaction must have at least **IS** lock on the parent node.
- To get **X**, **IX**, or **SIX** on a node, the transaction must have at least **IX** lock on the parent node.

Example 1

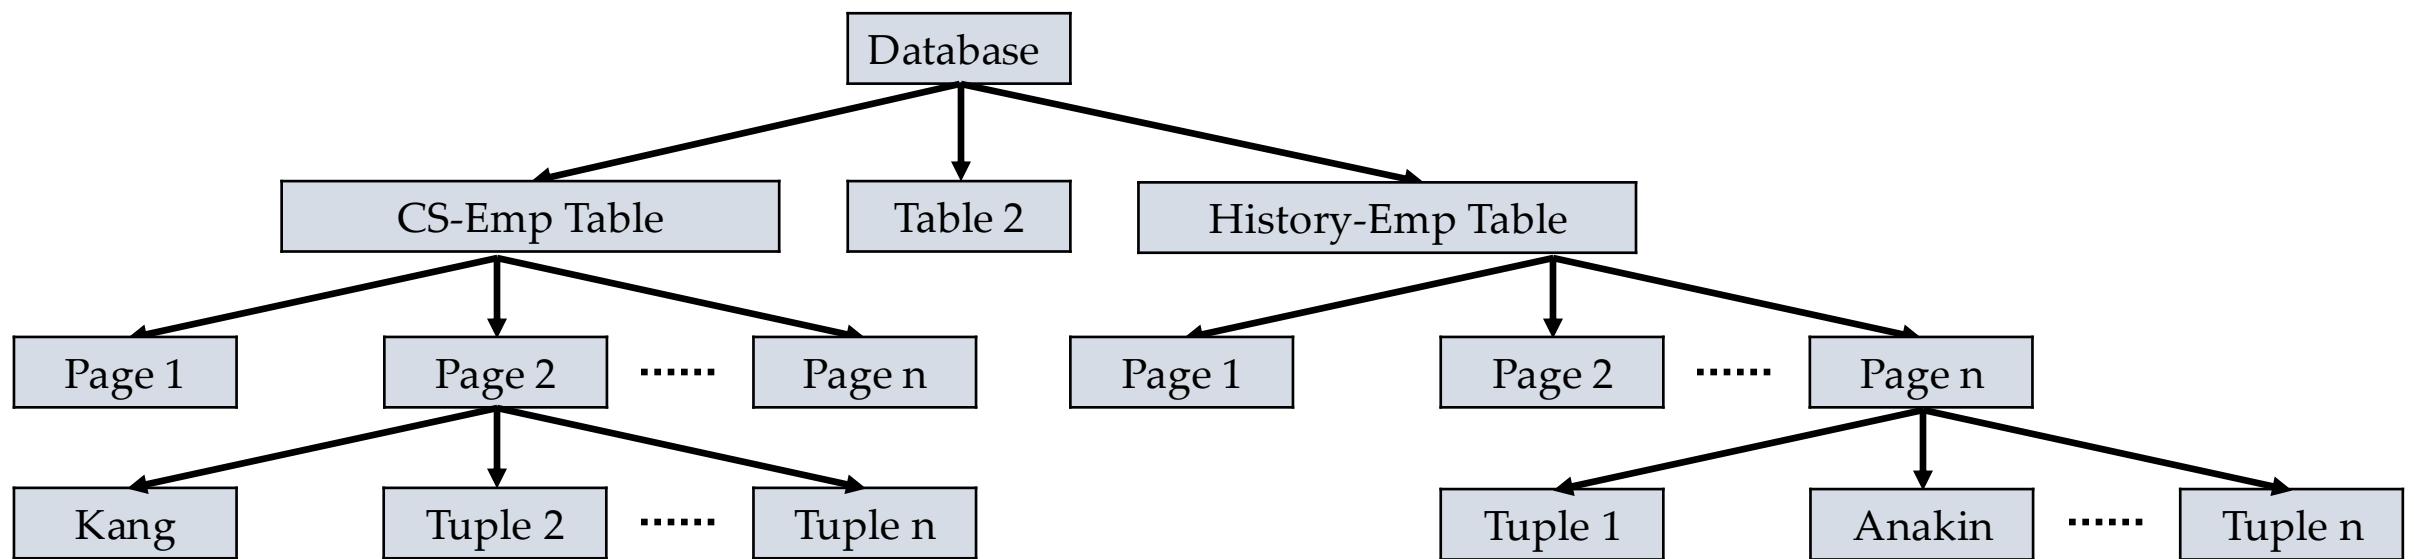
- T1 → Read a CS employee Kang's salary.
- T2 → Increase a History employee Anakin's salary by \$100.
- How can we acquire locks in this example?

Example 1



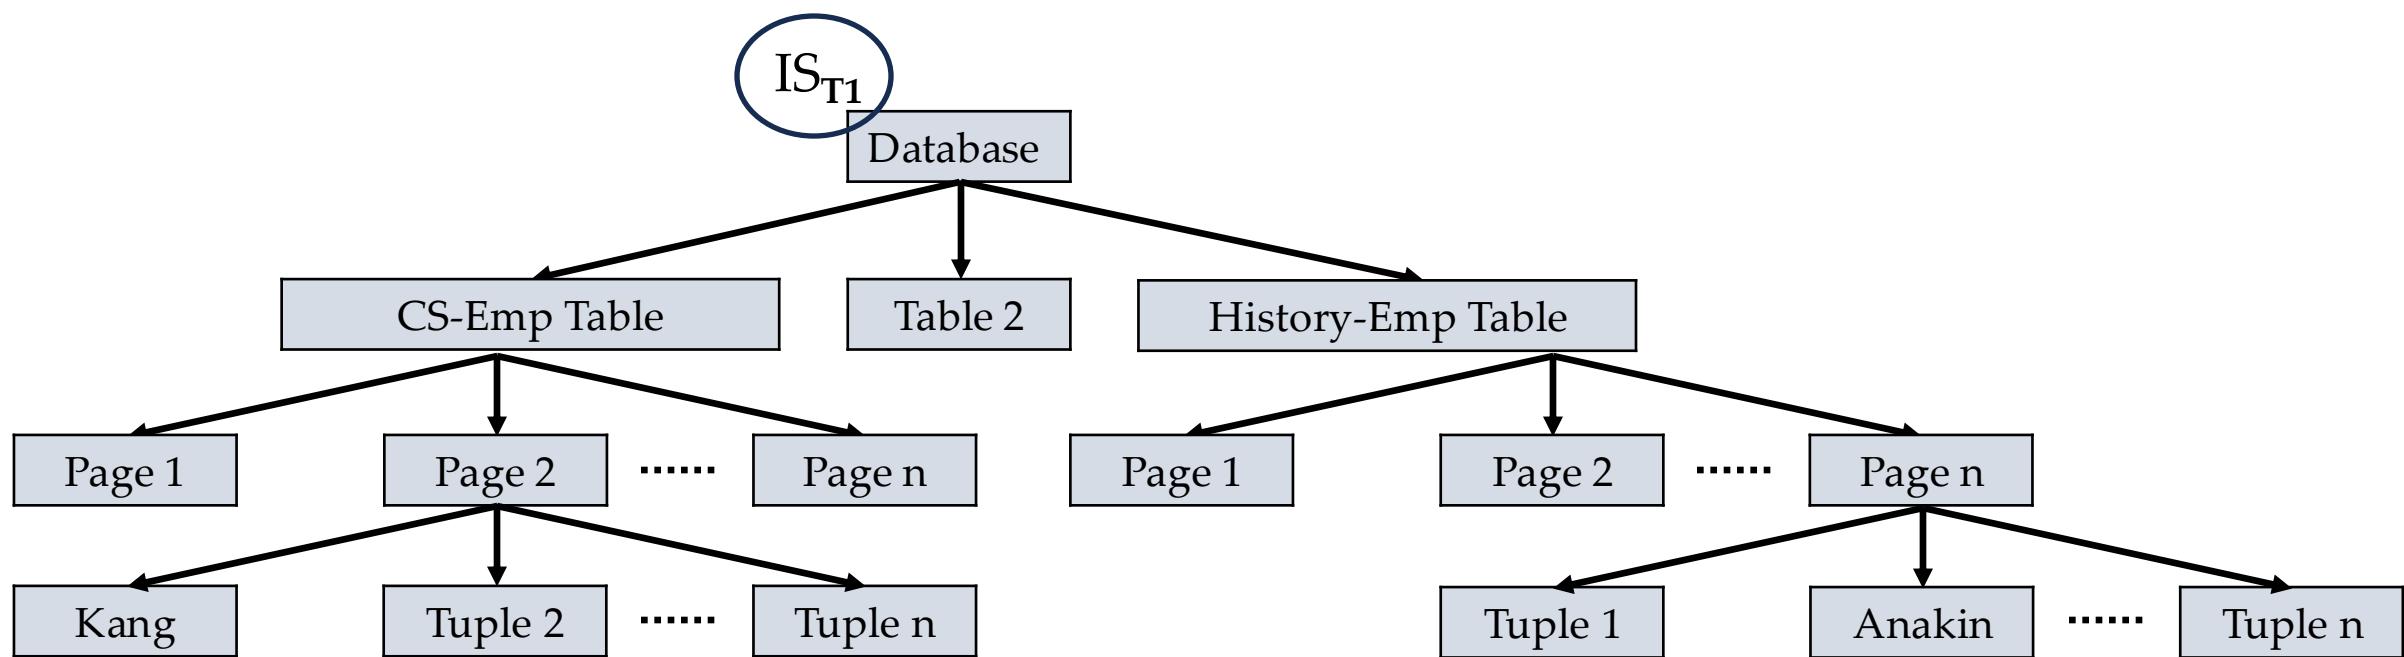
Example 1

T1 starts locking



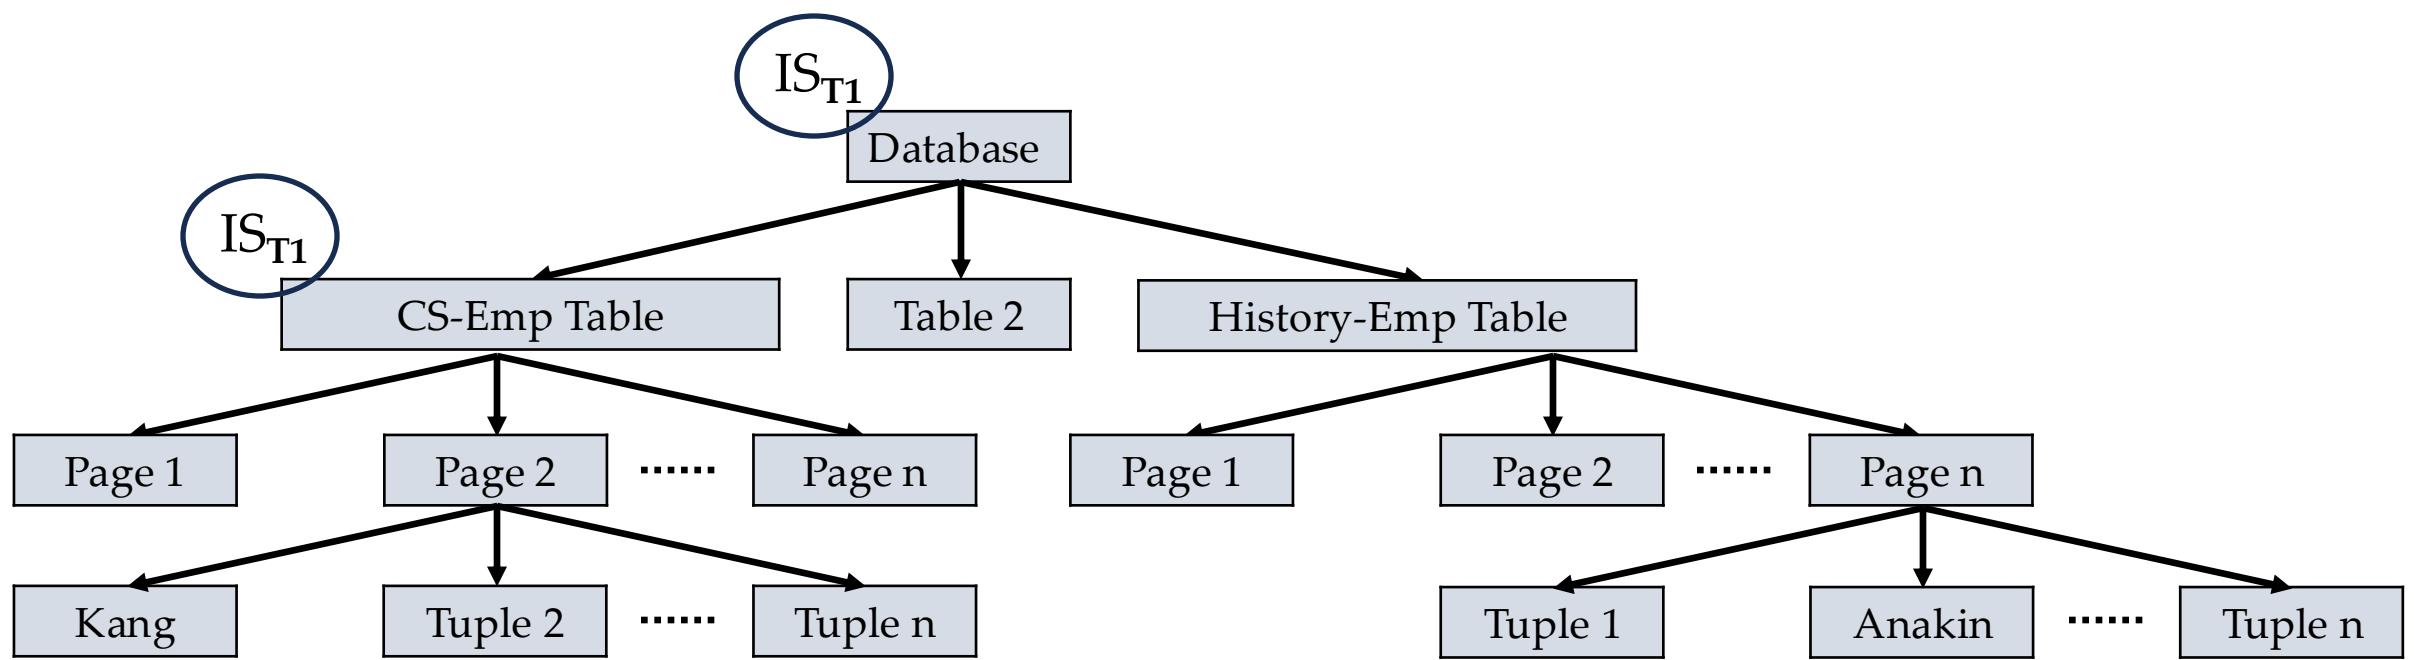
Example 1

T1 starts locking



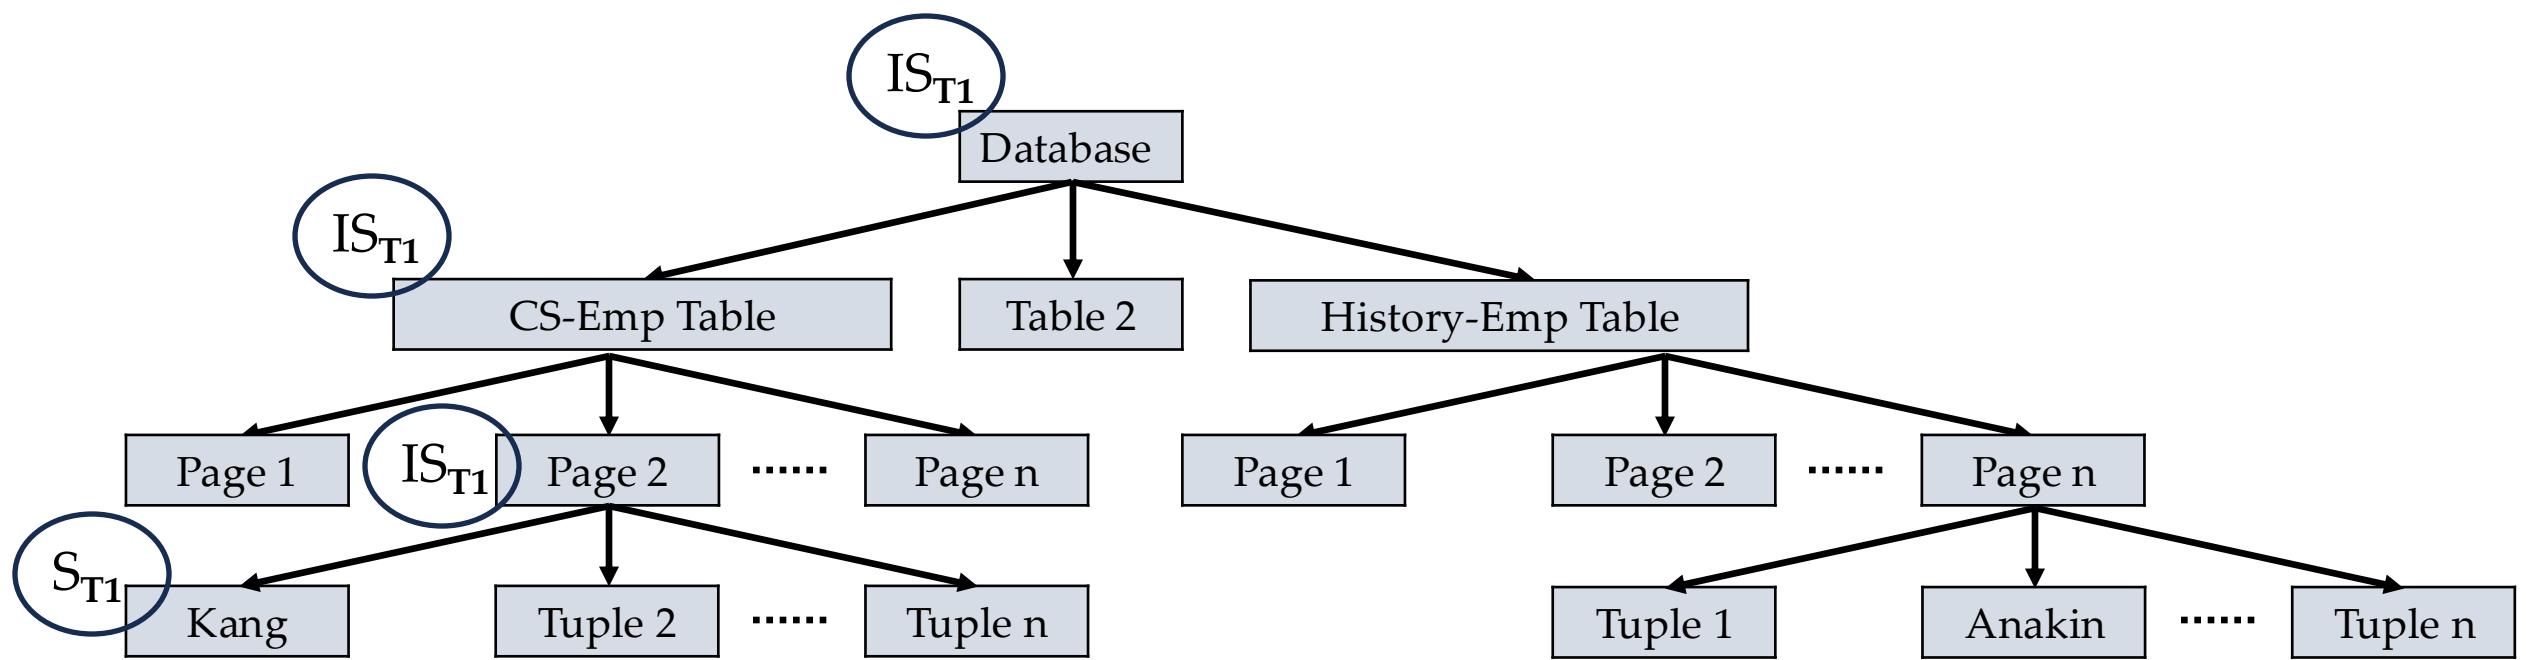
Example 1

T1 starts locking



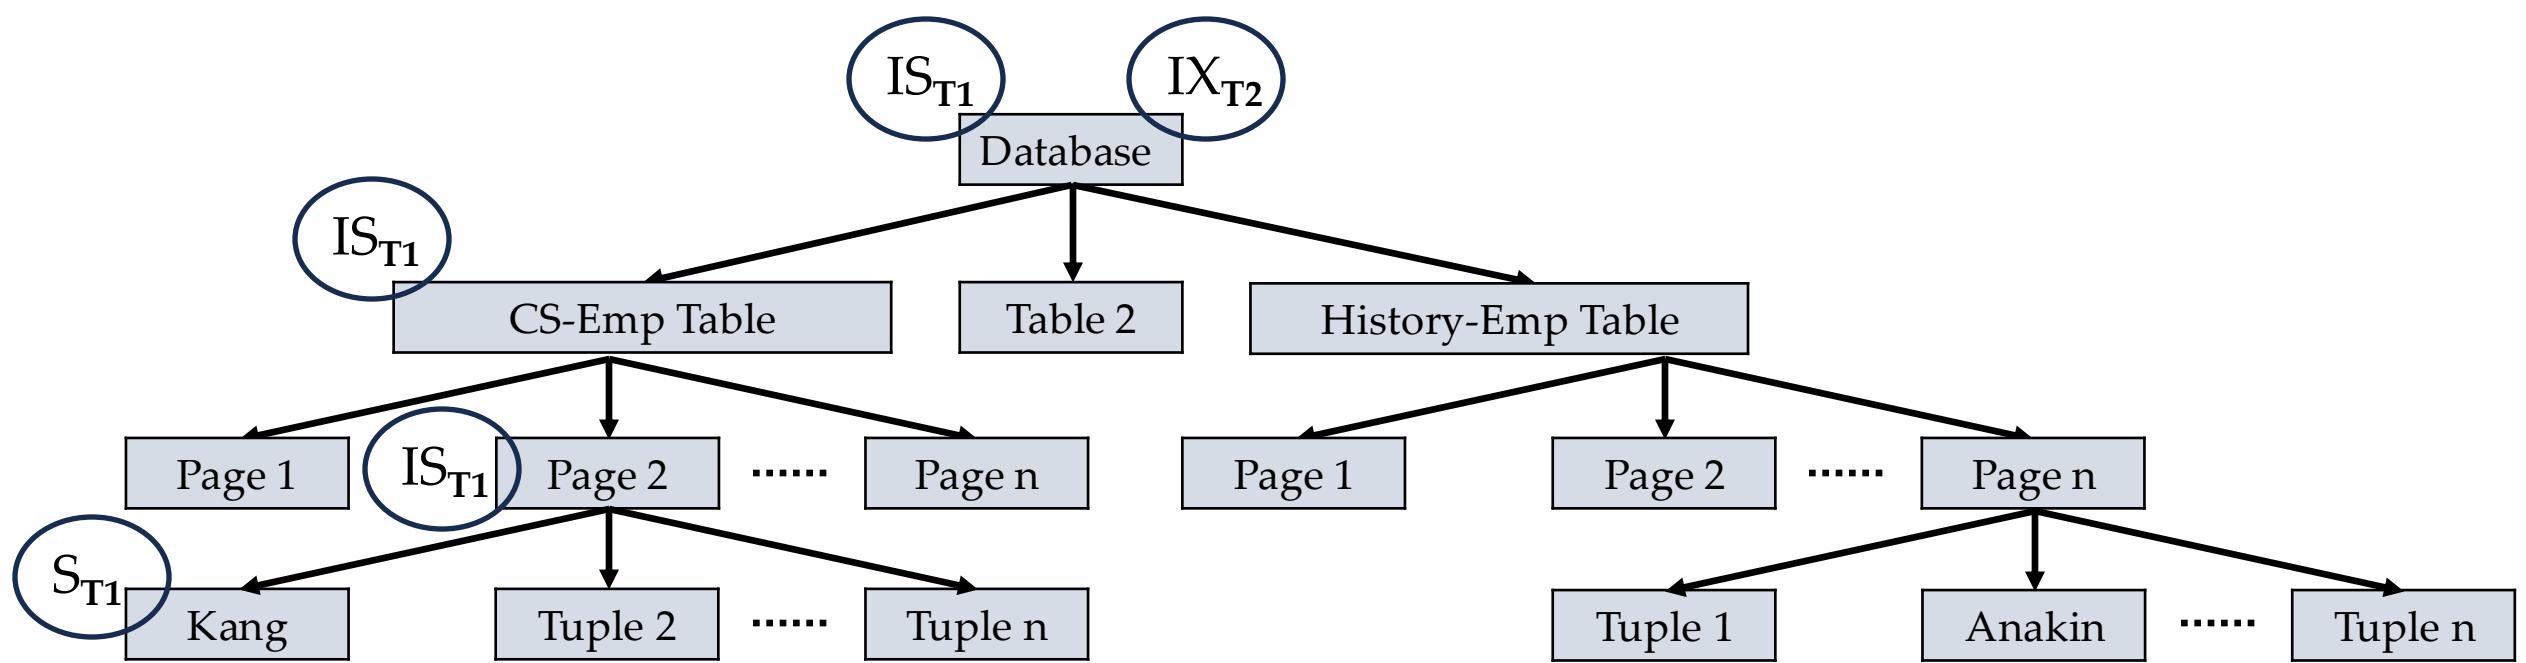
Example 1

T1 starts locking



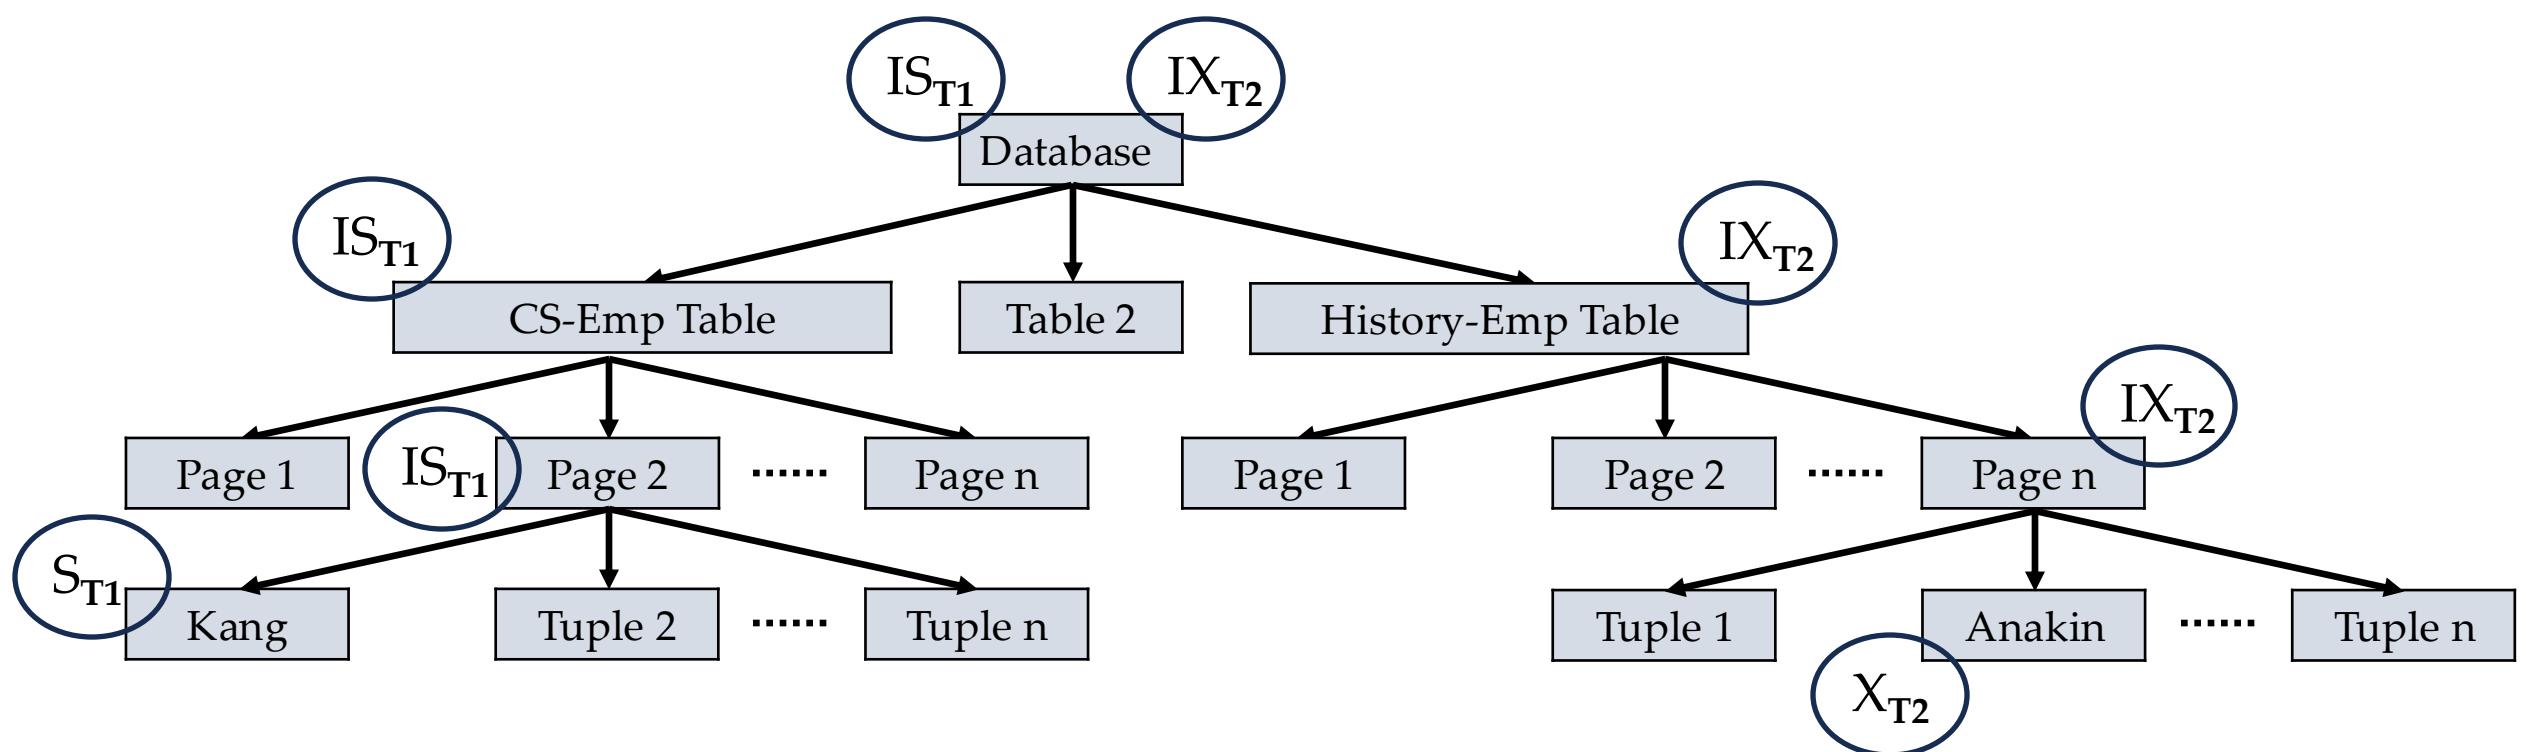
Example 1

T2 starts locking and IS and IX are compatible



Example 1

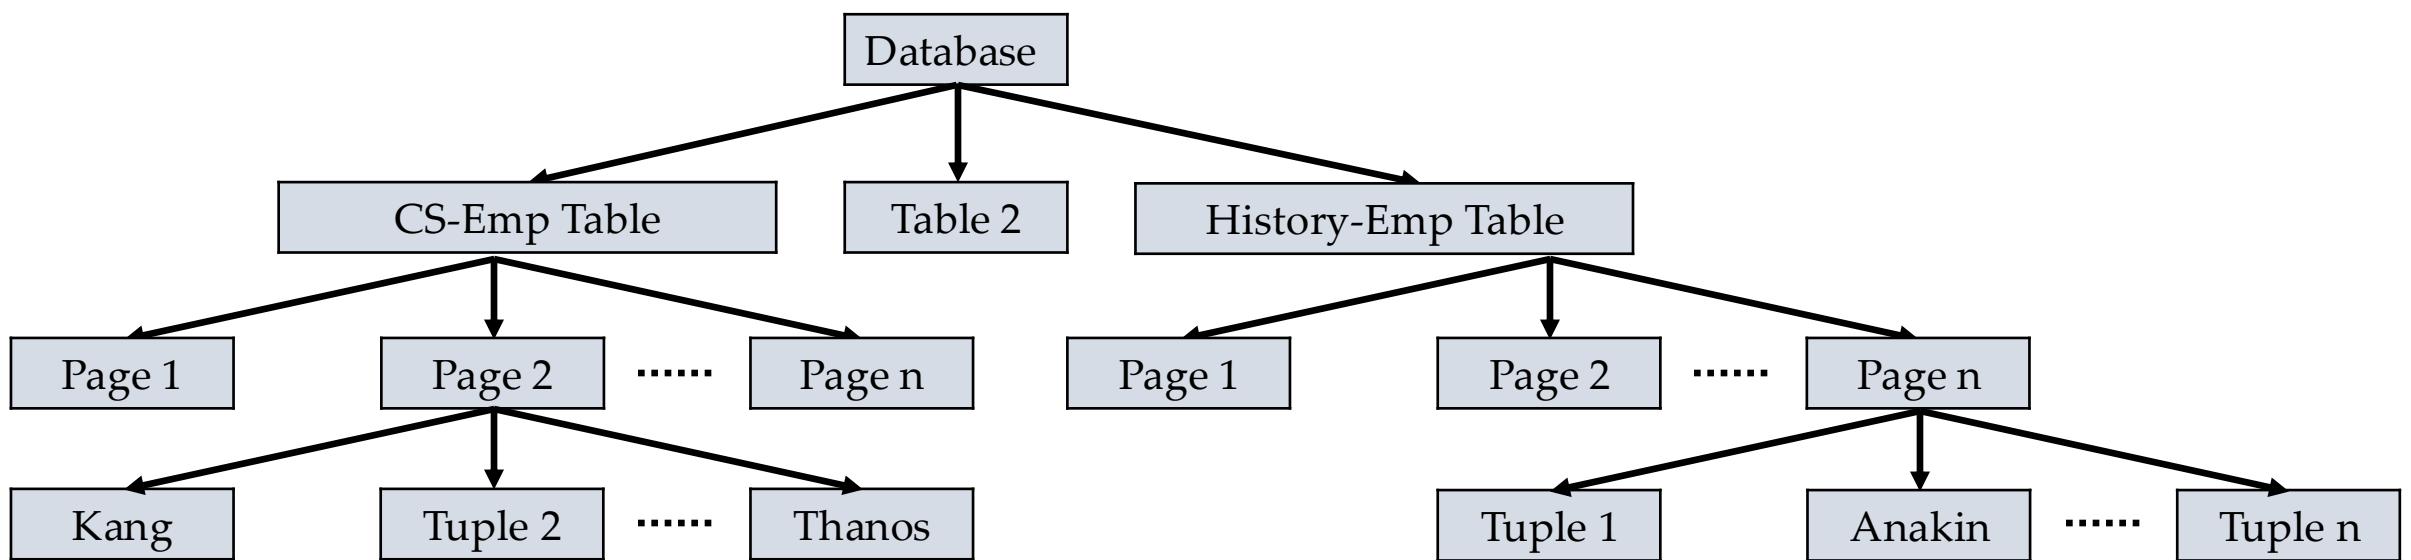
T2 starts locking



Example 2

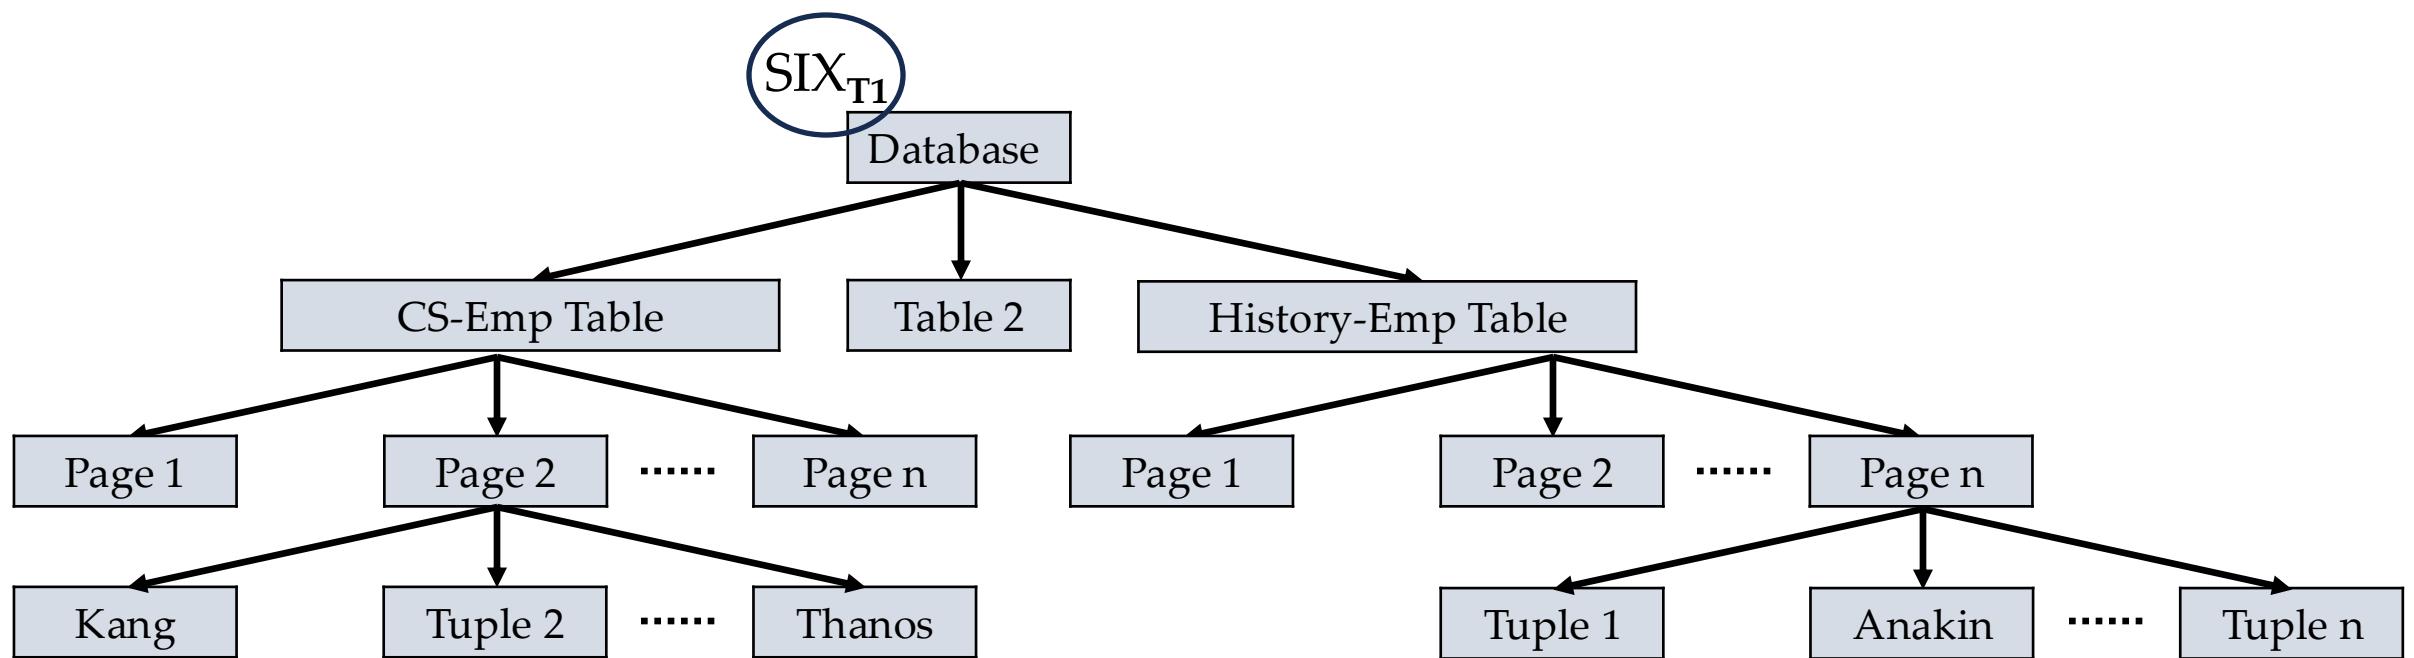
- T1 → Read all CS employees salary and increase Kang's salary by \$100.
- T2 → Read salary of CS employee Thanos.
- T3 → Read all CS employees salary .
- How can we acquire locks in this example?

Example 2



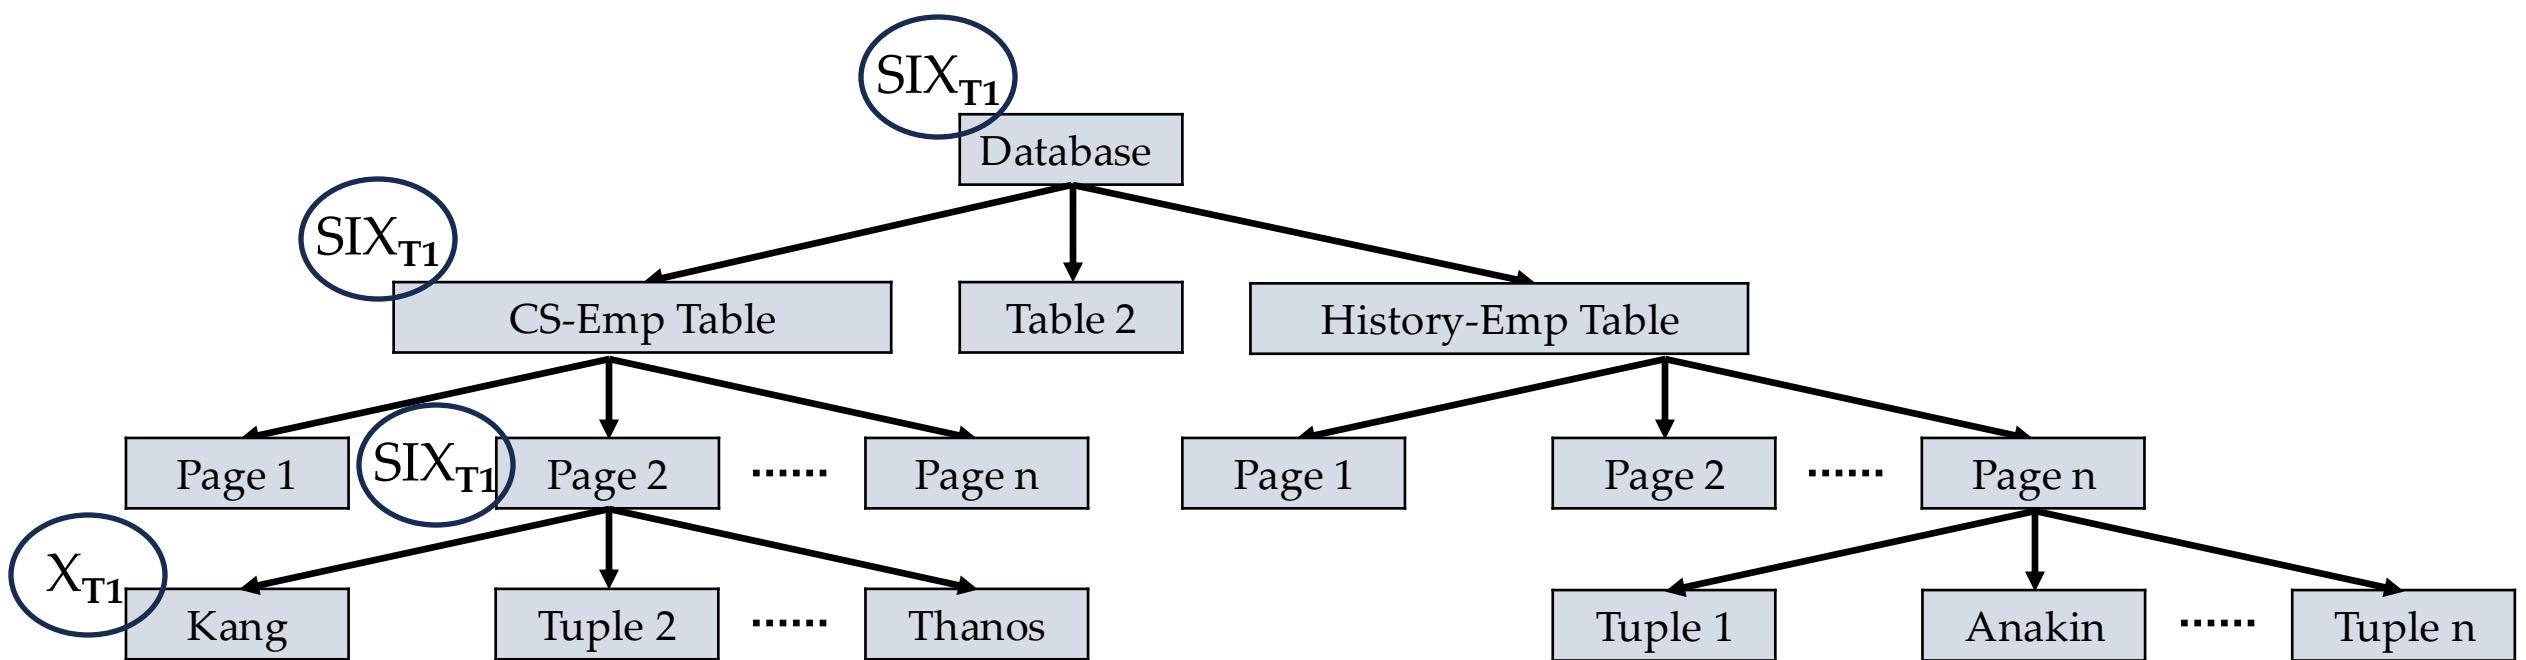
Example 2

T1 starts locking



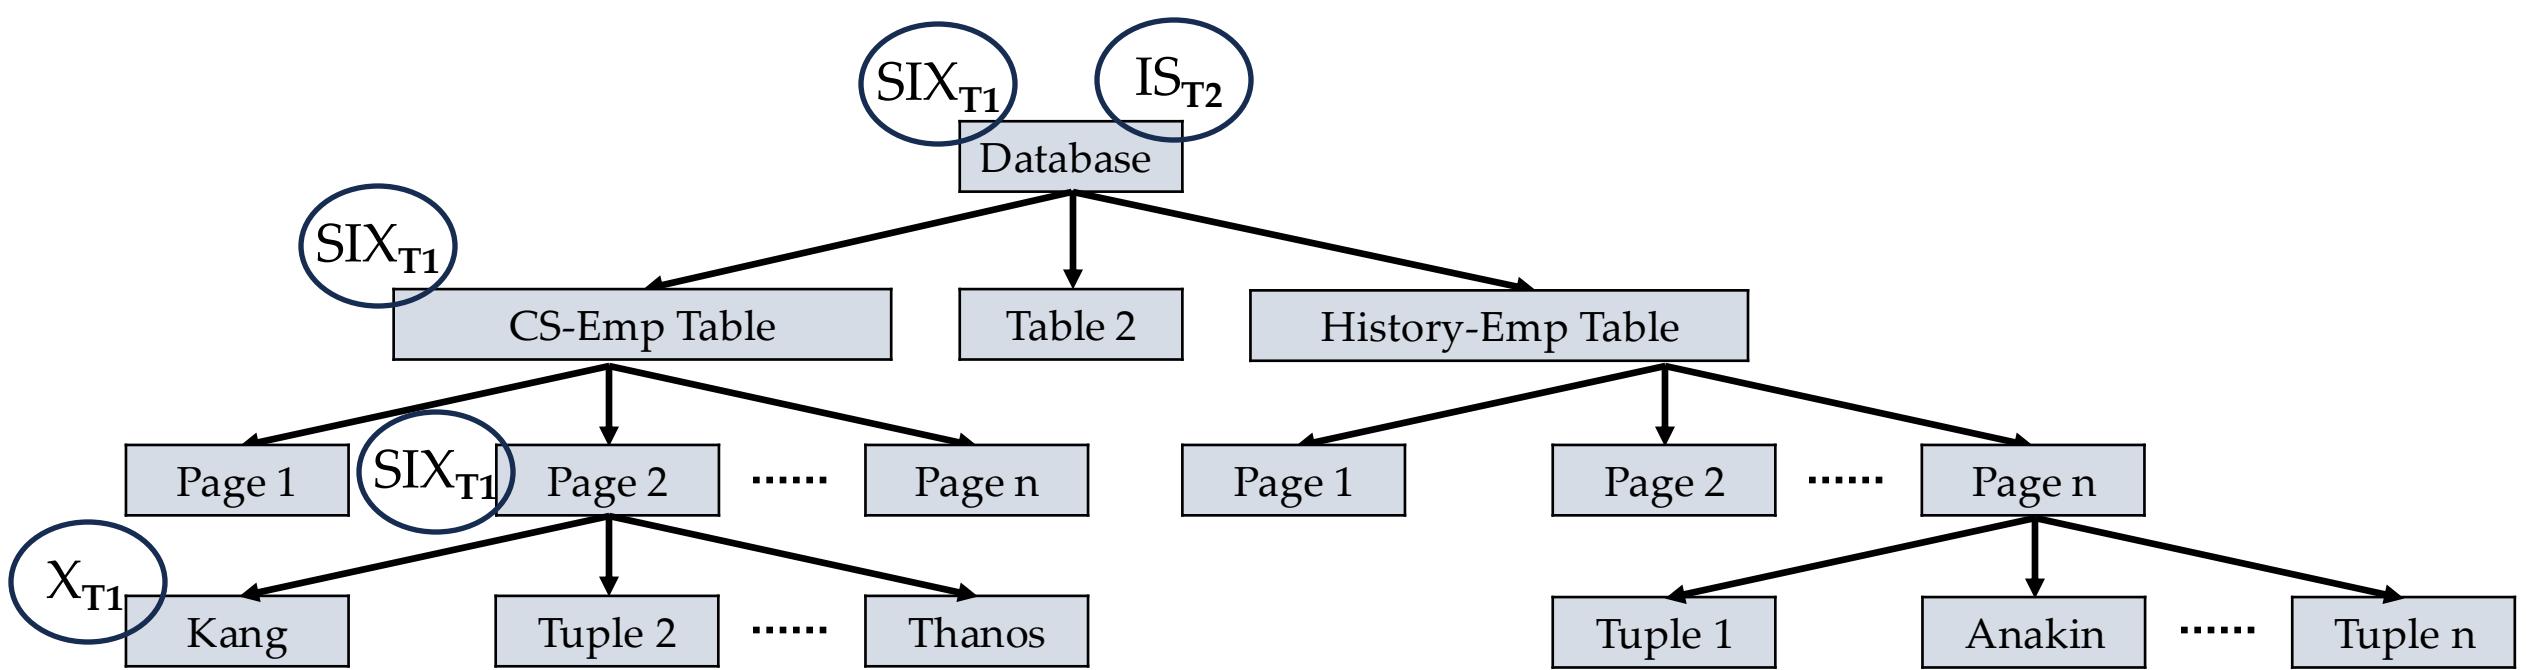
Example 2

T1 starts locking



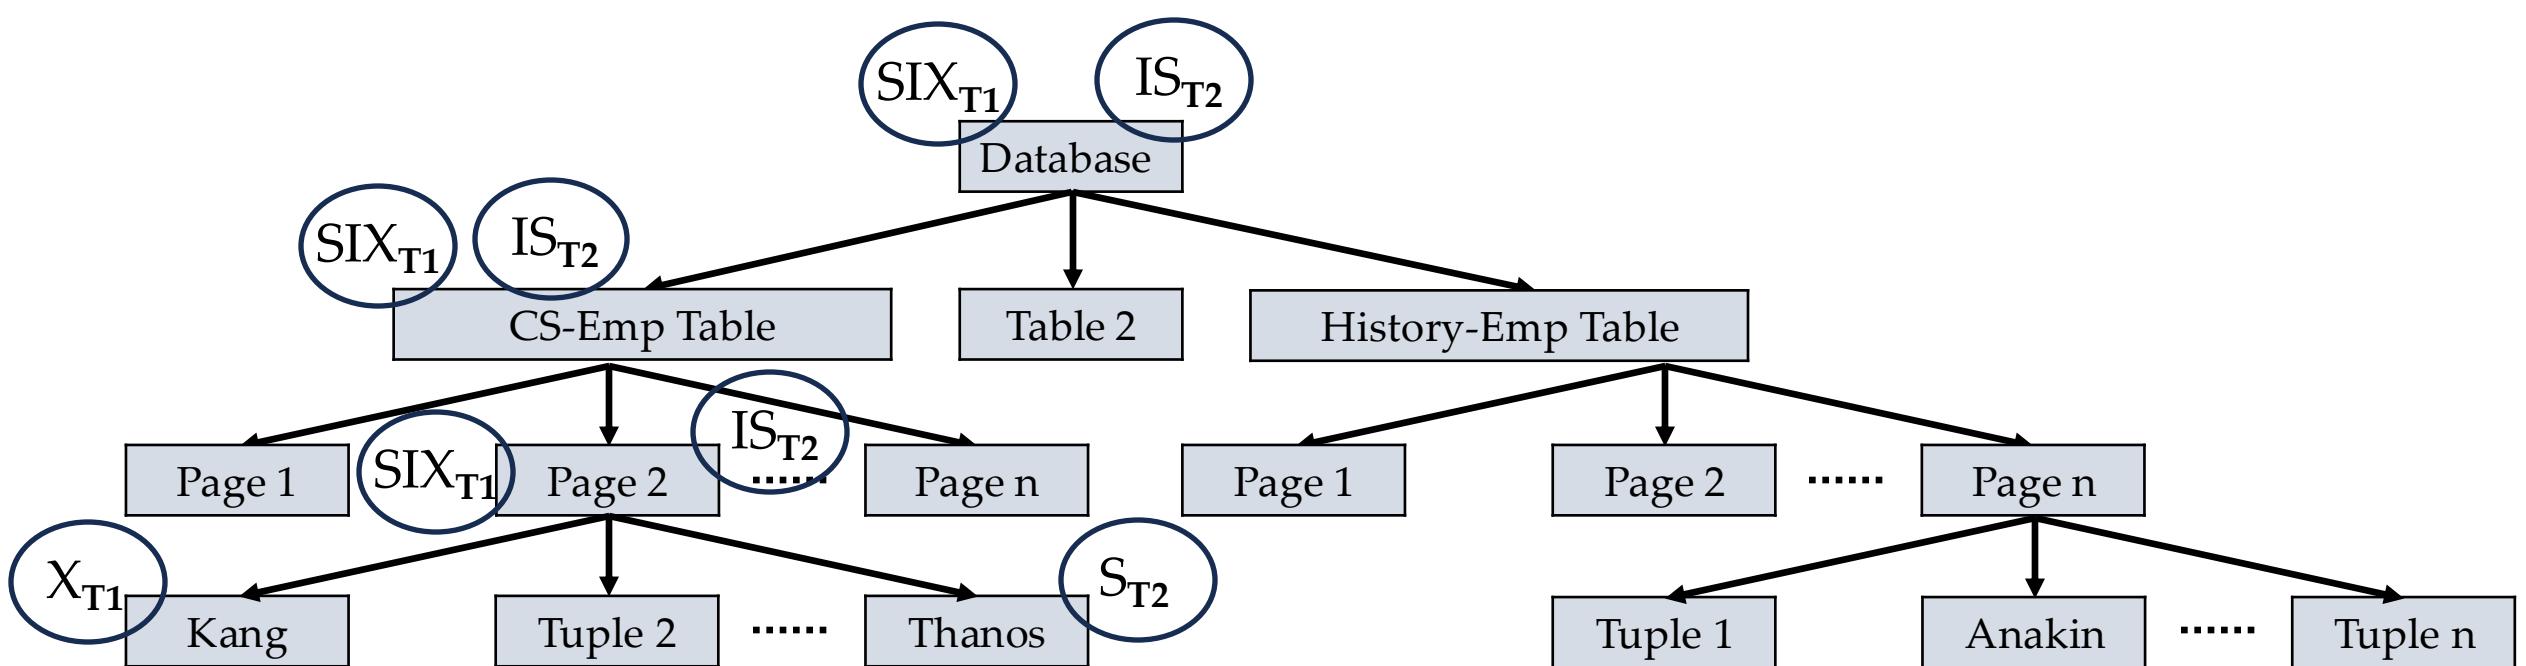
Example 2

T2 starts locking



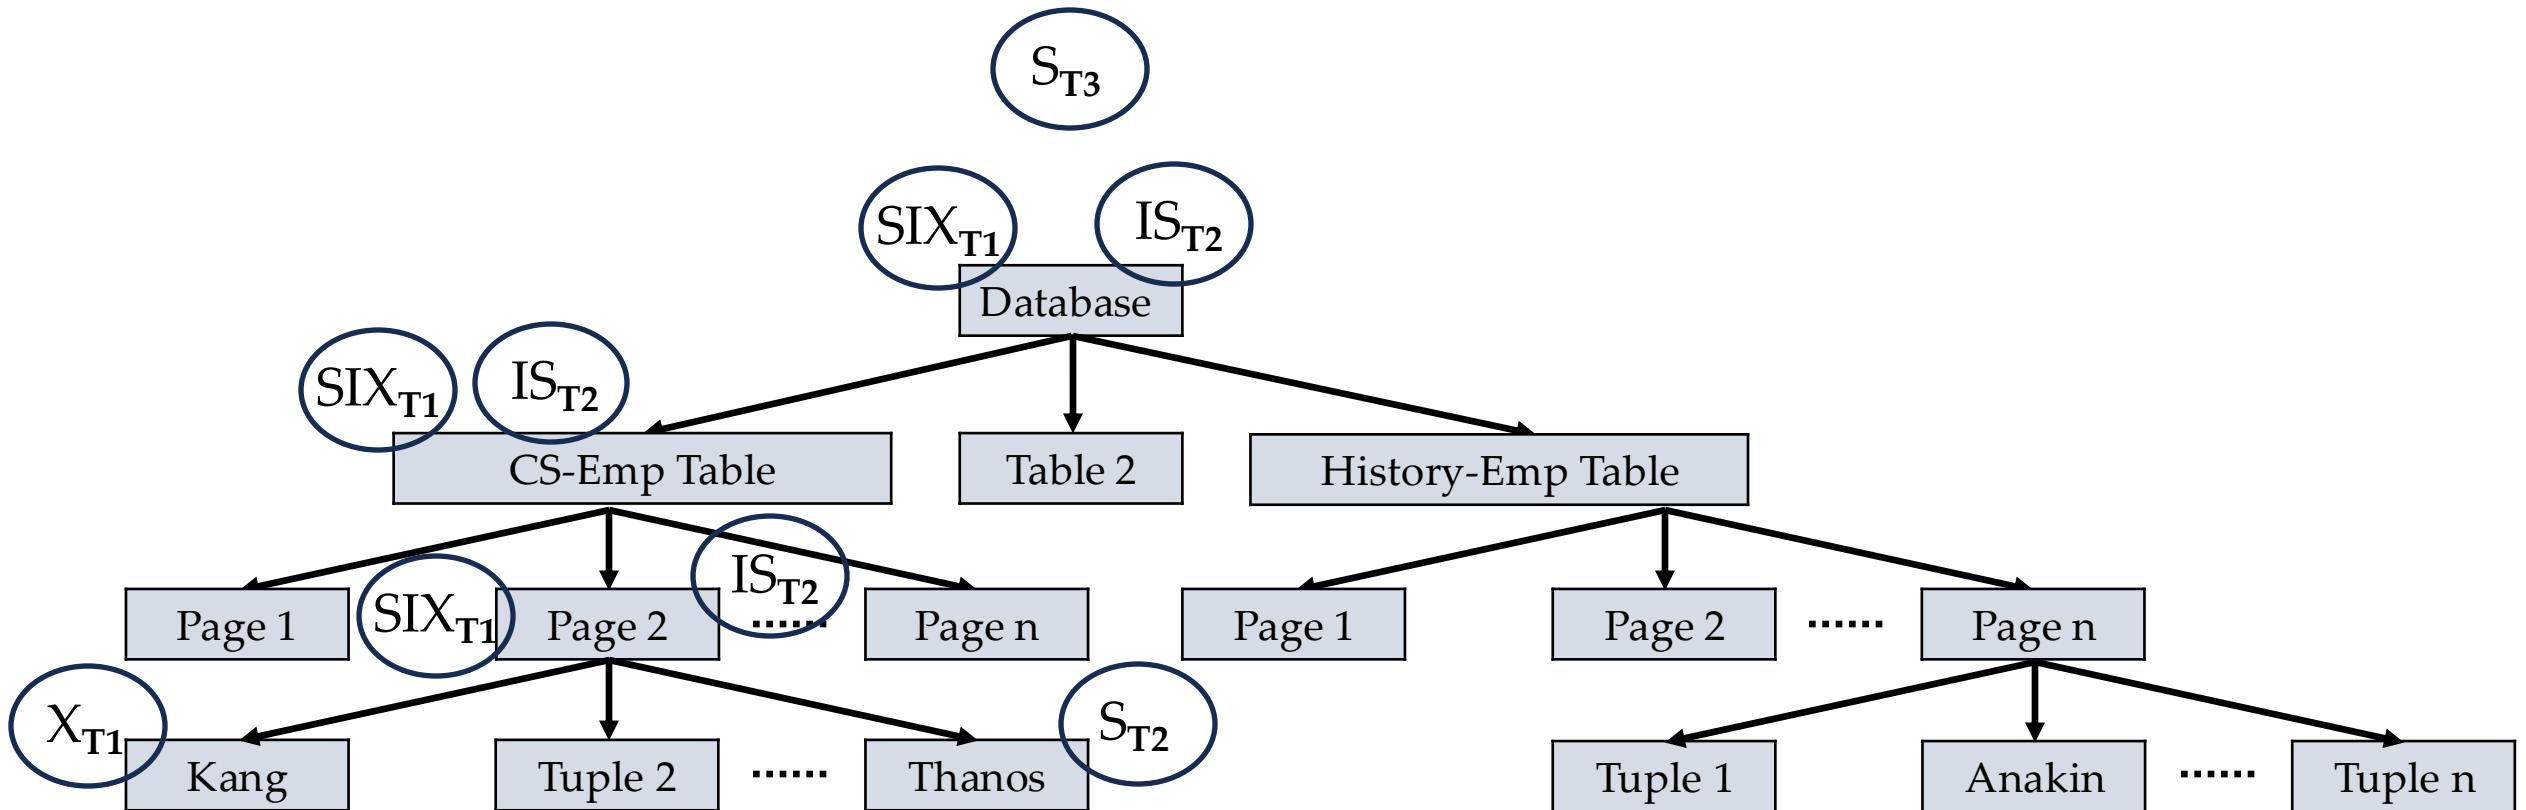
Example 2

T2 starts locking



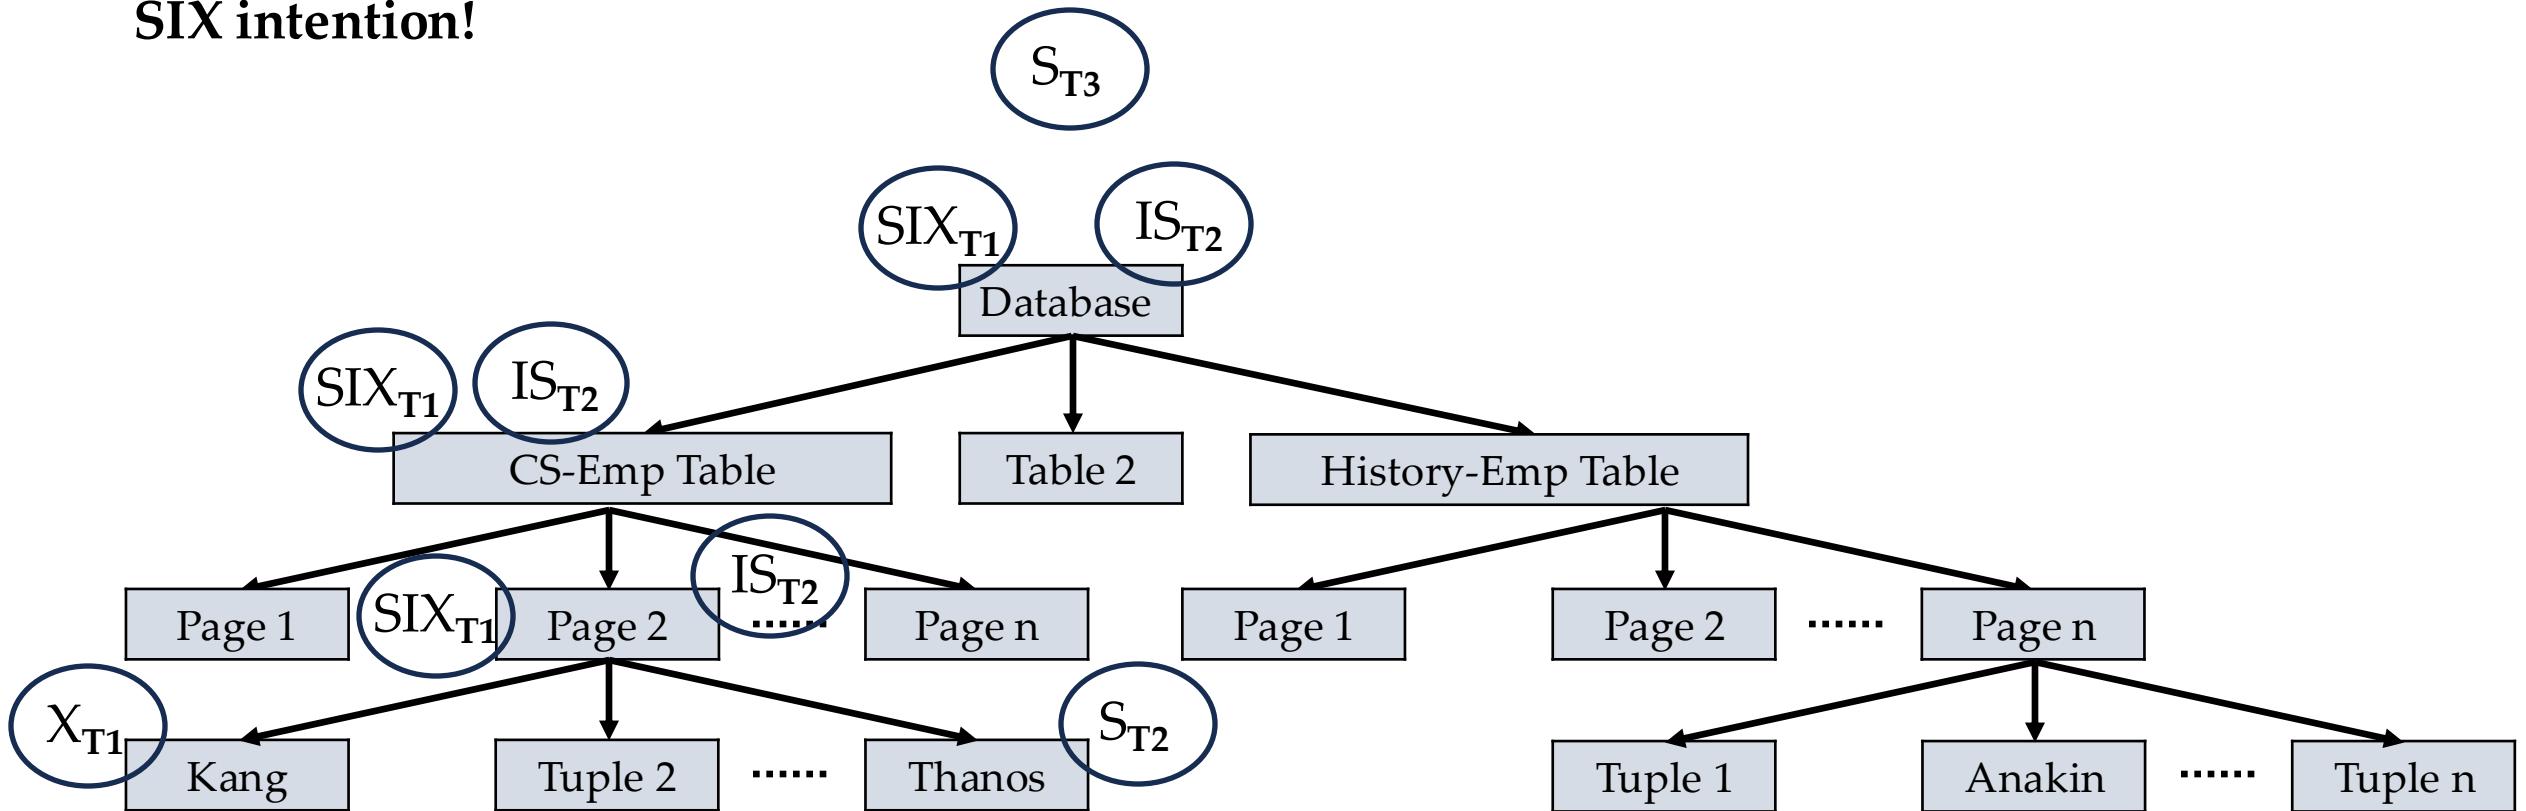
Example 2

T3 starts locking



Example 2

T3 is not allowed to lock until T1 finishes because T3 wants to read something which has SIX intention!



Lock Escalation

- The DBMS automatically switches to coarser-grained locks when a transaction acquires too many finer-grained locks.
- Reduces the number of requests that the lock manager needs to process.

Discussion

- 2PL forces transactions to acquire locks!
- Strong Strict 2PL forces transactions to acquire locks early to prevent cascade aborts!
- These protocols take a pessimistic approach and assume that conflicts are common and transactions access a lot of data items!
- Can we do better?

Timestamp Ordering Concurrency Control

Timestamp Ordering Concurrency Control

- An optimistic concurrency control protocol.
- Assumption:
 - Conflicts between transactions are rare.
 - Transactions are short-lived.
- Optimized for the no-conflict cases.

Assigning Timestamps

- Each transaction T_i is assigned a **unique monotonically increasing** timestamp.
- Let $TS(T_i)$ be the timestamp allocated to transaction T_i .
- When to assign the timestamp → Depends on the design.
- How to generate a timestamp?

Assigning Timestamps

- Each transaction T_i is assigned a **unique monotonically increasing** timestamp.
- Let $TS(T_i)$ be the timestamp allocated to transaction T_i .
- When to assign the timestamp → Depends on the design.
- How to generate a timestamp?
 - Wall clock time / System time
 - Logical counter
 - Hybrid

Timestamp Ordering Concurrency Control

- Timestamps are used to determine the **serializability order of transactions**.
- For two transactions T_i and T_j , if $TS(T_i) < TS(T_j)$, then
 - The DBMS must ensure that the execution schedule for these transactions is equivalent to the serial schedule where T_i appears before T_j .
- Each database object (e.g., tuple) need to track the timestamps of that last accessed/modified them.

Optimistic Concurrency Control

- Timestamp Ordering (T/O) can be used to design an OCC protocol.
- In OCC using T/O, DBMS creates a **private workspace** for each transaction.

Optimistic Concurrency Control

- Timestamp Ordering (T/O) can be used to design an OCC protocol.
- In OCC using T/O, DBMS creates a **private workspace** for each transaction.
 - Each object **read** is copied into workspace.

Optimistic Concurrency Control

- Timestamp Ordering (T/O) can be used to design an OCC protocol.
- In OCC using T/O, DBMS creates a **private workspace** for each transaction.
 - Each object **read** is copied into workspace.
 - Updates/Writes are applied to workspace.

Optimistic Concurrency Control

- Timestamp Ordering (T/O) can be used to design an OCC protocol.
- In OCC using T/O, DBMS creates a **private workspace** for each transaction.
 - Each object **read** is copied into workspace.
 - Updates/Writes are applied to workspace.
 - When a transaction commits, the DBMS checks if the workspace writes conflict with other transactions.
 - If there are no conflicts, the workspace write set is copied to the database.

OCC Phases

- How does OCC work?

OCC Phases

- How does OCC work?
- **Three Phases of OCC:**
 - Read Phase
 - Validation Phase
 - Write Phase

Read Phase

- Track the read/write sets of each transaction and store the writes of each transaction in a private workspace.
- DBMS copies every tuple that the transaction accesses from the database to its private workspace.

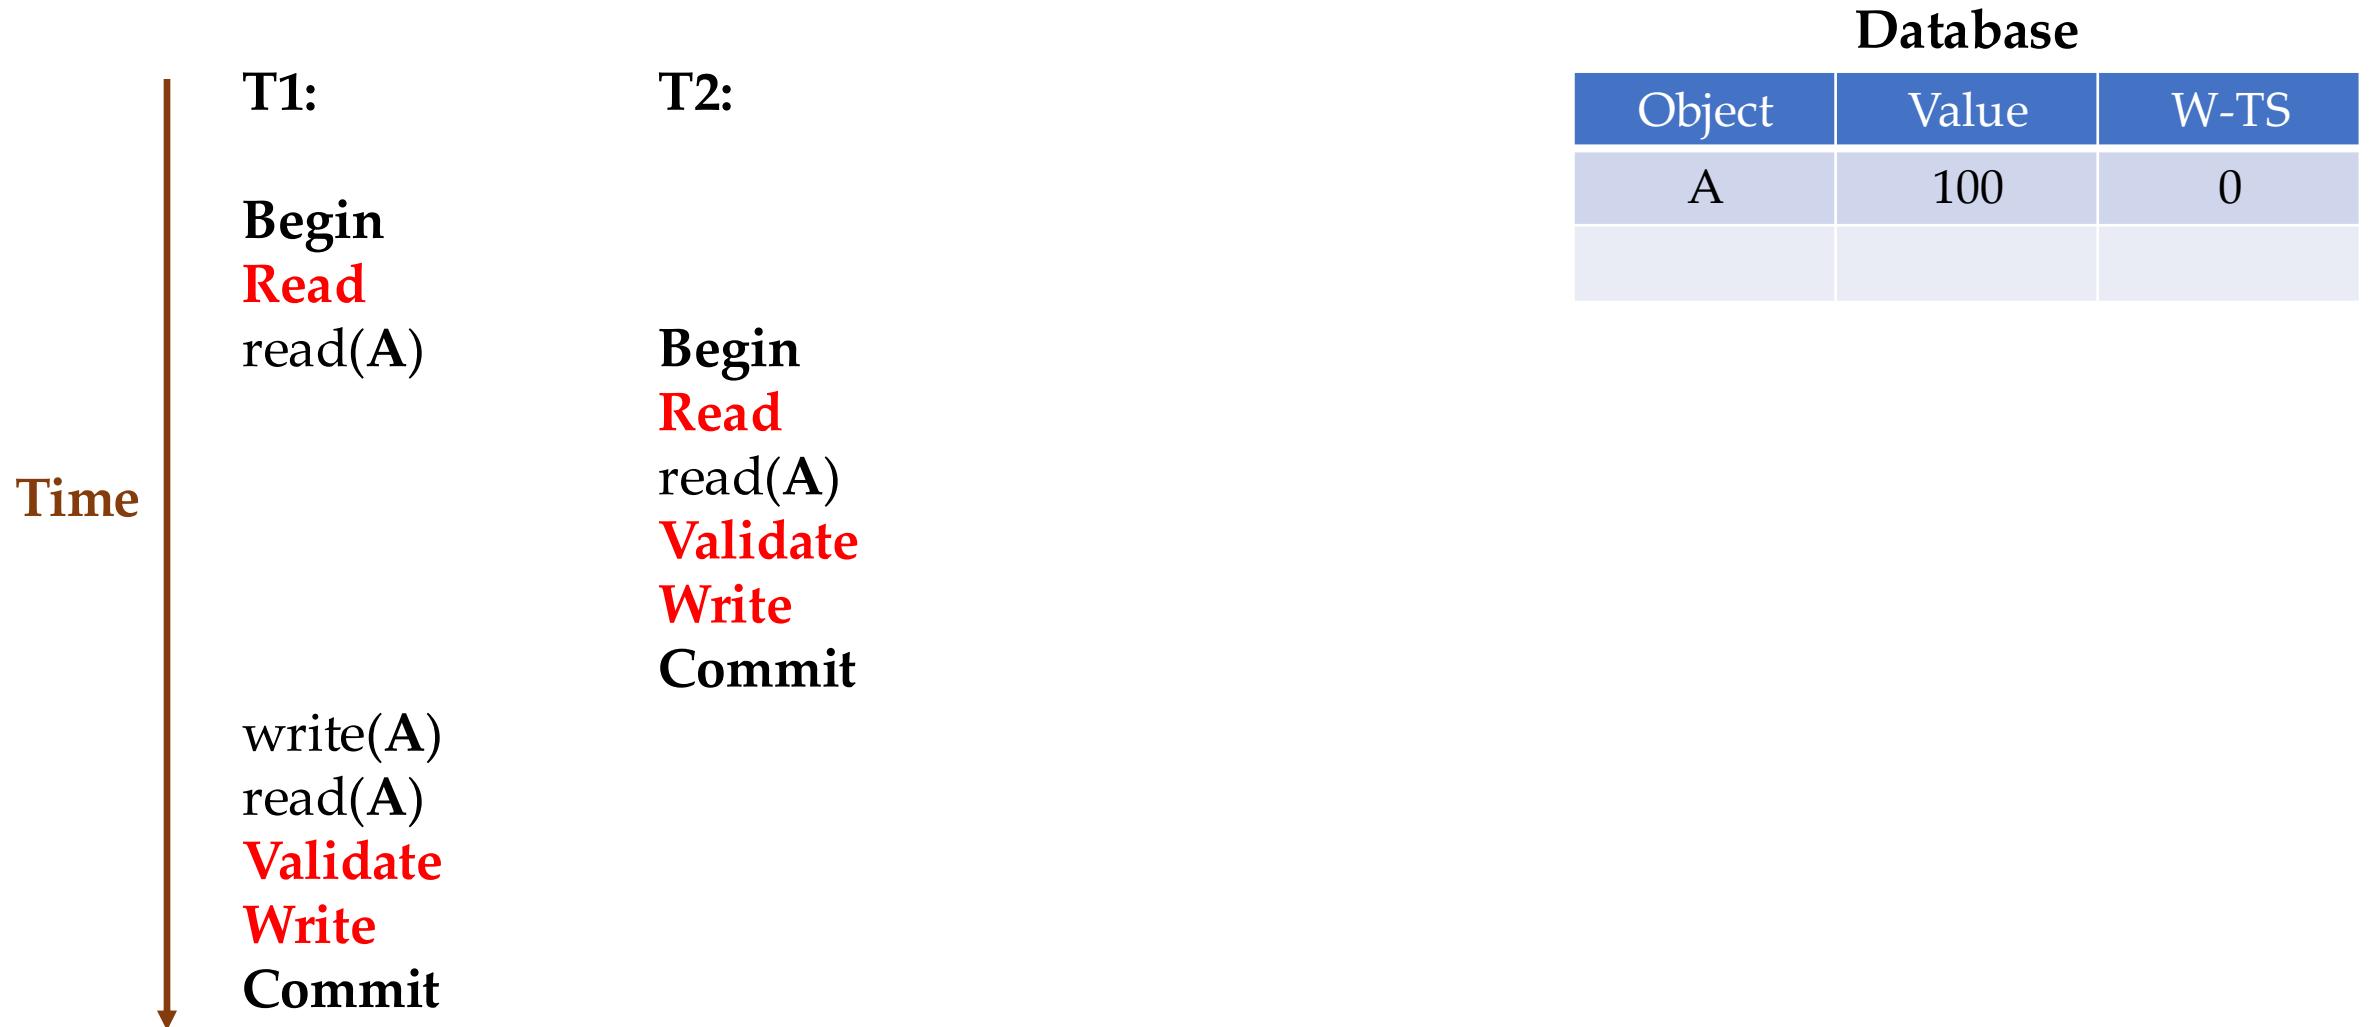
Validation Phase

- Assign the transaction a unique timestamp (TS) and then check whether it conflicts with other transactions.

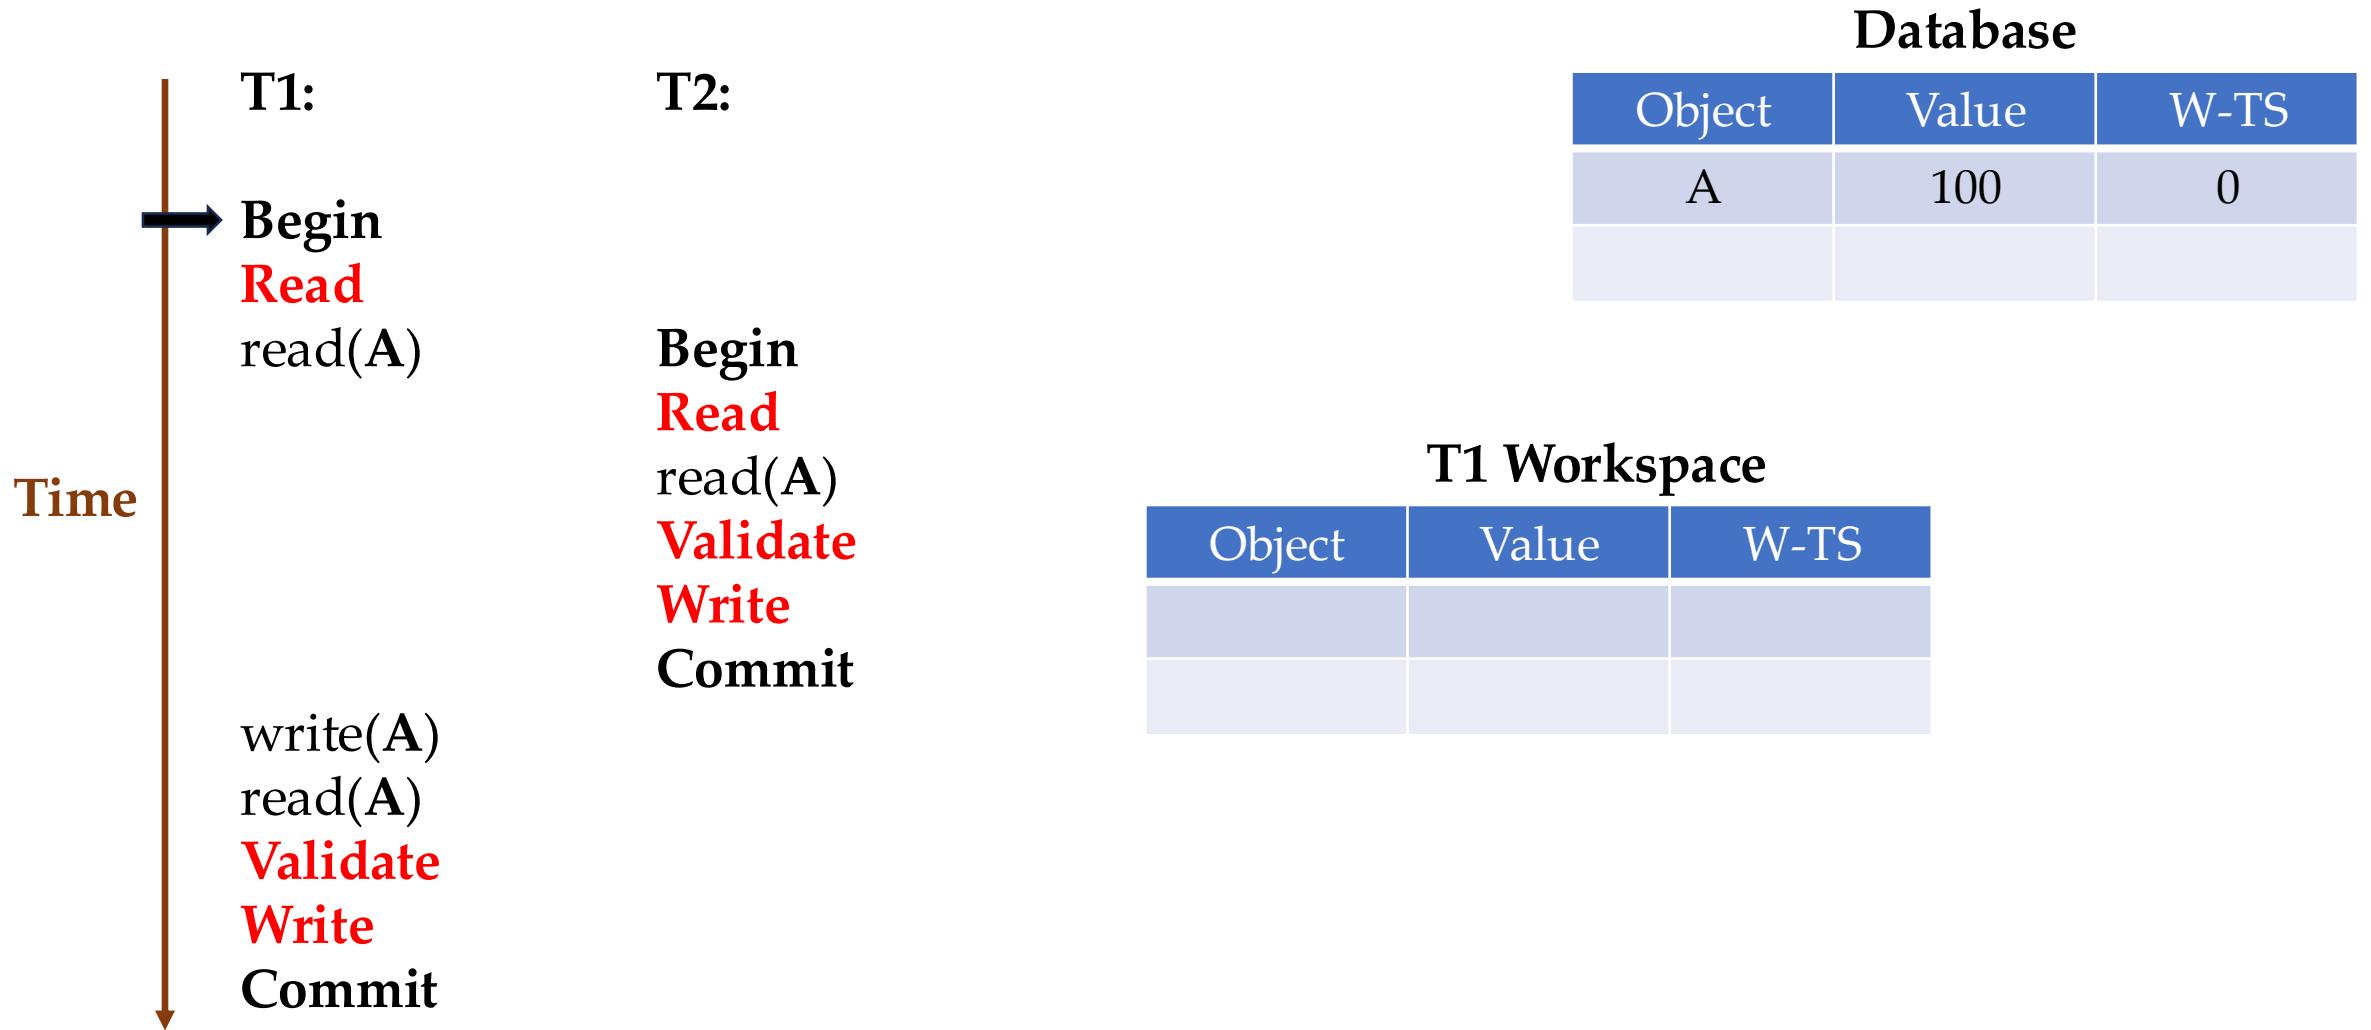
Write Phase

- If validation is successful, set the write timestamp (W-TS) for all the modified objects in private workspace to the validation timestamp.
 - Next, update the value and timestamp in the database.
- Otherwise abort transaction.

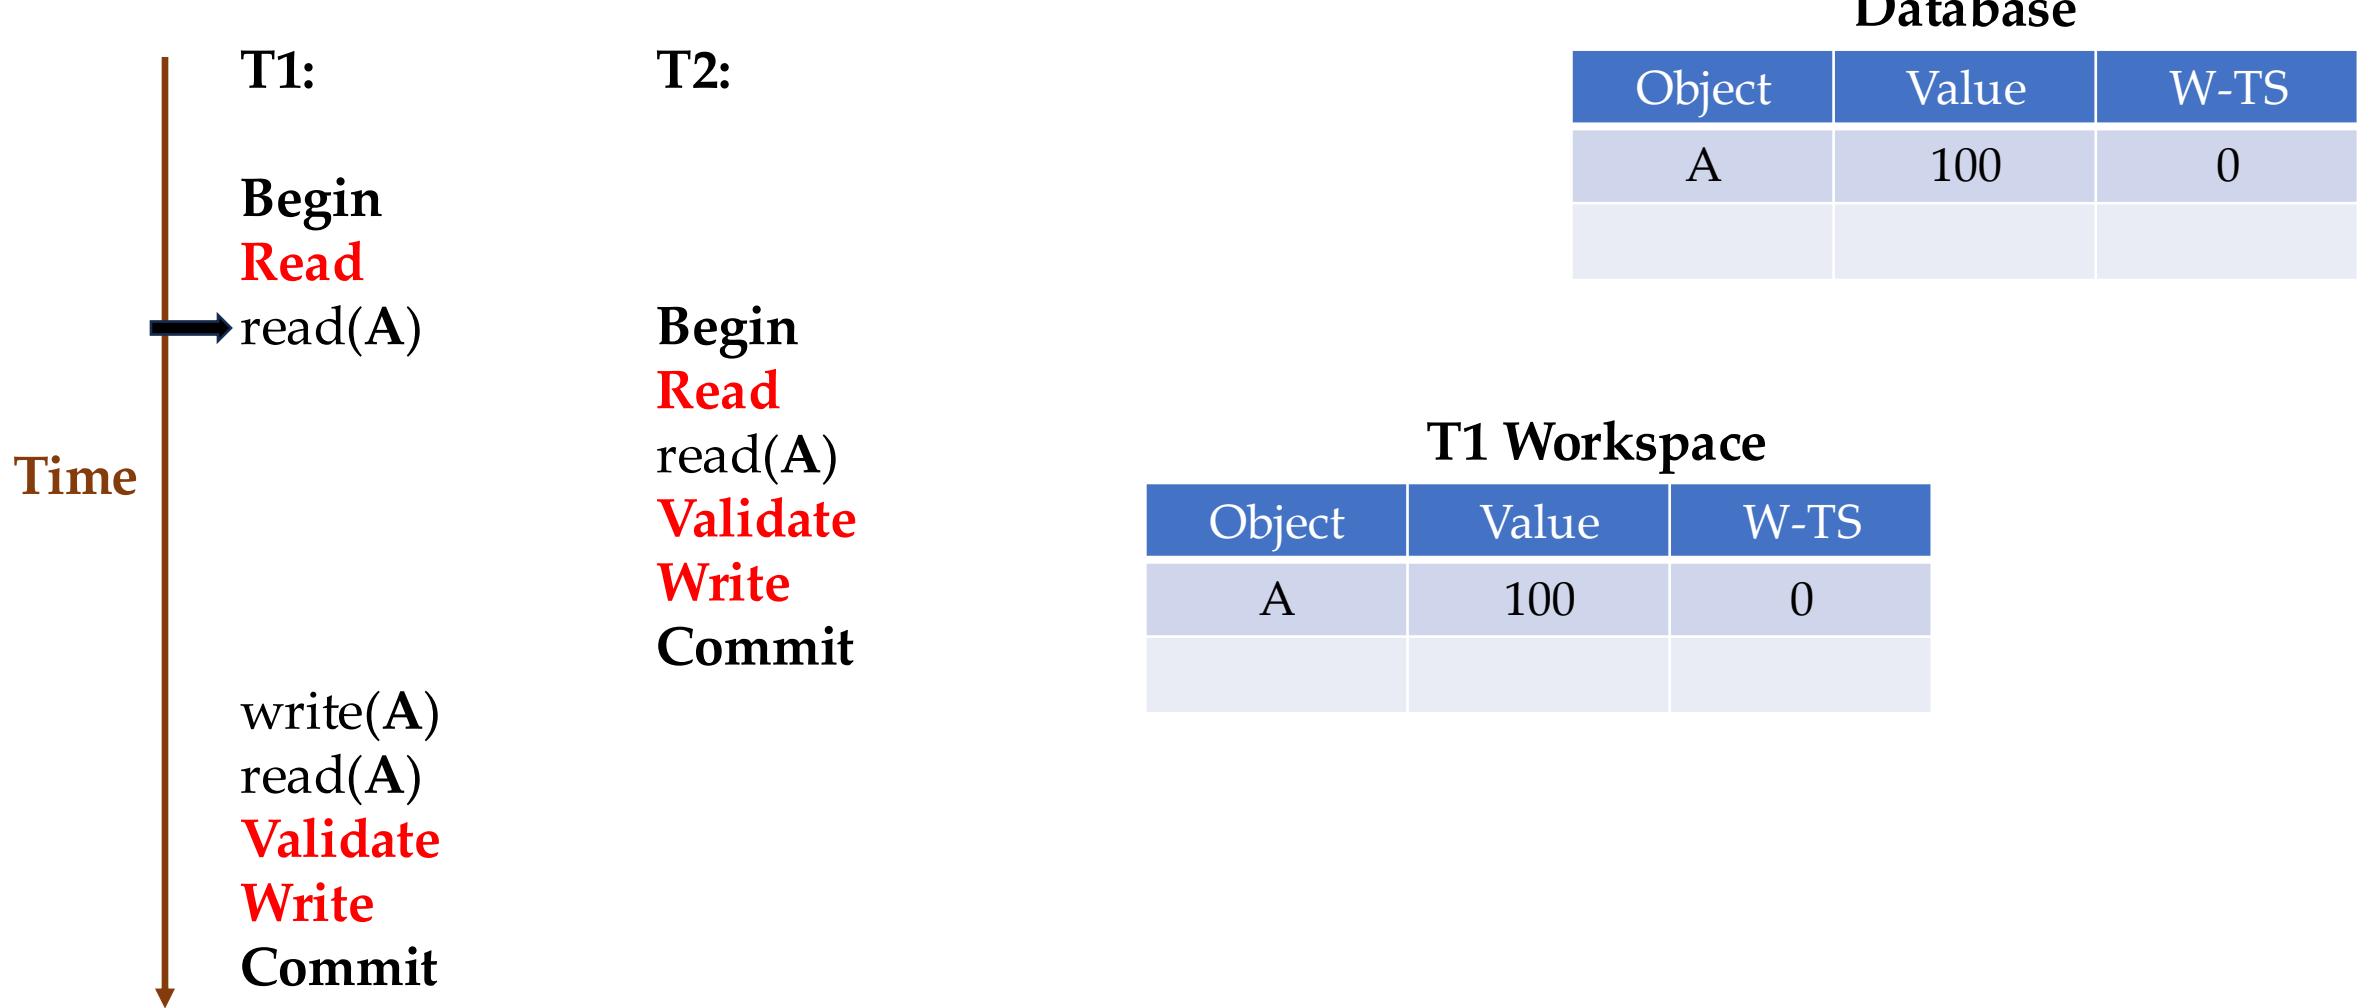
OCC Example I



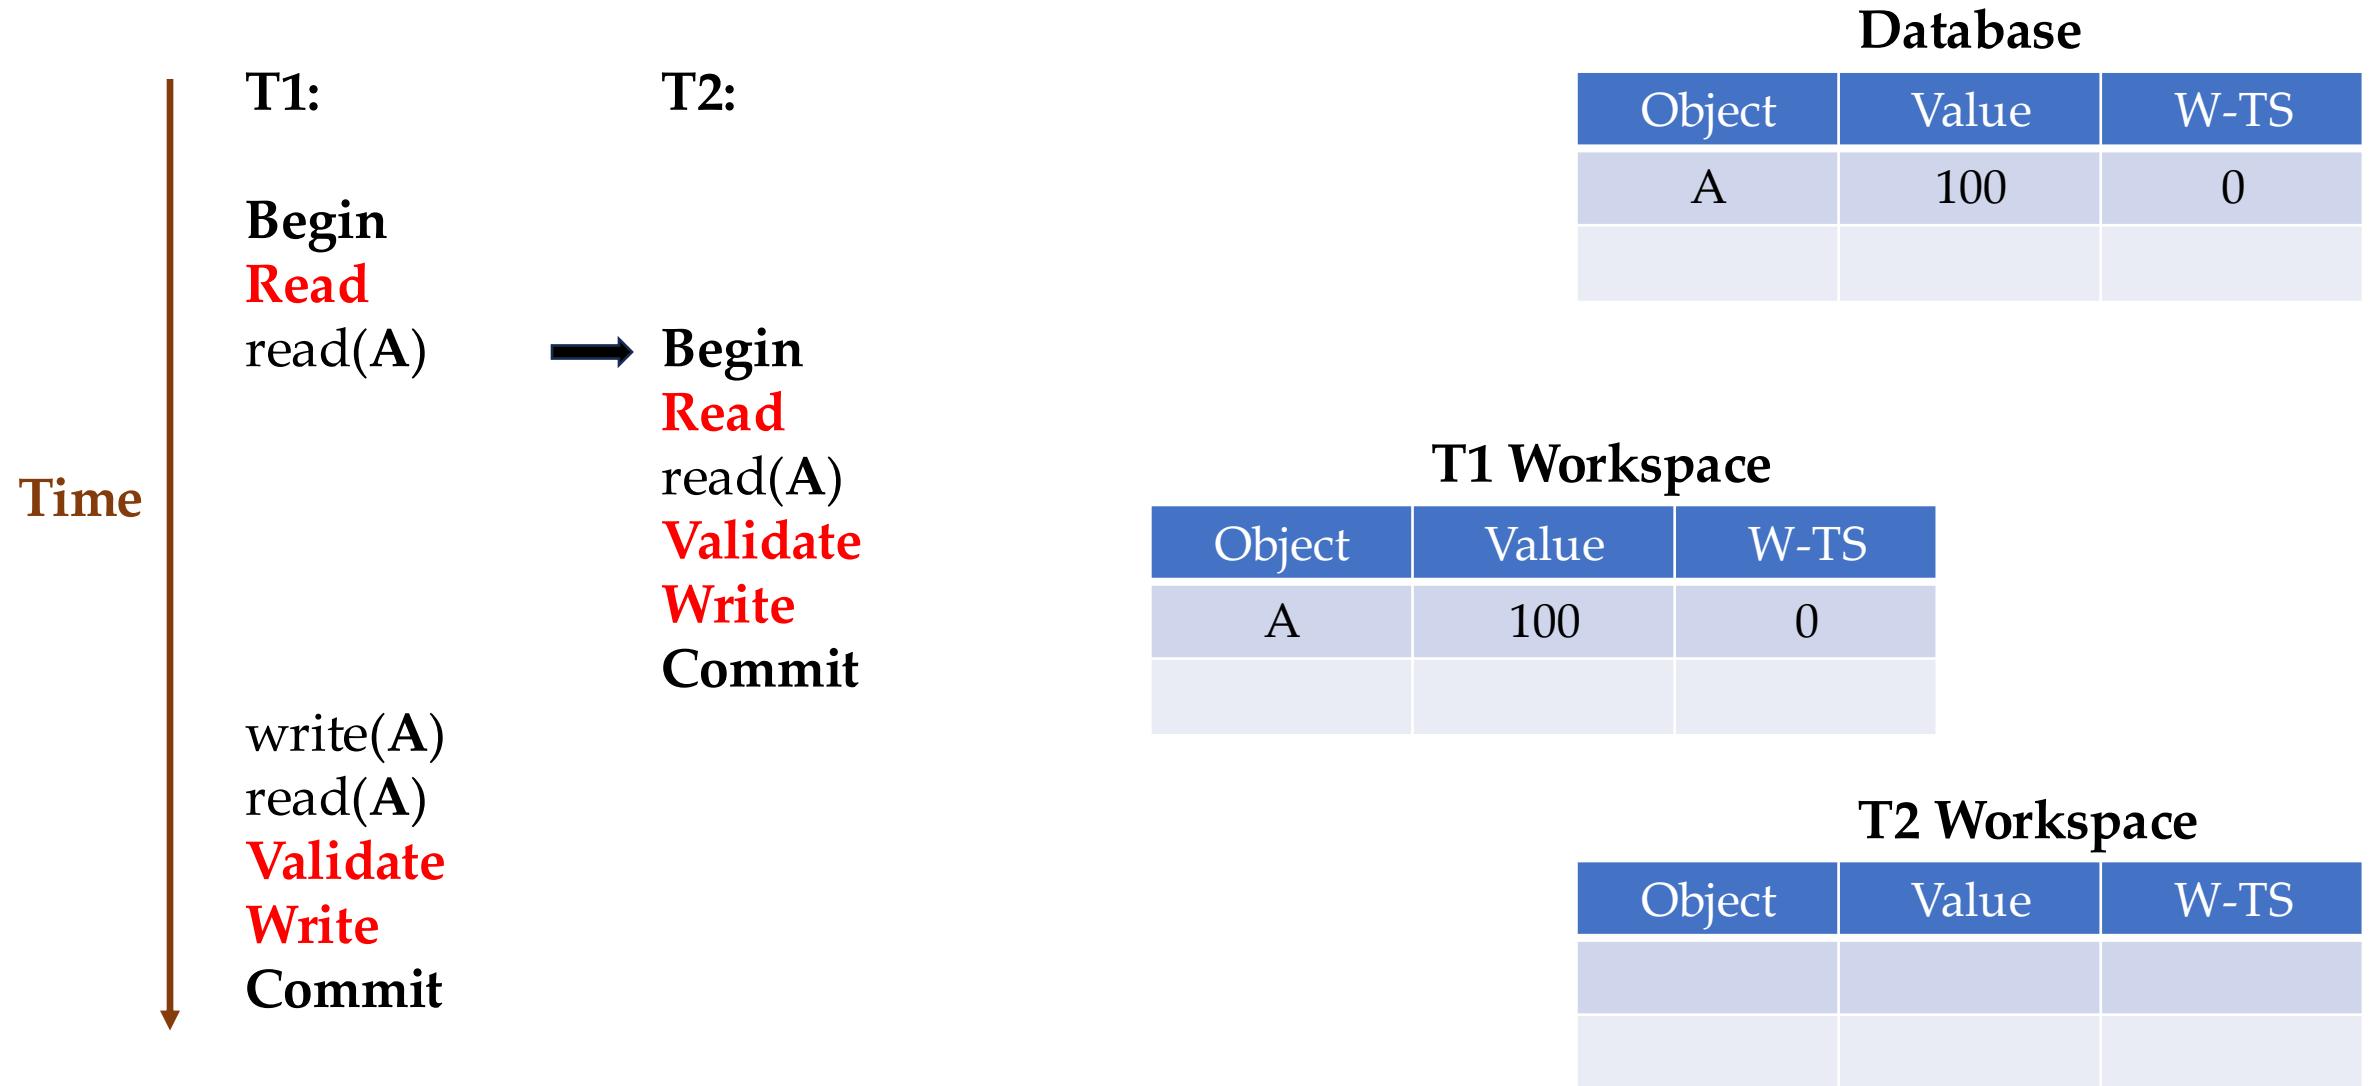
OCC Example I



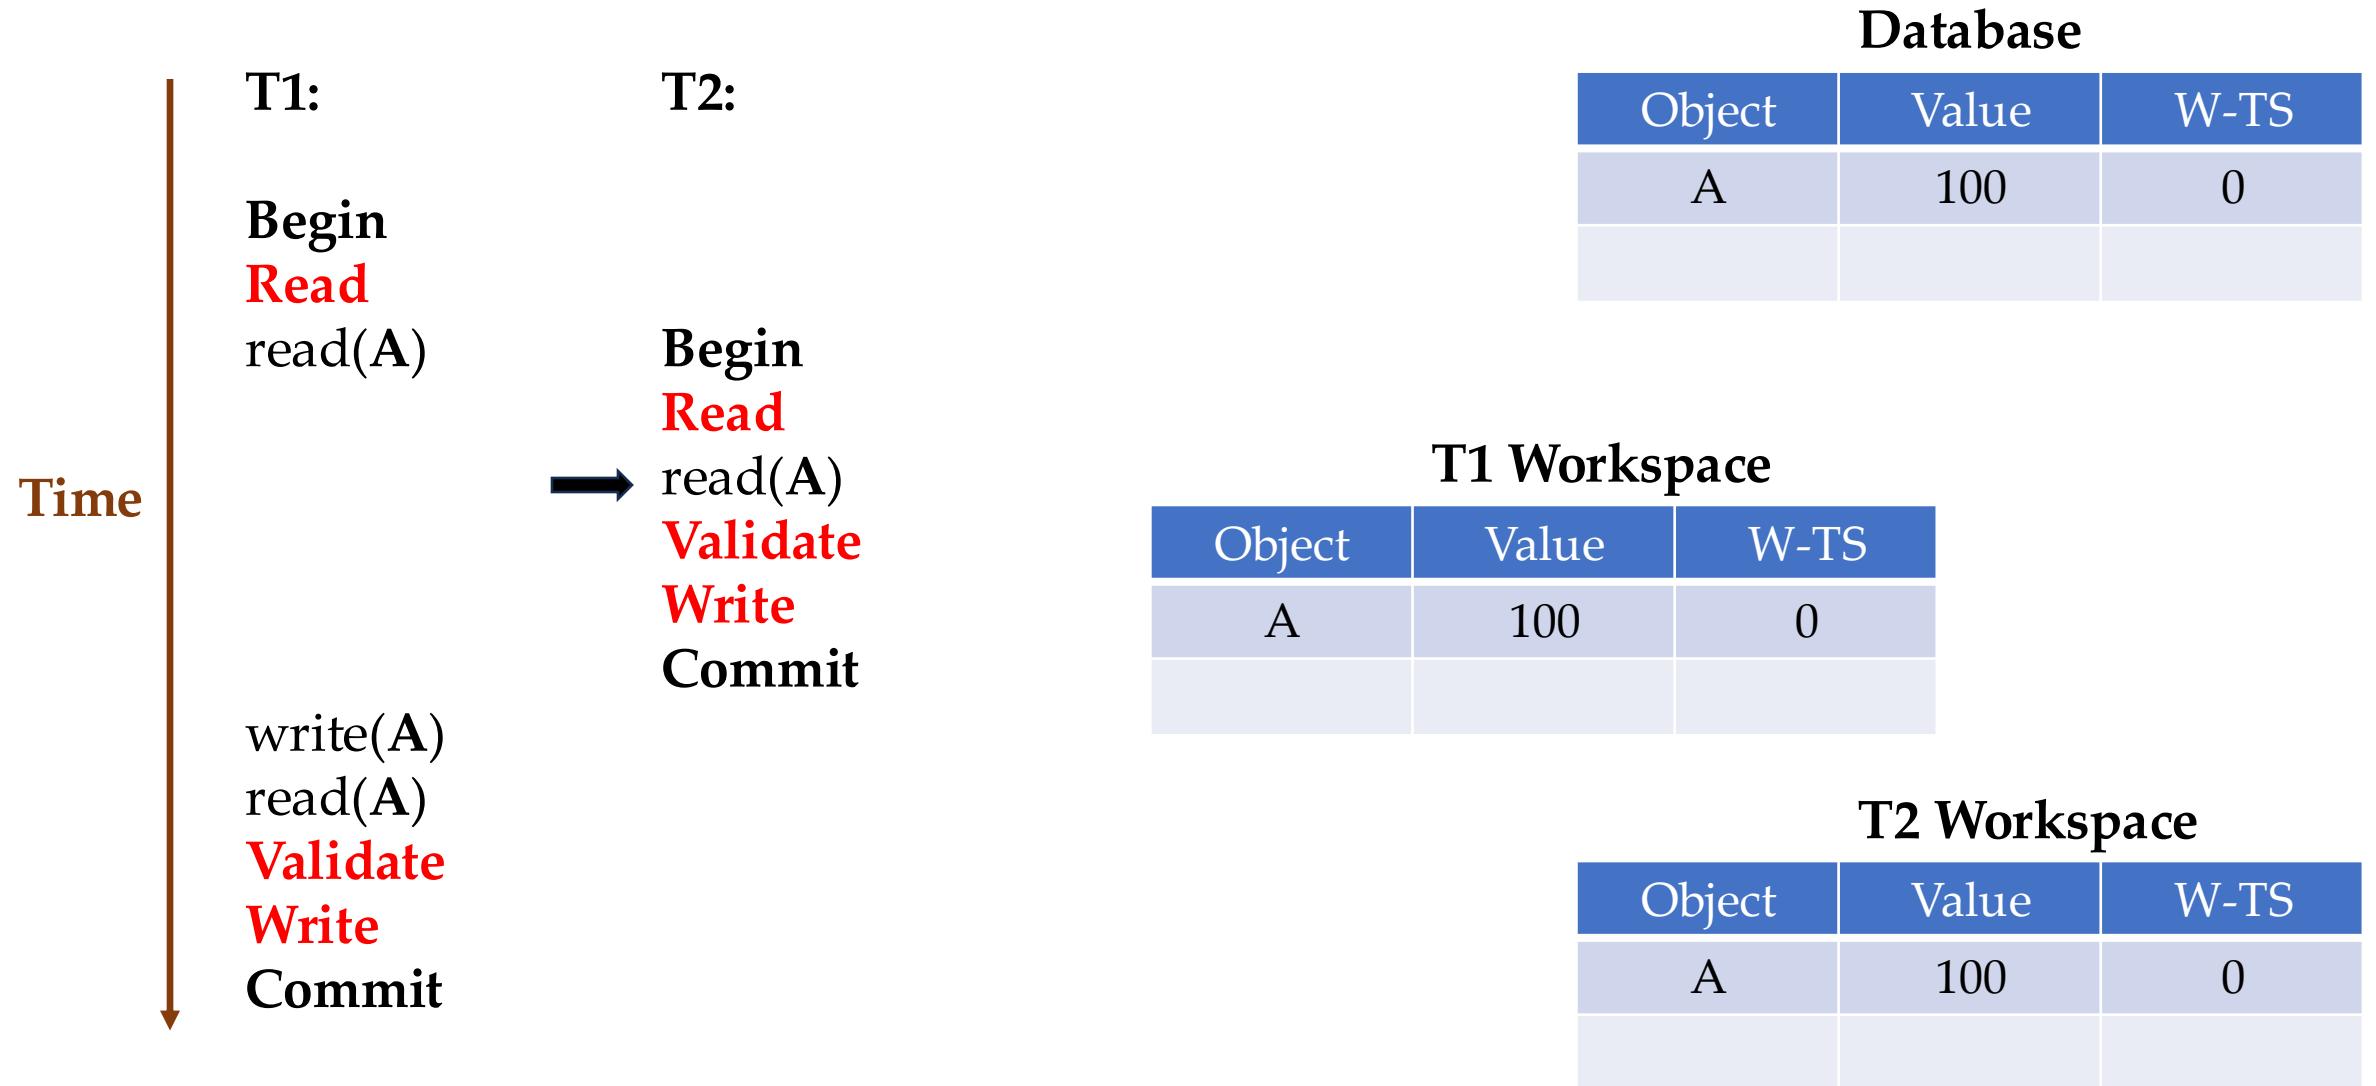
OCC Example I



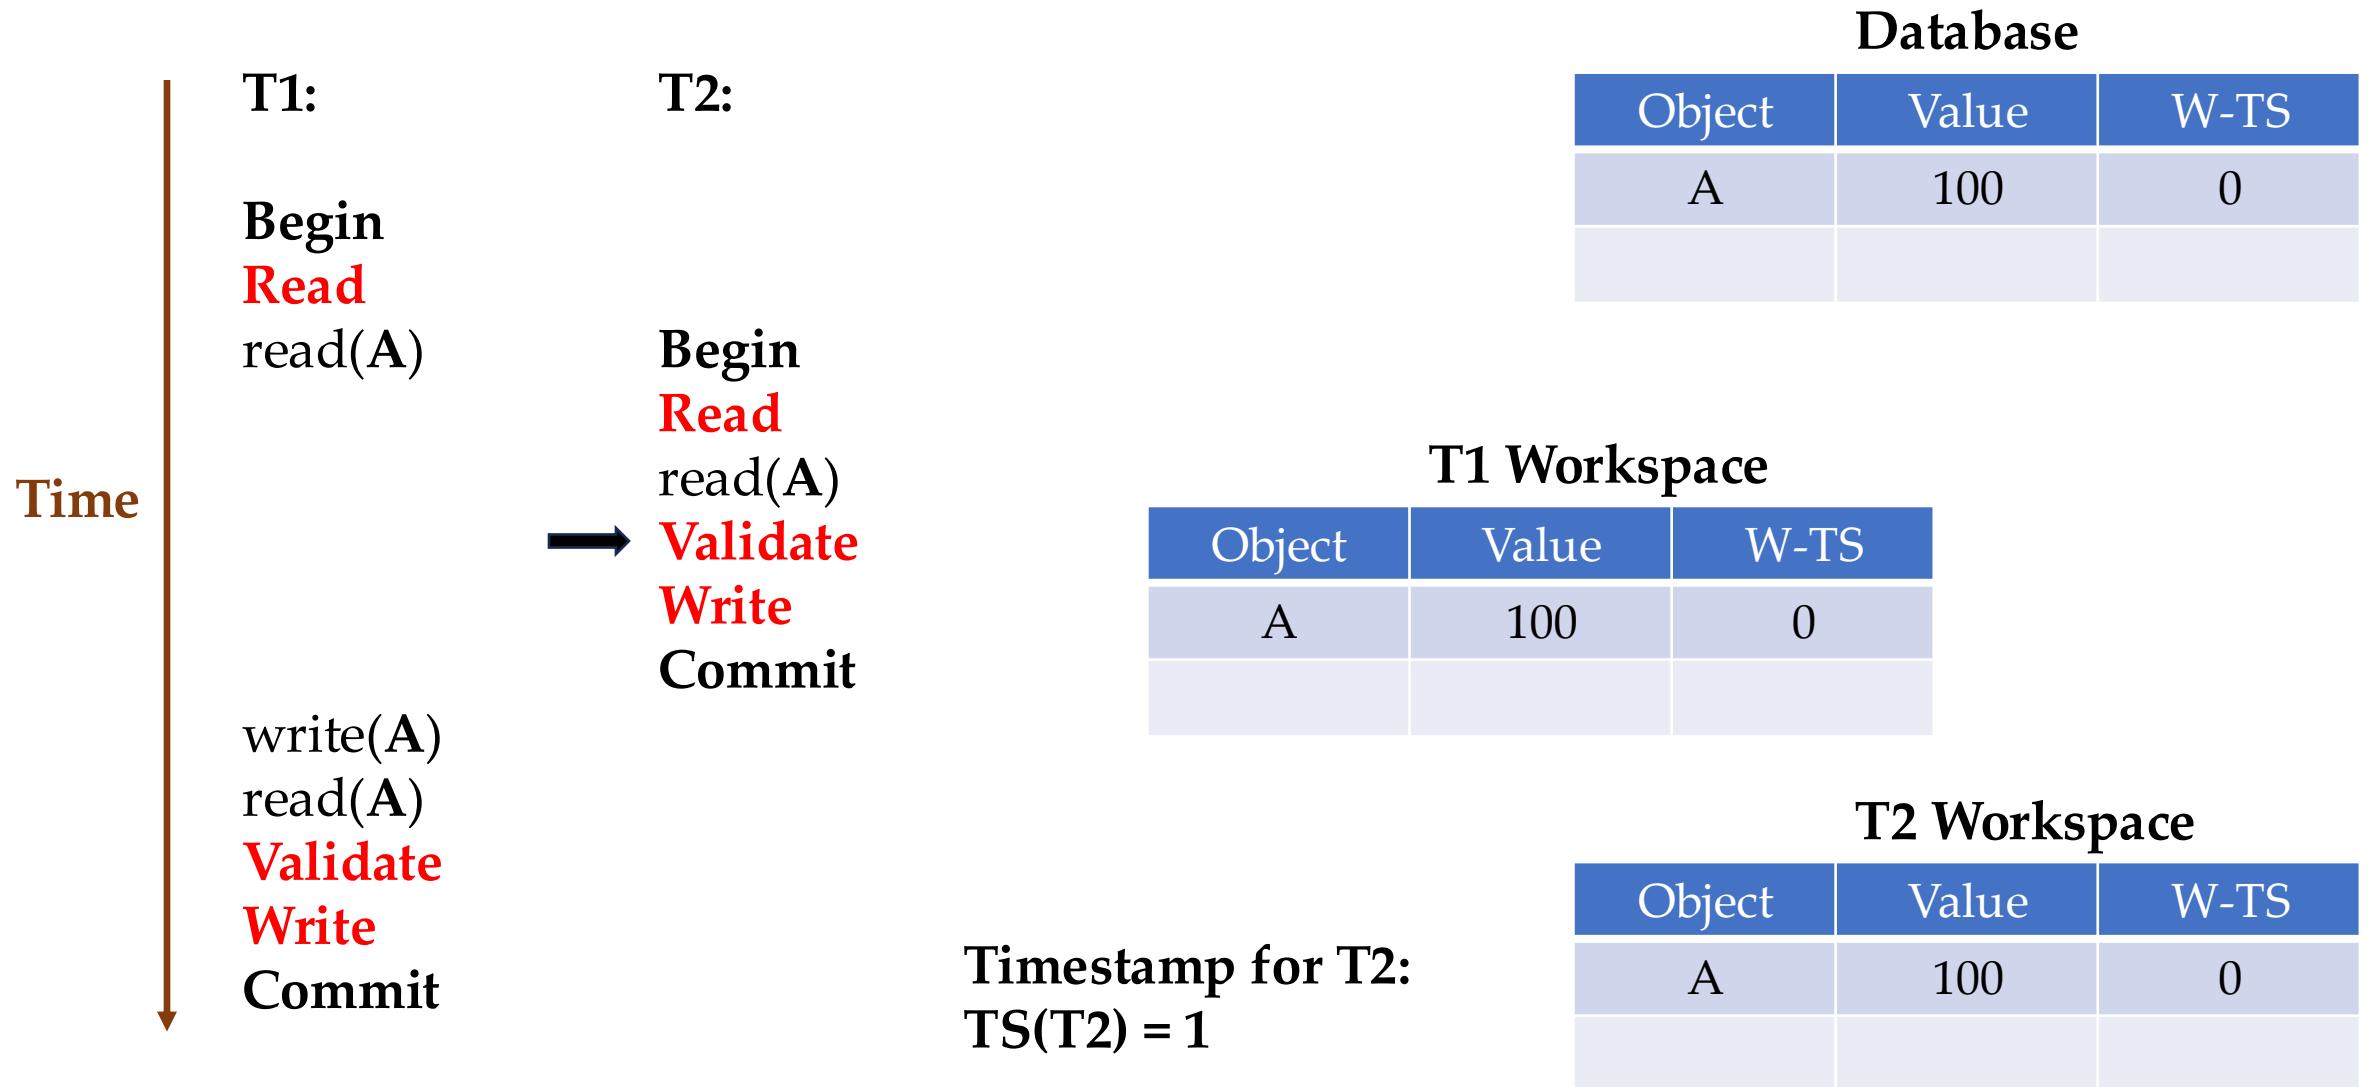
OCC Example I



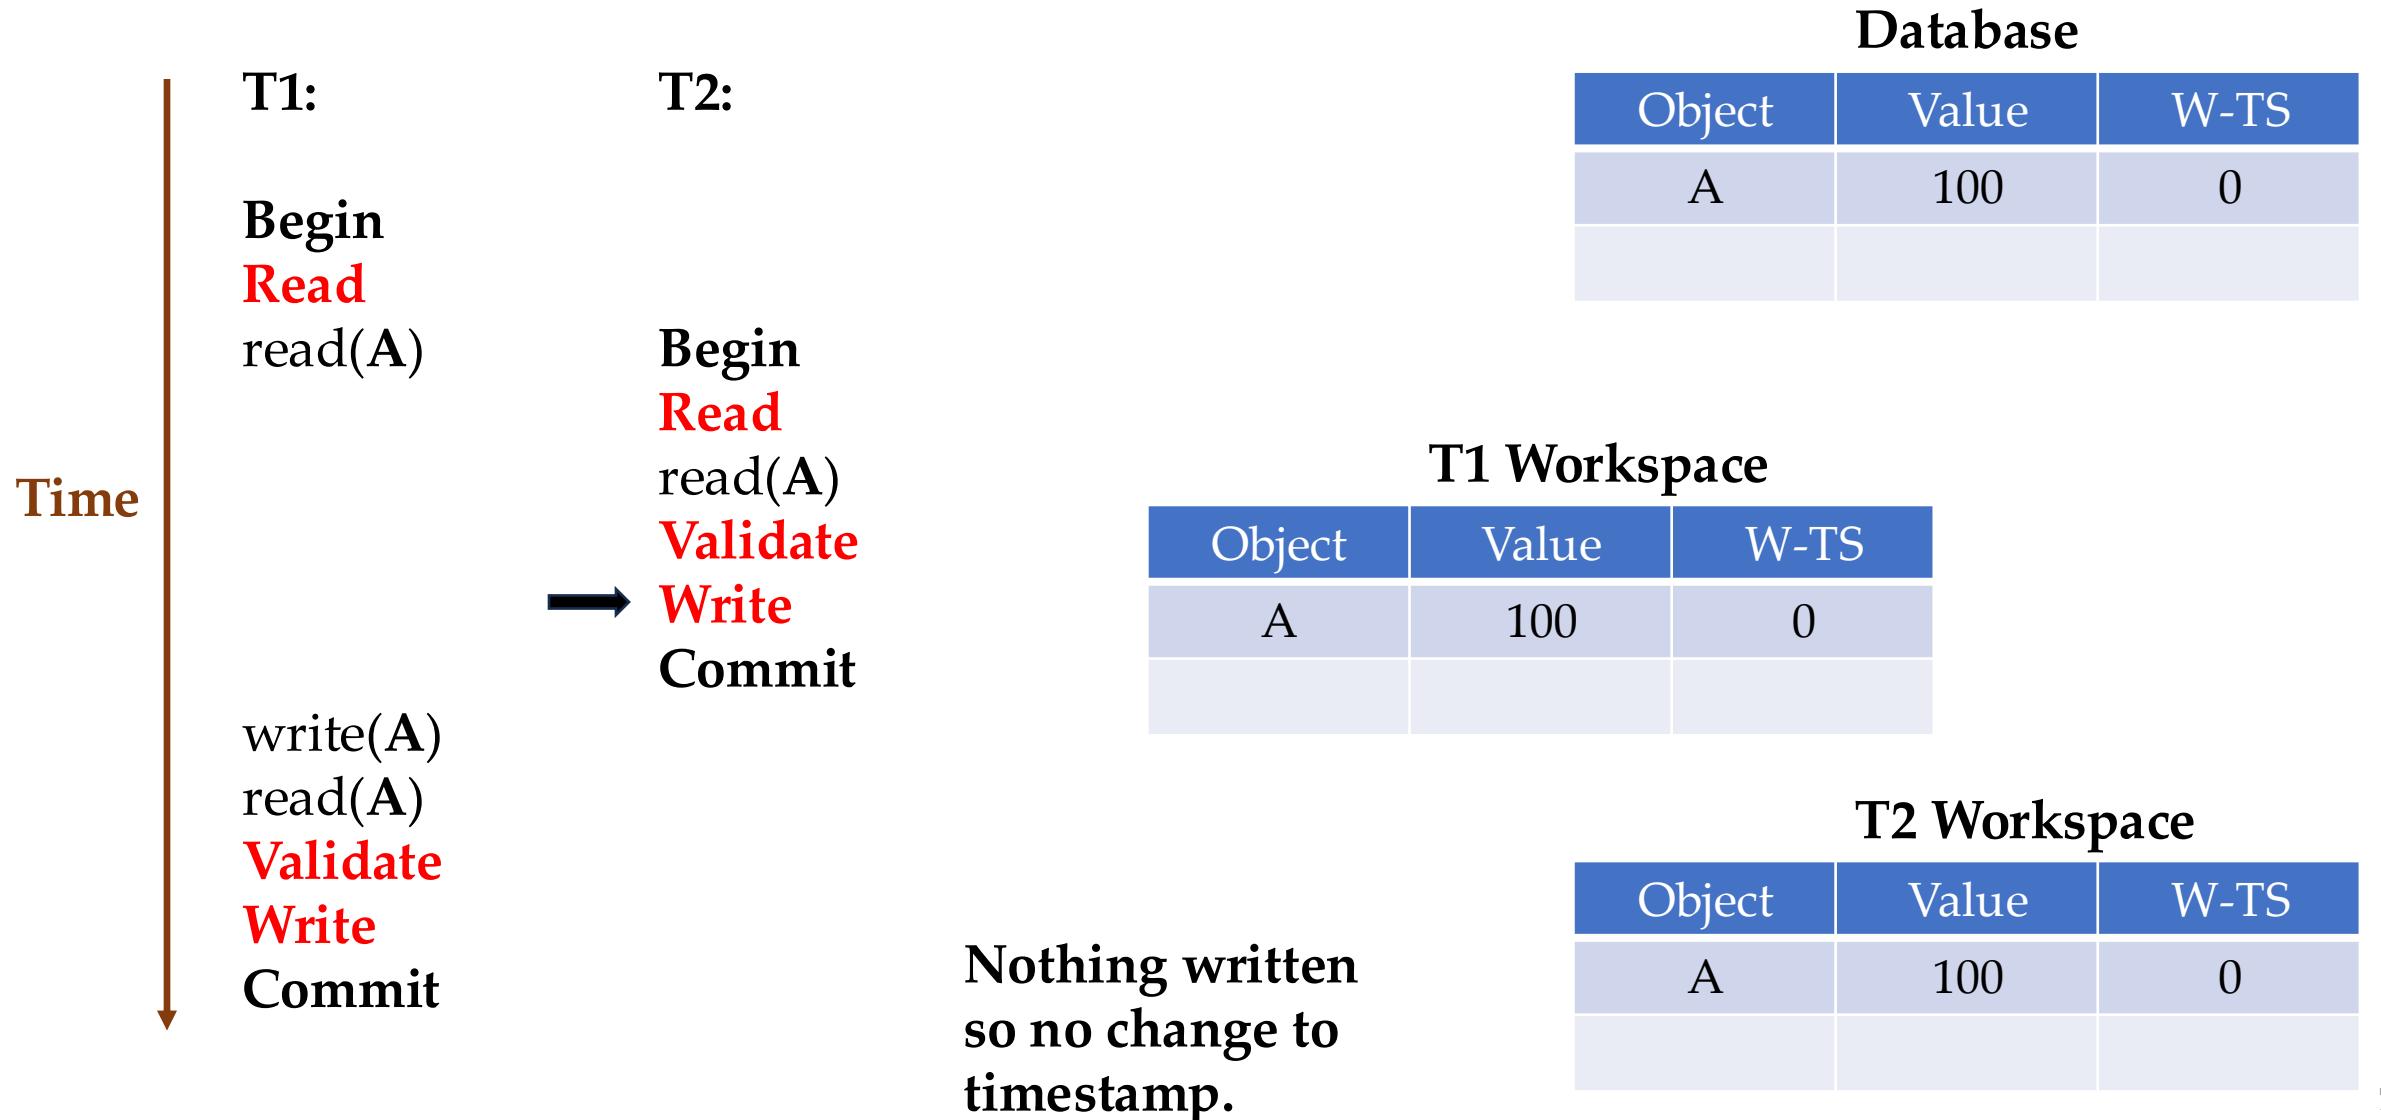
OCC Example I



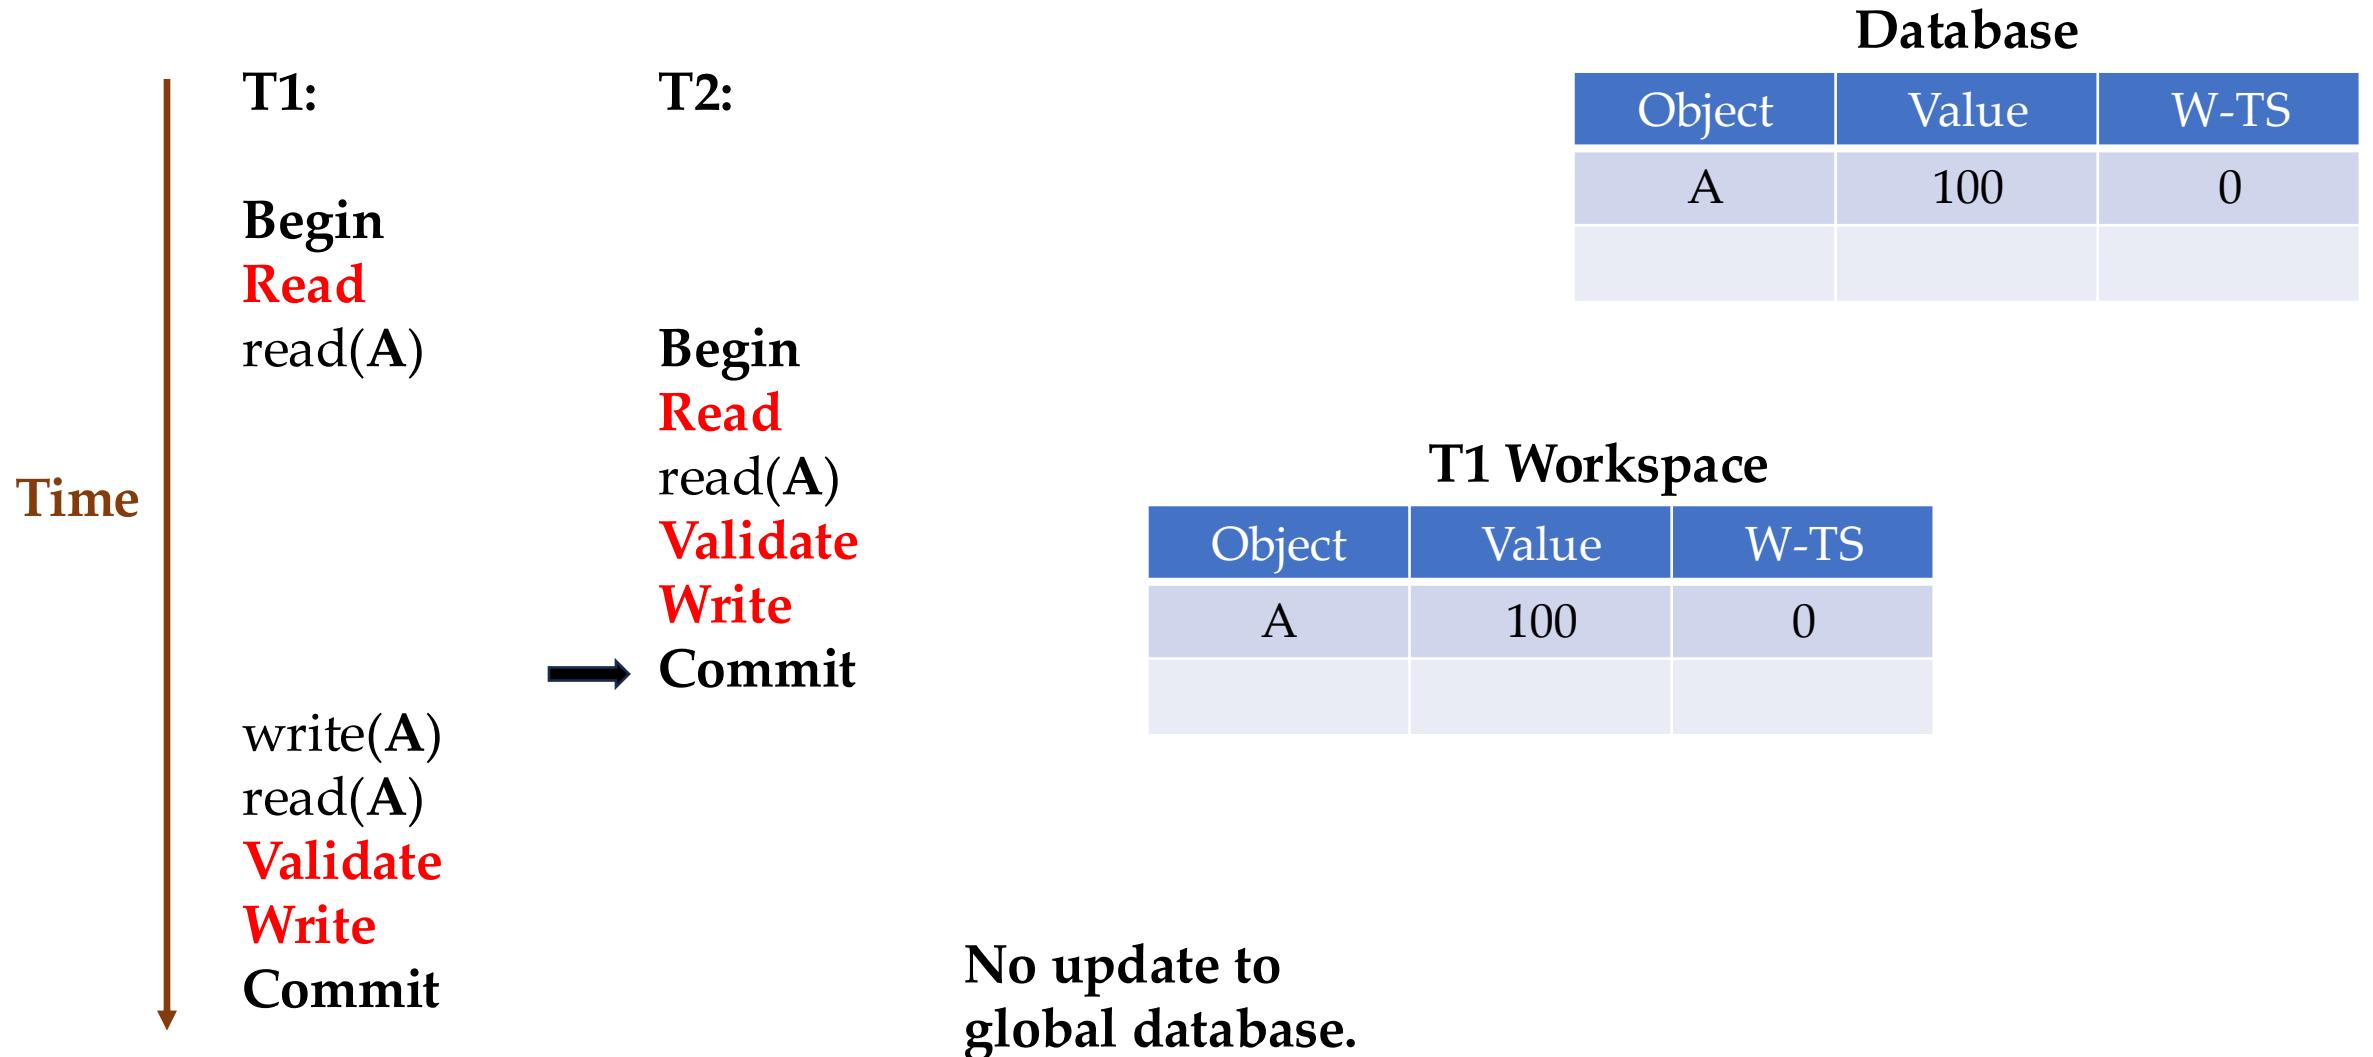
OCC Example I



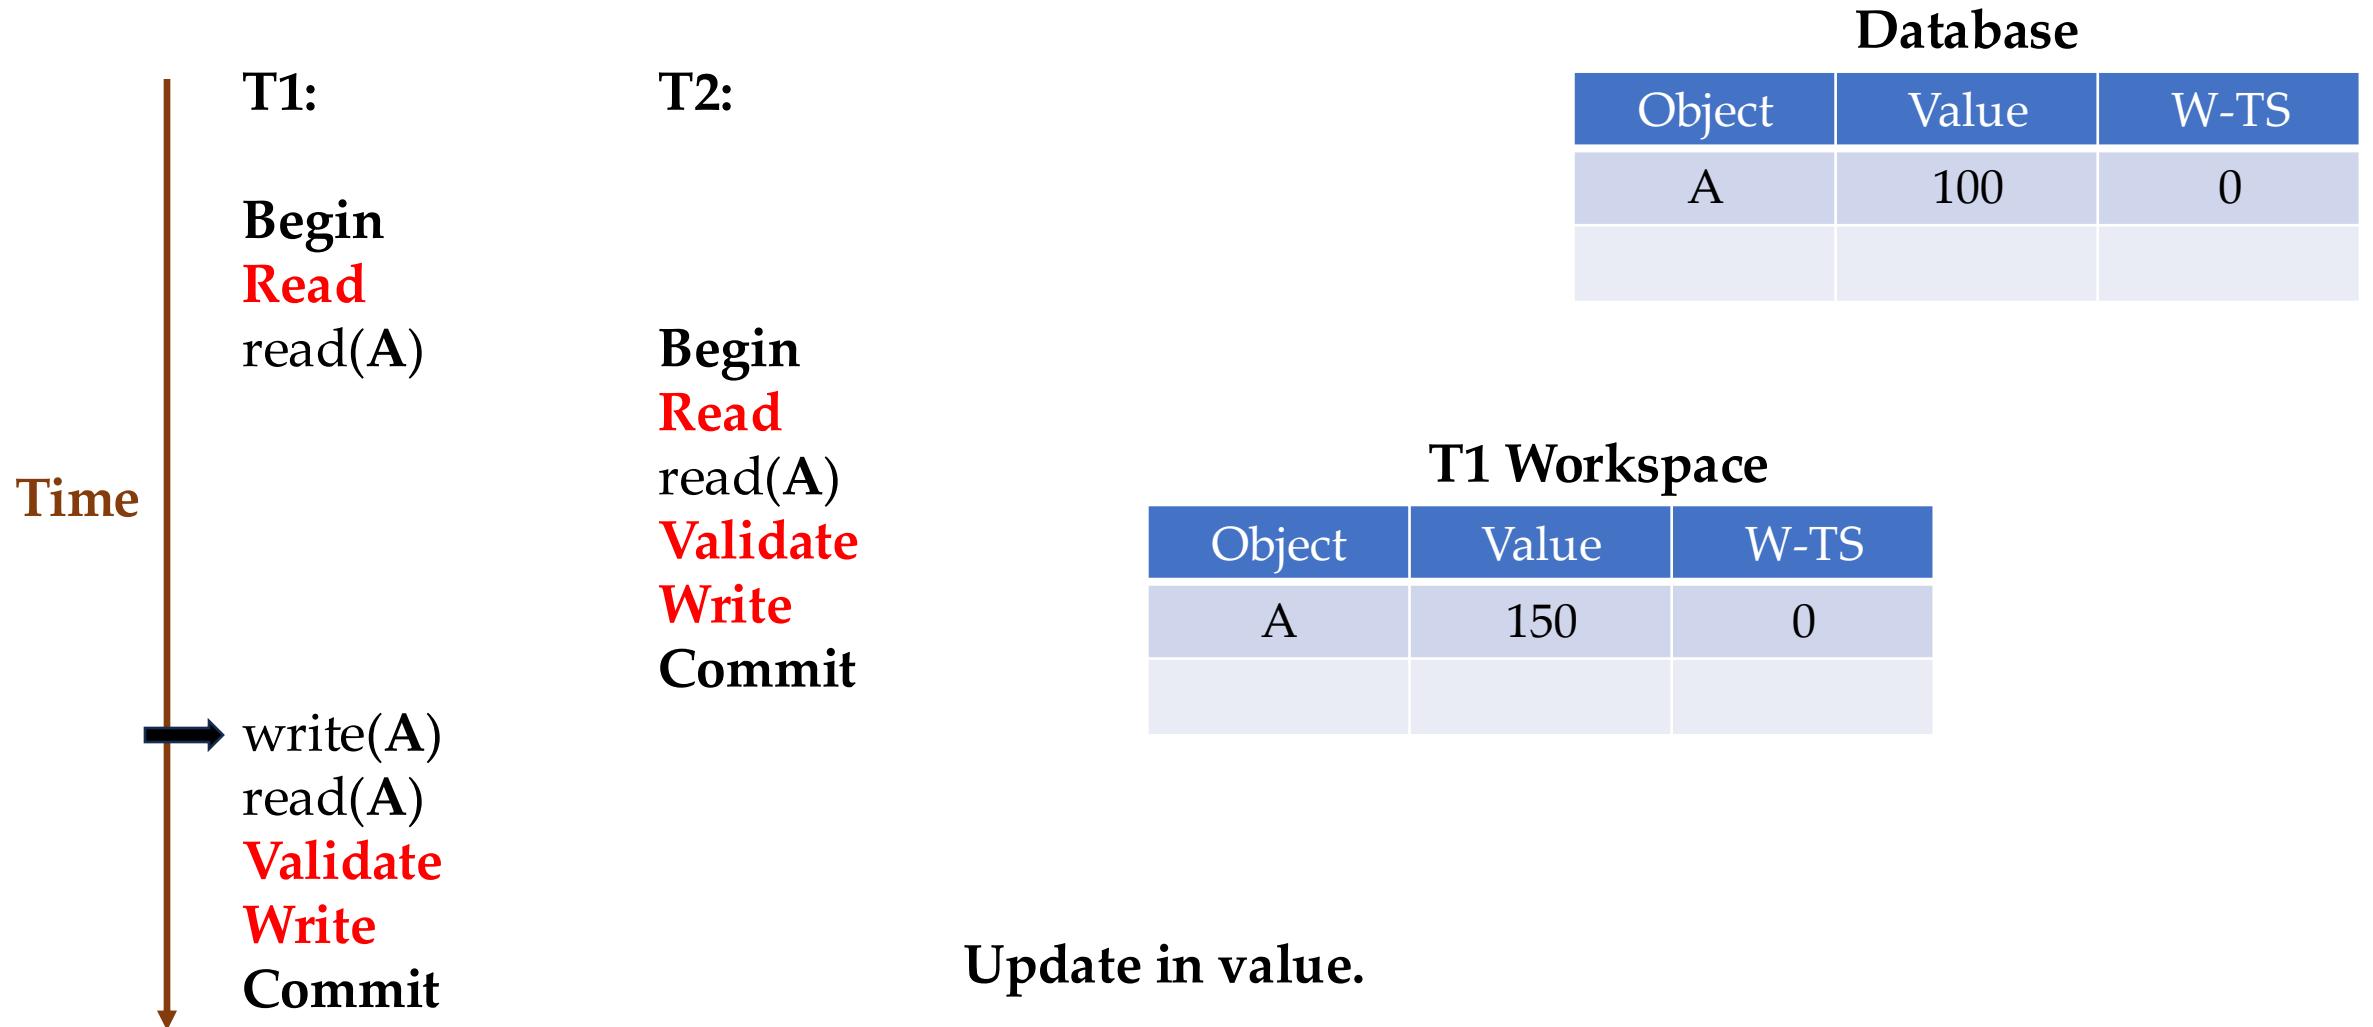
OCC Example I



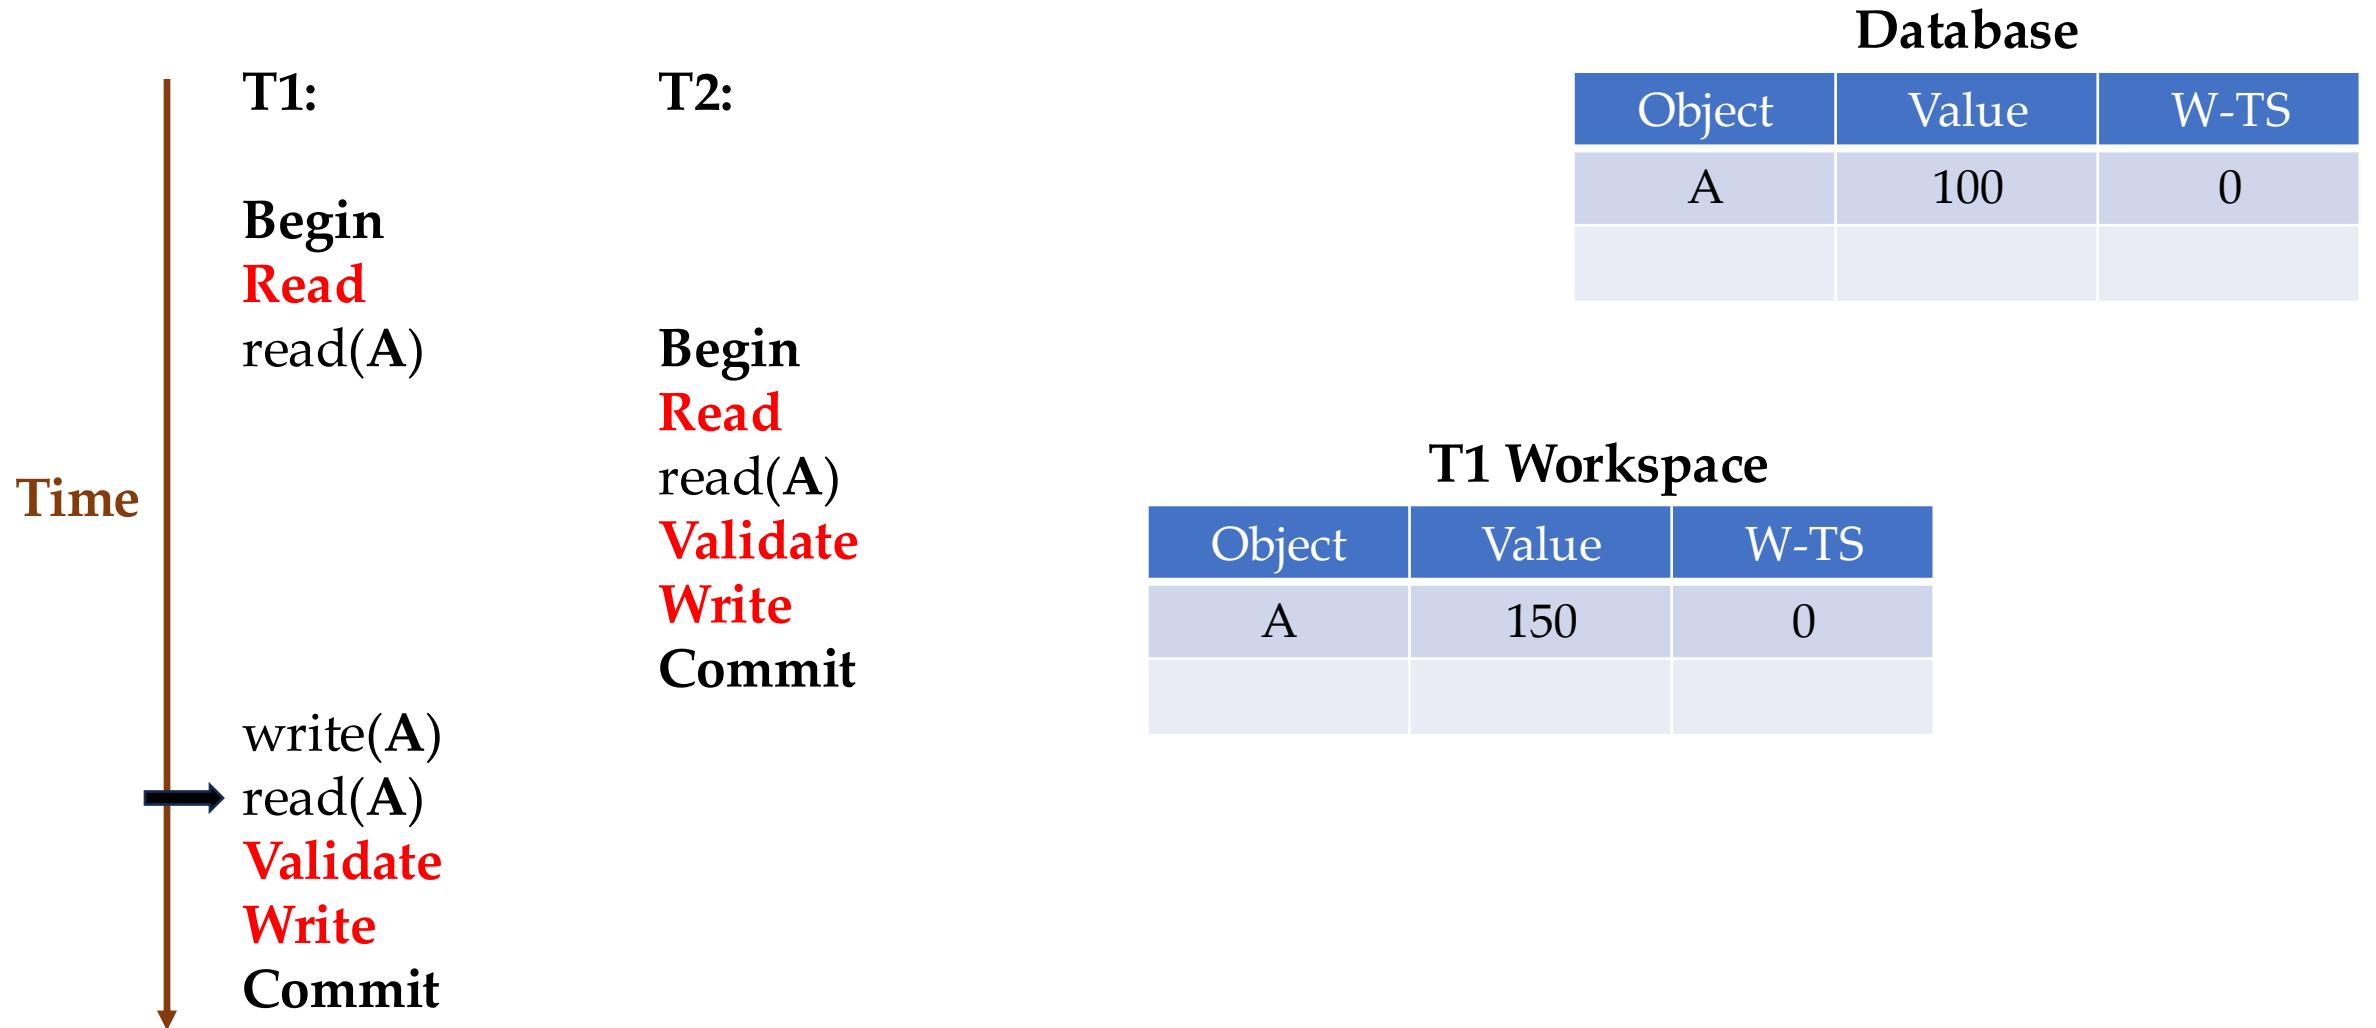
OCC Example I



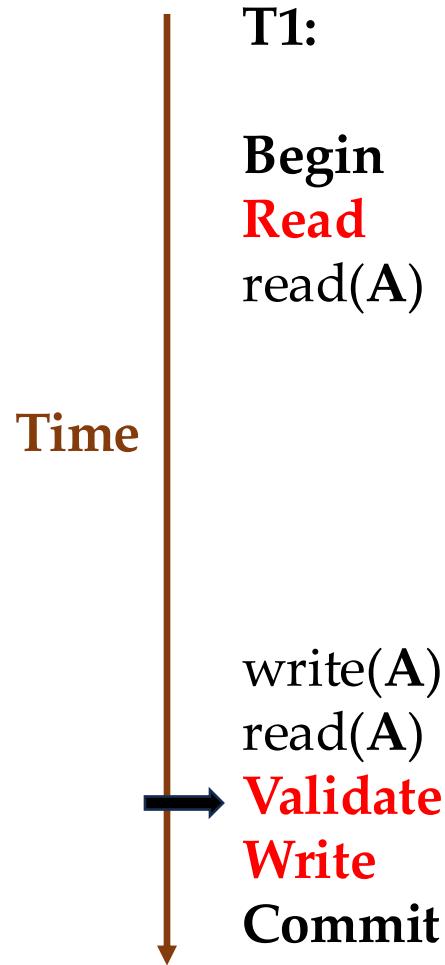
OCC Example I



OCC Example I



OCC Example I



T2:
Begin
Read
read(A)
Validate
Write
Commit

Database

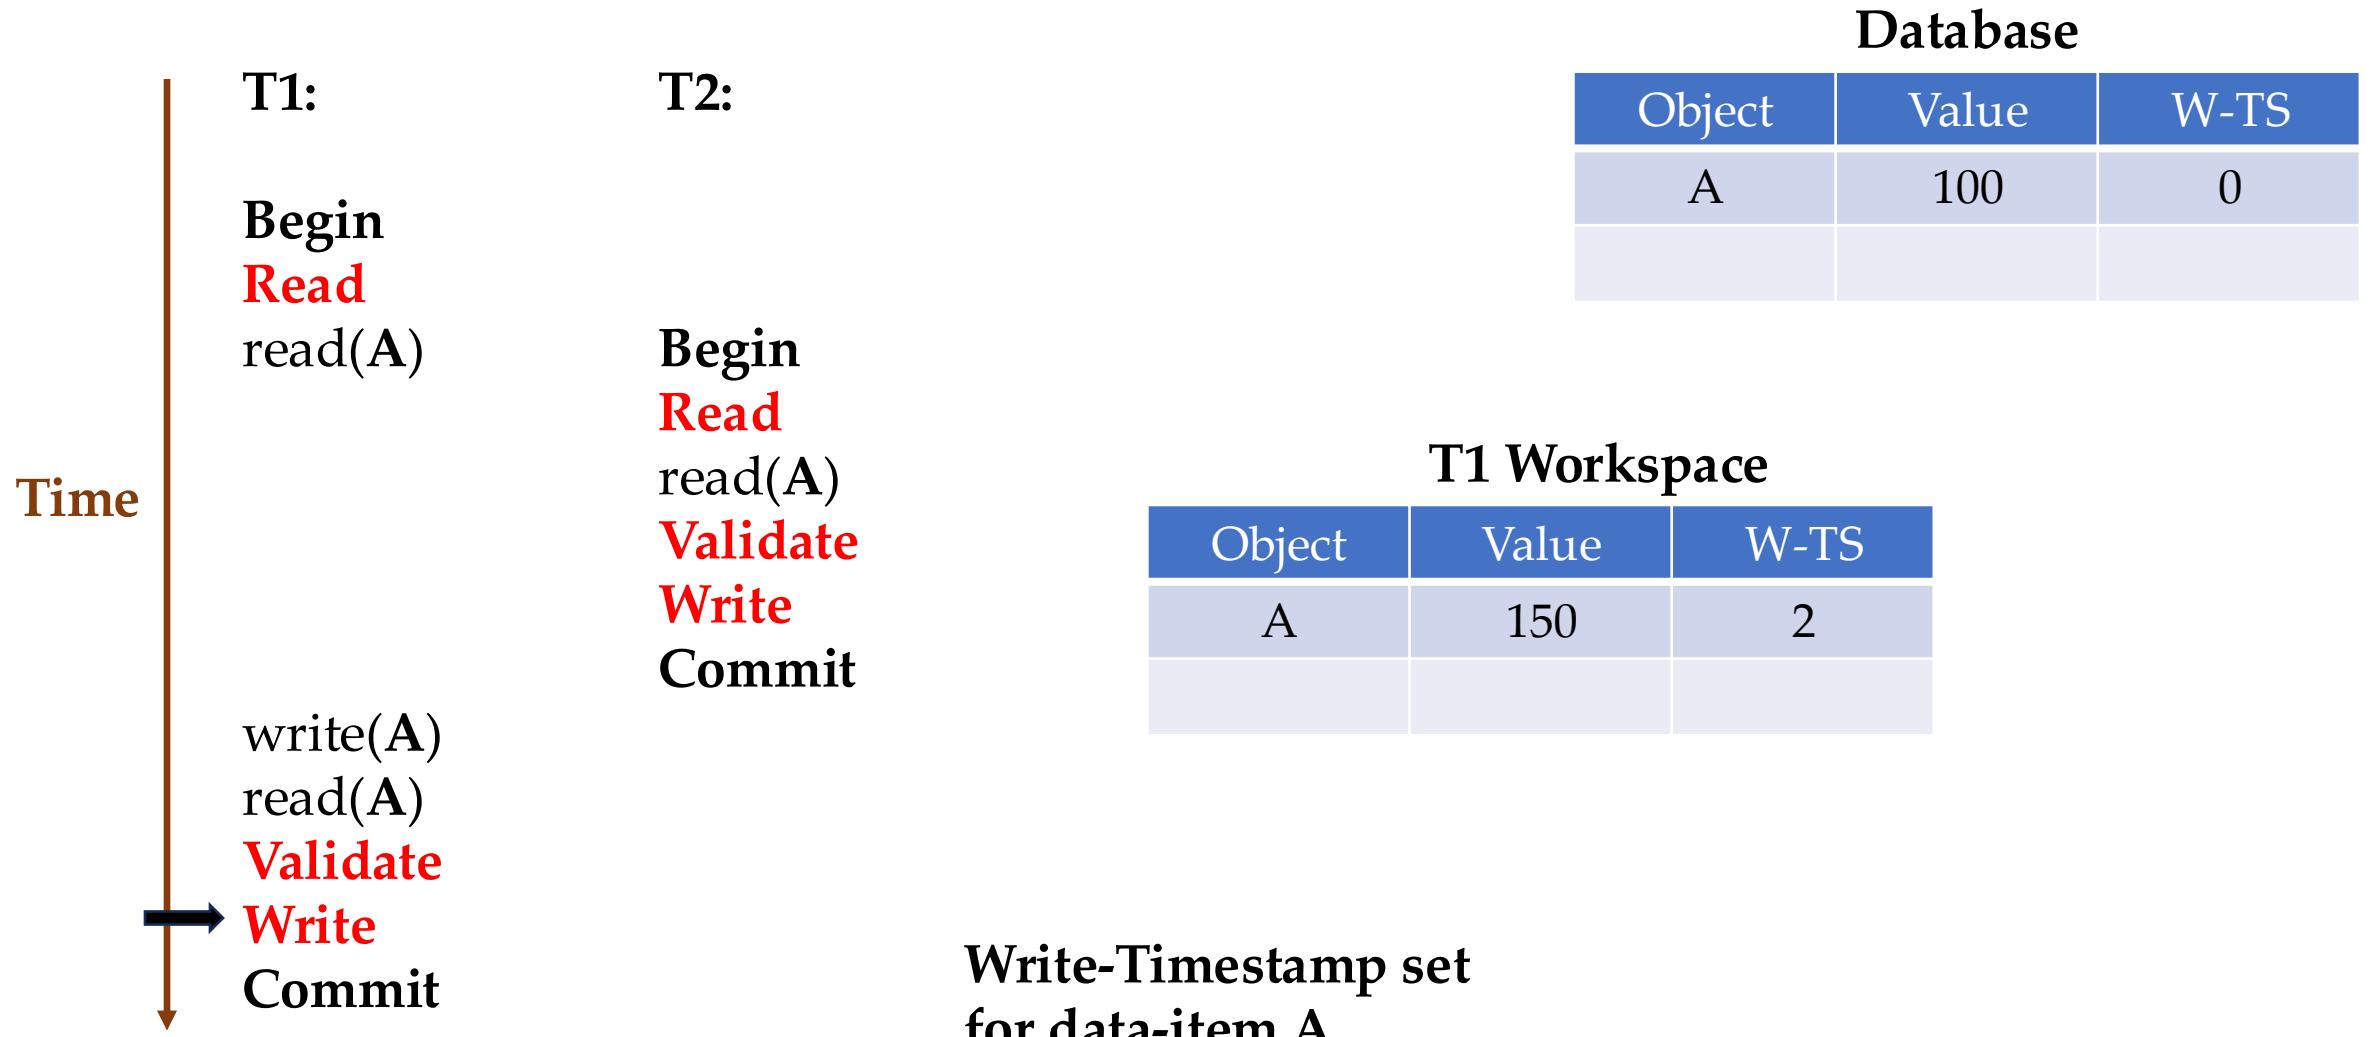
Object	Value	W-TS
A	100	0

T1 Workspace

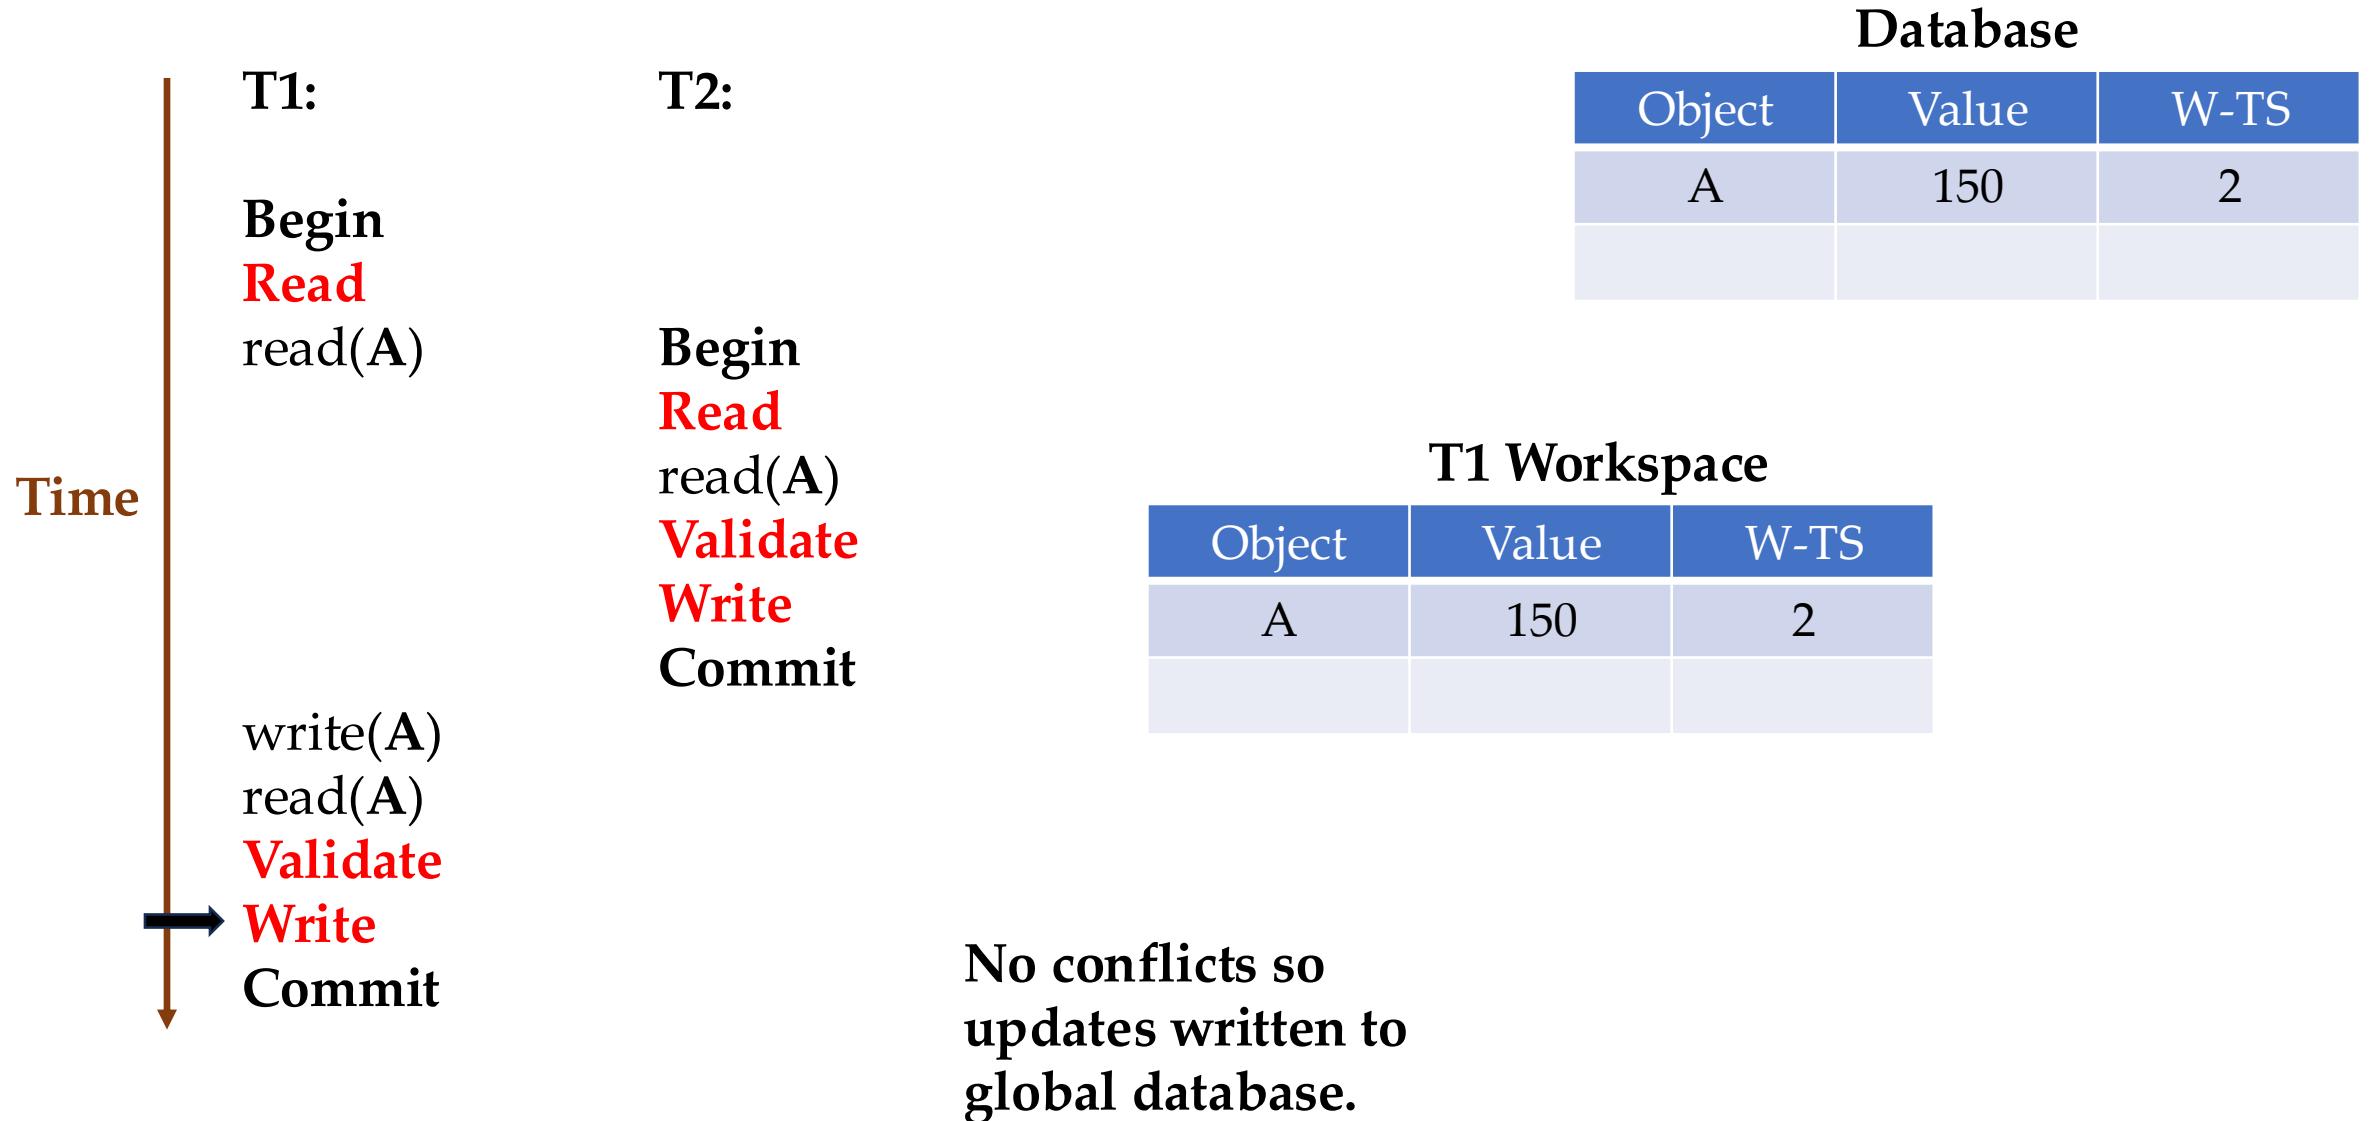
Object	Value	W-TS
A	150	0

Timestamp for T1:
 $TS(T1) = 2$

OCC Example I



OCC Example I



Validation Phase

- Assign the transaction a unique timestamp (TS) and then check whether it conflicts with other transactions.
- When transaction T_i invokes **Commit**, the DBMS checks if it conflicts with other transactions.
- Simplest mechanism → Use serial validation.
- How can DBMS guarantee only serializable schedules are permitted?

Validation Phase

- Assign the transaction a unique timestamp (TS) and then check whether it conflicts with other transactions.
- When transaction T_i invokes **Commit**, the DBMS checks if it conflicts with other transactions.
- Simplest mechanism → Use serial validation.
- How can DBMS guarantee only serializable schedules are permitted?
 - Forward Validation
 - Backward Validation

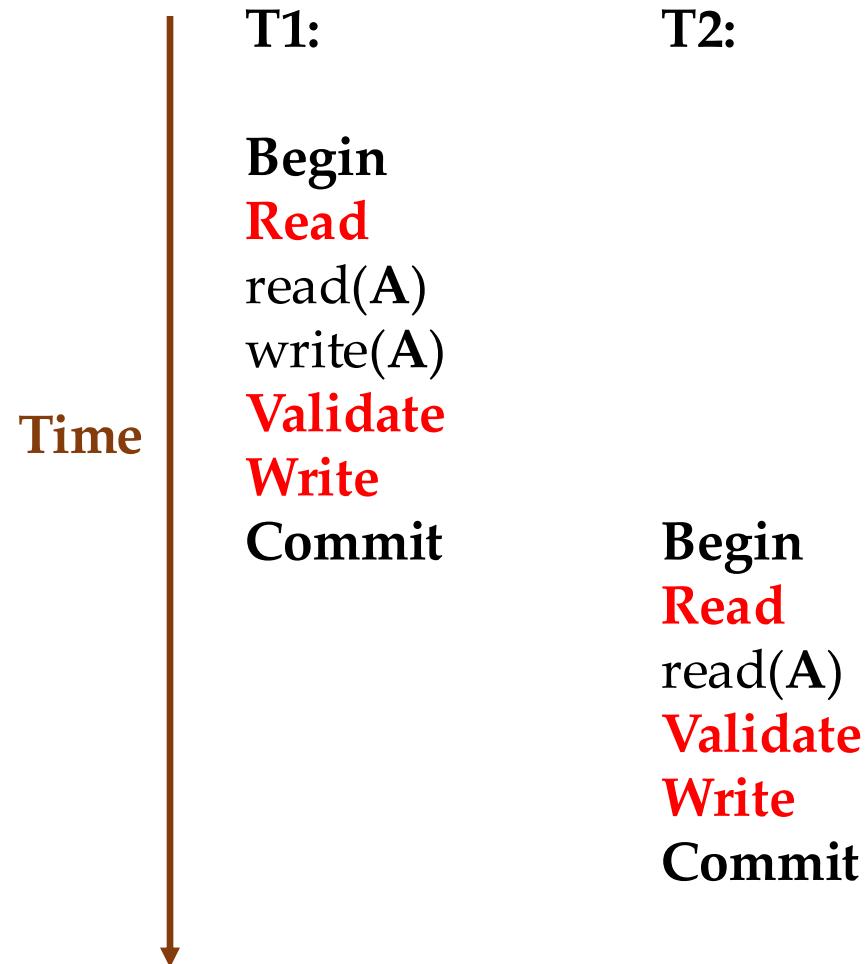
Forward Validation

- At the time of commit, each transaction checks if it conflicts with other **concurrently ongoing transactions** (yet to be committed).
- Each going to commit transaction (at the validation step), checks the timestamps and read/write sets of other ongoing transactions.
- There are three specific cases to satisfy:

Forward Validation: Case I

- For two transactions T1 and T2, say T1 is at the validation step ($T1 < T2$)
 - Check if T1 completes its Write phase before T2 begins its Read phase.
 - No conflict as all T1 's actions happen before T2 's.
 - Essentially, serial ordering.

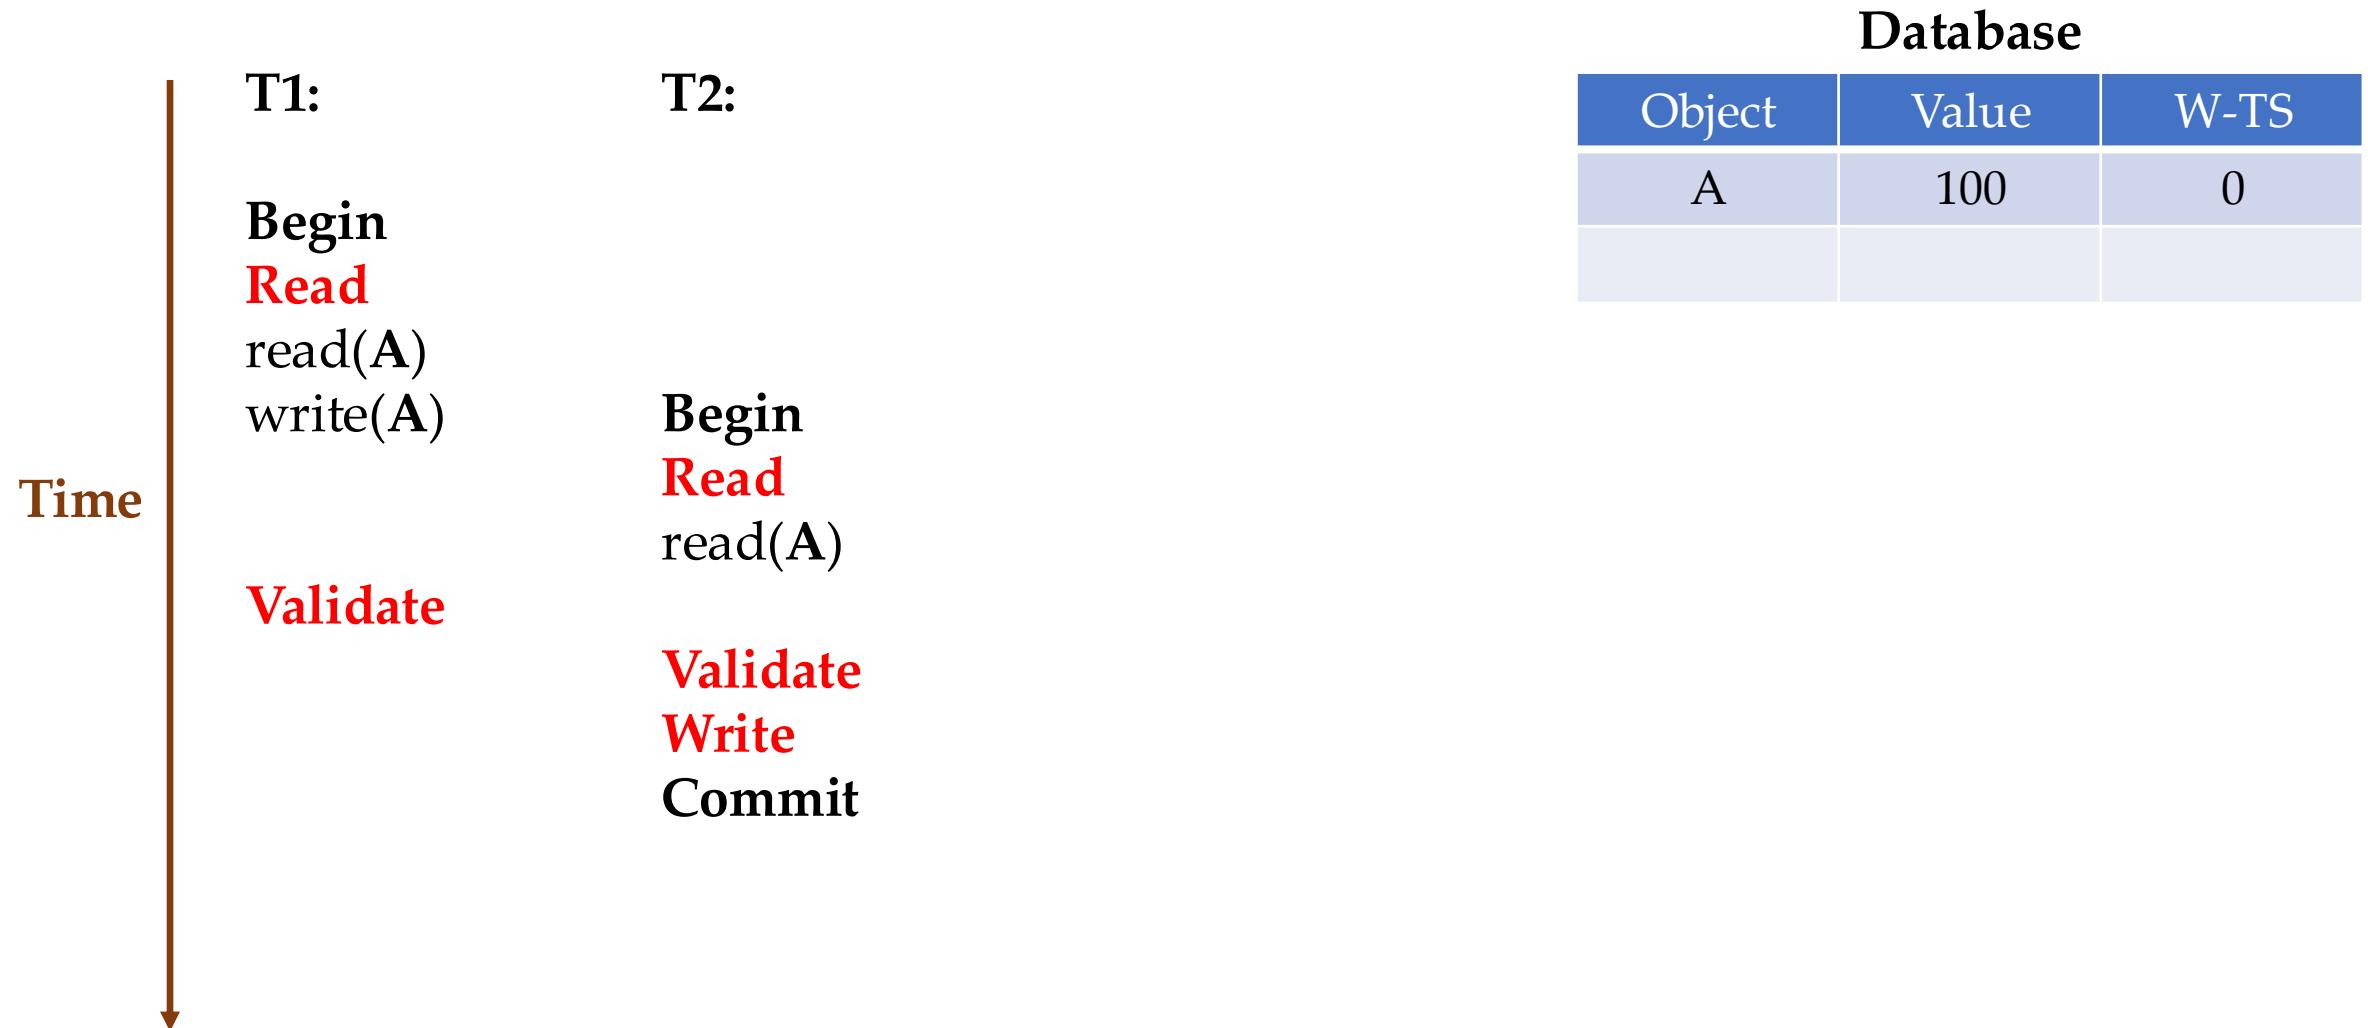
Forward Validation: Case I



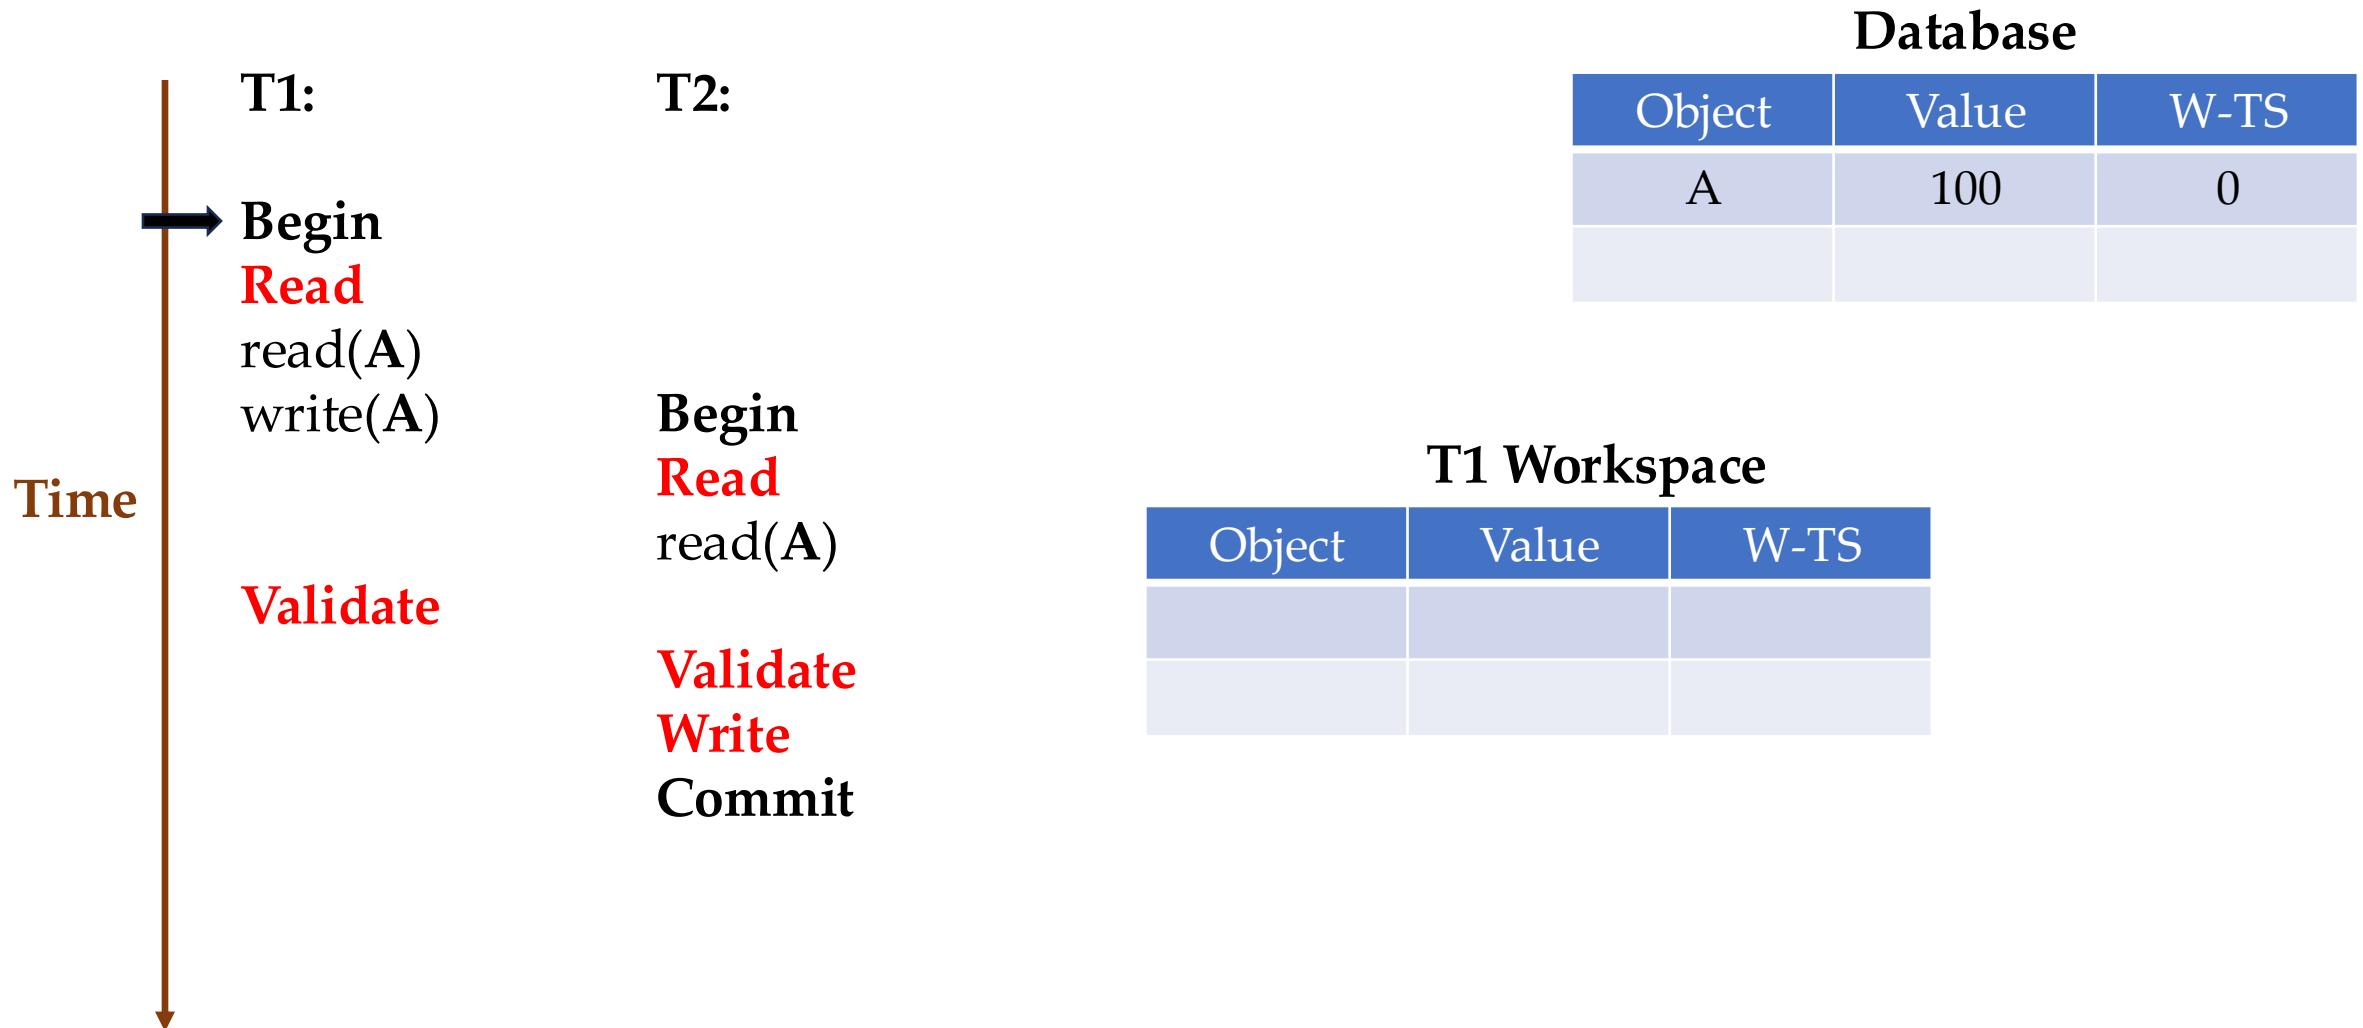
Forward Validation: Case II

- For two transactions T1 and T2, say T1 is at the validation step ($T1 < T2$)
 - Check if T1 completes its Write phase before T2 starts its Write phase.
 - T1 does not modify to any object read by T2.
 - $\text{WriteSet}(T1) \cap \text{ReadSet}(T2) = 0$

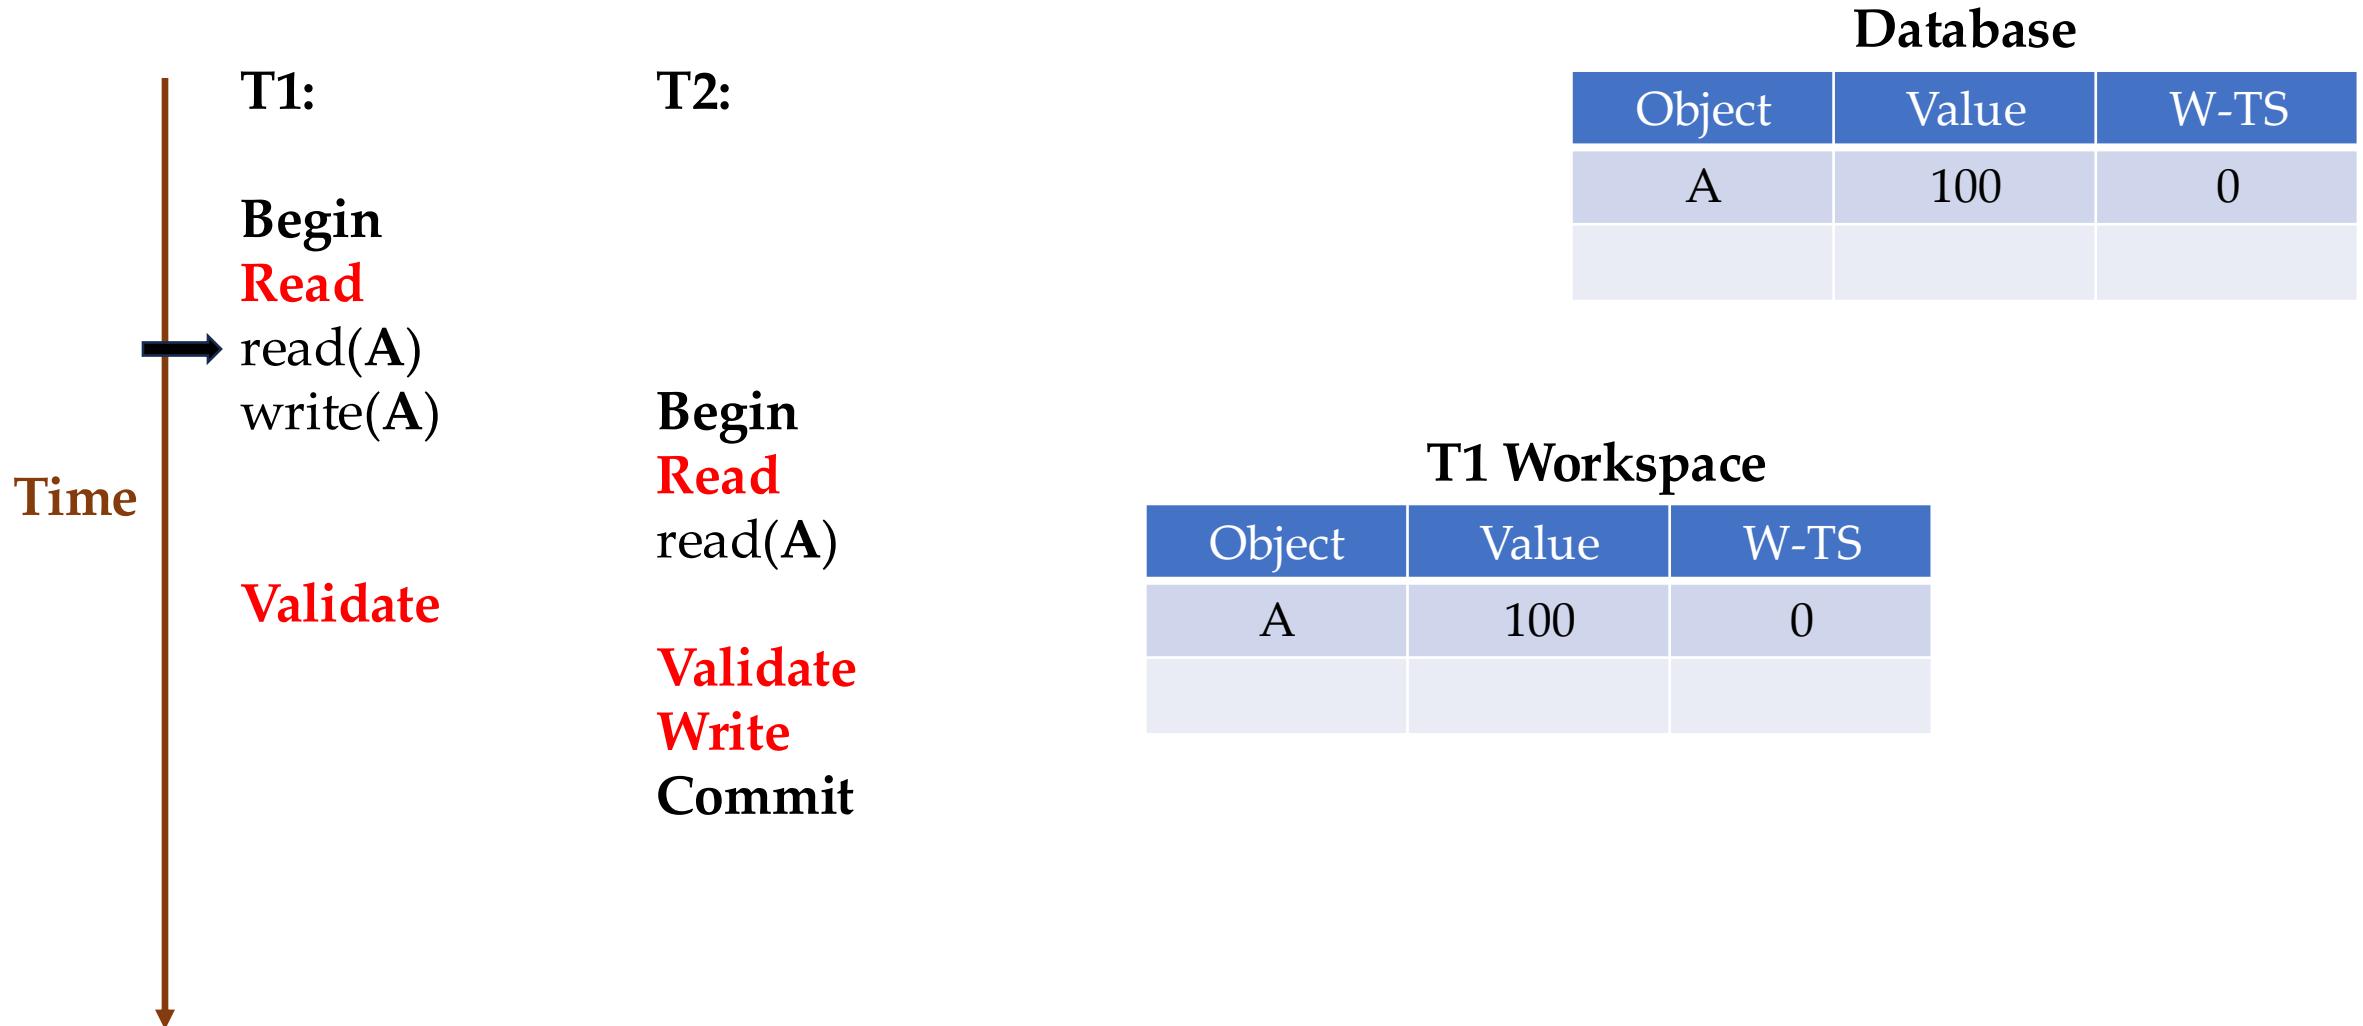
Forward Validation: Case II



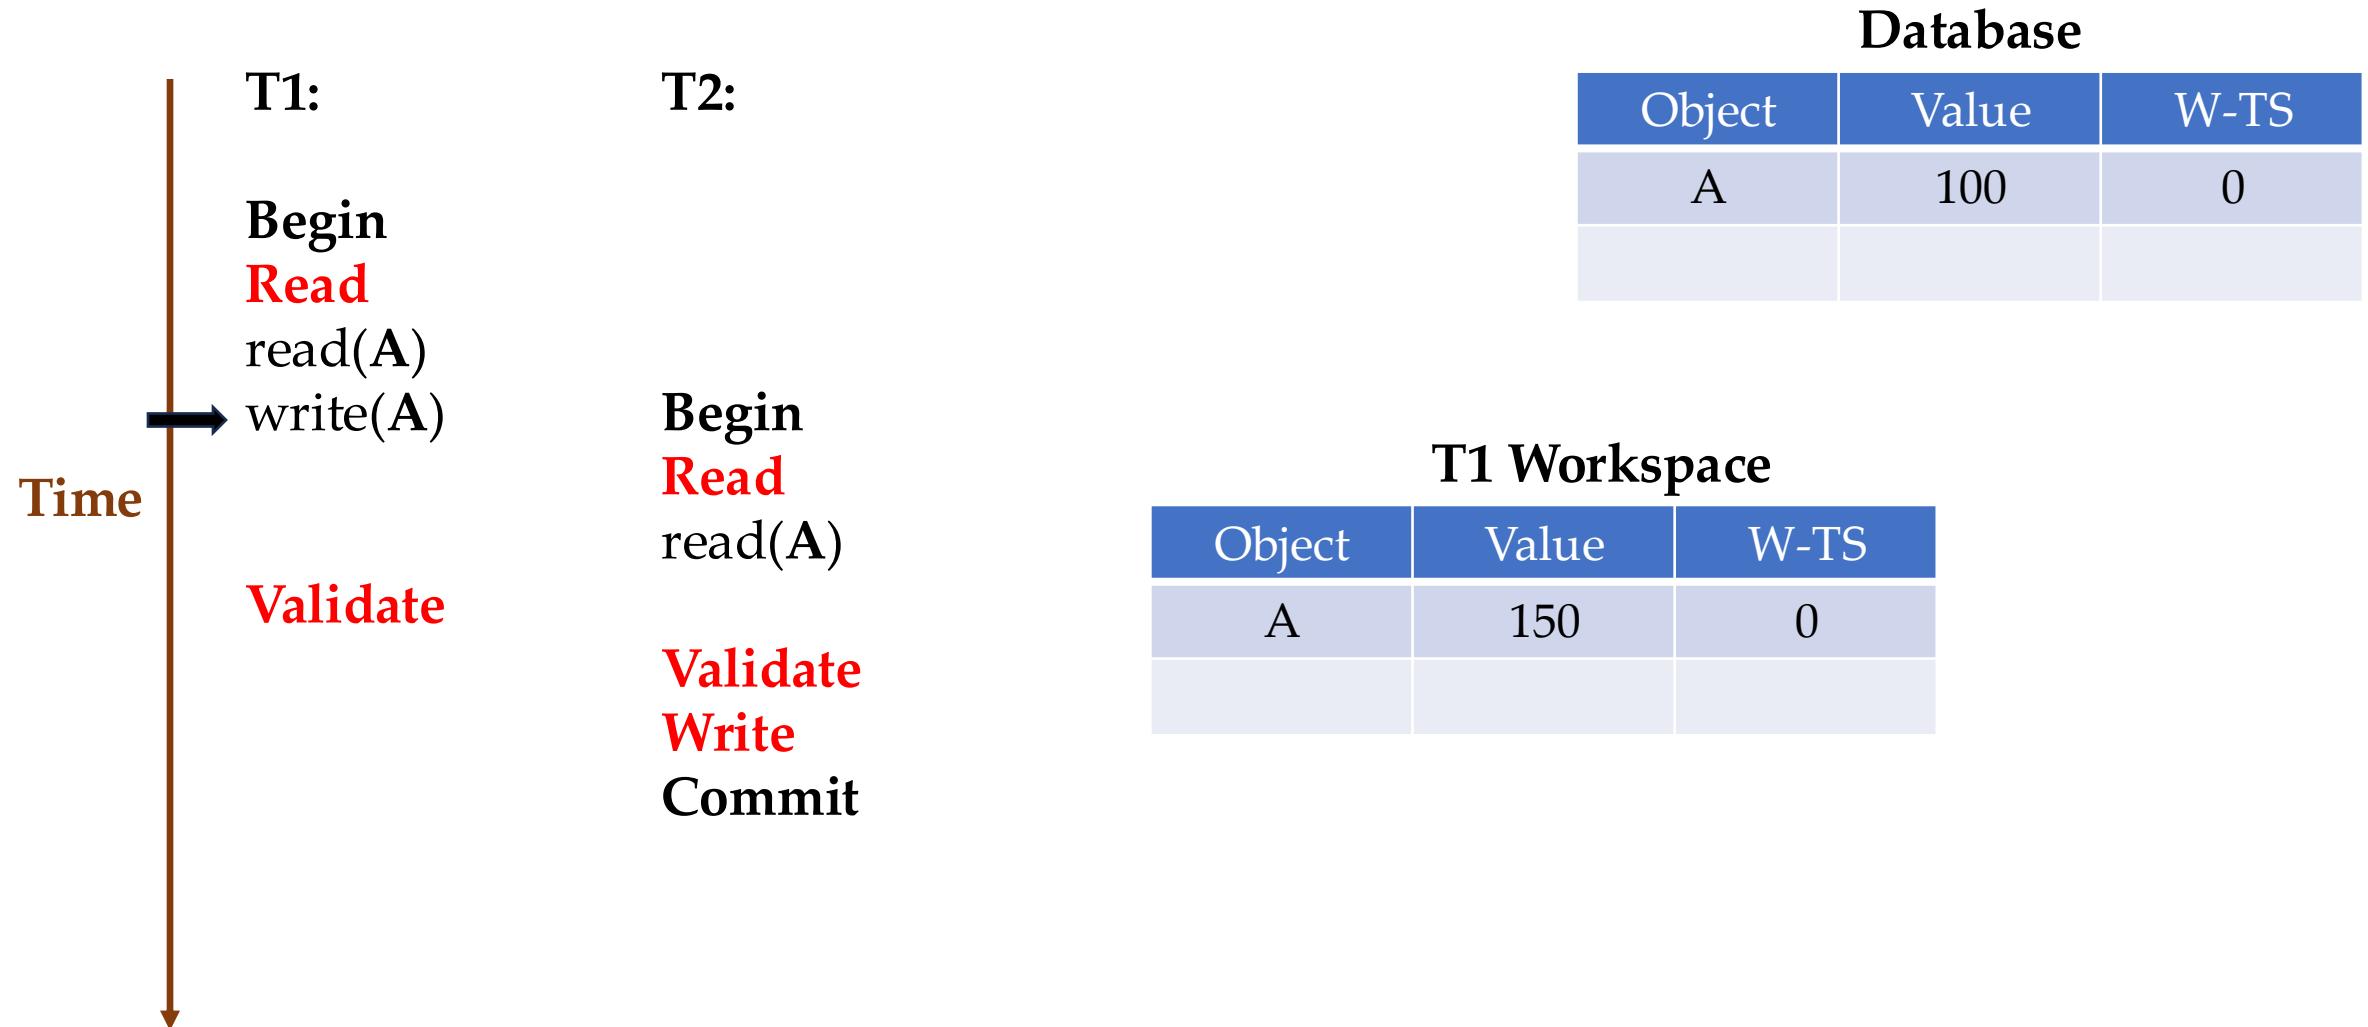
Forward Validation: Case II



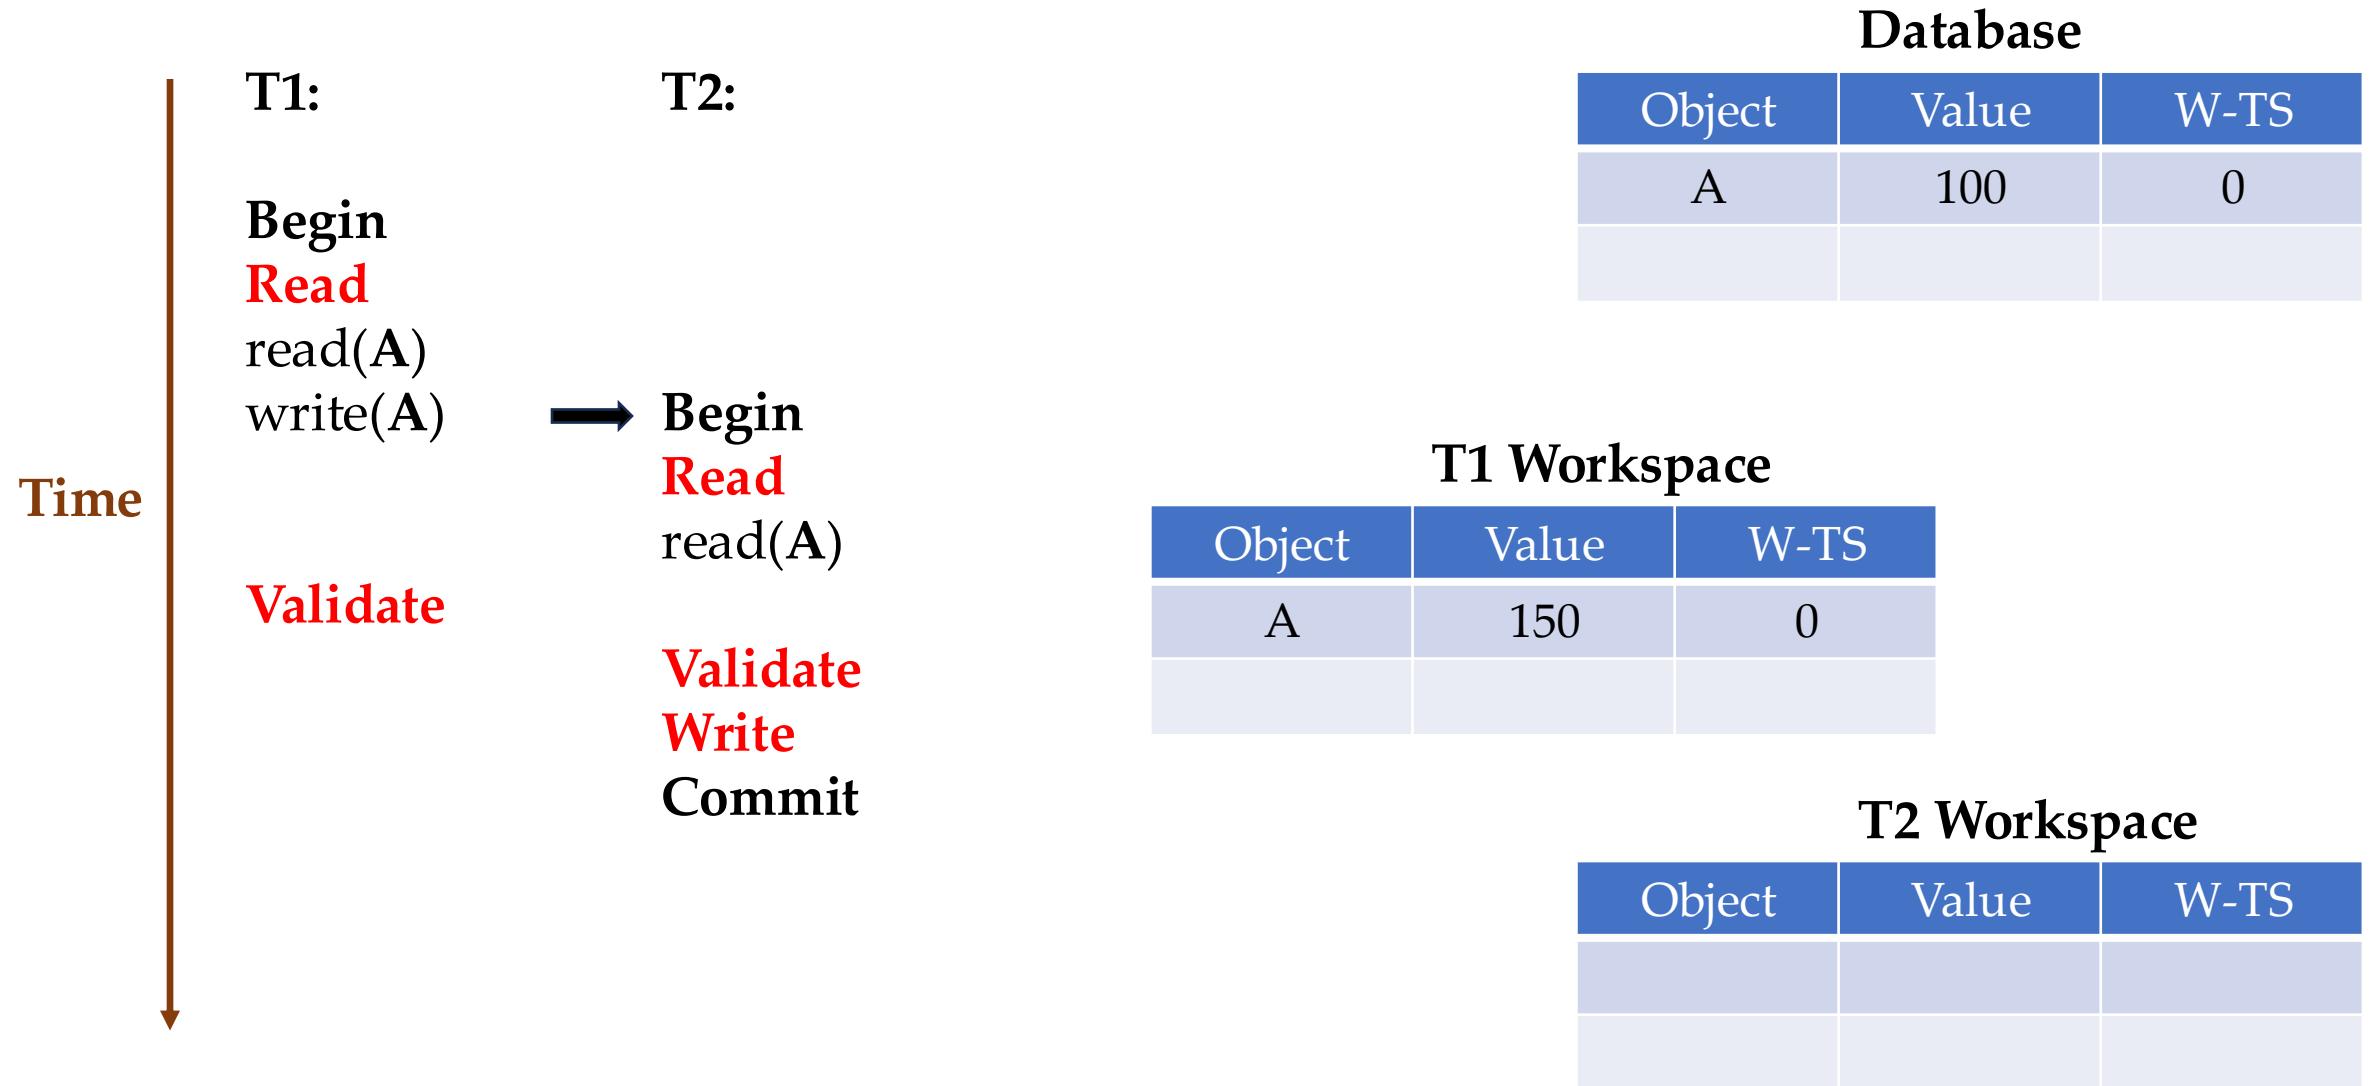
Forward Validation: Case II



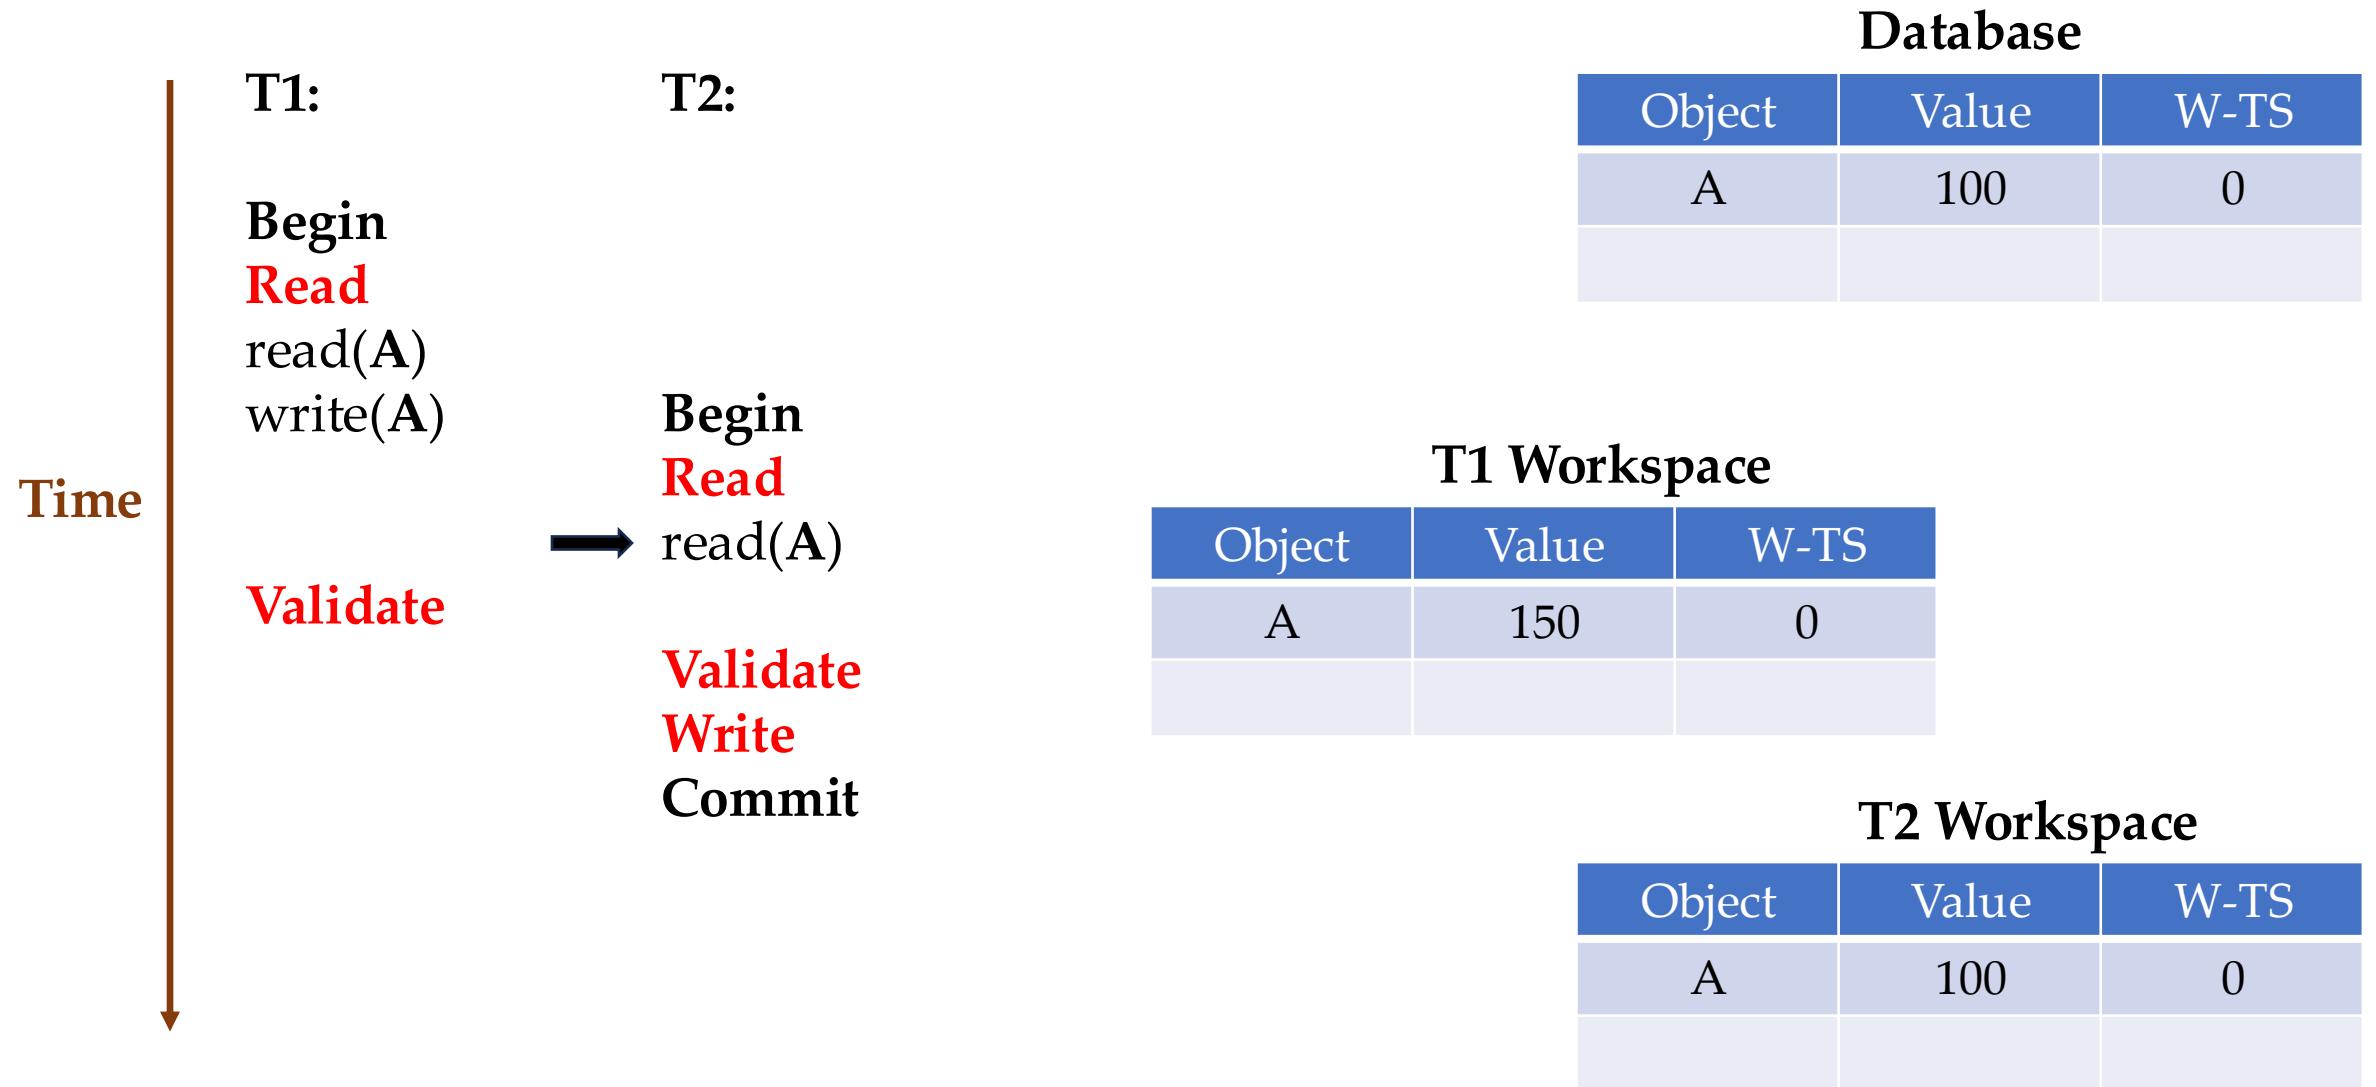
Forward Validation: Case II



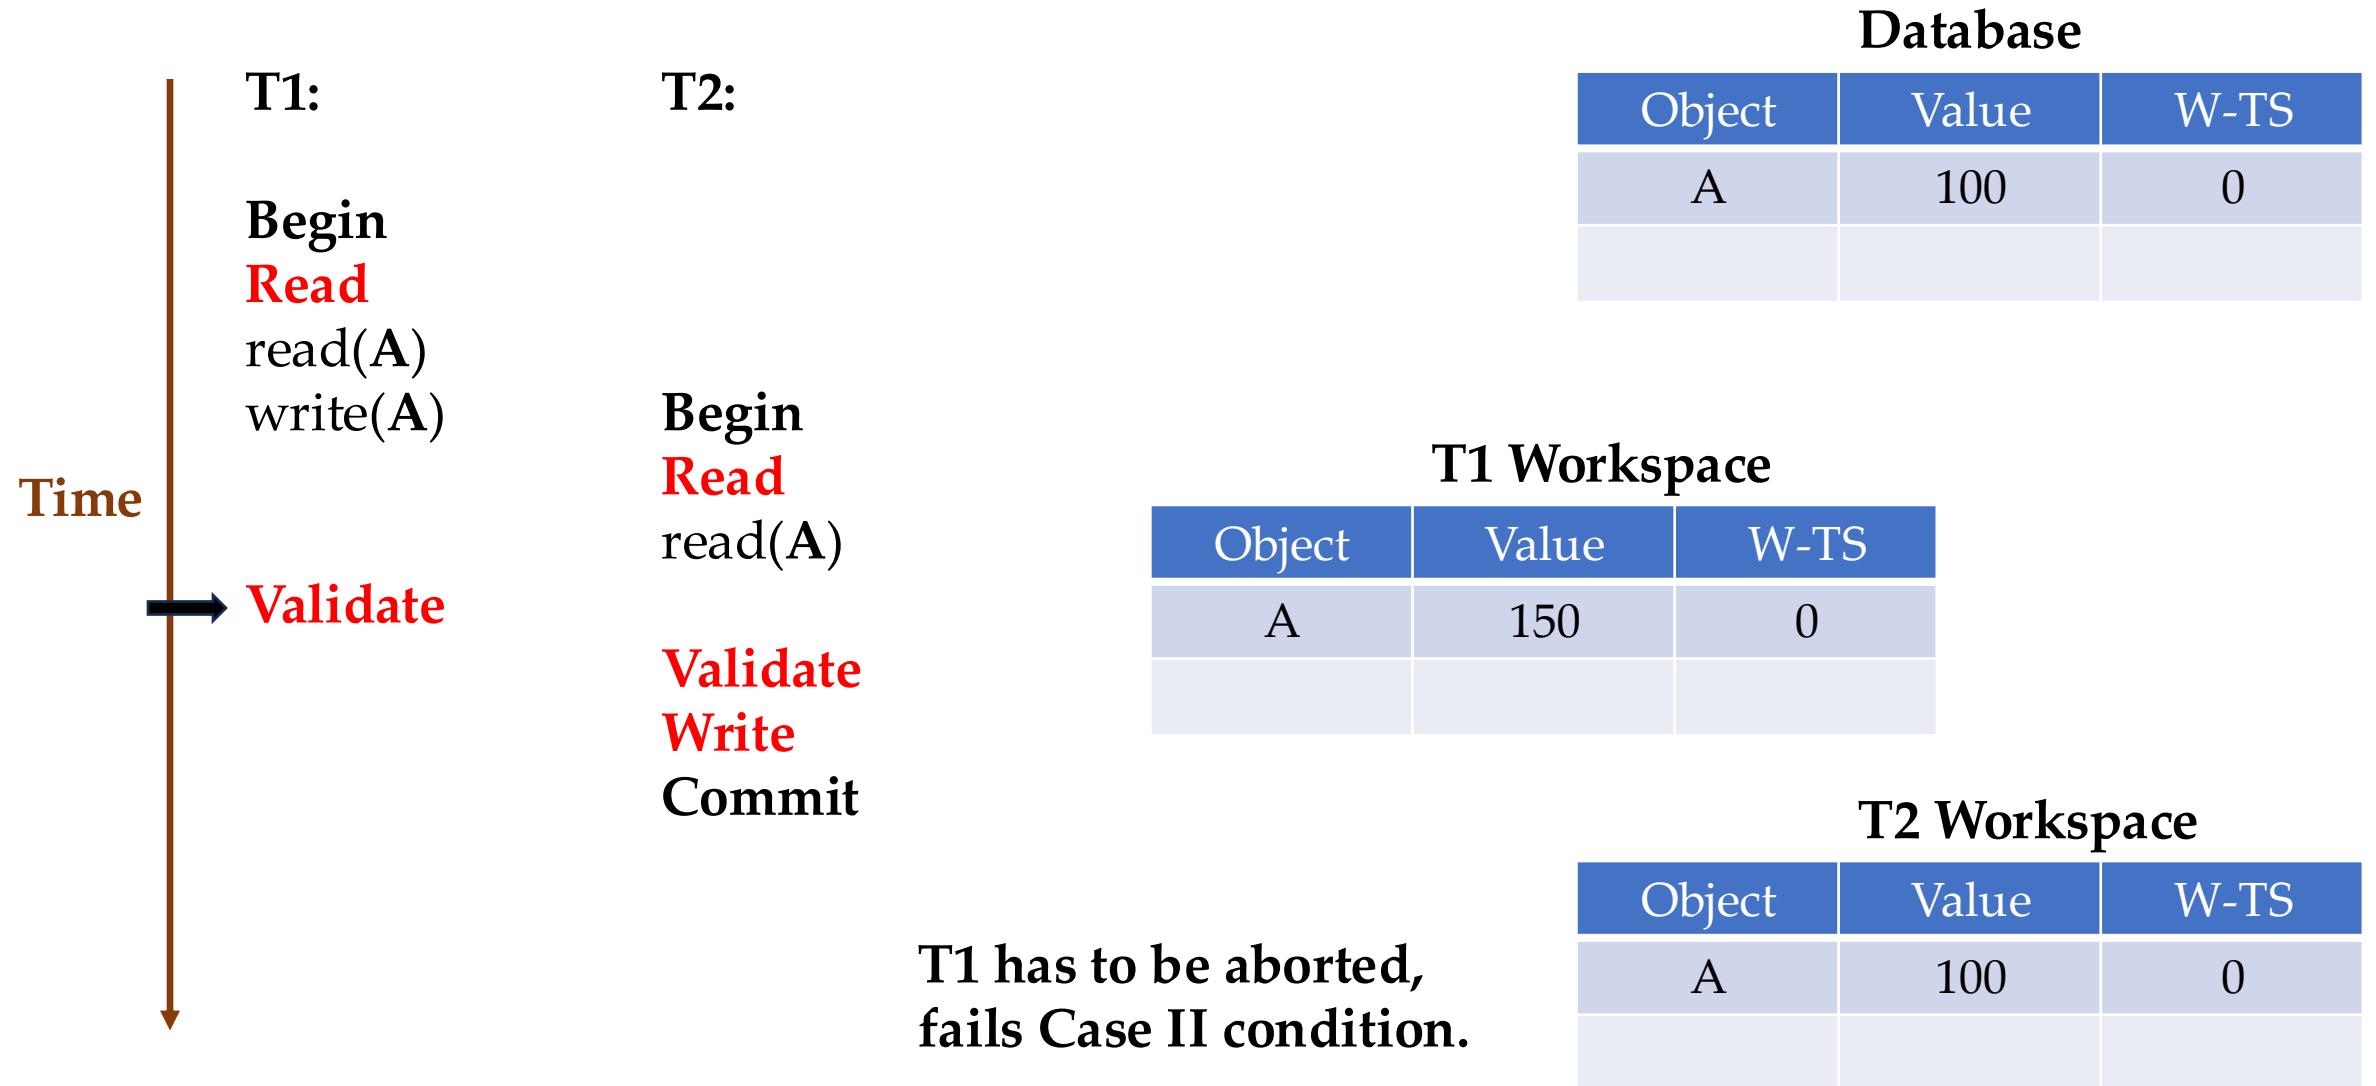
Forward Validation: Case II



Forward Validation: Case II

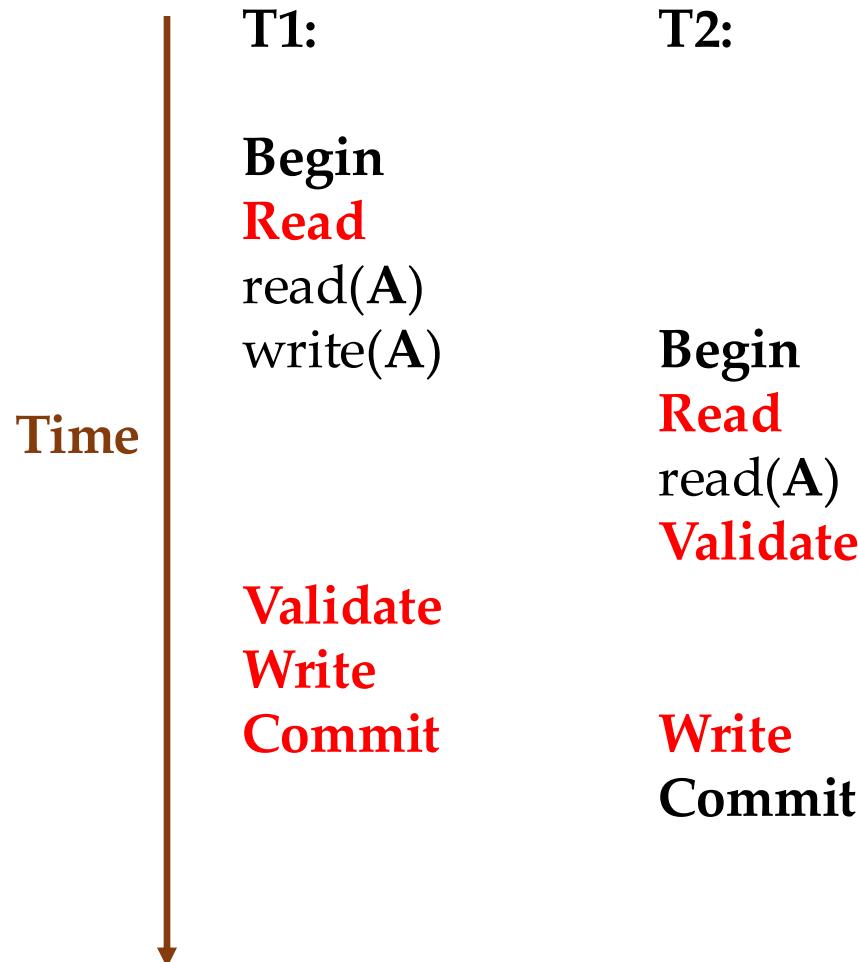


Forward Validation: Case II



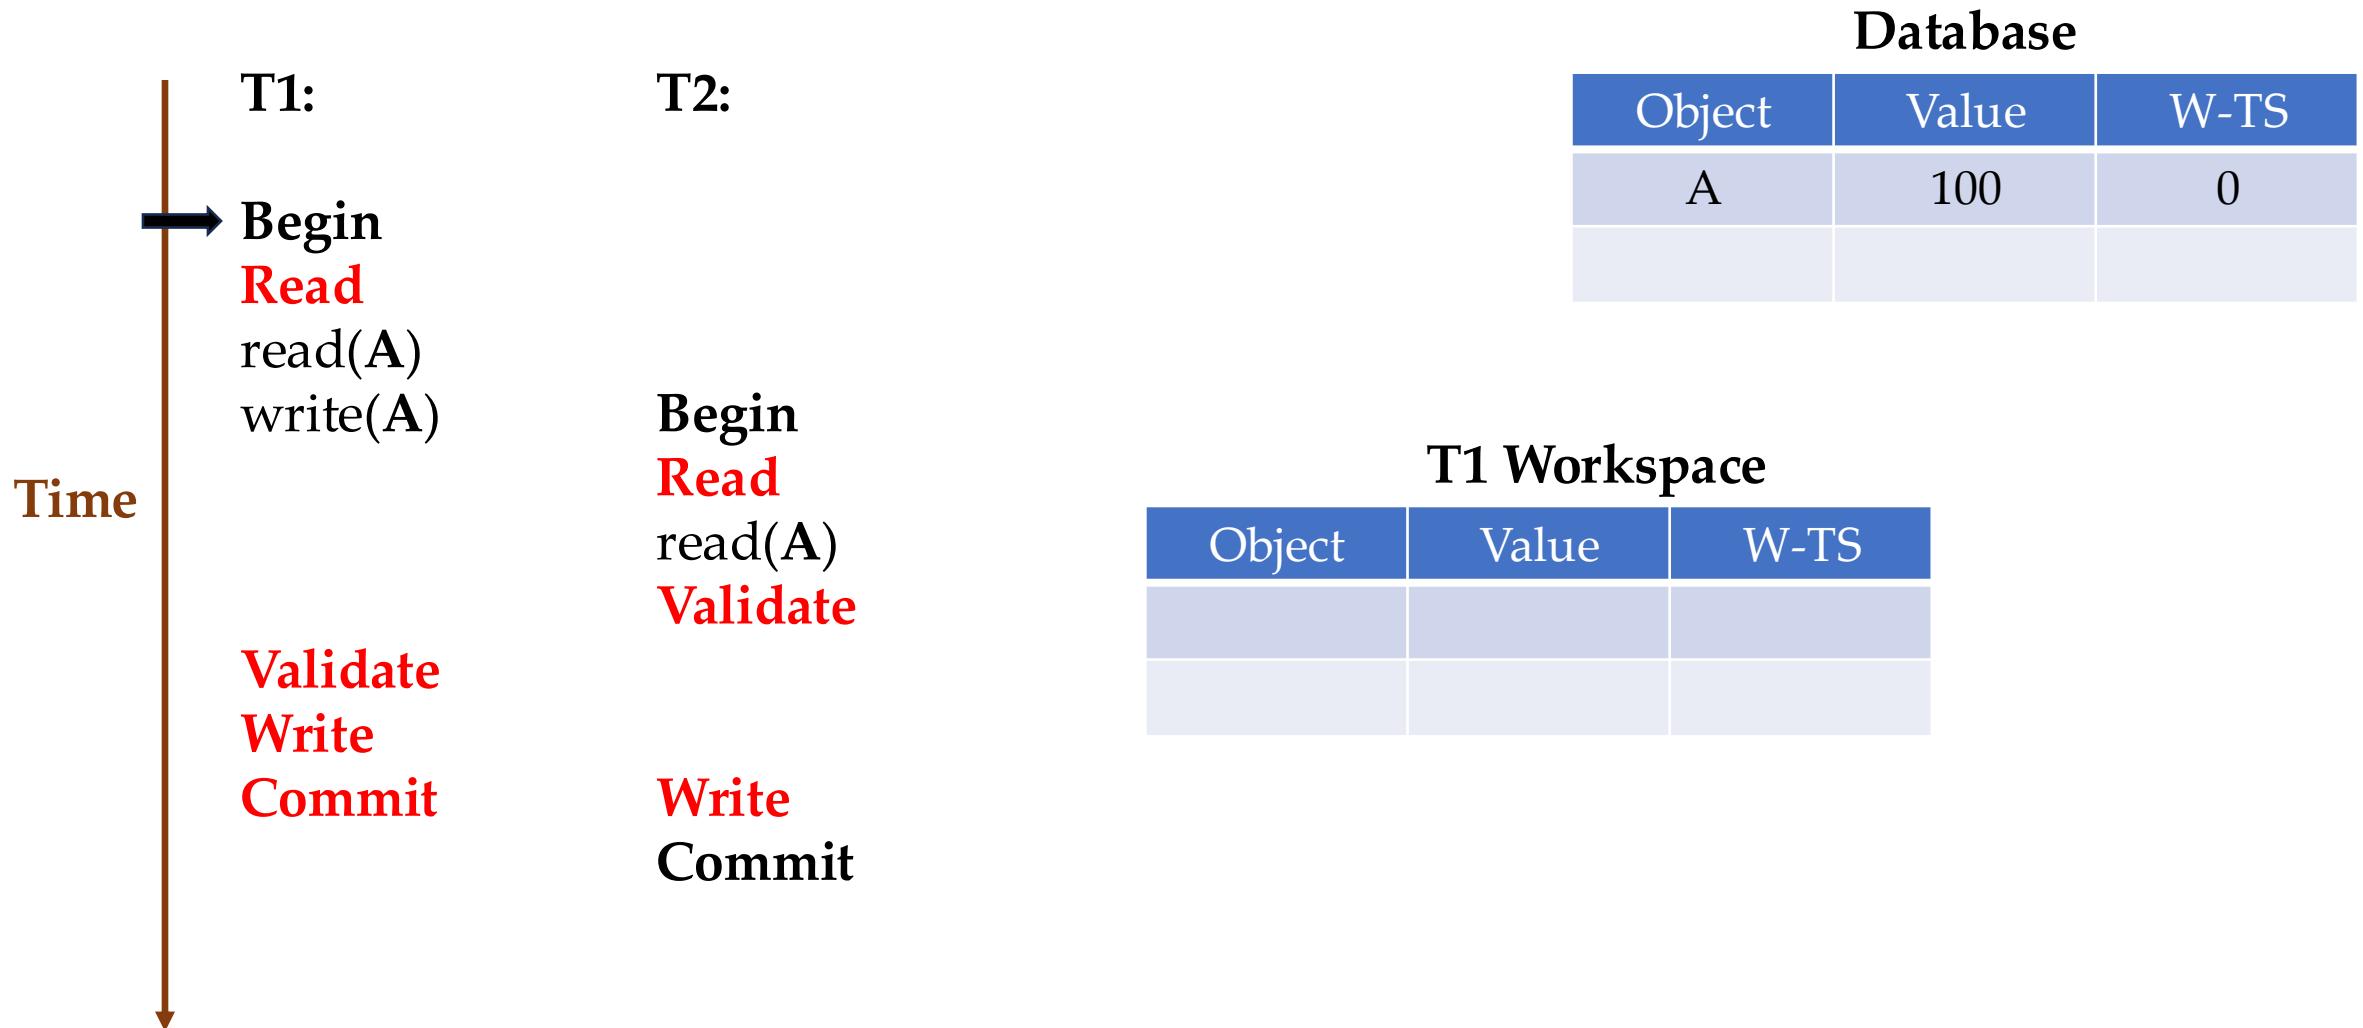
Forward Validation: Case II

How about this example?

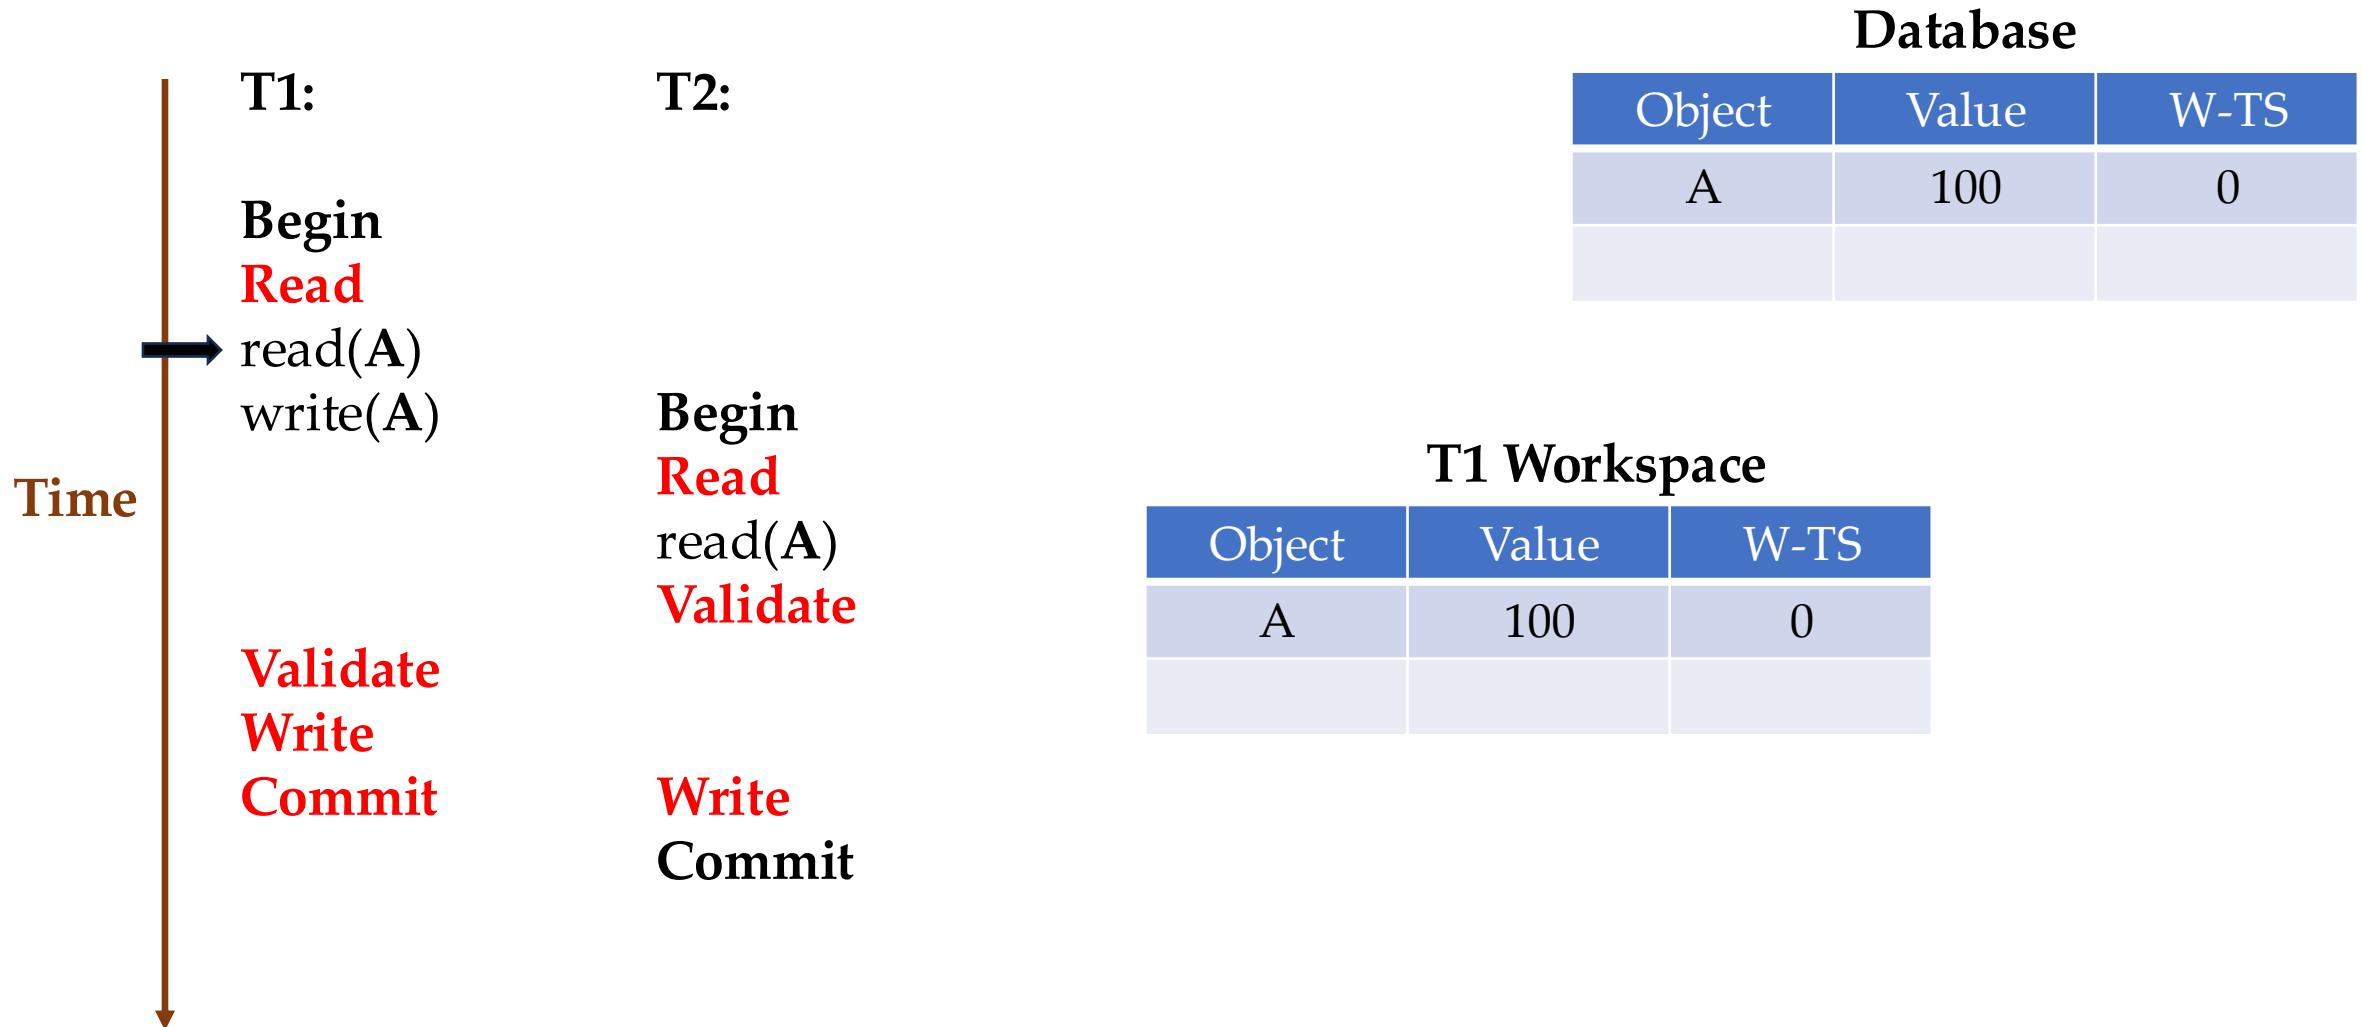


Database		
Object	Value	W-TS
A	100	0

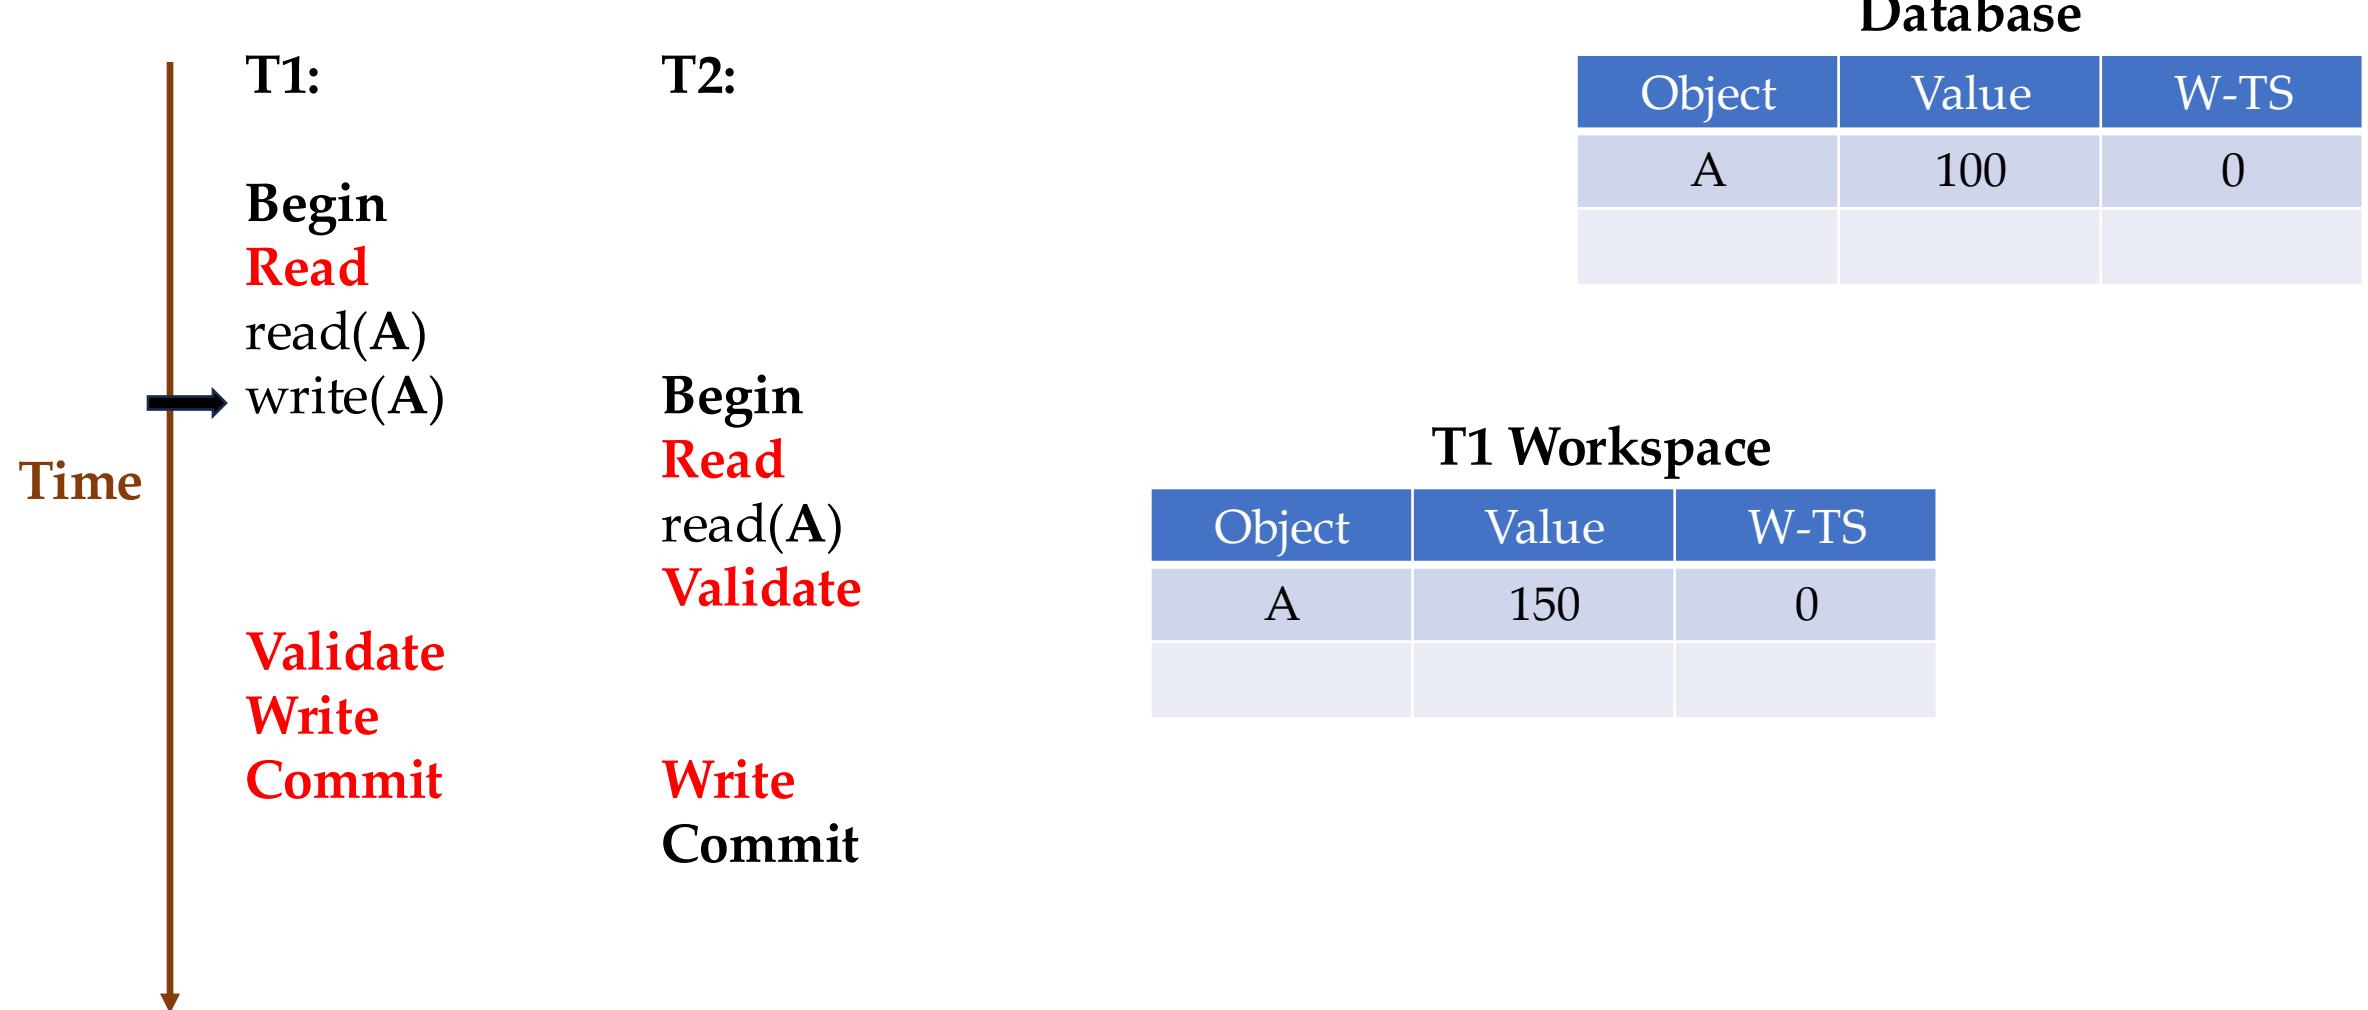
Forward Validation: Case II



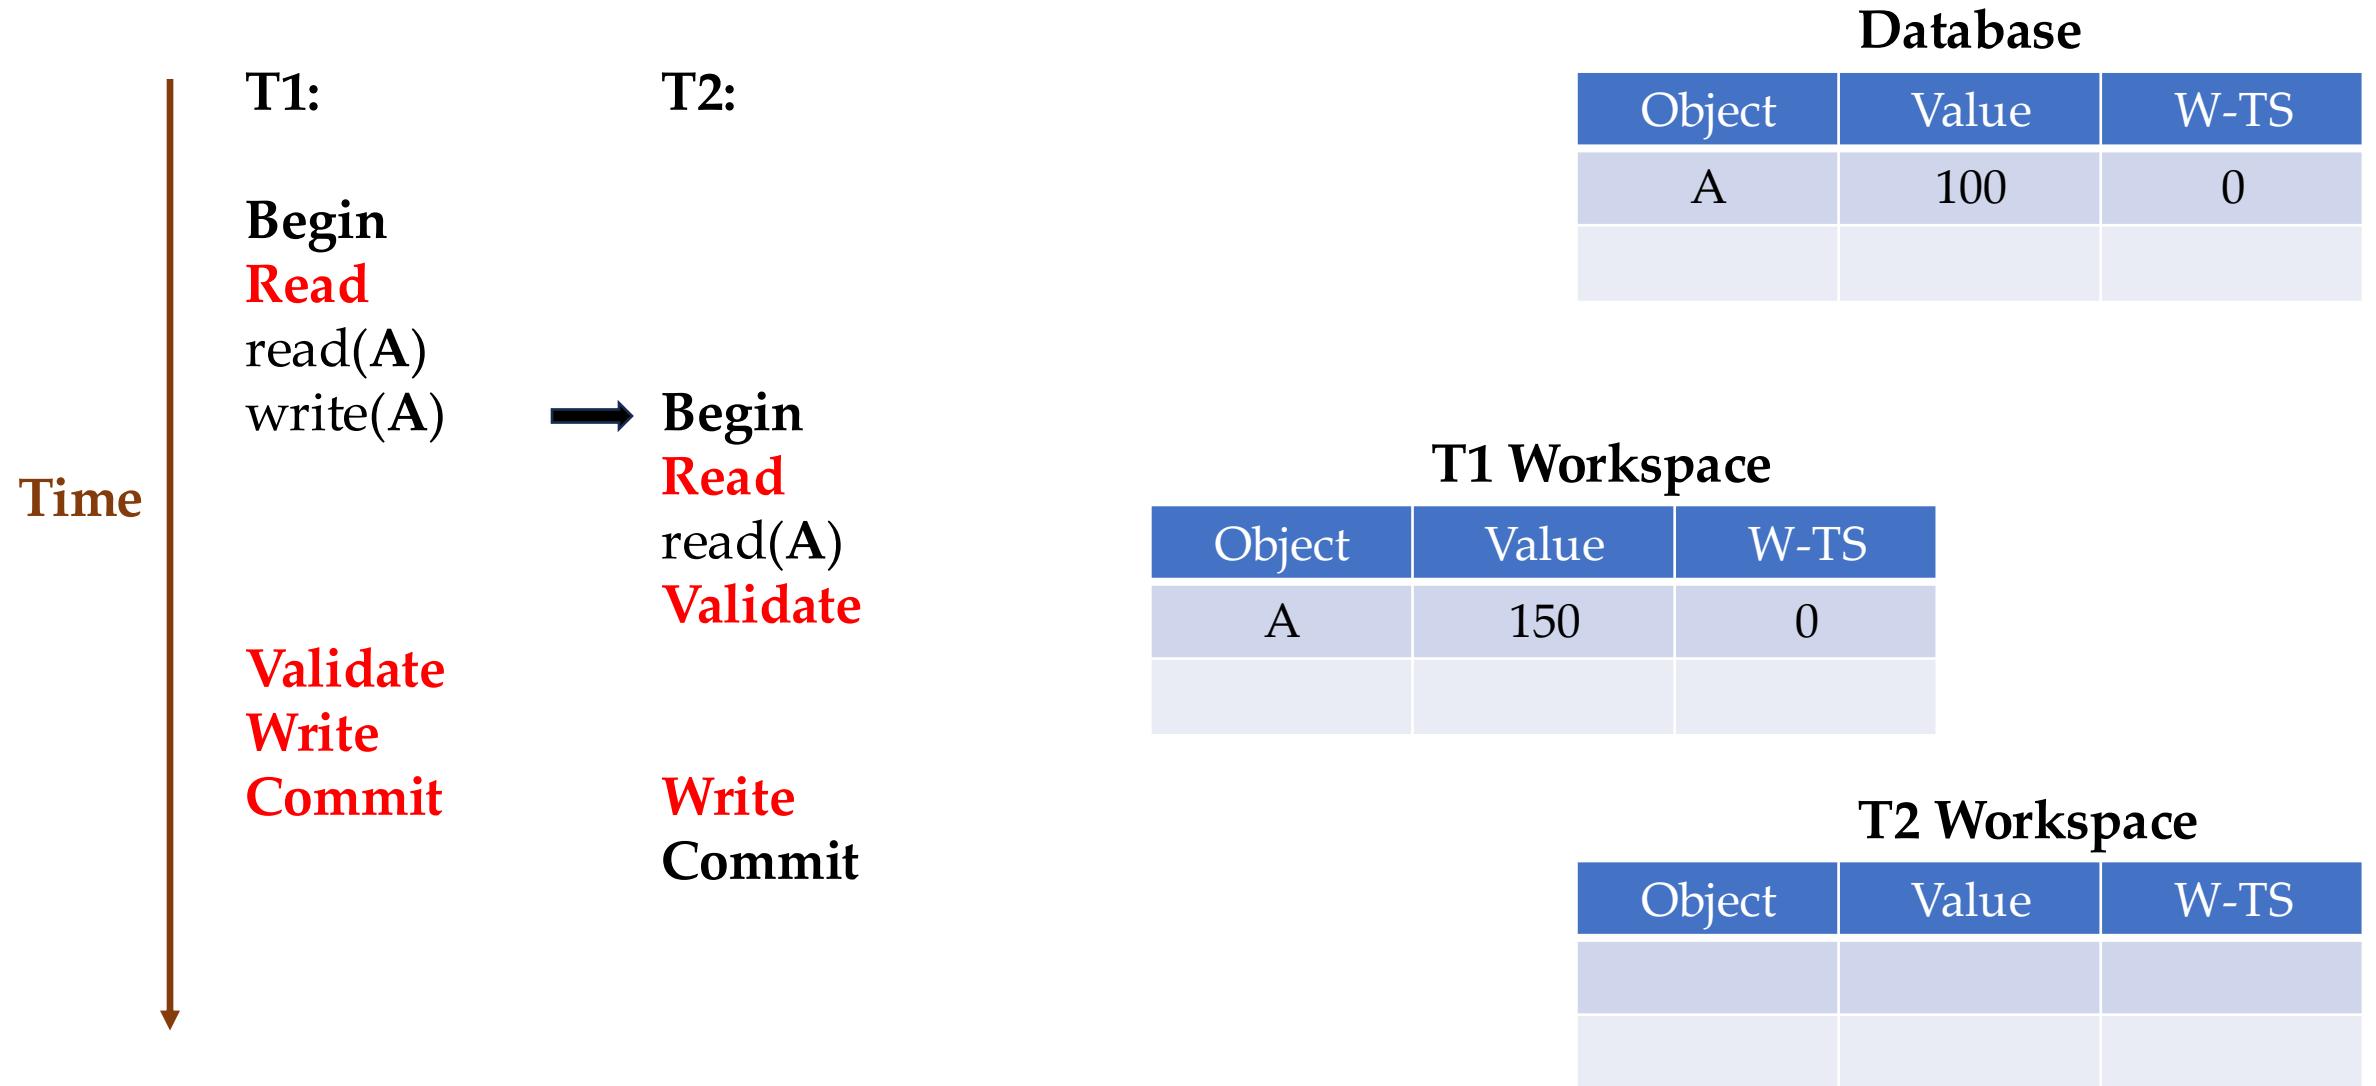
Forward Validation: Case II



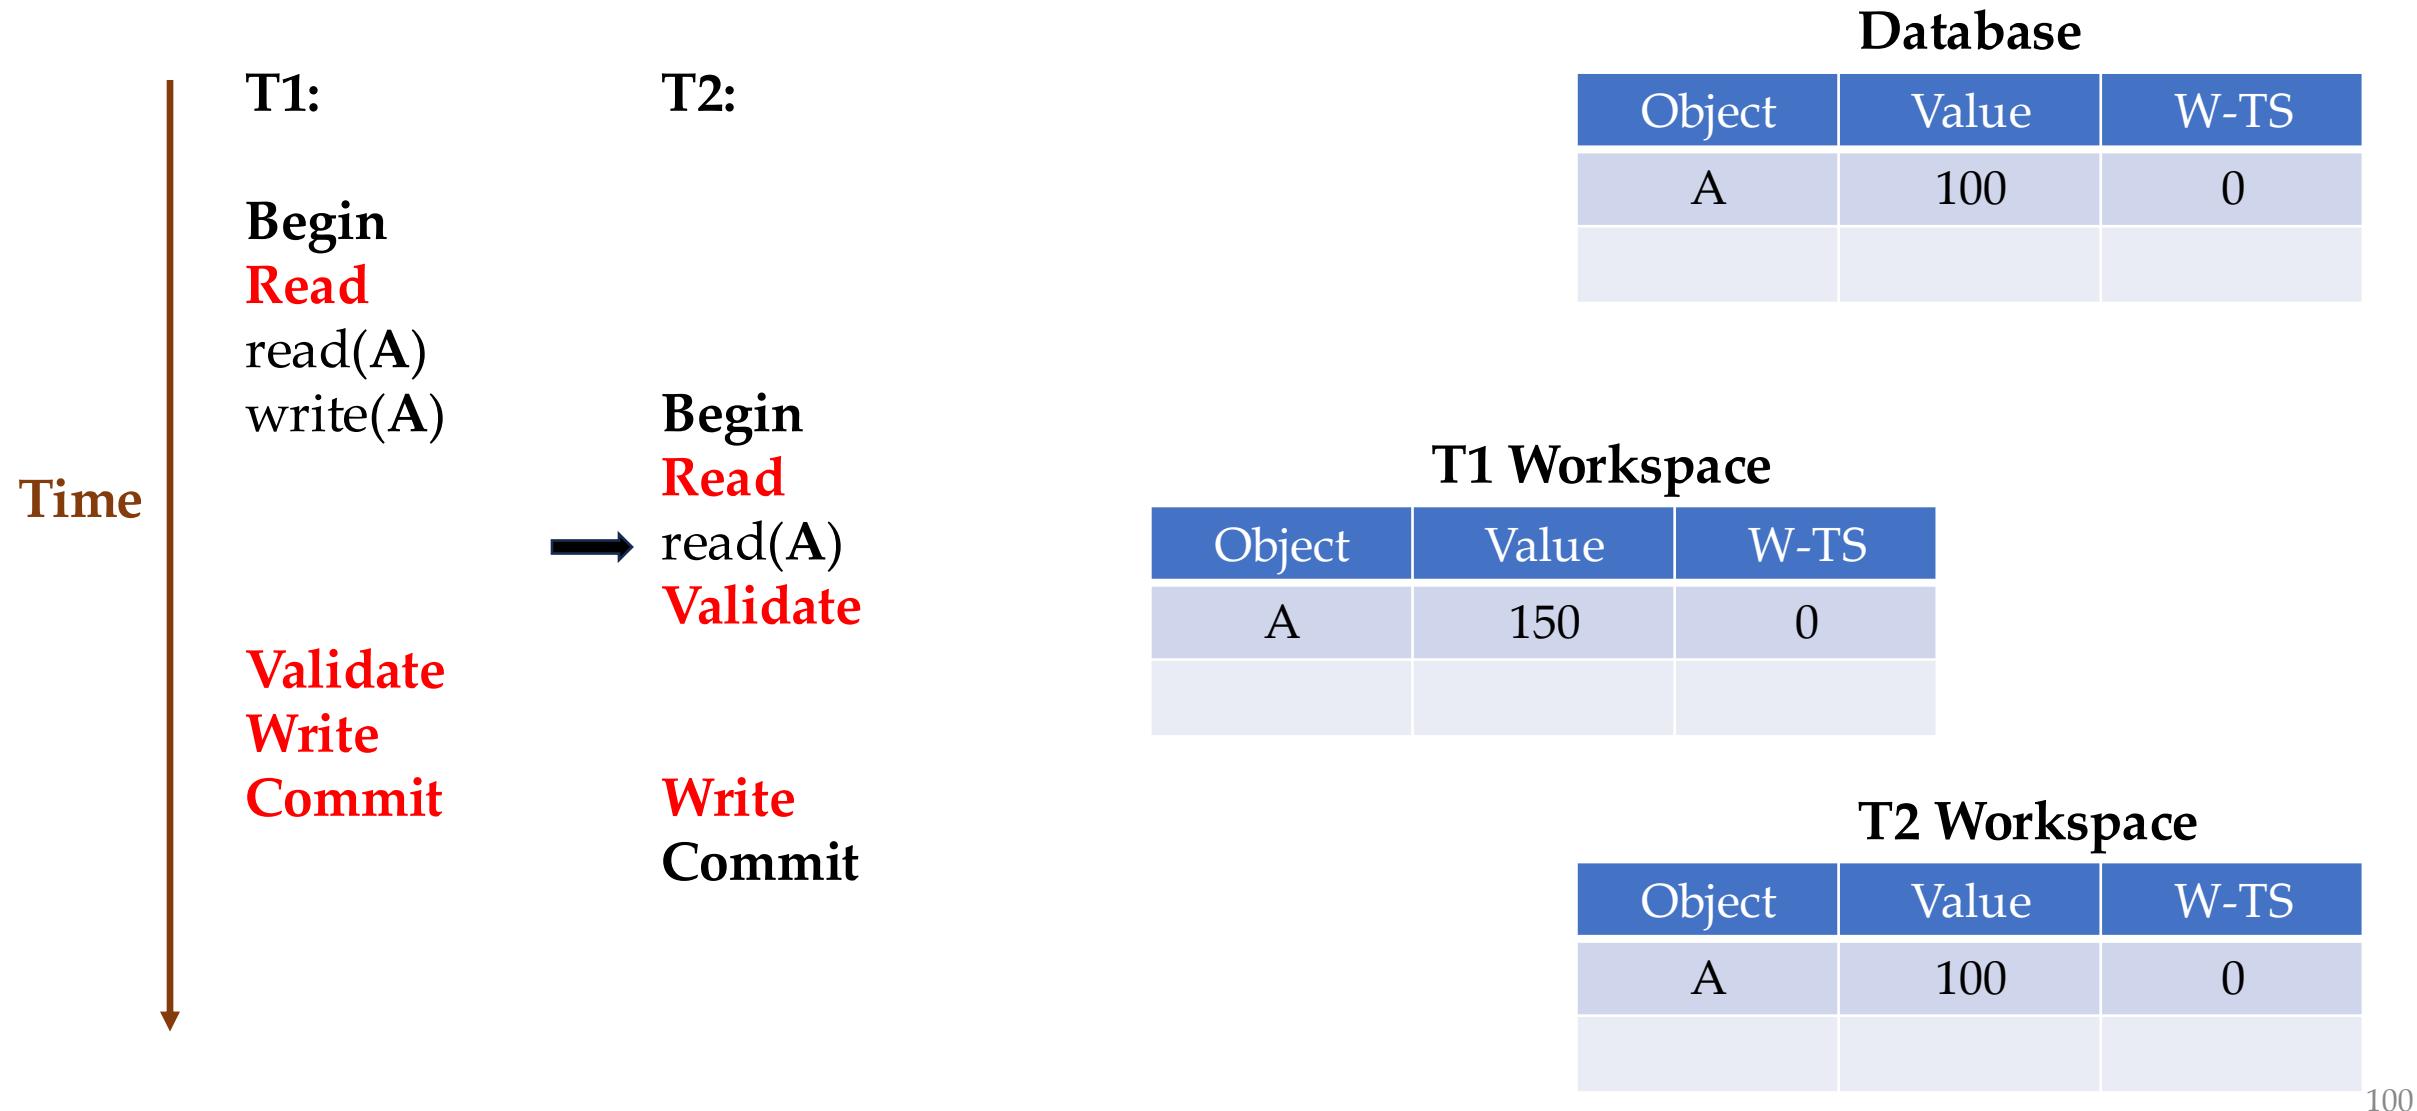
Forward Validation: Case II



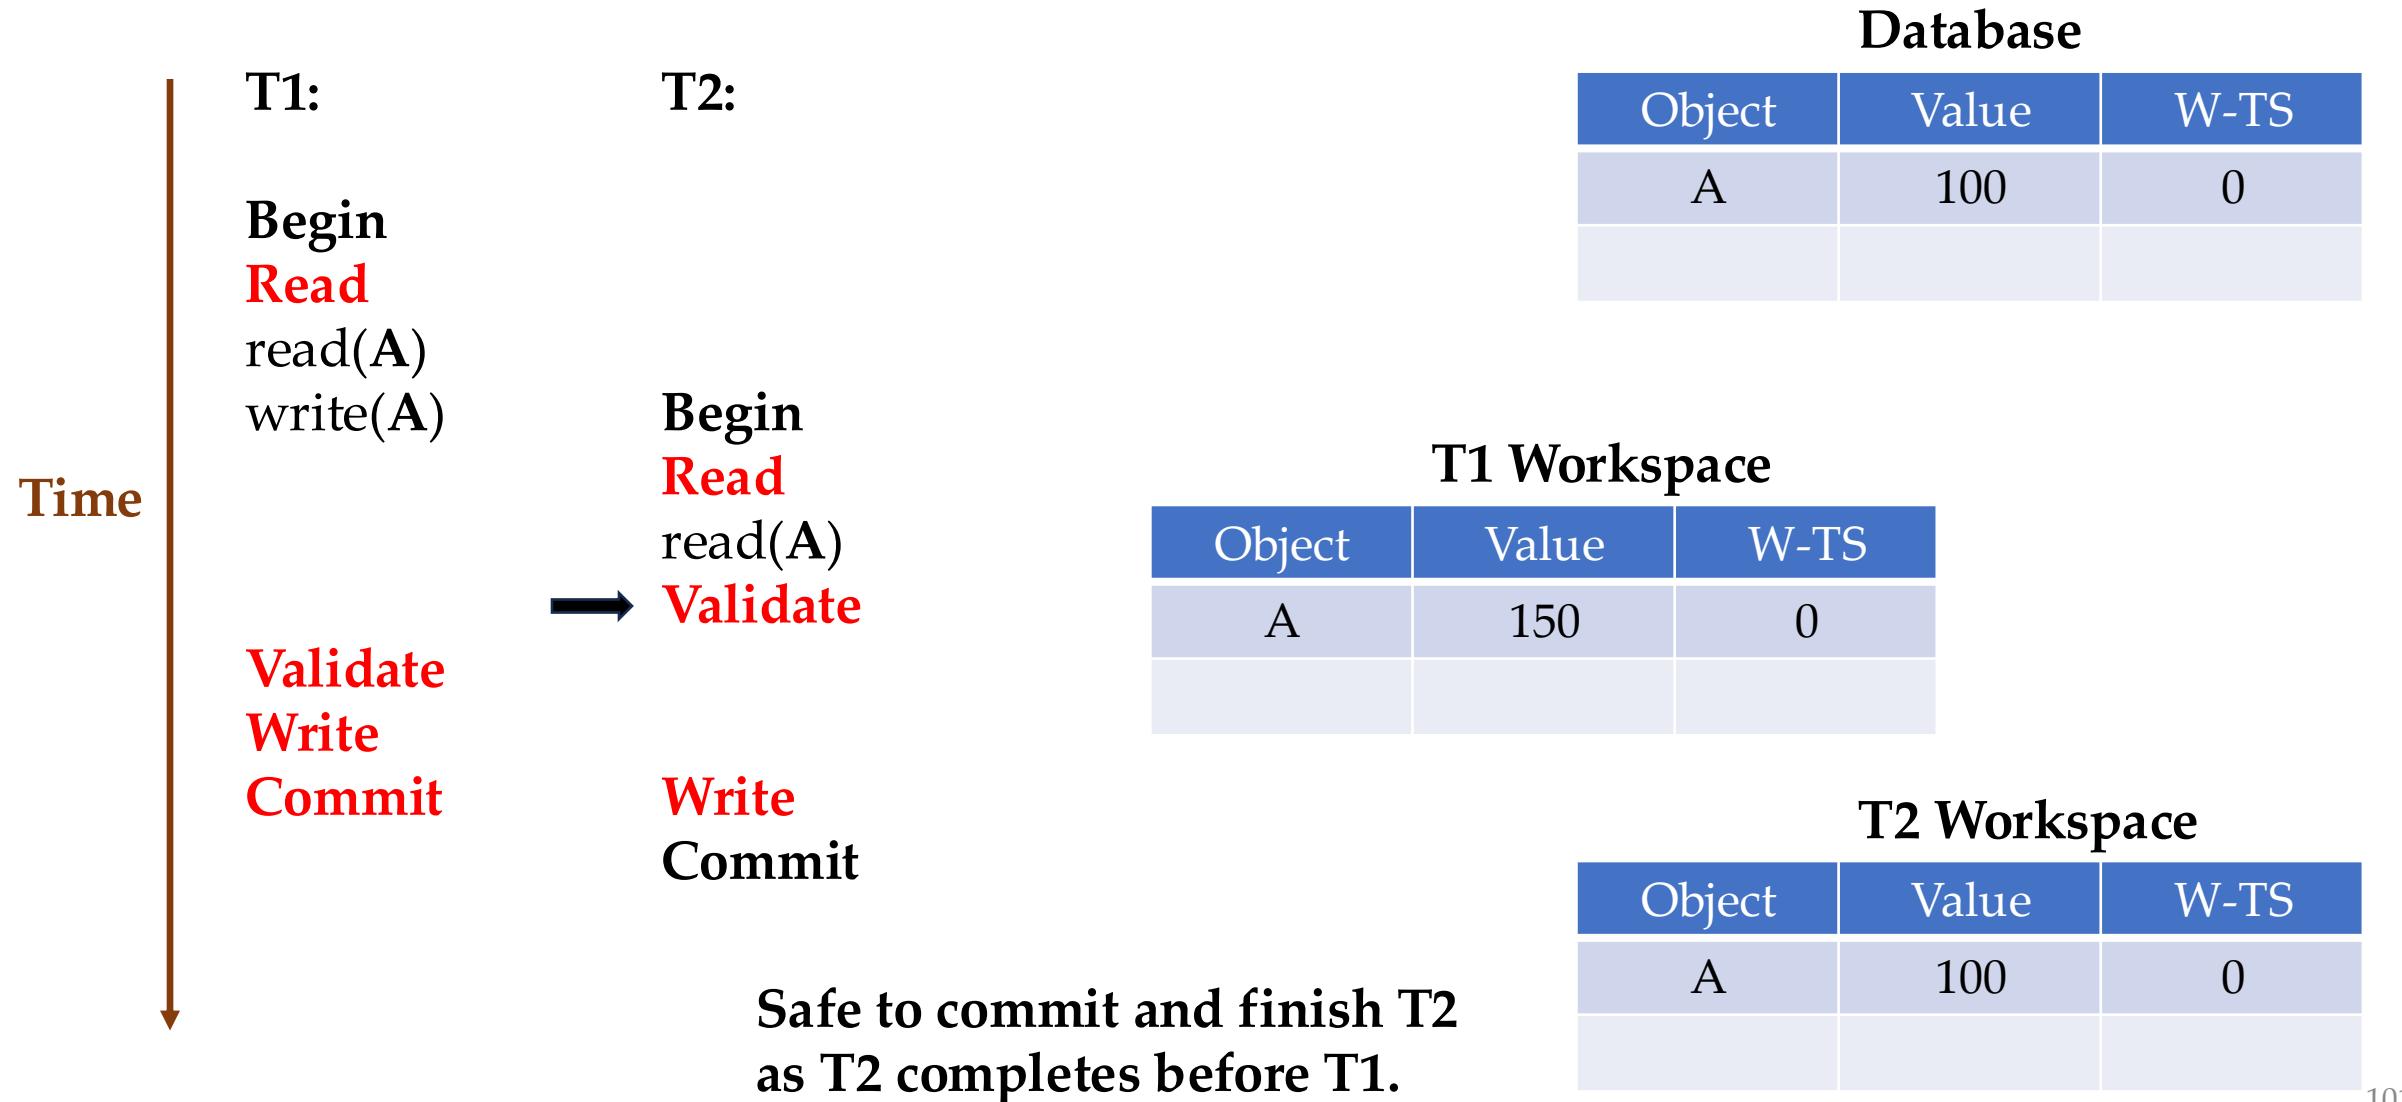
Forward Validation: Case II



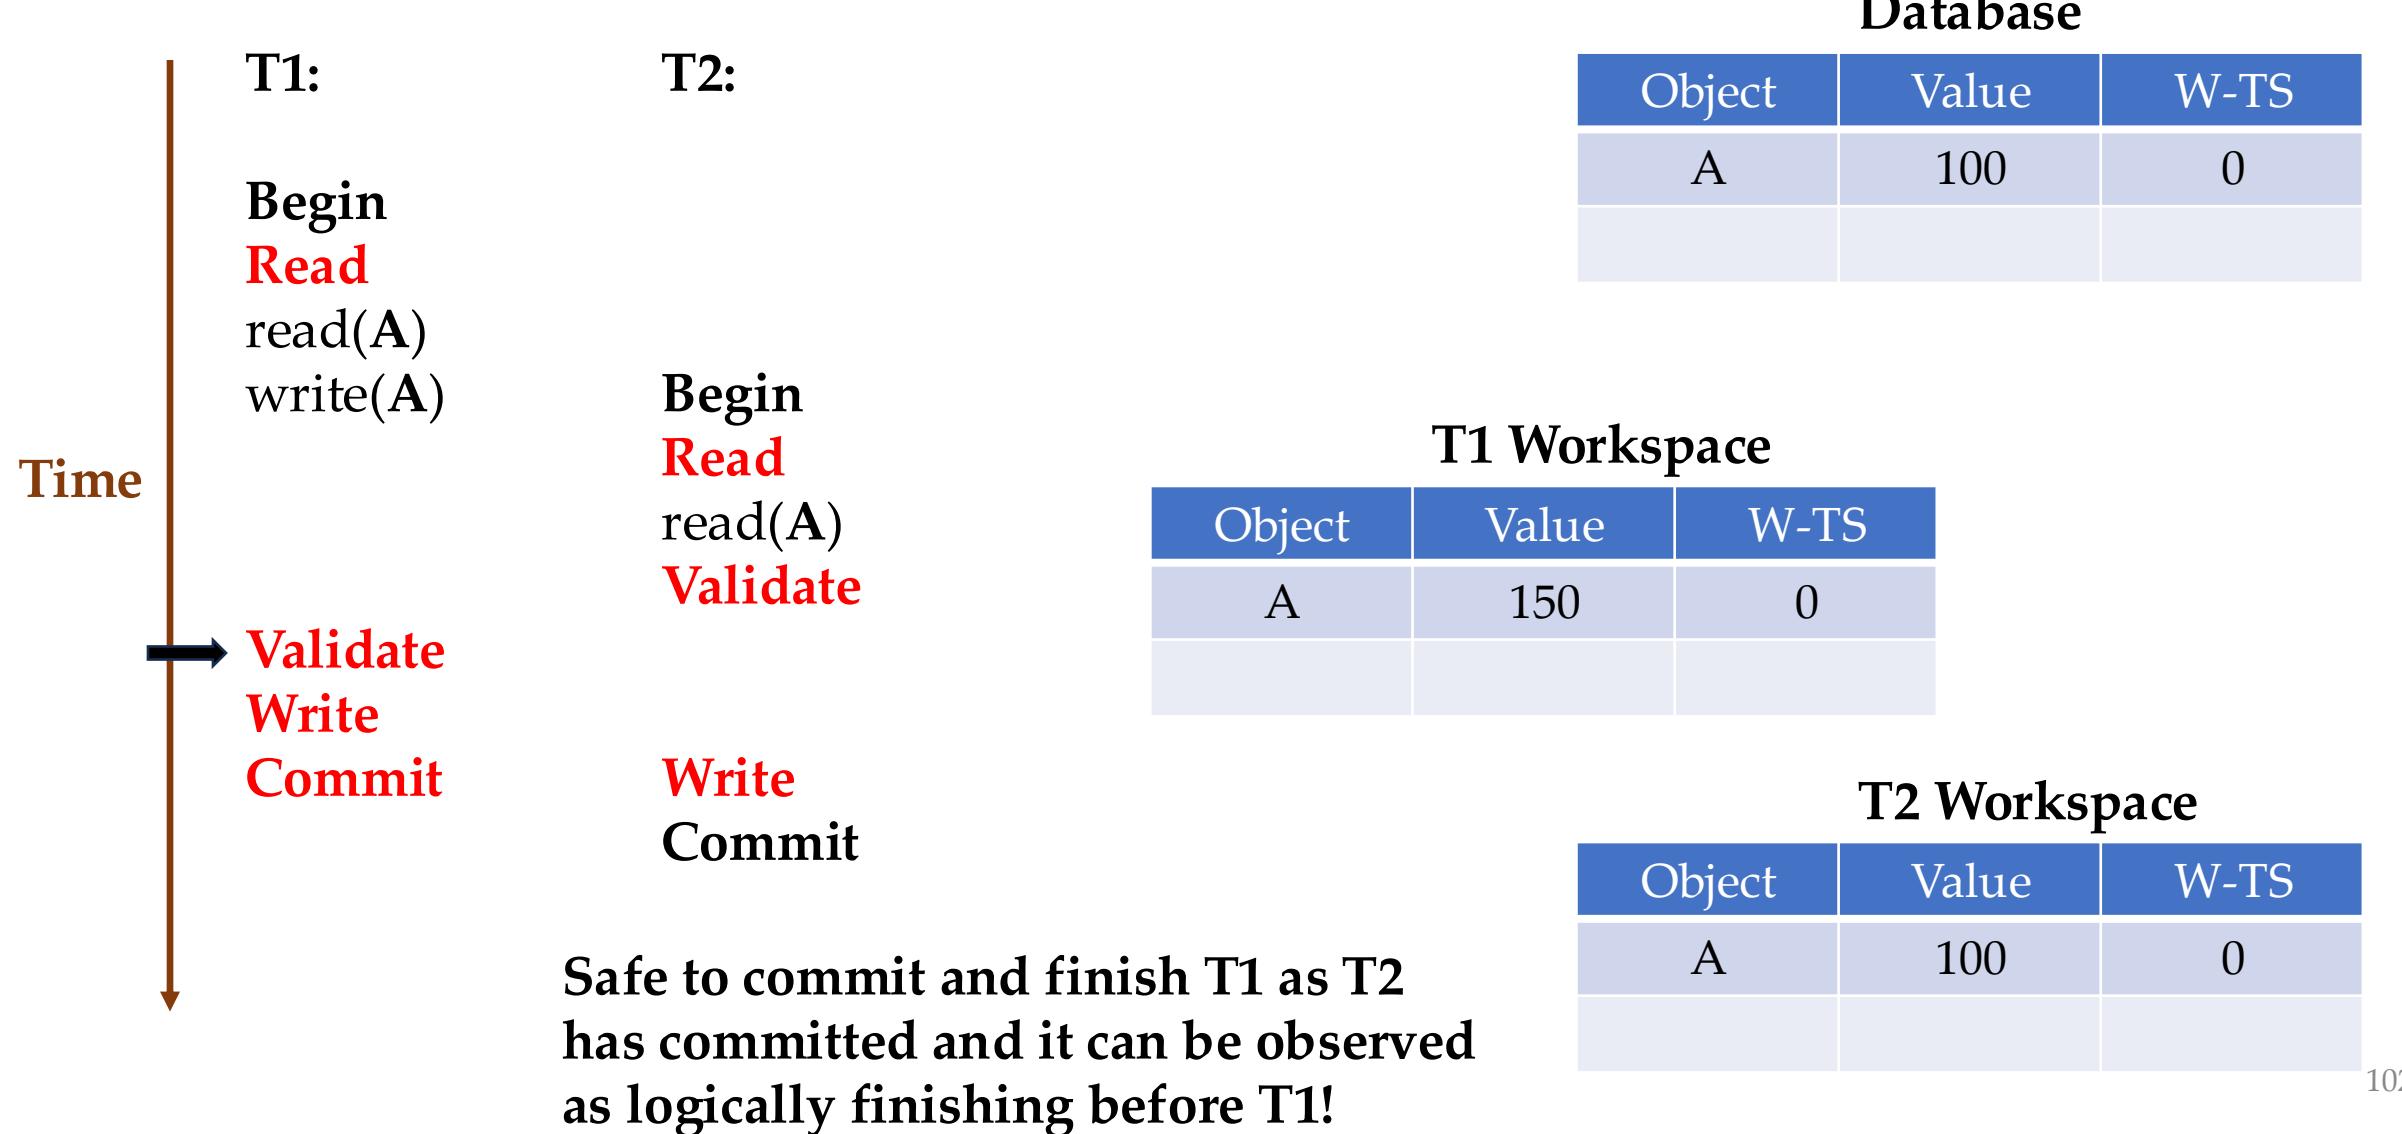
Forward Validation: Case II



Forward Validation: Case II



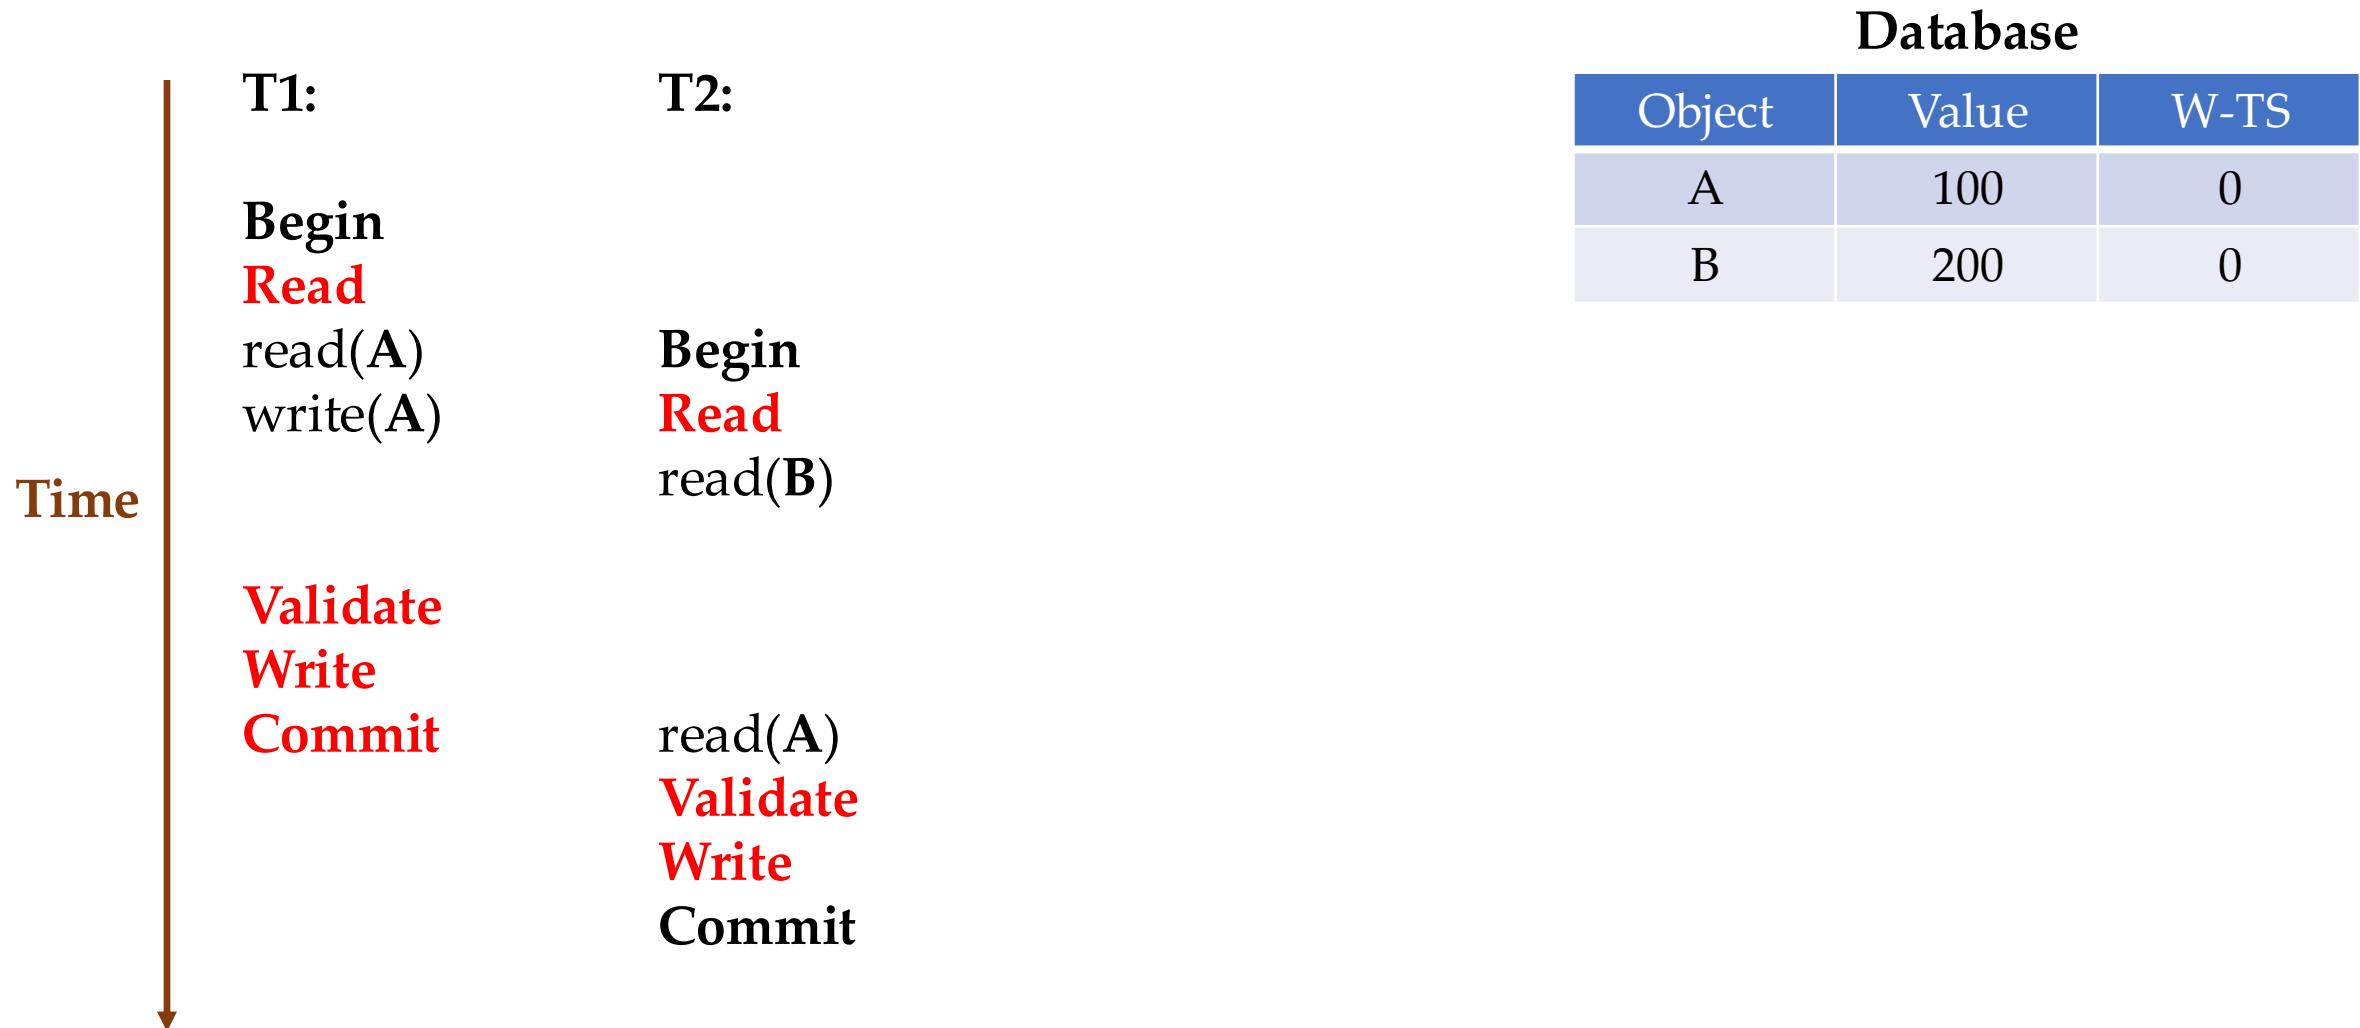
Forward Validation: Case II



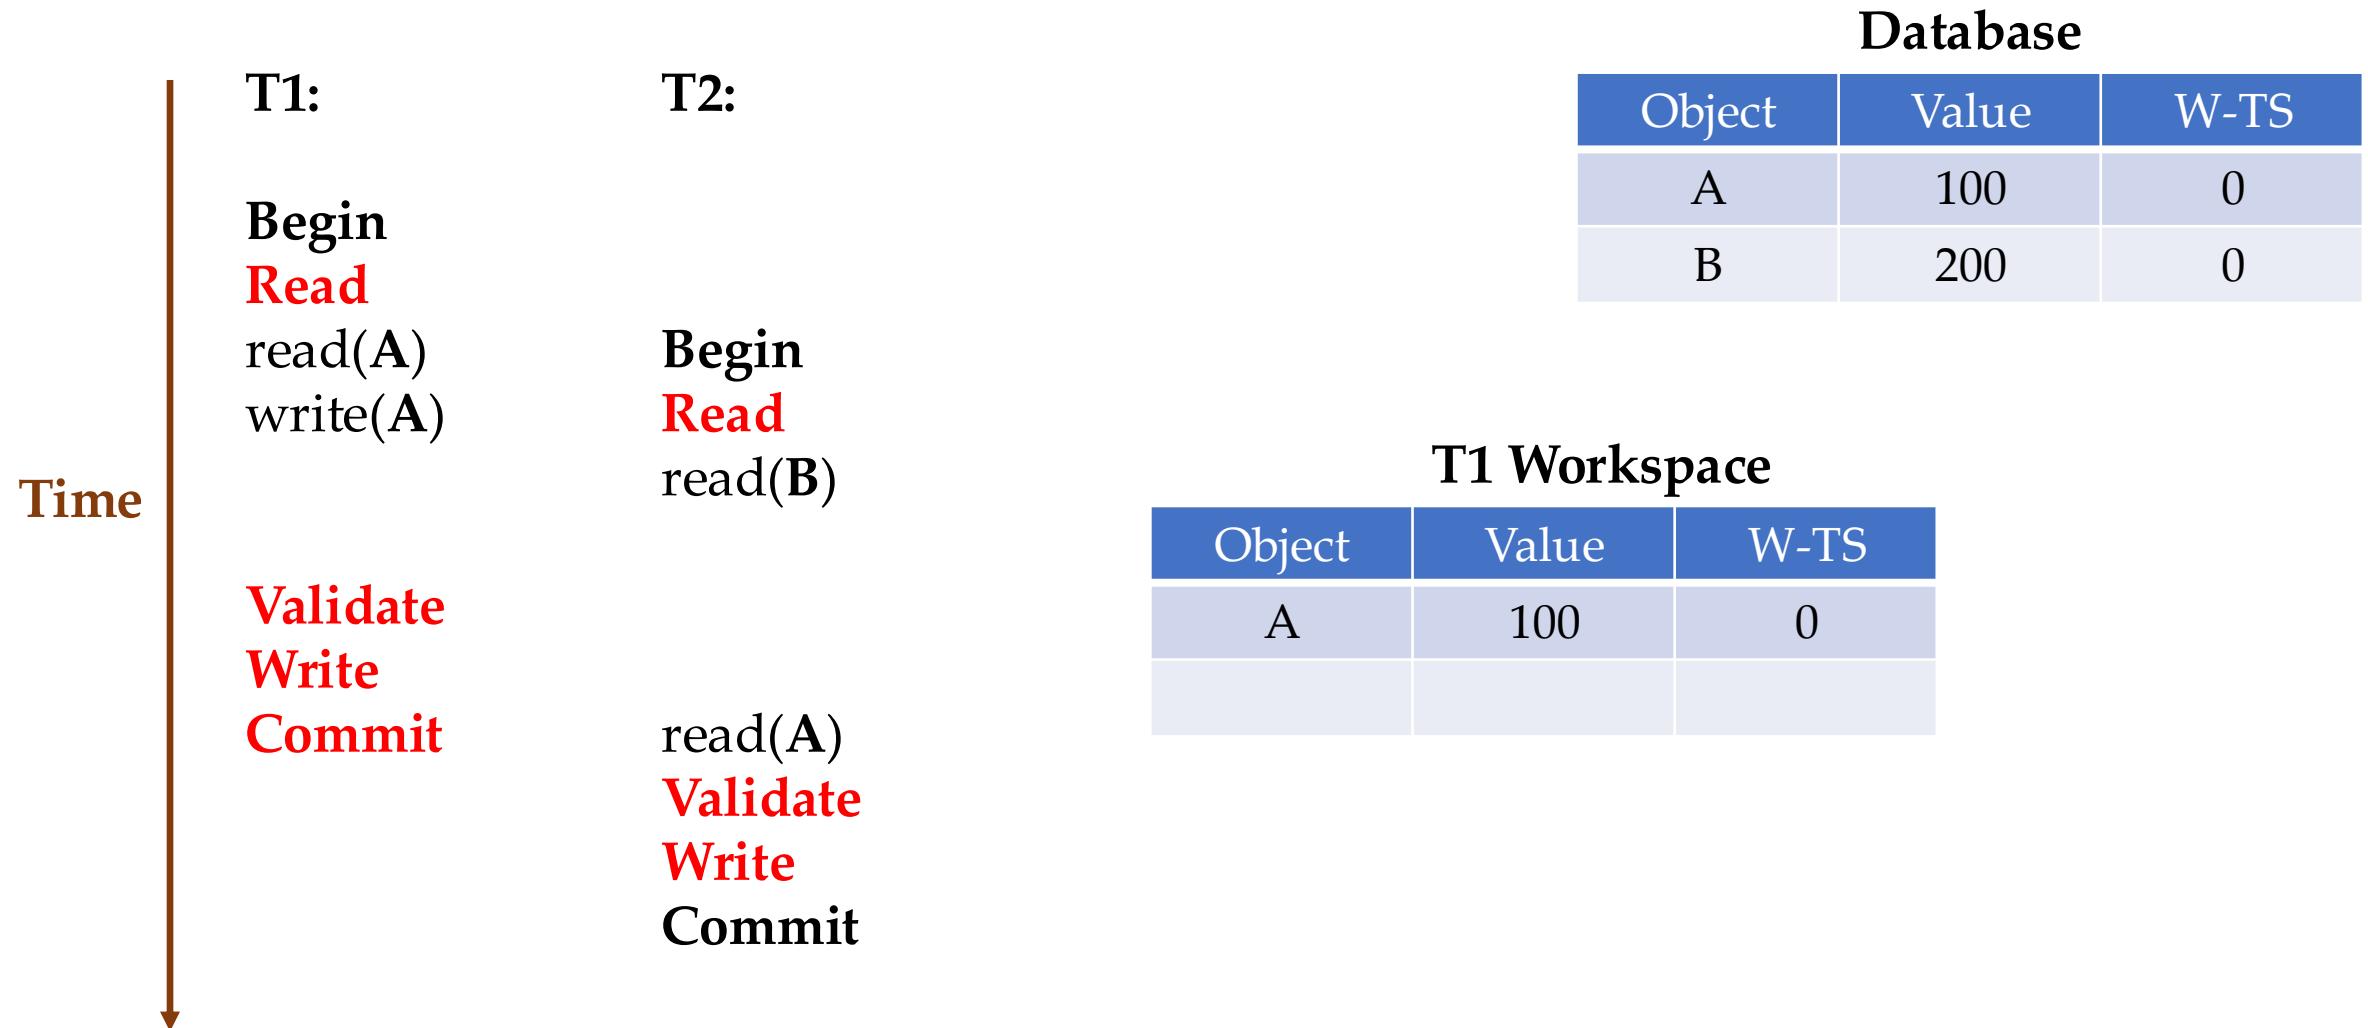
Forward Validation: Case III

- For two transactions T1 and T2, say T1 is at the validation step ($T1 < T2$)
 - Check if T1 completes its Read phase before T2 completes its Read phase.
 - T1 should not modify any object read or written by T2.
 - $\text{WriteSet}(T1) \cap \text{ReadSet}(T2) = 0$
 - $\text{WriteSet}(T1) \cap \text{WriteSet}(T2) = 0$

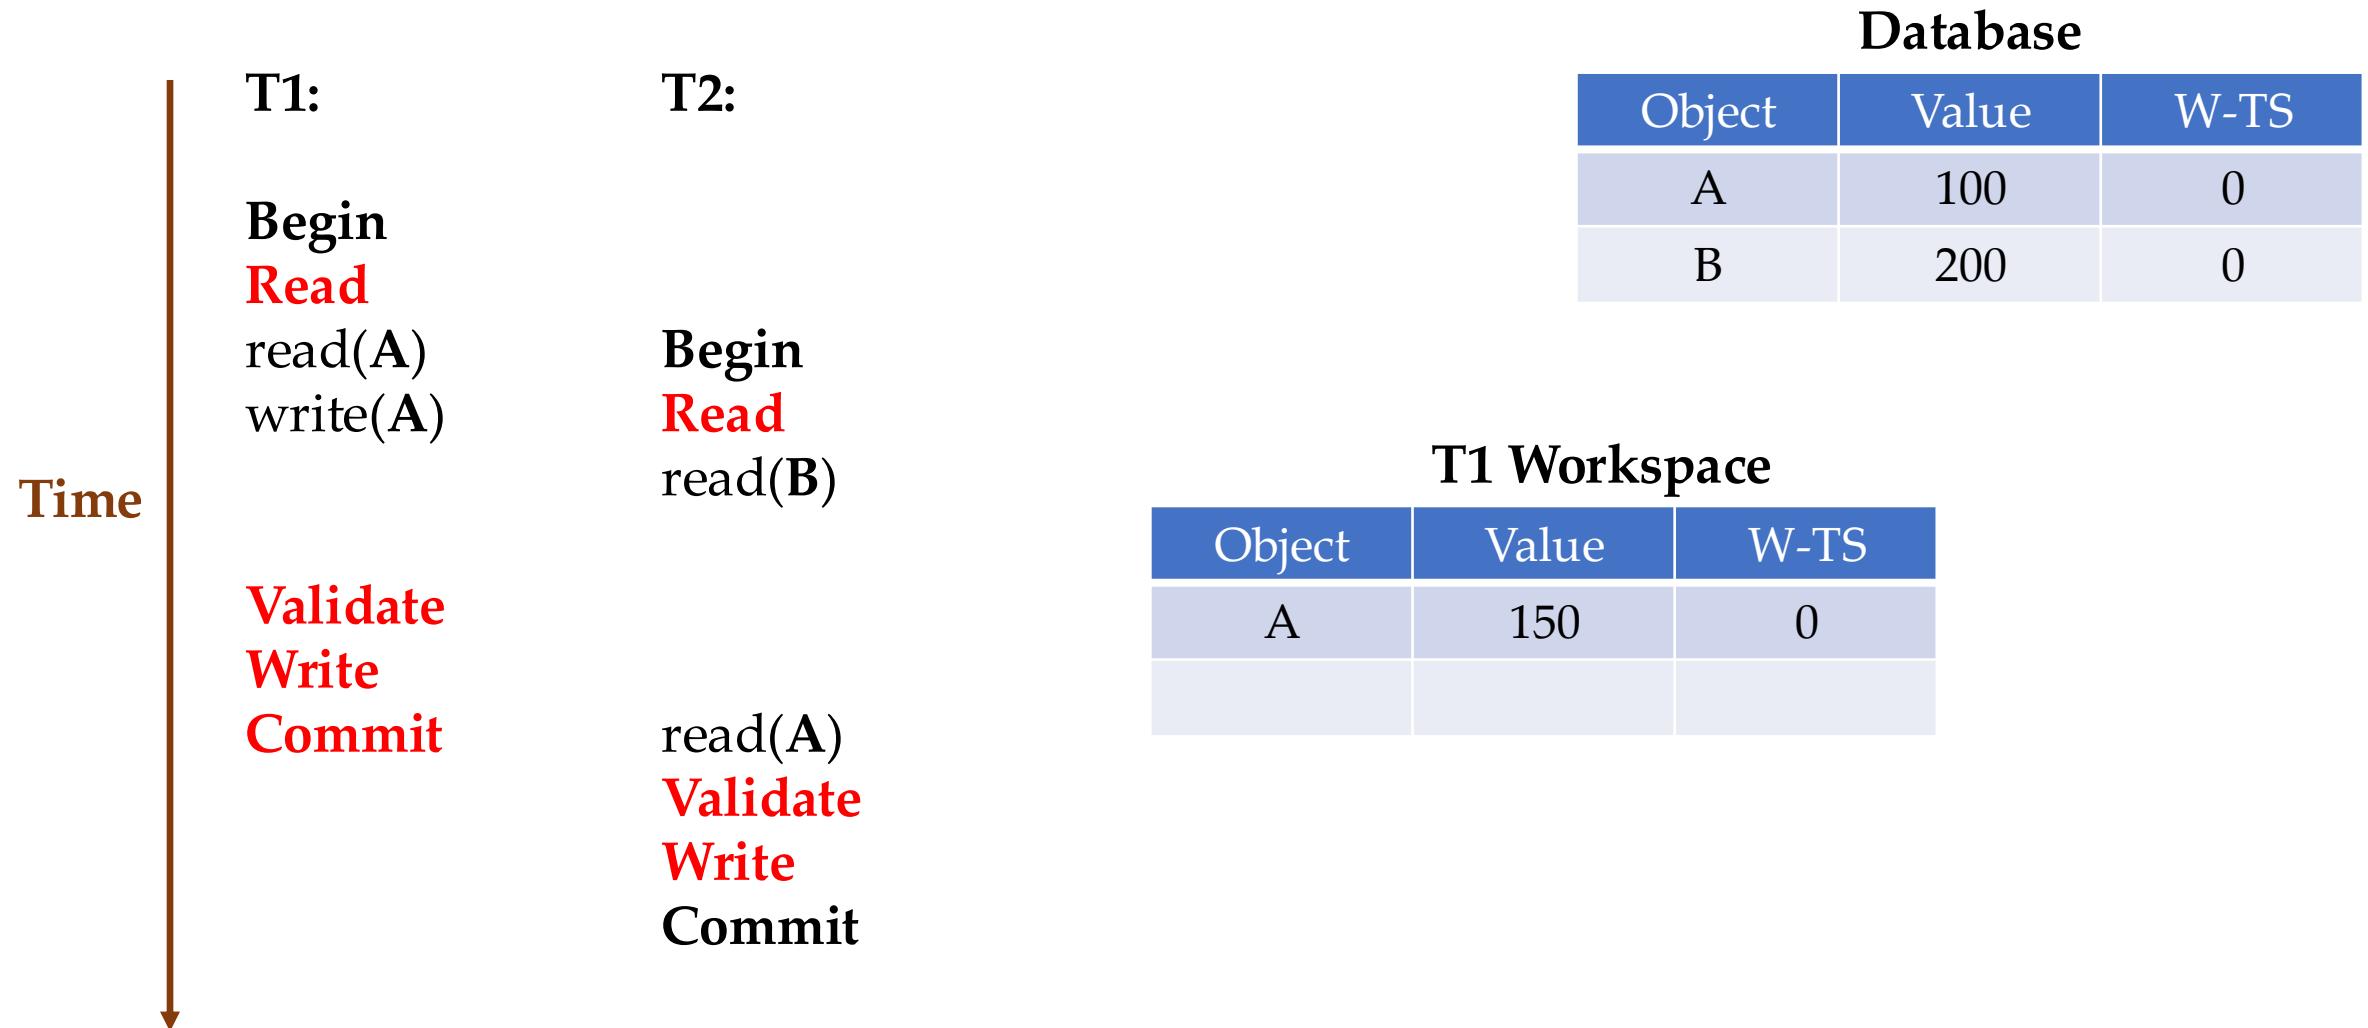
Forward Validation: Case III



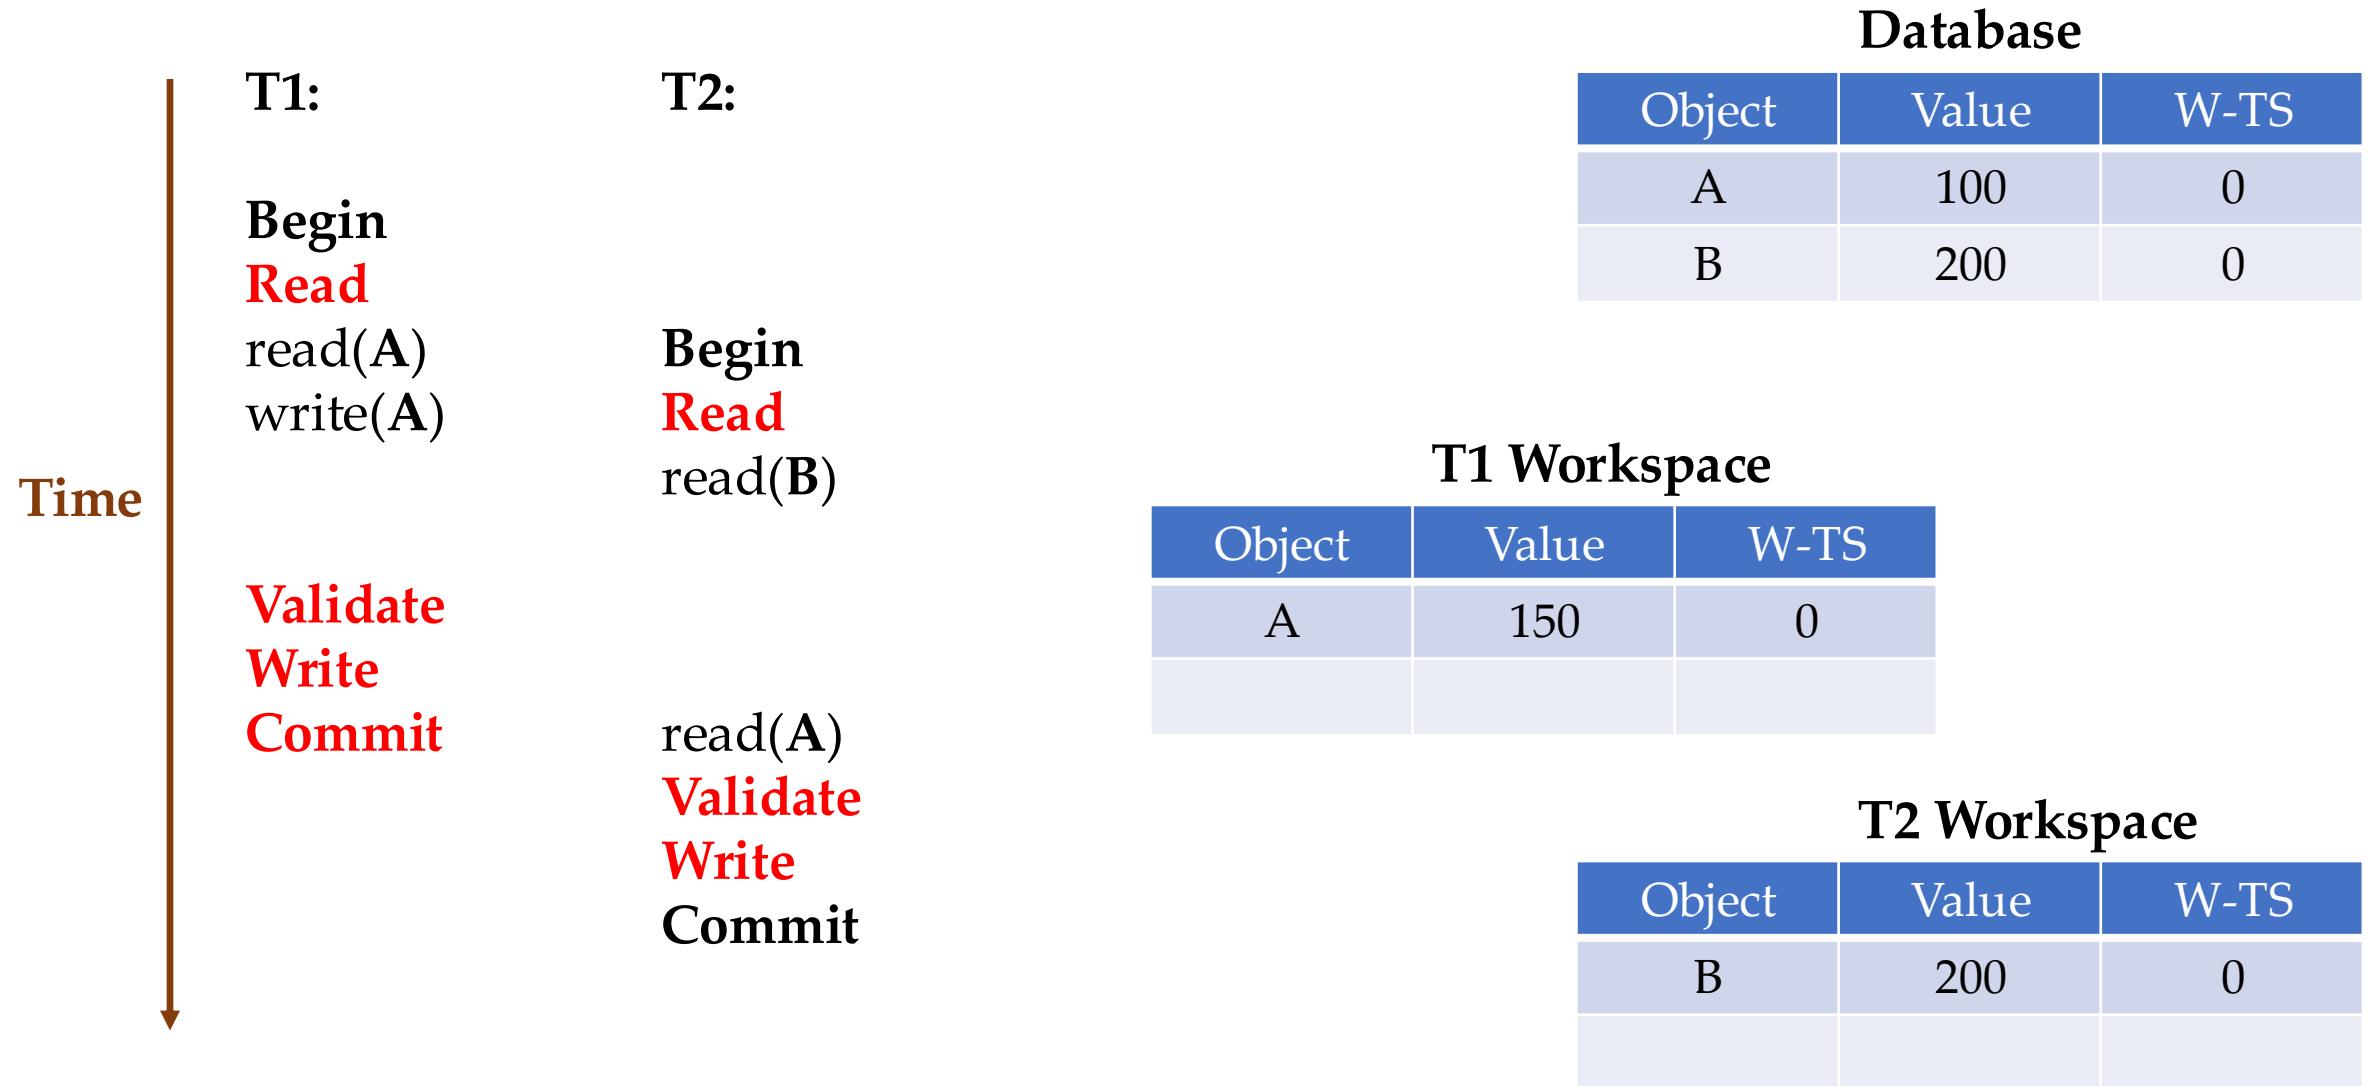
Forward Validation: Case III



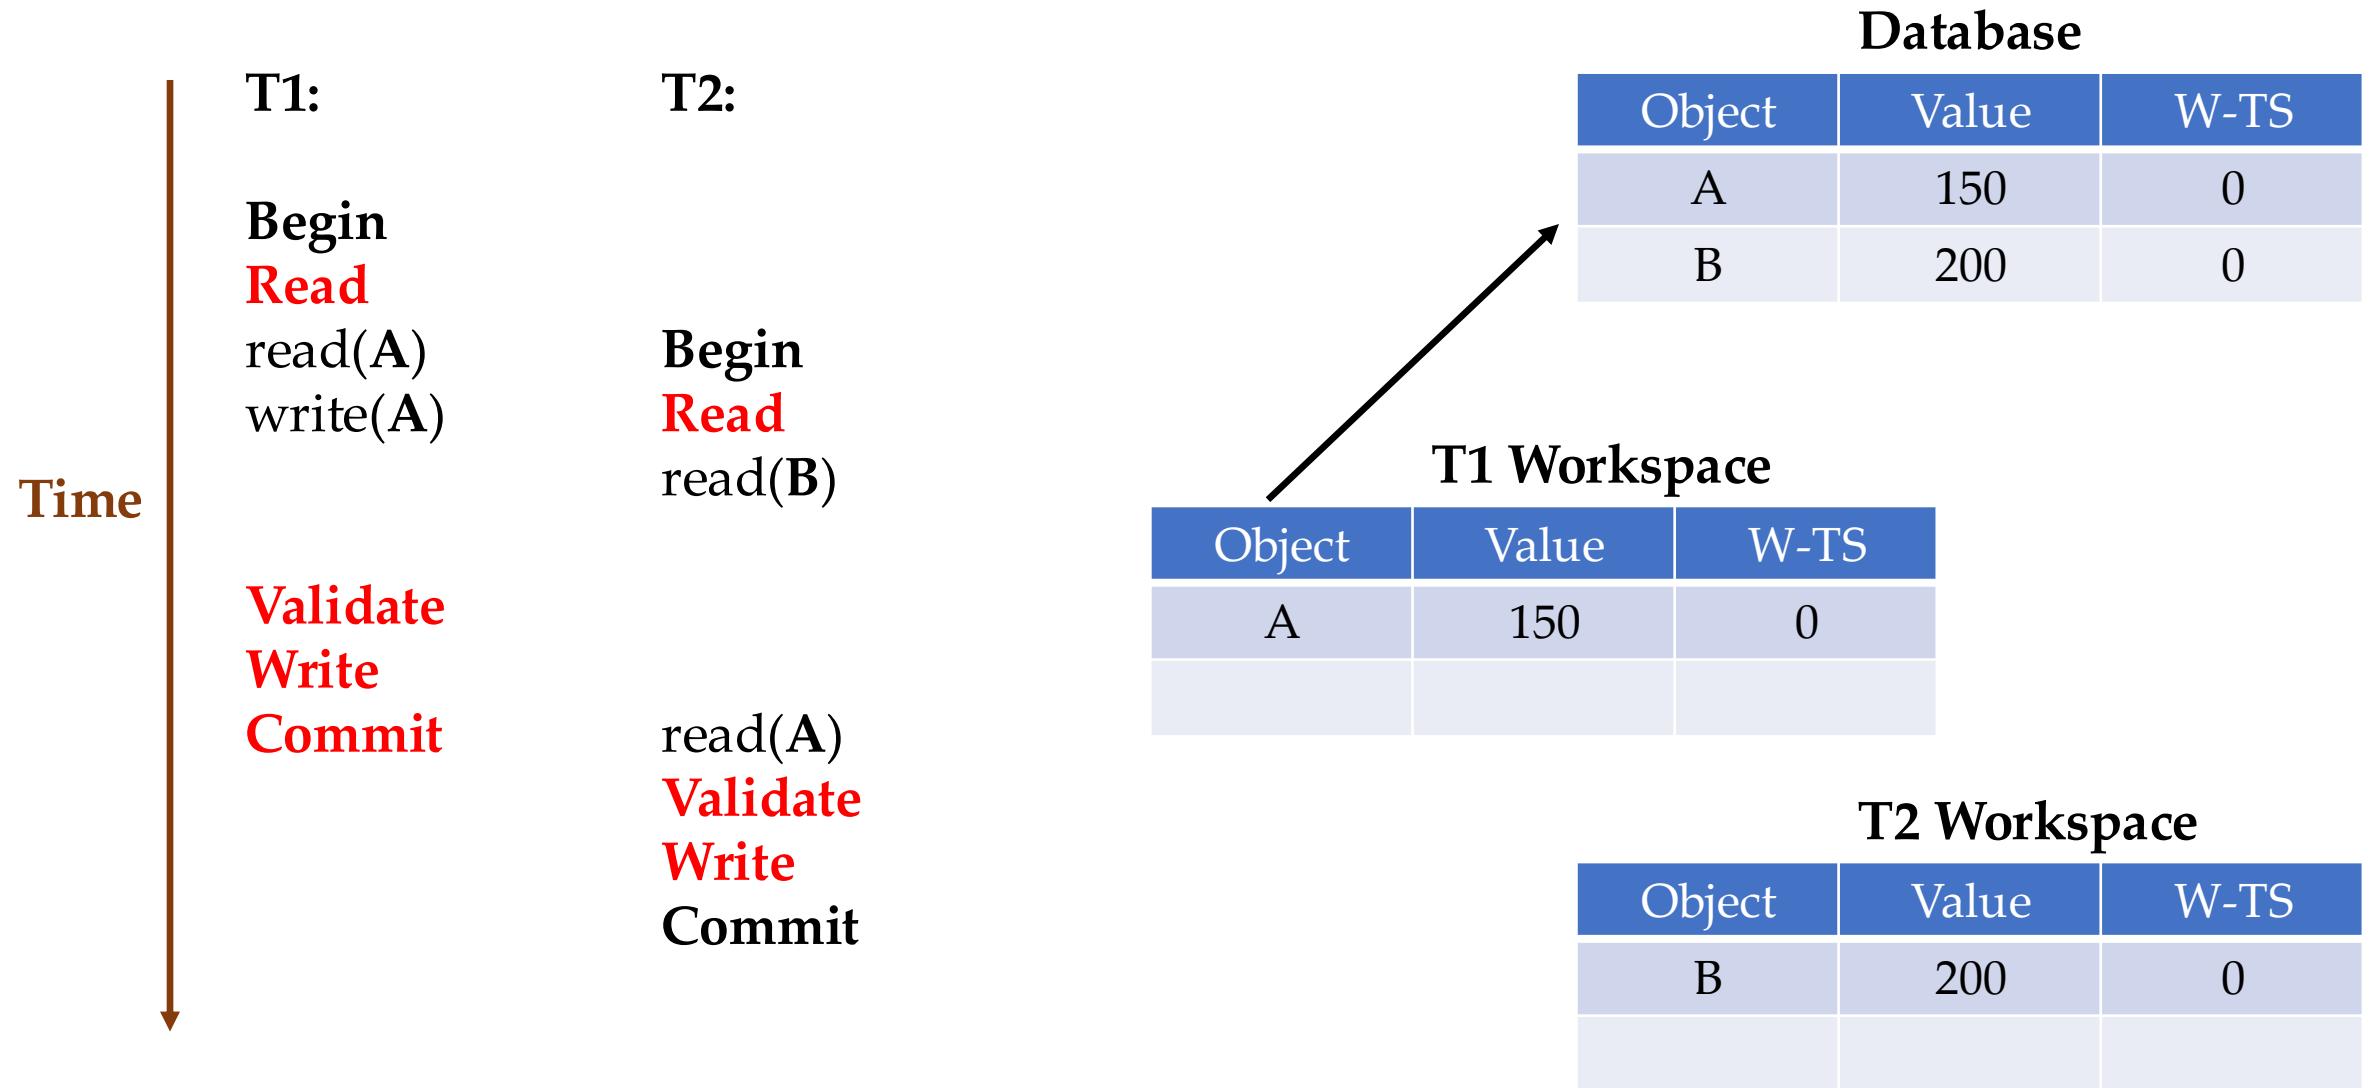
Forward Validation: Case III



Forward Validation: Case III



Forward Validation: Case III



Forward Validation: Case III

Time ↓

T1:

Begin
Read
read(A)
write(A)

Validate
Write
Commit

T2:

Begin
Read
read(B)

read(A)
Validate
Write
Commit

Database		
Object	Value	W-TS
A	150	0
B	200	0

→ T1 Workspace

Object	Value	W-TS
A	150	0

T2 Workspace

Object	Value	W-TS
B	200	0
A	150	0

Backward Validation

Backward Validation

- At the time of commit, each transaction checks if it conflicts with other **already committed transactions** (transactions which were concurrent and have committed).
- Each going to commit transaction (at the validation step), checks the timestamps and read/write sets of other committed transactions.
- There are three specific cases to satisfy:

OCC: Write Phase

- Propagate the changes in the transaction's private workspace (write set) to the database.
- The idea is to make the transaction's write-set visible to other transactions.
- **Serial Commits:** Use a global lock to limit a single transaction to be in the Validation/Write phases at a time.

OCC Disadvantages

OCC Disadvantages

- There is an overhead of copying data to private workspace.
 - More data to copy, more expensive!
- Validation/Write phase creates bottlenecks due to locking.
- Aborting a transaction is more expensive in OCC than in 2PL because it occurs after a transaction has already executed.