

Database Processing

CS 451 / 551

Lecture 14:
Isolation Levels and
Multi-Version Concurrency Control



Suyash Gupta

Assistant Professor

Distopia Labs and ONRG
Dept. of Computer Science

(E) suyash@uoregon.edu
(W) gupta-suyash.github.io



Assignment 3 is Out!
Deadline: Nov 30, 2025 at 11:59pm

Presentation Slots are Out.

Final Exam: Dec 8, 2025 at 8-10am

**Syllabus → Main focus on course not covered in Midterm,
but you should understand indexes and storage.**

Last Class

- We discussed Timestamp Ordering and Forward validation.
- In **Forward Validation**, at the time of commit, each transaction checks for conflicts with ongoing transactions.

Backward Validation

Backward Validation

- At the time of commit, each transaction checks if it conflicts with other **already committed transactions** (transactions which were concurrent and have committed).
- Each going to commit transaction (at the validation step), checks the timestamps and read/write sets of other committed transactions.

OCC Disadvantages

OCC Disadvantages

- There is an overhead of copying data to private workspace.
 - More data to copy, more expensive!
- Validation/Write phase creates bottlenecks due to locking.
- Aborting a transaction is more expensive in OCC than in 2PL because it occurs after a transaction has already executed.

Queries Isolation

Time

T1:

Begin

```
select count(*) as cnt  
from cs_cmemployees  
where age > 30
```

```
select count(*) as cnt  
from cs_cmemployees  
where age > 30
```

Commit

T2:

Begin


```
insert into cs_employees  
values (anakin, 70, 500)
```

Commit


```
create table cs_employees  
(  
    name    varchar(20),  
    age     int,  
    salary  int  
);
```

Is there any problem with this schedule?


Queries Isolation



Queries Isolation



Queries Isolation



```
create table cs_employees  
(  
    name  varchar(20),  
    age   int,  
    salary int  
);
```

Queries Isolation

Time

T1:

Begin

```
select count(*) as cnt  
from cs_cmemployees  
where age > 30
```

→ **select count(*) as cnt
from cs_cmemployees
where age > 30**

Commit

T2:

Begin

```
insert into cs_employees  
values (anakin, 70, 500)
```

Commit

Now, output = 11

```
create table cs_employees  
( name varchar(20),  
  age int,  
  salary int  
);
```

Queries Isolation

Time

T1:

Begin

```
select count(*) as cnt  
from cs_cmemployees  
where age > 30
```

→ **select count(*) as cnt
from cs_cmemployees
where age > 30**

Commit

T2:

Begin

```
insert into cs_employees  
values (anakin, 70, 500)
```

Commit

```
create table cs_employees  
(  
    name  varchar(20),  
    age   int,  
    salary int  
);
```

The output of the two queries changed!

Phantom Problem

Time

T1:

Begin

```
select count(*) as cnt  
from cs_cmemployees  
where age > 30
```

→ **select count(*) as cnt
from cs_cmemployees
where age > 30**

Commit

T2:

Begin

```
insert into cs_employees  
values (anakin, 70, 500)
```

Commit

```
create table cs_employees  
(  
    name  varchar(20),  
    age   int,  
    salary int  
);
```

This is also termed as a phantom problem.

Phantom Problem

Time

T1:

Begin

```
select count(*) as cnt  
from cs_cmemployees  
where age > 30
```

→ **select count(*) as cnt
from cs_cmemployees
where age > 30**

Commit

T2:

Begin

```
insert into cs_employees  
values (anakin, 70, 500)
```

Commit

```
create table cs_employees  
( name varchar(20),  
  age int,  
  salary int  
);
```

Violates our traditional definition of 2PL?
T1 cannot take a lock on something that
does not exist!

Why Phantom Problem?

Why Phantom Problem?

- We took read/write locks on existing records, and our locking scheme assumed a static system.
- But real-world databases are dynamic.
- Concurrent transactions can add new records, and our locking scheme does not consider insertions, deletions, and updates.

Solutions to Phantom Problem

Index Locking Schemes

- Index locking schemes can help eliminate phantom problem.
- Four key mechanisms in index locking schemes:
 - Key-Value Locks
 - Gap Locks
 - Key-Range Locks
 - Hierarchical Locking

Key-Value Locks

- Locks that cover a single key-value pair in an index → Standard Locks.
- For non-existent key-value pairs, we would need **virtual keys**.

B⁺-tree Leaf Nodes

6

8

10

12

Key-Value Locks

- Locks that cover a single key-value pair in an index → Standard Locks.
- For non-existent key-value pairs, we would need **virtual keys**.



Gap Locks

- Locks acquired on empty slots or gaps in the index.
- Gaps are like missing possible keys in the index.

B⁺-tree Leaf Nodes

6

8

10

12

Gap Locks

- Locks acquired on empty slots or gaps in the index.
- Gaps are like missing possible keys in the index.

B⁺-tree Leaf Nodes

6

8

10

12

**Say, we want to take a lock on
gap between 10-12.**

Gap Locks

- Locks acquired on empty slots or gaps in the index.
- Gaps are like missing possible keys in the index.
- Once a gap lock is taken, only the locking transaction can modify the gap.



**Say, we want to take a lock on
gap between 10-12.**

Key-Range Locks

- Locks that cover a key and a gap → Key-lock + Gap-lock.

B⁺-tree Leaf Nodes

6

8

10

12

Key-Range Locks

- Locks that cover a key and a gap → Key-lock + Gap-lock.

B⁺-tree Leaf Nodes

6

8

10

12

Say, we want to take a lock from 10-12.

Key-Range Locks

- Locks that cover a key and a gap → Key-lock + Gap-lock.



Say, we want to take a lock from 10-12.

Hierarchical Locks

- Allow a transaction to acquire key-range locks in a wider variety of modes.
- Remember the locking granularity matrix.

B⁺-tree Leaf Nodes

6

8

10

12

Hierarchical Locks


- Allow a transaction to acquire key-range locks in a wider variety of modes.
- Remember the locking granularity matrix.

B⁺-tree Leaf Nodes 6 8 10 12

Say, T1 wants to read all numbers from 6 to 12 (excluding 12) and update 10 to 12.

Hierarchical Locks


- Allow a transaction to acquire key-range locks in a wider variety of modes.
- Remember the locking granularity matrix.



Say, T1 wants to read all numbers from 6 to 12 (excluding 12) and update 10 to 12.

Hierarchical Locks

- Allow a transaction to acquire key-range locks in a wider variety of modes.
- Remember the locking granularity matrix.



Say, T1 wants to read all numbers from 6 to 12 (excluding 12) and update 10 to 12.

Say, T2 wants to read all numbers from 6 to 8.

Weaker Levels of Isolation

Weaker Levels of Isolation

- Serializability permits programmers to ignore concurrency issues.
- But enforcing serializability restricts opportunities for concurrency and limits performance.
- **Solution?** → Use a weaker level of consistency to improve scalability.

Weaker Levels of Isolation

- Isolation Levels control the extent to which a transaction is exposed to the actions of other concurrent transactions.
- Providing greater concurrency leads to several challenges:
 - Dirty Reads (W-R)
 - Unrepeatable Reads (R-W)
 - Lost Updates (W-W)
 - Phantom Reads

Weaker Levels of Isolation

Isolation
(High → Low)

- **Serializable:** no phantoms, all reads repeatable, no dirty reads.

Weaker Levels of Isolation

Isolation
(High → Low)

- 
- **Serializable:** no phantoms, all reads repeatable, no dirty reads.
 - **Repeatable Reads:** phantoms may happen.

Weaker Levels of Isolation

Isolation
(High → Low)

- 
- **Serializable:** no phantoms, all reads repeatable, no dirty reads.
 - **Repeatable Reads:** phantoms may happen.
 - **Read Committed:** phantoms, unrepeatable reads, and lost updates may happen.

Weaker Levels of Isolation

Isolation
(High → Low)

- **Serializable:** no phantoms, all reads repeatable, no dirty reads.
- **Repeatable Reads:** phantoms may happen.
- **Read Committed:** phantoms, unrepeatable reads, and lost updates may happen.
- **Read Uncommitted:** all anomalies may happen.

Weaker Levels of Isolation

Isolation
(High → Low)

- **Serializable:** Strong Strict 2PL with phantom protection (example through use of index locks)
- **Repeatable Reads:** Same as above, but without phantom protection.
- **Read Committed:** Same as above, but S-Locks are released immediately.
- **Read Uncommitted:** Same as above but allows dirty reads (no S-Locks).

Multi-Version Concurrency Control

Multi-Version Concurrency Control

- The DBMS maintains multiple physical versions of each record in the database.
- When a transaction reads a record, it reads the newest version that existed when the transaction started.
- When a transaction writes/updates a record, the DBMS creates a new version of that record.

Multi-Version Concurrency Control

- In MVCC,
 - Writers do not block readers.
 - Readers do not block writers.
- Read-only transactions can read from a **consistent snapshot** without acquiring locks.
- MVCC uses timestamps to determine visibility.
- MVCC provides support for **time-travel queries** if you skip doing garbage collection.
 - Run this query on the database state 2 weeks ago.


Multi-Version Concurrency Control

- How does MVCC work?

Multi-Version Concurrency Control


- How does MVCC work?
- For each transaction:
 - Create a new version on write.
 - Assign a begin timestamp and end timestamp.
 - End timestamp of previous version = begin timestamp of new version.
- Remember, we will still try to maintain isolation.
 - A concurrent transaction should not see uncommitted versions.
 - Concurrent transactions should read only committed versions.
 - View uncommitted versions as written in the local/logical space.

Example 1



Database			
Object	Value	Begin-TS	End-TS
A ₀	100		


Example 1



Object	Value	Begin-TS	End-TS
A_0	100	0	

At start of the transaction get a begin timestamp.

Example 1




$$TS(T_1) = 1$$

Database			
Object	Value	Begin-TS	End-TS
A ₀	100	0	

Next, read the latest version from the database.

Example 1




$$TS(T_1) = 1$$

$$TS(T_2) = 2$$

Object	Value	Begin-TS	End-TS
A ₀	100	0	

Begin timestamp for T2.

Example 1



$$TS(T_1) = 1$$


$$TS(T_2) = 2$$

Database

Object	Value	Begin-TS	End-TS
A ₀	100	0	
A ₁	200		

Create a new version for T2.


Example 1



Database			
Object	Value	Begin-TS	End-TS
A ₀	100	0	
A ₁	200	2	

Begin Timestamp for the new version of A is T2's begin timestamp.

Example 1




$TS(T_1) = 1$

$TS(T_2) = 2$

Database			
Object	Value	Begin-TS	End-TS
A ₀	100	0	2
A ₁	200	2	

Set end Timestamp for the previous.


Example 1



Database			
Object	Value	Begin-TS	End-TS
A ₀	100	0	2
A ₁	200	2	

But how would a future transaction know which version should it read or which is the committed version?

Example 1



$TS(T_1) = 1$


$TS(T_2) = 2$

Database			
Object	Value	Begin-TS	End-TS
A ₀	100	0	2
A ₁	200	2	

Transaction Status		
Object	Timestamp	Status
T ₁	1	Active
T ₂	2	Active

Maintain transaction status table!

Example 1




Database			
Object	Value	Begin-TS	End-TS
A ₀	100	0	2
A ₁	200	2	

Transaction Status		
Object	Timestamp	Status
T ₁	1	Active
T ₂	2	Active

Assume a new transaction T3 arrives at this moment and wants to read A, which version should it read?

Example 1



$TS(T_1) = 1$

$TS(T_2) = 2$

Database


Object	Value	Begin-TS	End-TS
A_0	100	0	2
A_1	200	2	

Transaction Status

Object	Timestamp	Status
T_1	1	Active
T_2	2	Active

If T3 is allowed to read only committed changes, then version A_0 , and if uncommitted changes are allowed then A_1 .

Example 1




Database			
Object	Value	Begin-TS	End-TS
A_0	100	0	2
A_1	200	2	

Transaction Status		
Object	Timestamp	Status
T_2	2	Active

T1 commits

Example 1



$TS(T_1) = 1$

$TS(T_2) = 2$

Database


Object	Value	Begin-TS	End-TS
A_0	100	0	2
A_1	200	2	

Transaction Status

Object	Timestamp	Status

T2 commits.

Example 2




$$TS(T_1) = 1$$

Database			
Object	Value	Begin-TS	End-TS
A ₀	100	0	

Transaction Status		
Object	Timestamp	Status
T ₁	1	Active

Example 2




$$TS(T_1) = 1$$

Database			
Object	Value	Begin-TS	End-TS
A ₀	100	0	

Transaction Status

Object	Timestamp	Status
T ₁	1	Active

Example 2




$$TS(T_1) = 1$$

Database			
Object	Value	Begin-TS	End-TS
A ₀	100	0	1
A ₁	200	1	

Transaction Status		
Object	Timestamp	Status
T ₁	1	Active

New version due to write operation.

Example 2




$TS(T_1) = 1$

$TS(T_2) = 2$

Database			
Object	Value	Begin-TS	End-TS
A ₀	100	0	1
A ₁	200	1	

Transaction Status		
Object	Timestamp	Status
T ₁	1	Active
T ₂	2	Active

Example 2



$TS(T_1) = 1$


$TS(T_2) = 2$

Database			
Object	Value	Begin-TS	End-TS
A ₀	100	0	1
A ₁	200	1	

Transaction Status		
Object	Timestamp	Status
T ₁	1	Active
T ₂	2	Active

What version will T2 read?

Example 2




Database			
Object	Value	Begin-TS	End-TS
A ₀	100	0	1
A ₁	200	1	

Transaction Status		
Object	Timestamp	Status
T ₁	1	Active
T ₂	2	Active

T₂ will read version A₀ because A₁ is not committed yet as T₁ is still Active!

Example 2




Database			
Object	Value	Begin-TS	End-TS
A ₀	100	0	1
A ₁	200	1	

Transaction Status		
Object	Timestamp	Status
T ₁	1	Active
T ₂	2	Active

What happens now?

Example 2




Database			
Object	Value	Begin-TS	End-TS
A ₀	100	0	1
A ₁	200	1	

Transaction Status		
Object	Timestamp	Status
T ₁	1	Active
T ₂	2	Active

This write operation should wait otherwise T2 will create a conflicting version.

Example 2




Database			
Object	Value	Begin-TS	End-TS
A ₀	100	0	1
A ₁	200	1	

Transaction Status		
Object	Timestamp	Status
T ₁	1	Active
T ₂	2	Active

Essentially, T2 is not allowed to create any version until T1 commits!

Example 2



$TS(T_1) = 1$

$TS(T_2) = 2$

Database


Object	Value	Begin-TS	End-TS
A_0	100	0	1
A_1	200	1	

Transaction Status

Object	Timestamp	Status
T_1	1	Active
T_2	2	Active

T1 reads version A_1 as it is local to T1.

Example 2




Database			
Object	Value	Begin-TS	End-TS
A ₀	100	0	1
A ₁	200	1	

Transaction Status		
Object	Timestamp	Status
T ₁	1	Committed
T ₂	2	Active

T1 commits and now we allow T2 to continue running.

Example 2




Database			
Object	Value	Begin-TS	End-TS
A ₀	100	0	1
A ₁	200	1	2
A ₂	200	2	

Transaction Status		
Object	Timestamp	Status
T ₁	1	Committed
T ₂	2	Active

T1 commits and now we allow T2 to continue running.

Example 2



$TS(T_1) = 1$

$TS(T_2) = 2$


Database			
Object	Value	Begin-TS	End-TS
A_0	100	0	1
A_1	200	1	2
A_2	200	2	

Transaction Status

Object	Timestamp	Status
T_1	1	Committed
T_2	2	Active

Is this serializable?

Example 2



Database			
Object	Value	Begin-TS	End-TS
A_0	100	0	1
A_1	200	1	2
A_2	200	2	

Transaction Status		
Object	Timestamp	Status
T_1	1	Committed
T_2	2	Active


Is this serializable? No, because T2 reads an older version.

Snapshot Isolation

- When a transaction starts, it sees a consistent snapshot of the database.
 - Snapshot of the database that existed when that the transaction started.
- No uncommitted writes from active transactions are visible.
- If two transactions update the same object, then the first writer does not wait.
- SI sometimes faces the **Write Skew Anomaly**.

Write Skew Anomaly


T1 → Turn all balls Black



T2 → Turn all balls White


Write Skew Anomaly

T1 → Turn all balls Black




T2 → Turn all balls White


Write Skew Anomaly




Write Skew Anomaly



Write Skew Anomaly




Write Skew Anomaly



But this is not serializable.

Write Skew Anomaly



This is serializable.

MVCC Design Decisions

- What do we need to consider while designing an MVCC scheme?

MVCC Design Decisions

- What do we need to consider while designing an MVCC scheme?
- Preventing Write Skew
- Version Storage
- Garbage Collection
- Index Management
- Deletes

Concurrency Control Protocol

- **Approach 1: Timestamp Ordering**
 - Assign transactions timestamps that determine serial order.
- **Approach 2: Optimistic Concurrency Control**
 - Three-phase protocol that we learnt in T/O lecture.
 - Use private workspace for new versions.
- **Approach 3: Two-Phase Locking**
 - Transactions acquire lock on physical version before they can read/write a logical tuple.

Version Storage

- How to store versions?

Version Storage

- **How to store versions?**
- The DBMS uses the record's pointer field to create a version chain per logical tuple.
- This allows the DBMS to find the version that is visible to a particular transaction at runtime.
- Indexes always point to the **head** of the chain.
- Different storage schemes determine where/what to store for each version.

Garbage Collection

- How to garbage collect old versions?

Garbage Collection

- How to garbage collect old versions?
- The DBMS needs to remove reclaimable physical versions from the database over time.
- No active transaction in the DBMS should be able to see a version going to be garbage collected.
 - For example: A version was created by an aborted transaction should be garbage collected.
- Two additional design decisions:
 - How to look for expired versions?
 - How to decide when it is safe to reclaim memory?